

Morten Hjorth-Jensen

Quantum mechanics for Many-particle Systems

From standard methods to quantum computing and
machine learning

July 26, 2026

Contents

Part I Mathematical basis and essential computational elements

1	Linear Algebra: Notation, Vectors, Matrices and Operations	3
1.1	Notations and conventions	3
1.2	Vectors	3
1.3	Matrices: definitions and algebraic operations	4
1.4	Special matrix types	4
1.5	Dirac bra-ket notation	5
1.6	Computational basis states and superposition	6
1.7	Pauli matrices	7
1.8	Orthonormal bases and the completeness relation	7
1.9	Unitary transformations and preservation of orthogonality	8
1.10	Tensor products	8
1.11	Spectral decomposition of Hermitian operators	9
1.12	From algebra to computation	10
1.13	Vector and matrix norms	10
1.14	Linear systems and Gaussian elimination	12
1.15	LU decomposition	14
1.16	Tridiagonal systems	16
1.17	Iterative methods for linear systems	16
1.18	The conjugate gradient method	17
1.19	Cubic spline interpolation	19
1.20	Applications: integral equations and matrix inversion	20
1.21	The algebraic eigenvalue problem	21
1.22	Similarity transformations	21
1.23	The power method	22
1.24	Householder's method for tridiagonalisation	22
1.25	Diagonalising a tridiagonal matrix	24
1.26	Lanczos' algorithm	26
1.27	Application: Schrödinger's equation through diagonalisation	29
1.28	The singular value decomposition	31
1.29	Rank, the pseudoinverse and ill-conditioned bases	33
1.30	Low-rank approximation and the Schmidt decomposition	34
1.31	Natural orbitals and occupation numbers	37
1.32	Low-rank factorisation of the two-body interaction	38
1.33	Summary and the programs	39
1.34	Exercises	40
	Warm-up exercises	40
	Exercise session, week 34	41

	Answers, week 34	42
2	From Linear Algebra to Many-Body Physics	45
2.1	The many-body Hamiltonian	45
2.2	The Pauli exclusion principle and antisymmetry	46
2.3	Determinants, and how they are actually computed	46
2.4	The Slater determinant	47
2.5	Single-particle bases and unitary transformations	48
2.6	Particle-hole notation and the Fermi level	49
2.7	The variational principle	50
2.8	Expectation values with a Slater determinant	50
2.9	The energy as a functional of the coefficients	52
2.10	The price of the two-body matrix elements	53
2.11	Evaluating Slater determinants in Monte Carlo simulations	54
2.12	Gradients, the Laplacian and the quantum force	56
2.13	Spin factorisation of the Slater determinant	57
2.14	Determinants, inverses and numerical stability	58
2.15	Quantum computing: the gate model and its linear-algebra foundation	60
2.16	The exponential-dimension problem and why it matters	60
2.17	Exercises	61
	Warm-up exercises	61
	Exercise session, week 35	63
	Answers, week 35	64
	Exercise session, week 36	67
	Answers, week 36	67
3	Second quantization	71
3.1	One-body operators in second quantization	74
3.2	Two-body operators in second quantization	76
3.3	Particle-hole formalism	79
3.4	Summarizing and defining a normal-ordered Hamiltonian	82
3.5	Normal ordering	83
3.6	Contractions	84
3.7	Two operators	85
3.8	Three operators	85
3.9	Four operators	85
3.10	Wick's theorem	87
3.11	Proof of Wick's theorem	87
3.12	Wick's generalised theorem	89
3.13	Applications: matrix elements between two-particle states	89
3.14	Operators in second quantization	91
4	Physical systems and models	95
4.1	The Lipkin model	95
4.2	The pairing model	98
4.3	The Fermi-Hubbard model	99
4.4	The Calogero-Sutherland model	101
4.5	What the four models have in common	102
4.6	Exercises	103
	Exercise 1: the Lipkin model	103
	Exercise 2: the pairing model	104
	Exercise 3: the Hubbard model	105
	Exercise 4: the Calogero ground state	106
	Exercise 5: numerical explorations	108

Part II Wave function based methods

5	Full configuration interaction theory	111
5.1	Slater determinants as a basis	111
5.2	The configuration-interaction expansion	112
5.3	The eigenvalue problem	113
5.4	The structure of the Hamiltonian matrix	113
5.5	The exponential wall	114
5.6	Full CI for the pairing model	116
5.7	Full CI for the Hubbard model	116
5.8	Truncating the expansion	117
5.9	Size consistency	118
5.10	Truncations and the other many-body methods	118
5.11	Practical considerations	119
5.12	Summary	120
5.13	Exercises	120
	Exercise 1: counting determinants	120
	Exercise 2: which blocks vanish	121
	Exercise 3: full CI for the pairing model	122
	Exercise 4: truncations and size consistency	123
	Exercise 5: the Hubbard ring by full CI	124
	Exercise 6: the exponential wall in practice	125

Part III Stochastic methods**Part IV Machine Learning based approaches****Part V Quantum Computing**

References	133
------------	-----

Index	135
--------------	-----

Part I
Mathematical basis and essential computational
elements

Chapter 1

Linear Algebra: Notation, Vectors, Matrices and Operations

1.1 Notations and conventions

Throughout these notes vectors, matrices and higher-order tensors are always set in boldface. Vectors are denoted by lower-case letters ($\mathbf{x}$, $\mathbf{v}$, ...) and matrices and tensors by upper-case letters ($\mathbf{A}$, $\mathbf{U}$, ...).

Unless stated otherwise the elements v_i of a vector $\mathbf{v}$ are assumed real. A real vector of length n satisfies $\mathbf{x} \in \mathbb{R}^n$; for complex vectors we write $\mathbf{x} \in \mathbb{C}^n$. A real $n \times n$ matrix satisfies $\mathbf{A} \in \mathbb{R}^{n \times n}$. Index counting starts at zero: matrix elements run from a_{00} to $a_{n-1,n-1}$.

We use the standard shorthand symbols: $\forall$ (for all), $\Rightarrow$ (implies), $\equiv$ (equivalent / defined as), with $\mathbb{R}$, $\mathbb{I}$, $\mathbb{C}$ for the real, integer and complex number fields.

1.2 Vectors

A column vector of length n and its transpose are

$$\mathbf{x} = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_{n-1} \end{bmatrix}, \quad \mathbf{x}^T = [x_0 \ x_1 \ \cdots \ x_{n-1}].$$

For complex vectors the *Hermitian conjugate* (adjoint) is

$$\mathbf{x}^\dagger = [x_0^* \ x_1^* \ \cdots \ x_{n-1}^*].$$

The *inner (dot) product* is the scalar

$$\mathbf{x}^T \mathbf{x} = \sum_{i=0}^{n-1} x_i x_i = x_0^2 + x_1^2 + \cdots + x_{n-1}^2,$$

and for complex vectors $\mathbf{x}^\dagger \mathbf{x} = \sum_{i=0}^{n-1} x_i^* x_i = \|\mathbf{x}\|^2$.

The *outer product* of two real vectors $\mathbf{y}$ (length n) and $\mathbf{z}$ (length m) is the $n \times m$ matrix

$$\mathbf{y} \mathbf{z}^T \implies (\mathbf{y} \mathbf{z}^T)_{ij} = y_i z_j.$$

For complex vectors one writes $\mathbf{y} \mathbf{z}^\dagger$ with $(\mathbf{y} \mathbf{z}^\dagger)_{ij} = y_i z_j^*$.

Important vector operations:

Addition / subtraction: $\mathbf{x} = \mathbf{y} \pm \mathbf{z} \Rightarrow x_i = y_i \pm z_i$.

Scalar multiplication: $\mathbf{x} = \gamma \mathbf{y} \Rightarrow x_i = \gamma y_i$.

Hadamard (element-wise) product: $\mathbf{x} = \mathbf{y} \circ \mathbf{z} \Rightarrow x_i = y_i z_i$.

1.3 Matrices: definitions and algebraic operations

A general $n \times n$ matrix and its column-vector representation:

$$\mathbf{A} = \begin{bmatrix} a_{00} & a_{01} & \cdots & a_{0,n-1} \\ a_{10} & a_{11} & \cdots & a_{1,n-1} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n-1,0} & \cdots & \cdots & a_{n-1,n-1} \end{bmatrix} = [\mathbf{a}_0 \ \mathbf{a}_1 \ \cdots \ \mathbf{a}_{n-1}].$$

The fundamental algebraic operations are:

Scalar mult.: $\mathbf{A} = \gamma \mathbf{B} \Rightarrow a_{ij} = \gamma b_{ij}$.

Addition: $\mathbf{A} = \mathbf{B} \pm \mathbf{C} \Rightarrow a_{ij} = b_{ij} \pm c_{ij}$.

Matrix-vector product: $\mathbf{y} = \mathbf{A} \mathbf{x} \Rightarrow y_i = \sum_{j=0}^{n-1} a_{ij} x_j$.

Matrix-matrix product: $\mathbf{A} = \mathbf{B} \mathbf{C} \Rightarrow a_{ij} = \sum_{k=0}^{n-1} b_{ik} c_{kj}$.

Transpose: $\mathbf{A} = \mathbf{B}^T \Rightarrow a_{ij} = b_{ji}$.

Inverse.

If $\mathbf{A}$ is square and nonsingular, its inverse satisfies $\mathbf{A}^{-1} \mathbf{A} = \mathbf{I}$, where $\mathbf{I}$ is the identity matrix.

Hermitian conjugate.

The Hermitian conjugate $\mathbf{A}^\dagger$ is obtained by transposing $\mathbf{A}$ and complex-conjugating every element: $(\mathbf{A}^\dagger)_{ij} = a_{ji}^*$.

Nonsingularity: equivalent characterisations.

For an $n \times n$ matrix $\mathbf{A}$ the following properties are all equivalent: (i) $\mathbf{A}^{-1}$ exists; (ii) $\mathbf{A} \mathbf{x} = \mathbf{0}$ implies $\mathbf{x} = \mathbf{0}$; (iii) the rows of $\mathbf{A}$ form a basis of $\mathbb{C}^n$; (iv) the columns of $\mathbf{A}$ form a basis of $\mathbb{C}^n$; (v) $\mathbf{A}$ is a product of elementary matrices; (vi) zero is not an eigenvalue of $\mathbf{A}$.

1.4 Special matrix types

Table 1.1 lists the matrix types that appear most frequently in quantum mechanics and many-body physics.

Structured (sparse) matrices important for numerical many-body work:

- **Diagonal:** $a_{ij} = 0$ for $i \neq j$.

Name	Defining relation	Element condition
Symmetric	$\mathbf{A} = \mathbf{A}^T$	$a_{ij} = a_{ji}$
Real orthogonal	$\mathbf{A} = (\mathbf{A}^T)^{-1}$	$\sum_k a_{ik} a_{jk} = \delta_{ij}$
Hermitian	$\mathbf{A} = \mathbf{A}^\dagger$	$a_{ij} = a_{ji}^*$
Unitary	$\mathbf{A} = (\mathbf{A}^\dagger)^{-1}$	$\sum_k a_{ik} a_{jk}^* = \delta_{ij}$
Real matrix	$\mathbf{A} = \mathbf{A}^*$	$a_{ij} = a_{ij}^*$

Table 1.1 Important special matrix types and their defining properties.

- **Upper/lower triangular:** $a_{ij} = 0$ for $i > j$ / $i < j$.
- **Upper/lower Hessenberg:** $a_{ij} = 0$ for $i > j + 1$ / $i < j + 1$.
- **Tridiagonal:** $a_{ij} = 0$ for $|i - j| > 1$.
- **Banded with bandwidth p :** $a_{ij} = 0$ for $|i - j| > p$.
- Block upper/lower triangular matrices, and more.

Many-body connection. Hamiltonian matrices in shell-model (full configuration-interaction, FCI) calculations are typically large, sparse, and often banded. Exploiting sparsity is essential to diagonalise them for systems with many single-particle states.

1.5 Dirac bra-ket notation

In quantum mechanics it is conventional to write state vectors using Dirac's bra-ket formalism. A ket $|x\rangle$ is a column vector

$$|x\rangle = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_{n-1} \end{bmatrix},$$

and the corresponding bra $\langle x|$ is the row vector (dual)

$$\langle x| = [x_0^* \ x_1^* \ \cdots \ x_{n-1}^*],$$

so $|x\rangle^\dagger = \langle x|$. The *inner product* (a complex number) is

$$\langle y|x\rangle = \sum_{i=0}^{n-1} y_i^* x_i.$$

Note that $\langle x|x\rangle \geq 0$ is always real and non-negative, and $\langle y|x\rangle^* = \langle x|y\rangle$.

The *outer product* $|x\rangle\langle y|$ is an $n \times n$ matrix with elements $(|x\rangle\langle y|)_{ij} = x_i y_j^*$:

$$|x\rangle\langle y| = \begin{bmatrix} x_0 y_0^* & x_0 y_1^* & \cdots & x_0 y_{n-1}^* \\ x_1 y_0^* & x_1 y_1^* & \cdots & x_1 y_{n-1}^* \\ \vdots & \vdots & \ddots & \vdots \\ x_{n-1} y_0^* & \cdots & \cdots & x_{n-1} y_{n-1}^* \end{bmatrix}.$$

Worked example.

Let $|x\rangle = \begin{bmatrix} 1-i \\ 2+i \end{bmatrix}$. Then

$$\langle x|x\rangle = (1+i)(1-i) + (2-i)(2+i) = 2+5 = 7 \in \mathbb{R}.$$

For two different vectors $|x\rangle = \begin{bmatrix} -1 \\ 2i \\ 1 \end{bmatrix}$, $|y\rangle = \begin{bmatrix} 1 \\ 0 \\ i \end{bmatrix}$: $\langle x|y\rangle = -1+i \neq \langle y|x\rangle = -1-i$, confirming the rule $\langle x|y\rangle^* = \langle y|x\rangle$.

1.6 Computational basis states and superposition

A *two-level system* (qubit) is described by two orthonormal basis states, which we label

$$|0\rangle = \begin{bmatrix} 1 \\ 0 \end{bmatrix}, \quad |1\rangle = \begin{bmatrix} 0 \\ 1 \end{bmatrix}, \quad \langle 0|1\rangle = 0, \quad \langle 0|0\rangle = \langle 1|1\rangle = 1.$$

These are the *computational basis* states. A general state is a *superposition*

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle,$$

where $\alpha = \langle 0|\psi\rangle$ and $\beta = \langle 1|\psi\rangle$ are the probability amplitudes, satisfying $|\alpha|^2 + |\beta|^2 = 1$.

The identity operator decomposes as a sum of outer products:

$$\mathbf{I} = \sum_{i=0}^1 |i\rangle\langle i| = \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix} + \begin{bmatrix} 0 & 0 \\ 0 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}.$$

Projection operators.

The outer products

$$\mathbf{P} = |0\rangle\langle 0| = \begin{bmatrix} 1 & 0 \\ 0 & 0 \end{bmatrix}, \quad \mathbf{Q} = |1\rangle\langle 1| = \begin{bmatrix} 0 & 0 \\ 0 & 1 \end{bmatrix}$$

are *projection operators*. They are *idempotent* ($\mathbf{P}^2 = \mathbf{P}$, $\mathbf{Q}^2 = \mathbf{Q}$), commute with each other ($[\mathbf{P}, \mathbf{Q}] = 0$), and satisfy $\mathbf{PQ} = 0$, $\mathbf{P} + \mathbf{Q} = \mathbf{I}$. Acting on a superposition state:

$$\mathbf{P}|\psi\rangle = \alpha|0\rangle, \quad \mathbf{Q}|\psi\rangle = \beta|1\rangle.$$

The measurement postulate states that a measurement on $|\psi\rangle$ yields $|0\rangle$ with probability $|\alpha|^2$ and $|1\rangle$ with probability $|\beta|^2$.

Many-body connection. In many-body theory, projection operators $\hat{P}$ (onto occupied states) and $\hat{Q}$ (onto unoccupied / virtual states) define the particle-hole decomposition that is fundamental to perturbation theory, Hartree-Fock theory, and coupled-cluster methods.

Density matrix.

The density matrix (density operator) of a pure state $|\psi\rangle$ is the outer product

$$\rho = |\psi\rangle\langle\psi| = \begin{bmatrix} \alpha\alpha^* & \alpha\beta^* \\ \beta\alpha^* & \beta\beta^* \end{bmatrix},$$

with $\text{tr} \rho = |\alpha|^2 + |\beta|^2 = 1$. Density matrices generalise to mixed states and subsystems; they are central to quantum information theory and to reduced-density-matrix many-body methods.

1.7 Pauli matrices

The three Pauli matrices are

$$\sigma_x = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}, \quad \sigma_y = \begin{bmatrix} 0 & -i \\ i & 0 \end{bmatrix}, \quad \sigma_z = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}.$$

Key properties (easy to verify directly):

- **Involutory:** $\sigma_x^2 = \sigma_y^2 = \sigma_z^2 = I$.
- **Unitary:** $\sigma_\alpha^\dagger = \sigma_\alpha^{-1}$ (Hermitian conjugate equals inverse).
- **Hermitian:** $\sigma_\alpha^\dagger = \sigma_\alpha$.
- $\det(\sigma_\alpha) = -1$ for each α .
- **Commutation:** $[\sigma_x, \sigma_y] = 2i\sigma_z$ (and cyclic permutations).

The Pauli matrices act on the computational basis as

$$\sigma_x|0\rangle = |1\rangle, \quad \sigma_x|1\rangle = |0\rangle; \quad \sigma_z|0\rangle = +|0\rangle, \quad \sigma_z|1\rangle = -|1\rangle.$$

Acting on a superposition $|\psi\rangle = \alpha|0\rangle + \beta|1\rangle$:

$$\sigma_x|\psi\rangle = \beta|0\rangle + \alpha|1\rangle, \quad \sigma_y|\psi\rangle = i\beta|0\rangle - i\alpha|1\rangle.$$

Quantum computing connection. The Pauli matrices are the elementary single-qubit gates X, Y, Z . Together with the Hadamard gate $H = (\sigma_x + \sigma_z)/\sqrt{2}$ and the phase gate T they form a universal single-qubit gate set.

Many-body connection. Via the Jordan-Wigner transformation, the creation/annihilation operators of second quantisation are mapped directly onto products (strings) of Pauli matrices, enabling the representation of fermionic Hamiltonians on qubit registers.

1.8 Orthonormal bases and the completeness relation

A set of states $\{|\phi_i\rangle\}_{i=0}^{n-1}$ is an *orthonormal basis* (ONB) of an n -dimensional Hilbert space if

$$\langle\phi_j|\phi_i\rangle = \delta_{ij} \quad \text{and} \quad \sum_{i=0}^{n-1} |\phi_i\rangle\langle\phi_i| = I \quad (\text{completeness / resolution of the identity}).$$

Any state can then be expanded as $|\psi\rangle = \sum_i \langle\phi_i|\psi\rangle |\phi_i\rangle$, with $\langle\phi_i|\psi\rangle$ the projection (overlap) of $|\psi\rangle$ onto $|\phi_i\rangle$.

Many-body connection. The natural ONB in many-body physics is the set of single-particle eigenstates of the one-body Hamiltonian:

$$\hat{h}_0|\phi_i\rangle = \epsilon_i|\phi_i\rangle, \quad \phi_i(\mathbf{r}) = \langle\mathbf{r}|\phi_i\rangle.$$

This single-particle basis $\{|\phi_i\rangle\}$ is the starting point for virtually all many-body methods: Hartree-Fock, MBPT, FCI, coupled-cluster, and quantum Monte Carlo.

1.9 Unitary transformations and preservation of orthogonality

An $n \times n$ matrix $\mathbf{U}$ is *unitary* if $\mathbf{U}^\dagger \mathbf{U} = \mathbf{U} \mathbf{U}^\dagger = \mathbf{I}$, which means $\mathbf{U}^{-1} = \mathbf{U}^\dagger$.

Orthogonality is preserved.

If $\{\mathbf{v}_i\}$ is an orthonormal set, $\mathbf{v}_j^T \mathbf{v}_i = \delta_{ij}$, then the transformed set $\mathbf{w}_i = \mathbf{U} \mathbf{v}_i$ is also orthonormal:

$$\mathbf{w}_j^T \mathbf{w}_i = (\mathbf{U} \mathbf{v}_j)^T \mathbf{U} \mathbf{v}_i = \mathbf{v}_j^T \underbrace{\mathbf{U}^T \mathbf{U}}_{=\mathbf{I}} \mathbf{v}_i = \mathbf{v}_j^T \mathbf{v}_i = \delta_{ij}.$$

In Dirac notation, if $|\psi_i\rangle = \sum_j u_{ij} |\phi_j\rangle$ with (u_{ij}) unitary, then $\langle \psi_j | \psi_i \rangle = \delta_{ij}$.

Why only unitary transformations?

For the inner product (and hence the norm) to be preserved under $|\phi\rangle = \mathbf{U} |\psi\rangle$ we need

$$\langle \phi | \phi \rangle = \langle \psi | \mathbf{U}^\dagger \mathbf{U} | \psi \rangle = \langle \psi | \psi \rangle = 1,$$

which requires $\mathbf{U}^\dagger \mathbf{U} = \mathbf{I}$. Since probabilities must sum to unity, this is a fundamental physical requirement. Unitary transformations are rotations (and reflections) in Hilbert space. The fact that $\mathbf{U}^{-1} = \mathbf{U}^\dagger$ also means that every unitary transformation is perfectly reversible.

Example: two-state rotation.

Given $|\psi\rangle = \alpha|0\rangle + \beta|1\rangle$ and a unitary $\mathbf{U}$, the rotated state $|\phi\rangle = \gamma|0\rangle + \delta|1\rangle = \mathbf{U} |\psi\rangle$ satisfies

$$\begin{bmatrix} \gamma \\ \delta \end{bmatrix} = \begin{bmatrix} u_{00} & u_{01} \\ u_{10} & u_{11} \end{bmatrix} \begin{bmatrix} \alpha \\ \beta \end{bmatrix}.$$

Since $|\alpha|^2 + |\beta|^2 = 1$ and $\mathbf{U}$ is unitary, $|\gamma|^2 + |\delta|^2 = 1$ automatically.

Many-body connection. In Hartree-Fock theory the self-consistent field (SCF) procedure generates a sequence of unitary rotations of the single-particle basis. The total HF energy is invariant under unitary rotations within the occupied or virtual subspace separately, which is the Thouless theorem.

Quantum computing connection. Every quantum gate is a unitary matrix. The evolution of a closed quantum system is described entirely by unitary operations; this is the quantum analogue of classical reversibility.

1.10 Tensor products

The *tensor (Kronecker) product* of two vectors $|x\rangle$ (length n) and $|y\rangle$ (length m) is the vector of length nm :

$$|x\rangle \otimes |y\rangle = |xy\rangle = \begin{bmatrix} x_0 y_0 \\ x_0 y_1 \\ x_1 y_0 \\ x_1 y_1 \end{bmatrix} \quad (\text{for } n = m = 2).$$

For the computational basis:

$$|0\rangle \otimes |0\rangle = |00\rangle = \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \quad |0\rangle \otimes |1\rangle = |01\rangle = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}, \quad |1\rangle \otimes |0\rangle = |10\rangle = \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}, \quad |1\rangle \otimes |1\rangle = |11\rangle = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}.$$

These four states form an ONB for the two-qubit Hilbert space $\mathbb{C}^4 = \mathbb{C}^2 \otimes \mathbb{C}^2$.

The most general state of N qubits is

$$|\Psi\rangle = \sum_{i_1 i_2 \dots i_N} \Psi_{i_1 i_2 \dots i_N} |i_1\rangle \otimes |i_2\rangle \otimes \dots \otimes |i_N\rangle,$$

requiring 2^N complex amplitudes. For $N \approx 300$ this exceeds the number of atoms in the observable universe — this is the *exponential complexity* at the heart of quantum many-body physics.

Many-body connection. The full N -particle Hilbert space is a tensor product of N single-particle spaces. The FCI method works in this exponentially large space; all approximate methods (HF, MBPT, coupled cluster, DMRG, ...) are strategies to identify the relevant low-dimensional submanifold.

Quantum computing connection. A quantum computer with N qubits *inherently* represents a 2^N -dimensional state vector. This is both the source of its potential power and the reason that quantum simulation of many-body systems is a leading application.

1.11 Spectral decomposition of Hermitian operators

A Hermitian operator $\mathbf{A} = \mathbf{A}^\dagger$ has the following key properties:

1. All eigenvalues λ_i are *real*.
2. Eigenvectors belonging to distinct eigenvalues are *orthogonal*.
3. $\mathbf{A}$ is diagonalisable (it is a *normal* operator, $[\mathbf{A}, \mathbf{A}^\dagger] = 0$).

The eigenvalue equation $\mathbf{A}|\psi\rangle = \lambda|\psi\rangle$ has non-trivial solutions only when

$$\det(\mathbf{A} - \lambda \mathbf{I}) = 0 \quad (\text{characteristic polynomial}).$$

Spectral decomposition.

Let $\{|i\rangle\}$ be the eigenbasis of $\mathbf{A}$ with eigenvalues λ_i . Any state expands as $|\psi\rangle = \sum_i \alpha_i |i\rangle$. Using projection operators $\mathbf{P}_i = |i\rangle\langle i|$ we obtain:

$$\mathbf{A}|\psi\rangle = \sum_i \alpha_i \lambda_i |i\rangle = \sum_i \lambda_i \mathbf{P}_i |\psi\rangle,$$

from which, since this holds for any $|\psi\rangle$,

$$\boxed{\mathbf{A} = \sum_{i=0}^{n-1} \lambda_i |i\rangle\langle i| = \sum_{i=0}^{n-1} \lambda_i \mathbf{P}_i.}$$

This is the *spectral decomposition*: the operator is expressed as a weighted sum of its projection operators.

Example for a 2×2 operator.

If $\mathbf{A}$ has eigenstates $|\psi_a\rangle = \alpha_0|0\rangle + \alpha_1|1\rangle$ and $|\psi_b\rangle = \beta_0|0\rangle + \beta_1|1\rangle$ with eigenvalues λ_a, λ_b , then

$$\mathbf{A} = \lambda_a \begin{bmatrix} |\alpha_0|^2 & \alpha_0\alpha_1^* \\ \alpha_1\alpha_0^* & |\alpha_1|^2 \end{bmatrix} + \lambda_b \begin{bmatrix} |\beta_0|^2 & \beta_0\beta_1^* \\ \beta_1\beta_0^* & |\beta_1|^2 \end{bmatrix}.$$

Many-body connection. The spectral decomposition is used every time one diagonalises a Hamiltonian matrix. In FCI the full N -body Hamiltonian is diagonalised in the many-determinant basis; in Hartree-Fock the Fock matrix (a single-particle Hermitian operator) is diagonalised to find the optimal orbital basis.

Quantum computing connection. Any unitary gate $\mathbf{U}$ can be written as $\mathbf{U} = e^{i\mathbf{H}}$ for a Hermitian generator $\mathbf{H}$. The spectral decomposition of $\mathbf{H}$ gives $\mathbf{U} = \sum_i e^{i\lambda_i} |i\rangle\langle i|$. This exponential relationship connects the physical Hamiltonian to the time-evolution operator $\hat{U} = e^{-i\hat{H}t/\hbar}$ and is the basis of the variational quantum eigensolver (VQE) and unitary coupled-cluster (UCC) ansatz.

1.12 From algebra to computation

Everything said so far is exact algebra. The moment we hand a matrix to a computer, however, two new questions arise. The first is one of *accuracy*: floating-point numbers carry a finite number of digits, and we need a way of measuring how large the resulting error is. The second is one of *cost*: a many-body Hamiltonian matrix may have a dimension of 10^6 or more, and an algorithm whose work grows as n^3 is then simply not available to us.

The remainder of this chapter is devoted to these two questions. We first introduce the norms with which errors are measured, then develop the direct (elimination) methods for linear systems, the iterative alternatives, and finally the algorithms for the algebraic eigenvalue problem $\mathbf{A}|x\rangle = \lambda|x\rangle$, which is the computational problem at the heart of essentially every method in this book. All programs are written in Python; the complete, runnable versions are collected in the `BookMaterial/Programs` directory.

Throughout we keep the convention of Section 1.1 that indices run from 0 to $n-1$, so that the mathematical formulae and the Python code carry the same indices.

1.13 Vector and matrix norms

A class of vector norms is given by the so-called p -norms

$$\|\mathbf{x}\|_p = (|x_0|^p + |x_1|^p + \cdots + |x_{n-1}|^p)^{1/p}, \quad p \geq 1. \quad (1.1)$$

The three cases we shall need are $p = 1$, $p = 2$ and the limit $p \rightarrow \infty$,

$$\|\mathbf{x}\|_1 = |x_0| + |x_1| + \cdots + |x_{n-1}|, \quad (1.2)$$

$$\|\mathbf{x}\|_2 = (|x_0|^2 + \cdots + |x_{n-1}|^2)^{1/2} = (\mathbf{x}^\dagger \mathbf{x})^{1/2}, \quad (1.3)$$

$$\|\mathbf{x}\|_\infty = \max_{0 \leq i \leq n-1} |x_i|. \quad (1.4)$$

The 2-norm is the one we have already been using implicitly: it is the square root of the inner product of Section 1.2, and in Dirac notation $\|\mathbf{x}\|_2^2 = \langle x|x \rangle$. A quantum state is normalised precisely when its 2-norm equals unity.

Two inequalities follow from these definitions and will be used repeatedly. The first is the Cauchy-Schwarz inequality, valid for any two vectors in an inner product space,

$$|\mathbf{x}^\dagger \mathbf{y}| \leq \|\mathbf{x}\|_2 \|\mathbf{y}\|_2, \quad (1.5)$$

with equality only when $\mathbf{x}$ and $\mathbf{y}$ are linearly dependent. In Dirac notation this reads $|\langle \mathbf{x} | \mathbf{y} \rangle|^2 \leq \langle \mathbf{x} | \mathbf{x} \rangle \langle \mathbf{y} | \mathbf{y} \rangle$, and it is the inequality behind the variational principle: the overlap of a trial state with the exact ground state can never exceed the product of their norms. The second is the triangle inequality

$$\|\mathbf{x} + \mathbf{y}\|_2 \leq \|\mathbf{x}\|_2 + \|\mathbf{y}\|_2, \quad (1.6)$$

which follows from Eq. (1.5). Proofs may be found in Golub and Van Loan [1].

For matrices the most frequently used norms are the Frobenius norm

$$\|\mathbf{A}\|_F = \sqrt{\sum_{i=0}^{m-1} \sum_{j=0}^{n-1} |a_{ij}|^2}, \quad (1.7)$$

which is simply the 2-norm of the matrix read as a long vector, and the induced or operator p -norms

$$\|\mathbf{A}\|_p = \max_{\mathbf{x} \neq 0} \frac{\|\mathbf{A}\mathbf{x}\|_p}{\|\mathbf{x}\|_p}. \quad (1.8)$$

The induced norm measures the largest factor by which $\mathbf{A}$ can stretch a vector. For a Hermitian matrix $\|\mathbf{A}\|_2$ is the largest absolute eigenvalue, a fact which follows immediately from the spectral decomposition of Section 1.11. A unitary matrix has $\|\mathbf{U}\|_2 = 1$; it rotates but never stretches, which is Section 1.9 restated in the language of norms.

Relative error.

Let $fl(\mathbf{x})$ denote the machine representation of a vector $\mathbf{x} \neq 0$. The relative error is

$$\varepsilon = \frac{\|fl(\mathbf{x}) - \mathbf{x}\|}{\|\mathbf{x}\|}, \quad (1.9)$$

and if we use the ∞ -norm the statement $\|fl(\mathbf{x}) - \mathbf{x}\|_\infty / \|\mathbf{x}\|_\infty \approx 10^{-l}$ says that the largest component of $fl(\mathbf{x})$ carries roughly l correct significant digits.

The condition number.

Suppose we solve $\mathbf{A}\mathbf{x} = \mathbf{b}$ but, because of rounding, actually solve a slightly perturbed problem with right-hand side $\mathbf{b} + \delta\mathbf{b}$. The resulting error $\delta\mathbf{x}$ satisfies $\mathbf{A}\delta\mathbf{x} = \delta\mathbf{b}$, hence $\|\delta\mathbf{x}\| \leq \|\mathbf{A}^{-1}\| \|\delta\mathbf{b}\|$, while $\|\mathbf{b}\| \leq \|\mathbf{A}\| \|\mathbf{x}\|$. Combining the two gives

$$\frac{\|\delta\mathbf{x}\|}{\|\mathbf{x}\|} \leq \kappa(\mathbf{A}) \frac{\|\delta\mathbf{b}\|}{\|\mathbf{b}\|}, \quad \kappa(\mathbf{A}) = \|\mathbf{A}\| \|\mathbf{A}^{-1}\|. \quad (1.10)$$

The quantity $\kappa(\mathbf{A})$ is the *condition number*. It tells us how much an input error can be amplified, and it is a property of the problem, not of the algorithm: no amount of cleverness can rescue an ill-conditioned system. In the 2-norm and for a Hermitian matrix

$$\kappa_2(\mathbf{A}) = \frac{\max_i |\lambda_i|}{\min_i |\lambda_i|}, \quad (1.11)$$

so a matrix is ill-conditioned exactly when it has eigenvalues of very different magnitude, that is, when it is close to being singular. As a rule of thumb, if $\kappa \approx 10^k$ one loses about k significant digits.

Many-body connection. Ill-conditioning is not an academic worry in many-body physics. Overlap matrices of non-orthogonal single-particle bases, such as the Gaussian basis sets of quantum chemistry, become nearly singular as the basis approaches linear dependence, and their condition numbers can reach 10^{10} and beyond. The standard cure is to diagonalise the overlap matrix and discard the eigenvectors with the smallest eigenvalues, which is the canonical orthogonalisation procedure met again in Hartree-Fock theory.

1.14 Linear systems and Gaussian elimination

We now turn to the solution of

$$\mathbf{A}\mathbf{x} = \mathbf{w}, \quad (1.12)$$

with $\mathbf{A}$ non-singular. Before describing the algorithms it is worth seeing why such systems appear at all.

Where linear systems come from.

Consider the boundary value problem

$$-\frac{d^2u(x)}{dx^2} = f(x), \quad x \in (a, b), \quad u(a) = u(b) = 0. \quad (1.13)$$

We subdivide the interval into $n+1$ subintervals by setting $x_i = a + ih$ with step $h = (b-a)/(n+1)$, and replace the second derivative by the three-point formula

$$u''(x_i) \approx \frac{u_{i+1} - 2u_i + u_{i-1}}{h^2} + \mathcal{O}(h^2). \quad (1.14)$$

For the interior points $i = 0, 1, \dots, n-1$ this turns Eq. (1.13) into

$$\frac{-u_{i+1} + 2u_i - u_{i-1}}{h^2} = f_i, \quad (1.15)$$

where the boundary values u_{-1} and u_n are known and equal to zero. Collecting the coefficients we obtain $\mathbf{A}\mathbf{u} = \mathbf{f}$ with the tridiagonal matrix

$$\mathbf{A} = \frac{1}{h^2} \begin{bmatrix} 2 & -1 & & & \\ -1 & 2 & -1 & & \\ & -1 & 2 & \ddots & \\ & & \ddots & \ddots & -1 \\ & & & -1 & 2 \end{bmatrix}. \quad (1.16)$$

A differential equation has become a matrix equation. This is the single most important lesson of this chapter, and it will be repeated in a many-body setting in Section 1.27: discretising a continuous problem, or expanding a state in a finite basis, always produces a matrix problem.

The elimination.

The idea of Gaussian elimination is to use the first equation to eliminate x_0 from the remaining $n-1$ equations, then the new second equation to eliminate x_1 from the remaining $n-2$, and so on. After $n-1$ such steps we are left with an upper triangular system

$$U\mathbf{x} = \mathbf{y}, \quad (1.17)$$

which is solved from the bottom upwards by *backward substitution*

$$x_m = \frac{1}{u_{mm}} \left(y_m - \sum_{k=m+1}^{n-1} u_{mk} x_k \right), \quad m = n-1, n-2, \dots, 0. \quad (1.18)$$

To eliminate x_0 we multiply the first equation by a_{j0}/a_{00} and subtract it from equation j , for $j = 1, \dots, n-1$. The elements of the remaining $(n-1) \times (n-1)$ block are updated according to

$$a_{jk}^{(1)} = a_{jk} - \frac{a_{j0} a_{0k}}{a_{00}}, \quad j, k = 1, \dots, n-1, \quad (1.19)$$

and the same formula, applied to successively smaller blocks, defines the whole algorithm. Counting operations, the elimination requires $2n^3/3 + \mathcal{O}(n^2)$ floating point operations, while the backward substitution costs only $\mathcal{O}(n^2)$. The elimination dominates, and this cubic scaling is the reason why direct methods become unusable for the large matrices of Section 1.26.

Pivoting.

The derivation assumed $a_{00} \neq 0$, and more generally that no pivot element vanishes. Even a small pivot is dangerous, since dividing by it amplifies rounding errors. The cure is *partial pivoting*: before eliminating column k we search the column for the element of largest absolute value and interchange rows so that this element becomes the pivot. Row interchanges do not alter the solution, and they keep all multipliers bounded by unity. In practice one always pivots.

The following class implements the algorithm.

```
class GaussianElimination(DirectSolver):
    """Gaussian elimination with partial pivoting."""

    def solve(self, b):
        n = self.n
        M = self.A.copy()           # the elimination destroys its input
        y = np.asarray(b, dtype=float).copy()

        for k in range(n - 1):      # forward elimination
            # partial pivoting: use the largest element in the column
            p = k + np.argmax(np.abs(M[k:, k]))
            if abs(M[p, k]) < 1.0e-14:
                raise np.linalg.LinAlgError("matrix is singular")
            if p != k:
                M[[k, p]] = M[[p, k]]
                y[k], y[p] = y[p], y[k]
            for i in range(k + 1, n):
                factor = M[i, k] / M[k, k]
                M[i, k:] -= factor * M[k, k:]
                y[i] -= factor * y[k]

        return self._back_substitute(M, y)

    @staticmethod
    def _back_substitute(U, y):
        """Solve  $Ux = y$  for an upper triangular  $U$ ."""
        n = U.shape[0]
        x = np.zeros(n)
        for m in range(n - 1, -1, -1):
            x[m] = (y[m] - U[m, m+1:] @ x[m+1:]) / U[m, m]
        return x
```

1.15 LU decomposition

Gaussian elimination as written above throws away the multipliers once it has used them. If we keep them instead, we obtain a factorisation of $\mathbf{A}$ itself, and this is what makes the method reusable. The LU decomposition writes

$$\mathbf{A} = \mathbf{L}\mathbf{U}, \quad (1.20)$$

with $\mathbf{L}$ unit lower triangular and $\mathbf{U}$ upper triangular; for a 4×4 matrix

$$\begin{bmatrix} a_{00} & a_{01} & a_{02} & a_{03} \\ a_{10} & a_{11} & a_{12} & a_{13} \\ a_{20} & a_{21} & a_{22} & a_{23} \\ a_{30} & a_{31} & a_{32} & a_{33} \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ l_{10} & 1 & 0 & 0 \\ l_{20} & l_{21} & 1 & 0 \\ l_{30} & l_{31} & l_{32} & 1 \end{bmatrix} \begin{bmatrix} u_{00} & u_{01} & u_{02} & u_{03} \\ 0 & u_{11} & u_{12} & u_{13} \\ 0 & 0 & u_{22} & u_{23} \\ 0 & 0 & 0 & u_{33} \end{bmatrix}. \quad (1.21)$$

The factorisation exists whenever $\mathbf{A}$ is non-singular and, with the normalisation $l_{ii} = 1$, it is unique. Multiplying out the first column gives $a_{00} = u_{00}$ and $a_{j0} = l_{j0}u_{00}$, which determines u_{00} and the multipliers l_{j0} . The second column then gives u_{01} , u_{11} and l_{j1} , and so on: at every stage the unknowns are determined by quantities already computed. In practice $\mathbf{L}$ and $\mathbf{U}$ are stored in the same array as $\mathbf{A}$, since the unit diagonal of $\mathbf{L}$ need not be kept.

With partial pivoting the factorisation reads $\mathbf{PA} = \mathbf{LU}$ with $\mathbf{P}$ a permutation matrix, which in the program is recorded as a permutation array rather than as a matrix.

The determinant.

Because $\det(\mathbf{L}) = 1$ and the determinant of a triangular matrix is the product of its diagonal elements,

$$|\mathbf{A}| = \pm u_{00}u_{11} \cdots u_{n-1,n-1}, \quad (1.22)$$

where the sign is $(-1)^r$ with r the number of row interchanges. This is the only sensible way of computing a determinant numerically; expanding in minors would cost $\mathcal{O}(n!)$ operations.

Solving the system.

Writing $\mathbf{Ax} = \mathbf{LUx} = \mathbf{w}$ and introducing the intermediate vector $\mathbf{y} = \mathbf{Ux}$, the problem splits into two triangular systems,

$$\mathbf{Ly} = \mathbf{Pw} \quad (\text{forward substitution}), \quad \mathbf{Ux} = \mathbf{y} \quad (\text{backward substitution}), \quad (1.23)$$

each costing only $\mathcal{O}(n^2)$ operations. This is the decisive advantage of LU over plain elimination: the expensive $\mathcal{O}(n^3)$ factorisation is performed once, after which any number of right-hand sides can be treated cheaply.

The inverse.

Column j of $\mathbf{A}^{-1}$ is the solution of $\mathbf{Ax} = \mathbf{e}_j$, where $\mathbf{e}_j$ is the j th unit vector. The inverse therefore costs one factorisation plus n substitutions. It is worth stressing that one should almost never compute an inverse: if the aim is to solve a linear system, solving it directly is both faster and more accurate than forming $\mathbf{A}^{-1}$ and multiplying.

```
class LUdecomposition(DirectSolver):
    """Doolittle LU factorisation with partial pivoting, P A = L U."""

    def _factorize(self):
```

```

n = self.n
LU = self.A.copy()
perm = np.arange(n)
sign = 1.0

for k in range(n):
    p = k + np.argmax(np.abs(LU[k:, k]))
    if abs(LU[p, k]) < 1.0e-14:
        raise np.linalg.LinAlgError("matrix is singular")
    if p != k:
        LU[[k, p]] = LU[[p, k]]
        perm[[k, p]] = perm[[p, k]]
        sign = -sign
    # the multipliers l_ik are stored in place, below the diagonal
    LU[k+1:, k] /= LU[k, k]
    # rank-one update of the trailing submatrix
    LU[k+1:, k+1:] -= np.outer(LU[k+1:, k], LU[k, k+1:])

self.LU, self.perm, self.sign = LU, perm, sign

def solve(self, b):
    """Solve A x = b in two steps: L y = P b, then U x = y."""
    y = np.asarray(b, dtype=float)[self.perm].copy()
    n = self.n
    for i in range(1, n):
        y[i] -= self.LU[i, :i] @ y[:i]
    x = np.zeros(n)
    for i in range(n-1, -1, -1):
        x[i] = (y[i] - self.LU[i, i+1:] @ x[i+1:]) / self.LU[i, i]
    return x

def determinant(self):
    return self.sign * np.prod(np.diag(self.LU))

```

Cholesky's factorisation.

If $\mathbf{A}$ is real, symmetric and positive definite it admits the special factorisation

$$\mathbf{A} = \mathbf{L}\mathbf{L}^T, \quad (1.24)$$

with $\mathbf{L}$ lower triangular. The elements follow from

$$L_{ii} = \left(A_{ii} - \sum_{k=0}^{i-1} L_{ik}^2 \right)^{1/2}, \quad (1.25)$$

$$L_{ji} = \frac{1}{L_{ii}} \left(A_{ji} - \sum_{k=0}^{i-1} L_{ik} L_{jk} \right), \quad j = i+1, \dots, n-1. \quad (1.26)$$

The algorithm costs half of a general LU factorisation, and since the pivots L_{ii} are guaranteed positive there is no need for pivoting at all. The square root in Eq. (1.25) also provides the cheapest available test of positive definiteness: if the argument turns negative, the matrix is not positive definite.

Many-body connection. Positive definite matrices are everywhere in many-body theory. Overlap matrices of linearly independent basis functions are positive definite by construction, and the Cholesky factor of the overlap matrix is the standard tool for transforming a generalised eigenvalue problem $\mathbf{H}\mathbf{c} = \mathbf{E}\mathbf{S}\mathbf{c}$ into an ordinary one. The Hessian of the variance in variational Monte Carlo is likewise

positive definite, which is what makes the conjugate gradient method of Section 1.18 applicable to the optimisation of trial wave functions.

1.16 Tridiagonal systems

The matrix of Eq. (1.16) is tridiagonal, and applying general Gaussian elimination to it would be an extravagant waste: almost all the operations would be spent on elements that are already zero. Writing the system as

$$a_i u_{i-1} + b_i u_i + c_i u_{i+1} = f_i, \quad i = 0, 1, \dots, n-1, \quad (1.27)$$

with a_0 and c_{n-1} not referenced, the elimination reduces to a single forward sweep followed by a single backward sweep. The cost is $\mathcal{O}(n)$ operations rather than $2n^3/3$, and the storage is three vectors rather than a full matrix. Whenever a problem produces a tridiagonal, banded or otherwise sparse matrix, exploiting that structure is not an optimisation but a necessity.

```
class TridiagonalSolver:
    """Solve a tridiagonal system in  $\mathcal{O}(n)$  operations."""

    def solve(self, f):
        n = self.n
        f = np.asarray(f, dtype=float)
        temp = np.zeros(n)
        u = np.zeros(n)

        btemp = self.b[0] # forward substitution
        u[0] = f[0] / btemp
        for i in range(1, n):
            temp[i] = self.c[i-1] / btemp
            btemp = self.b[i] - self.a[i] * temp[i]
            u[i] = (f[i] - self.a[i] * u[i-1]) / btemp

        for i in range(n-2, -1, -1): # backward substitution
            u[i] -= temp[i+1] * u[i+1]
        return u
```

Applied to $-u'' = 2$ on $(0, 1)$ with $u(0) = u(1) = 0$, whose exact solution is $u(x) = x(1-x)$, the solver reproduces the analytic answer to 8.6×10^{-16} with $n = 100$ interior points. The three-point formula is exact for polynomials of degree two, which is why the discretisation error vanishes here; for a general right-hand side the error would be $\mathcal{O}(h^2)$, as we shall see in Section 1.27.

1.17 Iterative methods for linear systems

The methods of the previous sections are *direct*: up to rounding they produce the exact answer in a finite, known number of operations. Their cost, however, grows as n^3 , and they require the matrix to be stored. Iterative methods take the opposite view. We start from a guess and improve it until the change falls below a tolerance. The matrix enters only through the product $\mathbf{A}\mathbf{x}$, so a matrix that is sparse, or that is never formed at all but only applied, costs nothing extra. This is the reason iterative methods, and not the direct ones, are what one uses for large-scale many-body calculations.

We split the matrix as

$$\mathbf{A} = \mathbf{D} + \mathbf{L} + \mathbf{U}, \quad (1.28)$$

with $\mathbf{D}$ diagonal, $\mathbf{L}$ strictly lower and $\mathbf{U}$ strictly upper triangular.

Jacobi's method.

Solving row i for x_i and evaluating the right-hand side with the previous iterate gives

$$x_i^{(k+1)} = \frac{1}{a_{ii}} \left(b_i - \sum_{j \neq i} a_{ij} x_j^{(k)} \right), \quad (1.29)$$

or, in matrix form,

$$\mathbf{x}^{(k+1)} = \mathbf{D}^{-1} \left(\mathbf{b} - (\mathbf{L} + \mathbf{U}) \mathbf{x}^{(k)} \right). \quad (1.30)$$

Every component of the new iterate is computed from the old ones only, so all n updates are independent and the method parallelises trivially. It converges whenever $\mathbf{A}$ is strictly diagonally dominant or symmetric positive definite.

Gauss-Seidel.

Nothing forces us to wait: the components $x_0^{(k+1)}, \dots, x_{i-1}^{(k+1)}$ are already available when we compute $x_i^{(k+1)}$. Using them immediately gives the Gauss-Seidel iteration

$$x_i^{(k+1)} = \frac{1}{a_{ii}} \left(b_i - \sum_{j < i} a_{ij} x_j^{(k+1)} - \sum_{j > i} a_{ij} x_j^{(k)} \right). \quad (1.31)$$

This is forward substitution applied to the splitting, and it typically halves the number of iterations. The price is that the sweep is now inherently sequential.

Successive over-relaxation.

Gauss-Seidel usually moves in the right direction but not far enough. Over-relaxation exaggerates the step by a factor ω ,

$$x_i^{(k+1)} = (1 - \omega) x_i^{(k)} + \frac{\omega}{a_{ii}} \left(b_i - \sum_{j < i} a_{ij} x_j^{(k+1)} - \sum_{j > i} a_{ij} x_j^{(k)} \right), \quad (1.32)$$

which in matrix form reads

$$\mathbf{x}^{(k+1)} = (\mathbf{D} + \omega \mathbf{L})^{-1} \left(\omega \mathbf{b} - [\omega \mathbf{U} + (\omega - 1) \mathbf{D}] \mathbf{x}^{(k)} \right). \quad (1.33)$$

For symmetric positive definite matrices convergence is guaranteed for $0 < \omega < 2$, and $\omega = 1$ returns Gauss-Seidel. A well chosen ω can reduce the iteration count by an order of magnitude, but the optimal value depends on the spectrum of $\mathbf{A}$ and is rarely known in advance.

1.18 The conjugate gradient method

The methods above treat the linear system as a fixed-point problem. The conjugate gradient method instead treats it as a minimisation. For a symmetric positive definite $\mathbf{A}$ the solution of $\mathbf{Ax} = \mathbf{b}$ is the unique minimum of the quadratic form

$$P(\mathbf{x}) = \frac{1}{2} \mathbf{x}^T \mathbf{Ax} - \mathbf{x}^T \mathbf{b}, \quad (1.34)$$

since $\nabla P = \mathbf{Ax} - \mathbf{b}$, which vanishes precisely at the solution. The residual

$$\mathbf{r} = \mathbf{b} - \mathbf{Ax} \quad (1.35)$$

is therefore the negative gradient of P , and steepest descent would move along it. Steepest descent is however notoriously slow, because successive steps undo part of the progress of their predecessors.

The remedy is to search along directions that are *conjugate*, meaning orthogonal with respect to the inner product defined by $\mathbf{A}$,

$$\mathbf{p}_i^T \mathbf{A} \mathbf{p}_j = 0, \quad i \neq j. \quad (1.36)$$

The eigenvectors of $\mathbf{A}$ are an example, since $\mathbf{v}_i^T \mathbf{A} \mathbf{v}_j = \lambda_j \mathbf{v}_i^T \mathbf{v}_j$ vanishes for $i \neq j$; but the whole point of the method is that we shall construct conjugate directions without knowing any eigenvectors. A set of n mutually conjugate vectors forms a basis of $\mathbb{R}^n$, so we may expand

$$\mathbf{x} = \sum_{i=0}^{n-1} \alpha_i \mathbf{p}_i. \quad (1.37)$$

Multiplying $\mathbf{A}\mathbf{x} = \mathbf{b}$ from the left by $\mathbf{p}_k^T$ and using Eq. (1.36) collapses the sum to a single term,

$$\alpha_k = \frac{\mathbf{p}_k^T \mathbf{b}}{\mathbf{p}_k^T \mathbf{A} \mathbf{p}_k}, \quad (1.38)$$

so each direction can be treated independently. In exact arithmetic the method therefore terminates after at most n steps, which makes it a direct method in disguise; what makes it useful is that a good approximation is normally reached in far fewer.

The directions are generated on the fly. Starting from $\mathbf{x}_0 = \mathbf{0}$ and $\mathbf{p}_0 = \mathbf{r}_0 = \mathbf{b}$, the new direction is the residual with its component along the previous direction projected out,

$$\mathbf{p}_{k+1} = \mathbf{r}_{k+1} - \frac{\mathbf{p}_k^T \mathbf{A} \mathbf{r}_{k+1}}{\mathbf{p}_k^T \mathbf{A} \mathbf{p}_k} \mathbf{p}_k, \quad (1.39)$$

and the residual itself is updated recursively. From $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$ we get

$$\mathbf{r}_{k+1} = \mathbf{b} - \mathbf{A}(\mathbf{x}_k + \alpha_k \mathbf{p}_k) = \mathbf{r}_k - \alpha_k \mathbf{A} \mathbf{p}_k, \quad (1.40)$$

which saves one matrix-vector product per iteration. Remarkably, the orthogonality relations that hold between the residuals allow Eq. (1.39) to be reduced to the two-term recurrence

$$\mathbf{p}_{k+1} = \mathbf{r}_{k+1} + \beta_k \mathbf{p}_k, \quad \beta_k = \frac{\mathbf{r}_{k+1}^T \mathbf{r}_{k+1}}{\mathbf{r}_k^T \mathbf{r}_k}, \quad (1.41)$$

so that only the current residual and the current direction need to be stored. This is what the implementation below uses.

```
class ConjugateGradient(IterativeSolver):
    """Conjugate gradient for symmetric positive definite A.

    Only the matrix-vector product A p is required, so a `matvec` callable
    may be supplied instead of a dense matrix.
    """

    def solve(self, b, x0=None):
        b = np.asarray(b, dtype=float)
        x = self._initial_guess(x0)
        r = b - self.matvec(x)          # residual
        p = r.copy()                    # first search direction
        rs = r @ r

        for k in range(self.max_iter):
            Ap = self.matvec(p)
            alpha = rs / (p @ Ap)
            x += alpha * p              # x_{k+1} = x_k + alpha_k p_k
```



```

r -= alpha * Ap                                # r_{k+1} = r_k - alpha_k A p_k
rs_new = r @ r
if np.sqrt(rs_new) < self.tol:
    break
p = r + (rs_new / rs) * p                      # beta_k = |r_{k+1}|^2 / |r_k|^2
rs = rs_new
return x

```

Table 1.2 compares the four methods on the discretised second derivative of Eq. (1.16) with $n = 50$ interior points, for a residual tolerance of 10^{-8} . The ordering is the one to remember: Gauss-Seidel is roughly twice as fast as Jacobi, a well chosen relaxation parameter gains an order of magnitude, and the conjugate gradient method gains another.

Method	Iterations	Error
Jacobi	11049	1.0×10^{-9}
Gauss-Seidel	5526	1.0×10^{-9}
Successive over-relaxation ($\omega = 1.9$)	220	2.8×10^{-10}
Conjugate gradient	25	1.7×10^{-16}

Table 1.2 Iterations needed to solve $-u'' = 2$ on $(0, 1)$ with $u(0) = u(1) = 0$, discretised on $n = 50$ interior points, to a residual tolerance of 10^{-8} . The error is measured against the analytic solution $u(x) = x(1 - x)$.

Many-body connection. The conjugate gradient method appears in two distinct roles in this book. As a linear solver it is used wherever a large sparse positive definite system arises. As a minimiser it is the workhorse of variational calculations: in variational Monte Carlo the parameters of a trial wave function are optimised by conjugate gradient steps on the energy or the variance, and the Hessian of the variance plays the role of $\mathbf{A}$. The non-linear generalisations of the method, discussed in the chapter on gradient methods, all rest on the linear theory developed here.

1.19 Cubic spline interpolation

Tridiagonal systems appear in a second, quite different context, and the solver of Section 1.16 can be reused verbatim. Suppose we are given $n + 1$ knots $x_0 < x_1 < \dots < x_n$ and the corresponding function values $y_i = f(x_i)$, and we want a smooth interpolant. A spline function s of degree k is a piecewise polynomial: on every subinterval $[x_i, x_{i+1})$ it is a polynomial of degree at most k , and it has $k - 1$ continuous derivatives on the whole interval. For $k = 1$ the spline is the broken line through the data points, with no continuous derivative at all. The useful choice is $k = 3$, the cubic spline, which is twice continuously differentiable and therefore looks smooth to the eye.

On each of the n subintervals we have a cubic

$$s_i(x) = a_{i0} + a_{i1}x + a_{i2}x^2 + a_{i3}x^3, \quad (1.42)$$

giving $4n$ unknown coefficients. Interpolation at both ends of every subinterval provides $2n$ conditions, continuity of s' at the interior knots $n - 1$ more, and continuity of s'' another $n - 1$, for a total of $4n - 2$. Two degrees of freedom remain, and they are fixed by a condition at the two ends; setting the second derivative to zero there gives the *natural* cubic spline.

The elegant way to construct the spline is to take the second derivatives $f_i = s''(x_i)$ as the unknowns. Since s'' is piecewise linear,

$$s_i''(x) = \frac{f_i}{x_{i+1} - x_i}(x_{i+1} - x) + \frac{f_{i+1}}{x_{i+1} - x_i}(x - x_i), \quad (1.43)$$

and integrating twice, fixing the integration constants by the interpolation conditions, gives $s_i(x)$ in terms of f_i and f_{i+1} . Demanding that s' be continuous at the interior knots then produces $n - 1$ equations for the $n - 1$ unknown interior second derivatives, and those equations are tridiagonal and diagonally dominant. A cubic spline through n points therefore costs $\mathcal{O}(n)$ operations, which is why splines rather than global polynomials are used to interpolate tabulated potentials, form factors and matrix elements.

1.20 Applications: integral equations and matrix inversion

Two further examples show how naturally physical problems turn into the matrix problems of this chapter.

A linear integral equation.

The scattering problem in quantum mechanics leads to the Lippmann-Schwinger equation, an integral equation of the form

$$R(k, k') = V(k, k') + \frac{2}{\pi} \mathcal{P} \int_0^\infty dq q^2 \frac{V(k, q) R(q, k')}{(k_0^2 - q^2)/m}, \quad (1.44)$$

where $\mathcal{P}$ denotes the Cauchy principal value. If we approximate the integral by a quadrature rule with N points q_j and weights w_j , the unknown function R is replaced by the $N + 1$ numbers $R(q_i, k_0)$, and Eq. (1.44) becomes

$$\sum_j A_{ij} R_j = V_i, \quad A_{ij} = \delta_{ij} - \frac{2}{\pi} u_j V(k_i, q_j), \quad (1.45)$$

with the weights u_j containing the quadrature weights and the propagator. The principal value is handled by the standard subtraction trick, which removes the singularity at $q = k_0$ before the quadrature is applied. What began as an integral equation is now a linear system of the type treated in Section 1.15: the physics is entirely in the construction of the matrix, and the solution is one call to an LU solver.

An ill-conditioned inverse.

Not every matrix inversion is benign. Fitting a polynomial of degree $n - 1$ through n points $x_0, \dots, x_{n-1}$ requires inverting the Vandermonde matrix

$$\mathbf{V} = \begin{bmatrix} 1 & x_0 & x_0^2 & \cdots & x_0^{n-1} \\ 1 & x_1 & x_1^2 & \cdots & x_1^{n-1} \\ \vdots & \vdots & \vdots & & \vdots \\ 1 & x_{n-1} & x_{n-1}^2 & \cdots & x_{n-1}^{n-1} \end{bmatrix}. \quad (1.46)$$

Its inverse is known in closed form, but its condition number grows exponentially with n : for equidistant points on $[0, 1]$ and $n = 20$ it exceeds 10^{20} , so that not a single significant digit survives in double precision. This is the numerical statement of a fact familiar from interpolation theory, namely that high-degree polynomial interpolation on a uniform grid is a bad idea, and it is the practical reason for preferring the piecewise construction of Section 1.19.

1.21 The algebraic eigenvalue problem

We return to the eigenvalue problem introduced in Section 1.11, this time asking how it is actually solved. The eigenvalues of a matrix $\mathbf{A}$ of dimension n are defined by

$$\mathbf{A}\mathbf{x}^{(v)} = \lambda^{(v)}\mathbf{x}^{(v)}, \quad (1.47)$$

where, unless stated otherwise, eigenvector means right eigenvector; the left eigenvectors satisfy $\mathbf{x}_L^{(v)}\mathbf{A} = \lambda^{(v)}\mathbf{x}_L^{(v)}$. Rewriting Eq. (1.47) as $(\mathbf{A} - \lambda^{(v)}\mathbf{I})\mathbf{x}^{(v)} = \mathbf{0}$, a non-trivial solution exists only if the matrix is singular, that is only if

$$P(\lambda) = \det(\lambda\mathbf{I} - \mathbf{A}) = 0. \quad (1.48)$$

The characteristic polynomial $P(\lambda)$ has degree n , and its roots

$$P(\lambda) = \prod_{i=0}^{n-1} (\lambda_i - \lambda) \quad (1.49)$$

form the *spectrum* $\lambda(\mathbf{A})$. Two useful identities follow at once from the factorisation,

$$|\mathbf{A}| = \lambda_0\lambda_1 \cdots \lambda_{n-1}, \quad \text{Tr}(\mathbf{A}) = \lambda_0 + \lambda_1 + \cdots + \lambda_{n-1}. \quad (1.50)$$

It is tempting to conclude that one should find eigenvalues by computing the characteristic polynomial and looking for its roots. This is almost always a bad idea: the roots of a polynomial are extremely sensitive to its coefficients, so the detour through $P(\lambda)$ throws away accuracy that the original matrix still contained. The one case where the polynomial route is respectable is a tridiagonal matrix, for which the leading principal minors obey the recursion

$$P_k(\lambda) = (d_k - \lambda)P_{k-1}(\lambda) - e_{k-1}^2 P_{k-2}(\lambda), \quad (1.51)$$

with $P_0(\lambda) = d_0 - \lambda$ and $P_{-1} = 1$. The sequence $\{P_k(\lambda)\}$ is a Sturm sequence, and counting its sign changes tells us how many eigenvalues lie below λ , which is a robust way of isolating selected eigenvalues.

The standard approach is entirely different. We apply a sequence of similarity transformations that bring $\mathbf{A}$ either to diagonal form directly, or first to tridiagonal form and then to diagonal form. The second route is the one we follow, because diagonalising a tridiagonal matrix is vastly cheaper than diagonalising a full one.

1.22 Similarity transformations

Let $\mathbf{A} \in \mathbb{R}^{n \times n}$ be real and symmetric. Then there exists a real orthogonal matrix $\mathbf{S}$ such that

$$\mathbf{S}^T \mathbf{A} \mathbf{S} = \text{diag}(\lambda_0, \lambda_1, \dots, \lambda_{n-1}) \equiv \mathbf{D}, \quad (1.52)$$

and the j th column of $\mathbf{S}$ is the eigenvector belonging to λ_j [1]. We say that $\mathbf{B}$ is a similarity transform of $\mathbf{A}$ if

$$\mathbf{B} = \mathbf{S}^T \mathbf{A} \mathbf{S}, \quad \mathbf{S}^T \mathbf{S} = \mathbf{S}^{-1} \mathbf{S} = \mathbf{I}. \quad (1.53)$$

The importance of such a transformation is that it leaves the eigenvalues unchanged while altering the eigenvectors in a controlled way. The proof is one line. Multiply $\mathbf{A}\mathbf{x} = \lambda\mathbf{x}$ from the left by $\mathbf{S}^T$ and insert $\mathbf{S}\mathbf{S}^T = \mathbf{I}$ between $\mathbf{A}$ and $\mathbf{x}$,

$$(\mathbf{S}^T \mathbf{A} \mathbf{S})(\mathbf{S}^T \mathbf{x}) = \lambda (\mathbf{S}^T \mathbf{x}), \quad (1.54)$$

that is $\mathbf{B}(\mathbf{S}^T \mathbf{x}) = \lambda(\mathbf{S}^T \mathbf{x})$. So λ is an eigenvalue of $\mathbf{B}$ as well, with eigenvector $\mathbf{S}^T \mathbf{x}$.

The strategy is therefore either to apply transformations until

$$\mathbf{S}_N^T \cdots \mathbf{S}_1^T \mathbf{A} \mathbf{S}_1 \cdots \mathbf{S}_N = \mathbf{D}, \quad (1.55)$$

which is Jacobi's method of successive plane rotations, or to reduce $\mathbf{A}$ to tridiagonal form in a finite number of steps and then attack the tridiagonal matrix. We follow the second route, which is what all modern library routines do.

Many-body connection. Similarity transformations are not merely a numerical device. The whole of coupled-cluster theory rests on the similarity-transformed Hamiltonian $\tilde{H} = e^{-T} \hat{H} e^T$, which has the same eigenvalues as $\hat{H}$ but is not Hermitian, since e^T is not unitary. The in-medium similarity renormalisation group instead uses a continuous sequence of *unitary* transformations to drive the Hamiltonian towards diagonal form, which is Eq. (1.55) carried out with infinitesimal steps.

1.23 The power method

The simplest of all eigenvalue algorithms is worth understanding, because everything that follows can be seen as a refinement of it. Assume $\mathbf{A}$ can be diagonalised, with eigenvalues $\lambda_0, \dots, \lambda_{n-1}$ and eigenvectors $\mathbf{v}_0, \dots, \mathbf{v}_{n-1}$, and assume there is a dominant eigenvalue, $|\lambda_0| > |\lambda_j|$ for $j > 0$. An arbitrary starting vector expands as

$$\mathbf{b}_0 = c_0 \mathbf{v}_0 + c_1 \mathbf{v}_1 + \cdots + c_{n-1} \mathbf{v}_{n-1}, \quad (1.56)$$

and applying $\mathbf{A}$ repeatedly gives

$$\mathbf{A}^k \mathbf{b}_0 = c_0 \lambda_0^k \left(\mathbf{v}_0 + \sum_{j>0} \frac{c_j}{c_0} \left(\frac{\lambda_j}{\lambda_0} \right)^k \mathbf{v}_j \right). \quad (1.57)$$

Every ratio $|\lambda_j/\lambda_0|$ is smaller than one, so the bracket tends to $\mathbf{v}_0$: repeated multiplication filters out everything except the dominant eigenvector. Normalising at every step, $\mathbf{b}_k = \mathbf{A}^k \mathbf{b}_0 / \|\mathbf{A}^k \mathbf{b}_0\|$, and estimating the eigenvalue by the Rayleigh quotient

$$\mu_k = \frac{\mathbf{b}_k^\dagger \mathbf{A} \mathbf{b}_k}{\mathbf{b}_k^\dagger \mathbf{b}_k}, \quad (1.58)$$

we obtain a method that converges geometrically with ratio $|\lambda_1/\lambda_0|$, where λ_1 is the second largest eigenvalue in magnitude. When the two are close, convergence is correspondingly slow.

The power method has two obvious weaknesses: it gives only one eigenpair, and it discards all the information contained in the intermediate vectors $\mathbf{b}_0, \mathbf{A}\mathbf{b}_0, \dots, \mathbf{A}^{k-1}\mathbf{b}_0$. Removing the second weakness by working in the whole space spanned by those vectors, the *Krylov space*, leads directly to the Lanczos algorithm of Section 1.26. Applying the power method to $\mathbf{A}^{-1}$ instead gives inverse iteration, which converges to the eigenvalue of smallest magnitude and, with a shift, to any eigenvalue one can bracket.

1.24 Householder's method for tridiagonalisation

The first of the two steps is to construct an orthogonal $\mathbf{S}$ that brings $\mathbf{A}$ to tridiagonal form. It is built as a product of $n-2$ elementary transformations,

$$\mathbf{S} = \mathbf{S}_0 \mathbf{S}_1 \cdots \mathbf{S}_{n-3}, \quad (1.59)$$

each of which clears one column and one row. Only $n-2$ steps are needed, since the last two rows are already tridiagonal. After the first transformation

$$\mathbf{S}_0^T \mathbf{A} \mathbf{S}_0 = \begin{bmatrix} a_{00} & e_0 & 0 & \cdots & 0 \\ e_0 & a'_{11} & a'_{12} & \cdots & a'_{1,n-1} \\ 0 & a'_{21} & a'_{22} & \cdots & a'_{2,n-1} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & a'_{n-1,1} & \cdots & \cdots & a'_{n-1,n-1} \end{bmatrix}, \quad (1.60)$$

where the primed block is a matrix of dimension $n - 1$ which the next transformation will treat in the same way. Note that the effective size of the problem shrinks at every step; in Jacobi's method, by contrast, each rotation acts on the full matrix and previously created zeros are destroyed. After all $n - 2$ steps we are left with the diagonal elements $a_{00}, a'_{11}, a'_{22}, \dots$ and the off-diagonal elements $e_0, e_1, \dots, e_{n-2}$ of a tridiagonal matrix.

It remains to find the transformations. We take

$$\mathbf{S}_0 = \begin{bmatrix} 1 & \mathbf{0}^T \\ \mathbf{0} & \mathbf{P} \end{bmatrix}, \quad \mathbf{P} = \mathbf{I} - 2\mathbf{u}\mathbf{u}^T, \quad \mathbf{u}^T \mathbf{u} = 1, \quad (1.61)$$

where $\mathbf{u}$ is a column vector of length $n - 1$. The matrix $\mathbf{P}$, whose elements are $P_{ij} = \delta_{ij} - 2u_i u_j$, is a *Householder reflector*: it is symmetric and satisfies $\mathbf{P}^2 = \mathbf{I}$, hence $\mathbf{P}^T = \mathbf{P}^{-1}$, so it is orthogonal. Geometrically it reflects a vector in the hyperplane perpendicular to $\mathbf{u}$.

Applying $\mathbf{S}_0$ gives

$$\mathbf{S}_0^T \mathbf{A} \mathbf{S}_0 = \begin{bmatrix} a_{00} & (\mathbf{P}\mathbf{v})^T \\ \mathbf{P}\mathbf{v} & \mathbf{A}' \end{bmatrix}, \quad \mathbf{v}^T = \{a_{10}, a_{20}, \dots, a_{n-1,0}\}, \quad (1.62)$$

so we must choose $\mathbf{u}$ such that $\mathbf{P}\mathbf{v}$ has only one non-zero component,

$$\mathbf{P}\mathbf{v} = \mathbf{v} - 2\mathbf{u}(\mathbf{u}^T \mathbf{v}) = \kappa \mathbf{e}, \quad \mathbf{e}^T = \{1, 0, \dots, 0\}. \quad (1.63)$$

The constant κ follows from taking the scalar product of Eq. (1.63) with itself and using $\mathbf{P}^2 = \mathbf{I}$,

$$(\mathbf{P}\mathbf{v})^T \mathbf{P}\mathbf{v} = \kappa^2 = \mathbf{v}^T \mathbf{v} = \sum_{i=1}^{n-1} a_{i0}^2, \quad (1.64)$$

so $\kappa = \pm \|\mathbf{v}\|_2$. Rewriting Eq. (1.63) as $\mathbf{v} - \kappa \mathbf{e} = 2\mathbf{u}(\mathbf{u}^T \mathbf{v})$ shows that $\mathbf{u}$ is proportional to $\mathbf{v} - \kappa \mathbf{e}$, and normalising gives

$$\mathbf{u} = \frac{\mathbf{v} - \kappa \mathbf{e}}{\|\mathbf{v} - \kappa \mathbf{e}\|_2}. \quad (1.65)$$

The sign of κ is free, and here numerical care is essential: if $v_0 = a_{10}$ and κ have the same sign, the leading component of $\mathbf{v} - \kappa \mathbf{e}$ is a difference of nearly equal numbers and precision is lost. We therefore always choose

$$\kappa = -\text{sign}(v_0) \|\mathbf{v}\|_2, \quad (1.66)$$

which makes the subtraction an addition.

Finally we need the transformed block. Writing $\mathbf{p} = \mathbf{A}'\mathbf{u}$ and $\mathbf{K} = \mathbf{u}^T \mathbf{p}$, and introducing $\mathbf{z} = \mathbf{p} - \mathbf{K}\mathbf{u}$, a short calculation gives the symmetric rank-two update

$$\mathbf{P}\mathbf{A}'\mathbf{P} = \mathbf{A}' - 2\mathbf{z}\mathbf{u}^T - 2\mathbf{u}\mathbf{z}^T, \quad (1.67)$$

which costs $\mathcal{O}(n^2)$ operations rather than the $\mathcal{O}(n^3)$ of forming the products explicitly. Summed over the $n - 2$ steps, the whole reduction costs $\frac{4}{3}n^3$ operations, or twice that if the transformation matrix is accumulated.

```
def tridiagonalize(self):
    """Return (d, e, S) with S^T A S tridiagonal."""
    n = self.n
    M = self.A.copy()
```

```

S = np.eye(n)

for k in range(n - 2):
    v = M[k+1:, k] # the column to be reduced
    norm_v = np.linalg.norm(v)
    if norm_v < self.tol:
        continue
    # the sign is chosen to avoid cancellation in v[0] - kappa
    kappa = -np.copysign(norm_v, v[0])
    w = v.copy()
    w[0] -= kappa
    u = w / np.linalg.norm(w) # unit vector, u^T u = 1

    # similarity transformation of the trailing block,
    # P A P = A - 2 z u^T - 2 u z^T, z = A u - (u^T A u) u
    block = M[k+1:, k+1:]
    p = block @ u
    z = p - (u @ p) * u
    M[k+1:, k+1:] = block - 2.0*np.outer(z, u) - 2.0*np.outer(u, z)

    M[k+1:, k] = 0.0 # the reduced column and row
    M[k+1, k] = kappa
    M[k, k+1:] = 0.0
    M[k, k+1] = kappa

    S[:, k+1:] -= 2.0 * np.outer(S[:, k+1:] @ u, u) # S <- S P

d = np.diag(M).copy()
e = np.diag(M, -1).copy()
return d, e, S

```

For a random symmetric 8×8 matrix the routine returns an $\mathbf{S}$ with $\|\mathbf{S}^T \mathbf{S} - \mathbf{I}\| = 1.1 \times 10^{-15}$ and reproduces the original matrix, $\|\mathbf{S}^T \mathbf{S}^T - \mathbf{A}\| = 3.1 \times 10^{-15}$. Orthogonality is preserved to machine precision, which is exactly why reflectors are preferred to the more obvious Gaussian-elimination-like schemes.

1.25 Diagonalising a tridiagonal matrix

The second step diagonalises the tridiagonal matrix

$$\mathbf{T} = \begin{bmatrix} d_0 & e_0 & & \\ e_0 & d_1 & e_1 & \\ & e_1 & d_2 & \ddots \\ & & \ddots & \ddots \end{bmatrix}. \quad (1.68)$$

The basic observation is that if any off-diagonal element e_i vanishes, the matrix splits into two independent blocks; in particular if $e_0 = 0$ then d_0 is already an eigenvalue. The algorithm therefore works by making the off-diagonal elements vanish one after another.

Francis' original scheme, and the QL and QR algorithms that grew out of it, apply a sequence of plane rotations

$$\mathbf{S}_i = \begin{bmatrix} \cos \theta & \cdots & \sin \theta \\ \vdots & & \vdots \\ -\sin \theta & \cdots & \cos \theta \end{bmatrix}, \quad (1.69)$$

choosing θ at each step so as to annihilate one off-diagonal element. Two refinements turn this idea into the algorithm actually used. The first is *shifting*: instead of working with $\mathbf{T}$ one works with $\mathbf{T} - \sigma \mathbf{I}$, where σ is an estimate of the eigenvalue being sought. The rate of convergence of the element e_{m-1} to zero is

governed by the ratio of the distances of σ to the two nearest eigenvalues, so a good shift produces cubic convergence and typically two or three iterations per eigenvalue. The Wilkinson shift, obtained from the trailing 2×2 submatrix, is the standard choice. The second refinement is to apply the shift *implicitly*, by chasing an unwanted element down the subdiagonal with successive rotations, which avoids ever forming $T - \sigma I$ explicitly and so avoids the associated cancellation.

The resulting routine is traditionally called `tqli`. It takes the diagonal d and the off-diagonal e produced by the Householder step, together with the accumulated transformation matrix, and returns the eigenvalues in d and the eigenvectors as the columns of that matrix.

```
@staticmethod
def tqli(d, e, z=None, max_sweeps=50):
    """QL algorithm with implicit shifts for a symmetric tridiagonal matrix.

    Pass the S of the Householder step as z to obtain the eigenvectors of
    the original matrix A.
    """
    d = np.asarray(d, dtype=float).copy()
    n = len(d)
    e = np.concatenate([np.asarray(e, dtype=float), [0.0]])
    z = np.eye(n) if z is None else np.asarray(z, dtype=float).copy()
    eps = np.finfo(float).eps

    for l in range(n):
        for _ in range(max_sweeps):
            # look for a negligible off-diagonal element to split off
            m = n - 1
            for i in range(l, n - 1):
                if abs(e[i]) <= eps * (abs(d[i]) + abs(d[i+1])):
                    m = i
                    break
            if m == l:
                # d[l] has converged
                break

            # Wilkinson shift, computed without cancellation
            g = (d[l+1] - d[l]) / (2.0 * e[l])
            r = np.hypot(g, 1.0)
            g = d[m] - d[l] + e[l] / (g + np.copysign(r, g))

            s = c = 1.0
            p = 0.0
            for i in range(m - 1, l - 1, -1):
                f = s * e[i]
                b = c * e[i]
                r = np.hypot(f, g)
                e[i+1] = r
                s, c = f / r, g / r
                g = d[i+1] - p
                r = (d[i] - g) * s + 2.0 * c * b
                p = s * r
                d[i+1] = g + p
                g = c * r - b
                col = z[:, i+1].copy()
                z[:, i+1] = s * z[:, i] + c * col
                z[:, i] = c * z[:, i] - s * col

            d[l] -= p
            e[l] = g
            e[m] = 0.0
    return d, z
```

Together, the two steps constitute a complete eigenvalue solver. On a random symmetric 8×8 matrix it reproduces the eigenvalues returned by `numpy.linalg.eigvalsh` to 2.2×10^{-15} , with $\|AV - VD\| =$

4.9×10^{-15} and $\|\mathbf{V}^T \mathbf{V} - \mathbf{I}\| = 2.7 \times 10^{-15}$. The total cost is $\mathcal{O}(n^3)$, dominated by the Householder reduction.

Many-body connection. The Householder and QL combination is what lies behind every call to a library eigenvalue routine, and it is entirely adequate as long as the matrix fits in memory. For the Hartree-Fock and random-phase-approximation matrices of the following chapters, whose dimensions are set by the number of single-particle states, this is the method of choice. It becomes useless the moment the dimension is set by the number of Slater determinants, which is the situation of the next section.

1.26 Lanczos' algorithm

The methods of the previous two sections transform the matrix, and therefore need to have it. In large-scale many-body calculations we do not. A full configuration-interaction Hamiltonian for a realistic system may have a dimension of 10^9 or more; it is sparse, it is generated on the fly from the occupation numbers of the Slater determinants, and it is never stored. What we can always do is compute the product $\mathbf{A}\mathbf{q}$ for a given vector $\mathbf{q}$. And we rarely want the whole spectrum: the ground state and a handful of excited states are enough.

The Lanczos algorithm is built for exactly this situation. It applies to real symmetric matrices and has the following properties:

- it generates a sequence of real tridiagonal matrices $\mathbf{T}_k$ of dimension $k \times k$ with $k \leq n$, whose extremal eigenvalues are progressively better estimates of the extremal eigenvalues of $\mathbf{A}$;
- the matrix enters only through the product $\mathbf{A}\mathbf{q}$;
- the underlying transformation is the similarity transformation $\mathbf{T} = \mathbf{Q}^T \mathbf{A} \mathbf{Q}$, with the first Lanczos vector $\mathbf{Q}\mathbf{e}_0 = \mathbf{q}_0$ chosen at will.

The three-term recurrence.

Write $\mathbf{Q}$ in terms of its column vectors, $\mathbf{Q} = [\mathbf{q}_0 \mathbf{q}_1 \cdots \mathbf{q}_{n-1}]$, and let

$$\mathbf{T} = \begin{bmatrix} \alpha_0 & \beta_0 & & & \\ \beta_0 & \alpha_1 & \beta_1 & & \\ & \beta_1 & \alpha_2 & \ddots & \\ & & \ddots & \ddots & \beta_{n-2} \\ & & & \beta_{n-2} & \alpha_{n-1} \end{bmatrix}. \quad (1.70)$$

Using $\mathbf{Q}\mathbf{Q}^T = \mathbf{I}$ we may rewrite $\mathbf{T} = \mathbf{Q}^T \mathbf{A} \mathbf{Q}$ as $\mathbf{Q}\mathbf{T} = \mathbf{A}\mathbf{Q}$, and equating column k on both sides gives the three-term recurrence

$$\mathbf{A}\mathbf{q}_k = \beta_{k-1}\mathbf{q}_{k-1} + \alpha_k\mathbf{q}_k + \beta_k\mathbf{q}_{k+1}, \quad (1.71)$$

with $\beta_{-1}\mathbf{q}_{-1} = \mathbf{0}$. This single equation contains the whole algorithm. Since the Lanczos vectors are orthonormal, multiplying Eq. (1.71) from the left by $\mathbf{q}_k^T$ gives

$$\alpha_k = \mathbf{q}_k^T \mathbf{A} \mathbf{q}_k, \quad (1.72)$$

and the remainder

$$\mathbf{r}_k = (\mathbf{A} - \alpha_k \mathbf{I})\mathbf{q}_k - \beta_{k-1}\mathbf{q}_{k-1} \quad (1.73)$$

determines the next vector. If $\mathbf{r}_k \neq \mathbf{0}$ then

$$\mathbf{q}_{k+1} = \mathbf{r}_k / \beta_k, \quad \beta_k = \pm \|\mathbf{r}_k\|_2. \quad (1.74)$$

If $\mathbf{r}_k$ does vanish we have found an invariant subspace, and the eigenvalues of $\mathbf{T}_k$ are exact eigenvalues of $\mathbf{A}$. The algorithm is therefore

$$\begin{aligned}
 &\mathbf{r}_0 = \mathbf{q}_0, \quad \beta_{-1} = 1, \quad \mathbf{q}_{-1} = \mathbf{0}, \quad k = 0 \\
 &\textbf{while } \beta_{k-1} \neq 0 \\
 &\quad \mathbf{q}_k = \mathbf{r}_{k-1} / \beta_{k-1} \\
 &\quad \alpha_k = \mathbf{q}_k^T \mathbf{A} \mathbf{q}_k \\
 &\quad \mathbf{r}_k = (\mathbf{A} - \alpha_k \mathbf{I}) \mathbf{q}_k - \beta_{k-1} \mathbf{q}_{k-1} \\
 &\quad \beta_k = \|\mathbf{r}_k\|_2, \quad k \leftarrow k + 1 \\
 &\textbf{end while}
 \end{aligned} \tag{1.75}$$

Krylov spaces and Ritz values.

After m steps the vectors $\mathbf{q}_0, \dots, \mathbf{q}_{m-1}$ span the Krylov space

$$\mathcal{K}_m(\mathbf{A}, \mathbf{q}_0) = \text{span}\{\mathbf{q}_0, \mathbf{A}\mathbf{q}_0, \mathbf{A}^2\mathbf{q}_0, \dots, \mathbf{A}^{m-1}\mathbf{q}_0\}, \tag{1.76}$$

which is precisely the sequence of vectors the power method of Section 1.23 generates and then throws away. Lanczos keeps them, orthonormalises them, and diagonalises $\mathbf{A}$ restricted to their span. The eigenvalues θ_i of the small matrix $\mathbf{T}_m$ are called *Ritz values* and the vectors $\mathbf{Q}\mathbf{y}_i$, with $\mathbf{y}_i$ the eigenvectors of $\mathbf{T}_m$, are the *Ritz vectors*. They are the best approximations to eigenpairs of $\mathbf{A}$ obtainable within the Krylov space, in the variational sense: the lowest Ritz value is the minimum of the Rayleigh quotient over $\mathcal{K}_m$, and is therefore an upper bound on the exact ground-state energy. This is the same variational principle that underlies Hartree-Fock theory and the variational Monte Carlo method.

Because $\mathbf{T}_m$ is tridiagonal and small, it is diagonalised with the QL algorithm of Section 1.25. The two algorithms are thus used together: Lanczos reduces the enormous problem to a tiny tridiagonal one, and `tqli` finishes the job.

```

def run(self, m, q0=None, tol=1.0e-12):
    """Perform m Lanczos steps, returning the tridiagonal matrix elements."""
    n = self.n
    q = np.random.default_rng(2026).normal(size=n) if q0 is None \
        else np.asarray(q0, dtype=float).copy()
    q /= np.linalg.norm(q)

    Q = np.zeros((n, m))
    alpha = np.zeros(m)
    beta = np.zeros(max(m - 1, 0))
    q_prev = np.zeros(n)
    b_prev = 0.0

    for k in range(m):
        Q[:, k] = q
        w = self.matvec(q) # the only use made of A
        alpha[k] = q @ w # alpha_k = q_k^T A q_k
        w = w - alpha[k]*q - b_prev*q_prev # r_k

        if self.reorthogonalize:
            for _ in range(2): # modified Gram-Schmidt
                w -= Q[:, :k+1] @ (Q[:, :k+1].T @ w)

        if k == m - 1:
            break
        b = np.linalg.norm(w)
        if b < tol: # invariant subspace found
            alpha, beta, Q = alpha[:k+1], beta[:k], Q[:, :k+1]
            break
        beta[k] = b

```

```

q_prev, b_prev = q, b
q = w / b

self.alpha, self.beta, self.Q = alpha, beta, Q
return alpha, beta, Q

```

Loss of orthogonality.

In exact arithmetic the three-term recurrence produces orthogonal vectors automatically. In floating-point arithmetic it does not. As soon as one Ritz value converges, the corresponding direction starts to reappear in the remainder r_k , orthogonality is lost, and the algorithm produces spurious extra copies of eigenvalues it has already found. These are the notorious *ghost* eigenvalues. The simplest cure, used in the code above, is full reorthogonalisation: at every step the new vector is explicitly orthogonalised against all previous ones. This requires storing all the Lanczos vectors and costs $\mathcal{O}(m^2n)$ operations, which is acceptable when m is a few hundred. For very long runs one uses selective or periodic reorthogonalisation instead.

A many-body example.

Table 1.3 shows the algorithm at work on the pairing model that runs through this book: twelve doubly degenerate single-particle levels, six pairs, single-particle spacing $\delta = 1$ and pairing strength $g = 0.5$. In the seniority-zero space a basis state is a choice of which levels carry a pair, so the dimension is $\binom{12}{6} = 924$. The ground-state energy is converged to ten digits after thirty matrix-vector products and to machine precision after forty, in a space of dimension 924.

m	Lowest Ritz value	Error
5	32.865169933329	8.0×10^0
10	25.140140936891	3.0×10^{-1}
20	24.839267115133	9.4×10^{-5}
30	24.839172748689	2.4×10^{-10}
40	24.839172748451	5.7×10^{-14}
exact	24.839172748451	

Table 1.3 Convergence of the lowest Ritz value for the pairing model with twelve levels and six pairs, $\delta = 1$, $g = 0.5$, in a space of dimension 924. The exact value is obtained by full diagonalisation.

Run without reorthogonalisation for 120 steps on the same matrix, the Lanczos vectors lose orthogonality completely, $\|Q^T Q - I\| = 4.9$, and three copies of the non-degenerate ground state appear among the Ritz values. With reorthogonalisation the deviation is 8.3×10^{-15} and the ground state appears exactly once.

Many-body connection. The Lanczos algorithm is the engine of large-scale shell-model and full configuration-interaction codes. Its combination of properties – only matrix-vector products are needed, the extremal eigenvalues converge first, and the lowest Ritz value is a variational upper bound – matches the structure of the many-body problem exactly: we want the ground state, the Hamiltonian is sparse, and the dimension is astronomical. The same Krylov idea reappears in quantum computing, where the quantum phase estimation and variational quantum eigensolver algorithms are, in different ways, attempts to extract the same information with quantum rather than classical resources.

1.27 Application: Schrödinger's equation through diagonalisation

We close the chapter by putting the machinery to work on a genuine quantum mechanical problem, and one that will return in a many-body guise later: solving Schrödinger's equation not as a differential equation, but as a matrix eigenvalue problem.

The radial equation for a particle in a central potential is

$$-\frac{\hbar^2}{2m} \left(\frac{1}{r^2} \frac{d}{dr} r^2 \frac{d}{dr} - \frac{l(l+1)}{r^2} \right) R(r) + V(r)R(r) = ER(r). \quad (1.77)$$

Substituting $R(r) = u(r)/r$ removes the first-derivative term,

$$-\frac{\hbar^2}{2m} \frac{d^2 u}{dr^2} + \left(V(r) + \frac{l(l+1)}{r^2} \frac{\hbar^2}{2m} \right) u(r) = Eu(r), \quad (1.78)$$

and introducing the dimensionless variable $\rho = r/\alpha$ with α a length makes the equation dimensionless. We specialise to one dimension, drop the centrifugal term, and take the harmonic oscillator potential $V(x) = \frac{1}{2}kx^2$. In units where $k = \hbar = m = \alpha = 1$ the equation becomes

$$-\frac{d^2 u}{dx^2} + x^2 u(x) = 2E u(x). \quad (1.79)$$

The factor two on the right is worth remembering: the program will return the values $2E_n = 2n + 1$, that is $1, 3, 5, 7, \dots$, rather than $E_n = n + \frac{1}{2}$.

We now repeat the discretisation of Section 1.14. Choose $R_{\min}$ and $R_{\max}$, a number of mesh points N_{step} , and the step $h = (R_{\max} - R_{\min})/N_{\text{step}}$. With $x_i = R_{\min} + ih$ and the three-point formula (1.14), Eq. (1.79) becomes

$$-\frac{u_{i+1} - 2u_i + u_{i-1}}{h^2} + V_i u_i = 2E u_i, \quad (1.80)$$

that is the tridiagonal eigenvalue problem

$$\begin{bmatrix} \frac{2}{h^2} + V_1 & -\frac{1}{h^2} & & \\ -\frac{1}{h^2} & \frac{2}{h^2} + V_2 & -\frac{1}{h^2} & \\ & \ddots & \ddots & \ddots \\ & & -\frac{1}{h^2} & \frac{2}{h^2} + V_{N_{\text{step}}-1} \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \\ \vdots \\ u_{N_{\text{step}}-1} \end{bmatrix} = 2E \begin{bmatrix} u_1 \\ u_2 \\ \vdots \\ u_{N_{\text{step}}-1} \end{bmatrix}. \quad (1.81)$$

The boundary conditions $u(R_{\min}) = u(R_{\max}) = 0$ are imposed simply by leaving the two endpoints out of the matrix, whose dimension is therefore $N_{\text{step}} - 1$. The eigenvectors are the wave functions on the grid, normalised with the trapezoidal rule,

$$\int_{R_{\min}}^{R_{\max}} |u(x)|^2 dx \longrightarrow h \sum_i u_i^2 = 1. \quad (1.82)$$

Since the matrix is already tridiagonal, the Householder step is unnecessary and we go straight to `qrli`. Alternatively, and this is the point of the exercise, we may pretend that the matrix is far too large to store and use the Lanczos algorithm with a matrix-free operator.

```
class SchrodingerDiagonalization:
    """Discretise -u'' + V(x) u = 2E u and diagonalise the result."""

    def __init__(self, potential, rmin=-10.0, rmax=10.0, nsteps=400):
        self.potential = potential
        self.h = (rmax - rmin) / nsteps
        # interior mesh points only; u vanishes at the two endpoints
        self.x = rmin + self.h * np.arange(1, nsteps)
        self.n = len(self.x)
```

```

@property
def diagonal(self):
    return 2.0 / self.h**2 + self.potential(self.x)

@property
def offdiagonal(self):
    return np.full(self.n - 1, -1.0 / self.h**2)

def matvec(self, v):
    """Apply the Hamiltonian without ever forming the matrix."""
    w = self.diagonal * v
    w[:-1] -= v[1:] / self.h**2
    w[1:] -= v[:-1] / self.h**2
    return w

def solve(self, k=5):
    """The k lowest eigenvalues and normalised eigenfunctions."""
    values, vectors = SymmetricEigenSolver.tqli(self.diagonal,
                                                self.offdiagonal)

    order = np.argsort(values)
    return (values[order][:k],
            self._normalise(vectors[:, order][:, :k]))

def solve_lanczos(self, k=5, m=None):
    """The same eigenvalues from the Lanczos algorithm."""
    lan = Lanczos(matvec=self.matvec, n=self.n)
    values, vectors = lan.ritz(10*k if m is None else m, k=k)
    return values, self._normalise(vectors)

```

Table 1.4 shows the result on $x \in [-10, 10]$. The eigenvalues converge to the exact values $1, 3, 5, \dots$ as h^2 : doubling N_{step} divides the error by four, exactly as the truncation error of the three-point formula demands. Running the same problem through the Lanczos algorithm with $m = 200$ steps reproduces the three lowest eigenvalues to 6×10^{-13} , and the computed wave functions agree with the analytic $u_0(x) = \pi^{-1/4} e^{-x^2/2}$ to 7×10^{-5} .

N_{step}	$2E_0$	$2E_1$	$2E_2$	Max error
100	0.9974937026	2.9874430342	4.9672772610	1.0×10^{-1}
200	0.9993746086	2.9968714733	4.9918612670	2.6×10^{-2}
400	0.9998437256	2.9992185301	4.9979678946	6.4×10^{-3}
800	0.9999609360	2.9998046738	4.9994921341	1.6×10^{-3}
exact	1	3	5	

Table 1.4 The one-dimensional harmonic oscillator obtained by diagonalising Eq. (1.81) on $x \in [-10, 10]$. The error decreases as h^2 .

Many-body connection. This calculation is the one-particle ancestor of everything that follows. Expanding a state in a finite basis and diagonalising the resulting matrix is exactly what full configuration interaction does, the only difference being that the basis consists of Slater determinants rather than grid points, and that the dimension grows exponentially with the number of particles rather than linearly with the number of mesh points. The eigenvectors computed here are the natural single-particle basis $\{|\phi_i\rangle\}$ of Section 1.8 from which those Slater determinants are built. Chapter 2 takes the first step in that direction.

1.28 The singular value decomposition

Every algorithm so far has assumed a square matrix, and most of them a symmetric one. Many of the objects we meet in many-body physics are neither. The coefficients of a state written across a partition of the system into two parts form a rectangular array; a set of sampled wave functions collected as columns forms a rectangular array; the two-body matrix elements arranged in a composite pair index form a matrix whose row and column dimensions have nothing to do with the dimension of the Hilbert space. For all of these we need a decomposition that does not require squareness, and preferably one that does not require diagonalisability either.

Recall that a square matrix can be diagonalised if and only if it is *normal*, that is if $\mathbf{X}\mathbf{X}^T = \mathbf{X}^T\mathbf{X}$ in the real case or $\mathbf{X}\mathbf{X}^\dagger = \mathbf{X}^\dagger\mathbf{X}$ in the complex one. Not every matrix satisfies this. The simplest counterexample is

$$\mathbf{X} = \begin{bmatrix} 1 & -1 \\ 1 & -1 \end{bmatrix}, \quad (1.83)$$

whose determinant vanishes and for which $\mathbf{X}\mathbf{X}^T \neq \mathbf{X}^T\mathbf{X}$. It is *defective*: it cannot be diagonalised, and it has no inverse.

The singular value decomposition has no such restriction.

The theorem.

Any $m \times n$ matrix $\mathbf{X}$, real or complex, square or rectangular, can be written

$$\mathbf{X} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T, \quad (1.84)$$

where $\mathbf{U}$ is an $m \times m$ orthogonal matrix, $\mathbf{U}\mathbf{U}^T = \mathbf{U}^T\mathbf{U} = \mathbf{I}_m$, where $\mathbf{V}$ is an $n \times n$ orthogonal matrix, $\mathbf{V}\mathbf{V}^T = \mathbf{V}^T\mathbf{V} = \mathbf{I}_n$, and where $\mathbf{\Sigma}$ is an $m \times n$ matrix which is zero except on the main diagonal, on which sit the *singular values*

$$\sigma_0 \geq \sigma_1 \geq \cdots \geq \sigma_{r-1} \geq 0, \quad r = \min(m, n). \quad (1.85)$$

The columns of $\mathbf{U}$ are the *left* singular vectors and the columns of $\mathbf{V}$ the *right* singular vectors. The decomposition always exists. For the defective matrix of Eq. (1.83) it reads

$$\mathbf{X} = \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix} \begin{bmatrix} 2 & 0 \\ 0 & 0 \end{bmatrix} \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & -1 \\ 1 & 1 \end{bmatrix} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T, \quad (1.86)$$

with $\sigma_0 = 2$ and $\sigma_1 = 0$. A matrix that cannot be diagonalised at all is decomposed without difficulty, and the vanishing singular value states plainly what the determinant only hinted at.

Geometry.

Equation (1.84) says that every linear map is a rotation, followed by a scaling along the new axes, followed by another rotation. The singular values are the scale factors, and $\|\mathbf{X}\|_2 = \sigma_0$, which is the statement of Section 1.13 that the induced 2-norm measures the largest stretching a matrix can produce.

Relation to the eigenvalue problem.

The SVD is not an independent construction; it is the symmetric eigenvalue problem of Section 1.21 in disguise. Using Eq. (1.84) and the orthogonality of $\mathbf{U}$,

$$\mathbf{X}^T\mathbf{X} = \mathbf{V}\mathbf{\Sigma}^T\mathbf{U}^T\mathbf{U}\mathbf{\Sigma}\mathbf{V}^T = \mathbf{V}\mathbf{\Sigma}^T\mathbf{\Sigma}\mathbf{V}^T \equiv \mathbf{V}\tilde{\mathbf{\Sigma}}^2\mathbf{V}^T, \quad (1.87)$$

where $\tilde{\Sigma}$ is the $n \times n$ diagonal matrix of singular values. Multiplying from the right by $\mathbf{V}$ gives

$$(\mathbf{X}^T \mathbf{X}) \mathbf{v}_i = \sigma_i^2 \mathbf{v}_i, \quad (1.88)$$

so the right singular vectors are the eigenvectors of $\mathbf{X}^T \mathbf{X}$ and the squared singular values its eigenvalues. In exactly the same way

$$\mathbf{X} \mathbf{X}^T = \mathbf{U} \Sigma \Sigma^T \mathbf{U}^T, \quad (\mathbf{X} \mathbf{X}^T) \mathbf{u}_i = \sigma_i^2 \mathbf{u}_i, \quad (1.89)$$

with the remaining $m - r$ eigenvalues equal to zero. Every singular value is therefore the non-negative square root of an eigenvalue of $\mathbf{X}^T \mathbf{X}$, and if $\mathbf{X}$ is itself symmetric the singular values are the absolute values of its eigenvalues. These two relations will do a great deal of work in the sections that follow: $\mathbf{X} \mathbf{X}^T$ and $\mathbf{X}^T \mathbf{X}$ will turn out to be the two reduced density matrices of a bipartite quantum state.

The economy-size decomposition.

If $m > n$ the last $m - n$ columns of $\mathbf{U}$ are multiplied by zeros in Σ and can never contribute. Removing them, and the corresponding rows of zeros in Σ , gives the economy-size decomposition, in which $\mathbf{U}$ is $m \times n$ and Σ is $n \times n$. If $n > m$ one keeps instead the first m columns of $\mathbf{V}$. Nothing is lost and both storage and work are reduced; this is the default in most library routines.

How not to compute it.

Equation (1.88) suggests an algorithm: form $\mathbf{X}^T \mathbf{X}$, diagonalise it with the Householder and QL machinery of Sections 1.24 and 1.25, take square roots. This is mathematically correct and numerically poor. From Eq. (1.11) the condition number of the product is the square of that of $\mathbf{X}$,

$$\kappa_2(\mathbf{X}^T \mathbf{X}) = \kappa_2(\mathbf{X})^2, \quad (1.90)$$

so half of the available significant digits are destroyed before the eigenvalue solver is even called. Table 1.5 shows the damage for a 60×12 matrix constructed with singular values spread logarithmically between 1 and 10^{-7} : the direct SVD returns all twelve to machine precision, the detour through $\mathbf{X}^T \mathbf{X}$ loses seven digits on the smallest.

Exact σ_i	From the SVD	From $\mathbf{X}^T \mathbf{X}$	Relative error
1.00000000×10^0	1.00000000×10^0	1.00000000×10^0	4.4×10^{-16}
$6.57933225 \times 10^{-4}$	$6.57933225 \times 10^{-4}$	$6.57933225 \times 10^{-4}$	2.0×10^{-11}
$1.87381742 \times 10^{-6}$	$1.87381742 \times 10^{-6}$	$1.87381889 \times 10^{-6}$	7.8×10^{-7}
$4.32876128 \times 10^{-7}$	$4.32876128 \times 10^{-7}$	$4.32893930 \times 10^{-7}$	4.1×10^{-5}
$1.00000000 \times 10^{-7}$	$1.00000000 \times 10^{-7}$	$9.99683051 \times 10^{-8}$	3.2×10^{-4}

Table 1.5 Singular values of a 60×12 matrix with $\kappa_2(\mathbf{X}) = 10^7$, computed directly and through the eigenvalues of $\mathbf{X}^T \mathbf{X}$, for which $\kappa_2 = 10^{14}$. Forming the product costs roughly half the significant digits.

The stable algorithm never forms the product. It reduces $\mathbf{X}$ to bidiagonal form by Householder reflectors applied alternately from the left and the right – the same reflectors as in Section 1.24 – and then runs an implicitly shifted QR sweep on the bidiagonal matrix. This is the Golub-Kahan algorithm, and it is what `numpy.linalg.svd` calls through LAPACK. We use the library routine in what follows: the interest here lies in what the decomposition is used for.

Many-body connection. Equations (1.88) and (1.89) are the reason the SVD keeps reappearing in many-body physics. Whenever a quantity is built as a product of a matrix with its own transpose – a density matrix, an overlap matrix, a two-body interaction written in a pair index – the SVD of the underlying rectangular object gives its spectrum directly, with better accuracy and often at lower cost than forming the product at all.

1.29 Rank, the pseudoinverse and ill-conditioned bases

The *rank* of a matrix is the number of linearly independent columns, and the SVD reads it off immediately: it is the number of non-zero singular values. In floating-point arithmetic one asks instead for the *numerical rank*, the number of singular values above a threshold, typically $\max(m, n) \varepsilon \sigma_0$ with ε the machine precision. A matrix is rank deficient when some columns are linear combinations of others, and near rank deficient – which is the practically dangerous case – when they nearly are.

The condition number of Section 1.13 is likewise a statement about singular values,

$$\kappa_2(\mathbf{X}) = \frac{\sigma_0}{\sigma_{r-1}}, \quad (1.91)$$

which is infinite exactly when $\mathbf{X}$ is singular. The SVD thus not only detects near-singularity, it quantifies it.

The Moore-Penrose pseudoinverse.

When $\mathbf{X}^{-1}$ does not exist, or exists but is useless because $\mathbf{X}$ is nearly singular, the SVD provides the natural substitute

$$\mathbf{X}^+ = \mathbf{V} \boldsymbol{\Sigma}^+ \mathbf{U}^T, \quad (1.92)$$

where $\boldsymbol{\Sigma}^+$ is obtained from $\boldsymbol{\Sigma}$ by transposing it and replacing every non-zero singular value by its reciprocal, *leaving the zeros as zeros*. In practice one treats as zero every singular value below some $\varepsilon_{\text{cut}} \sigma_0$. If $\mathbf{X}$ is invertible and well conditioned, $\mathbf{X}^+ = \mathbf{X}^{-1}$; otherwise $\mathbf{X}^+ \mathbf{b}$ is the least-squares solution of smallest norm. Inverting a singular value of 10^{-13} would amplify the rounding noise in the corresponding direction by 10^{13} ; discarding it simply declares that the problem contains no information about that direction, which is the truth.

Near-linearly-dependent basis sets.

This is not an abstract concern. Many-body calculations are routinely carried out in non-orthogonal single-particle bases – Gaussians in quantum chemistry, harmonic-oscillator states in different oscillator lengths, generator coordinates in nuclear physics – and such bases become linearly dependent long before one would like them to. The eigenvalue problem is then the generalised one,

$$\mathbf{H} \mathbf{c} = E \mathbf{S} \mathbf{c}, \quad S_{ij} = \langle \phi_i | \phi_j \rangle, \quad (1.93)$$

with $\mathbf{S}$ the overlap matrix. Diagonalising $\mathbf{S} = \mathbf{V} \mathbf{s} \mathbf{V}^T$ and forming

$$\mathbf{X} = \mathbf{V} \mathbf{s}^{-1/2} \quad (1.94)$$

turns Eq. (1.93) into the ordinary problem $(\mathbf{X}^T \mathbf{H} \mathbf{X}) \mathbf{c}' = E \mathbf{c}'$. The transformation involves $s_i^{-1/2}$, so a tiny overlap eigenvalue is precisely where the procedure breaks down. *Canonical orthogonalisation* is the fix: drop the eigenvectors whose eigenvalues lie below a threshold before forming Eq. (1.94).

Table 1.6 shows this for twelve Gaussians of width 1.4 placed on a uniform grid, a deliberately over-complete basis of the kind used in quantum chemistry. The overlap matrix has eigenvalues spanning thirteen orders of magnitude, from 6.9 down to 6.5×10^{-13} , and a condition number of 1.1×10^{13} : in double precision only three significant digits would survive. Discarding three basis directions removes six orders of magnitude of ill-conditioning at negligible cost in completeness.

Threshold	Functions kept	Effective $\kappa_2(\mathcal{S})$
none	12	1.1×10^{13}
10^{-12}	11	6.9×10^{10}
10^{-8}	9	2.1×10^7
10^{-5}	7	3.3×10^4

Table 1.6 Canonical orthogonalisation of a basis of twelve overlapping Gaussians. Eigenvectors of the overlap matrix with eigenvalue below the threshold, relative to the largest, are discarded.

```
class CanonicalOrthogonalization:
    """Orthogonalise a non-orthogonal basis with a near-singular overlap.

    The generalised problem  $H c = E S c$  is reduced to an ordinary one by
     $X = V s^{-1/2}$ , keeping only the eigenvalues of  $S$  above a threshold.
    """

    def __init__(self, S, threshold=1.0e-8):
        self.S = np.asarray(S, dtype=float)
        w, V = np.linalg.eigh(self.S)
        order = np.argsort(w)[::-1]
        self.eigenvalues = w[order]
        self.eigenvectors = V[:, order]
        keep = self.eigenvalues > threshold * self.eigenvalues[0]
        self.kept = int(np.sum(keep))
        self.X = self.eigenvectors[:, keep] / np.sqrt(self.eigenvalues[keep])

    def solve_generalized(self, H):
        """Solve  $H c = E S c$  through the orthogonalised problem."""
        values, vectors = np.linalg.eigh(self.X.T @ H @ self.X)
        return values, self.X @ vectors
```

Many-body connection. The threshold is a physical choice, not a numerical detail. Setting it too low keeps directions that carry only rounding noise and poisons the whole calculation; setting it too high throws away genuine variational freedom and raises the computed energy. The same trade-off appears in the generator-coordinate method, where the overlap kernel of non-orthogonal mean-field states is routinely rank deficient, and in explicitly correlated Gaussian expansions, where it is the principal practical obstacle to enlarging the basis.

1.30 Low-rank approximation and the Schmidt decomposition

We now come to the property that makes the SVD indispensable in many-body physics. Truncating the sum

$$\mathbf{X} = \sum_{k=0}^{r-1} \sigma_k \mathbf{u}_k \mathbf{v}_k^T \quad (1.95)$$

after χ terms gives a matrix $\mathbf{X}_\chi$ of rank χ . The Eckart-Young theorem states that this is the *best* rank- χ approximation to $\mathbf{X}$, in both the 2-norm and the Frobenius norm, and that the error it makes is

$$\|\mathbf{X} - \mathbf{X}_\chi\|_2 = \sigma_\chi, \quad \|\mathbf{X} - \mathbf{X}_\chi\|_F^2 = \sum_{k \geq \chi} \sigma_k^2. \quad (1.96)$$

No cleverer choice of χ vectors exists. If the singular values decay quickly, a matrix with mn entries is captured by $\chi(m+n+1)$ numbers.

The Schmidt decomposition.

Consider a quantum system divided into two parts, A and B , with orthonormal bases $\{|i\rangle_A\}$ and $\{|j\rangle_B\}$. Any state of the whole system can be written

$$|\Psi\rangle = \sum_{ij} C_{ij} |i\rangle_A \otimes |j\rangle_B, \quad (1.97)$$

and the coefficients form a matrix $\mathbf{C}$ which is rectangular whenever the two parts have different dimensions. Inserting its SVD $\mathbf{C} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$ and absorbing $\mathbf{U}$ and $\mathbf{V}$ into new basis states $|u_k\rangle_A = \sum_i U_{ik} |i\rangle_A$ and $|v_k\rangle_B = \sum_j V_{jk} |j\rangle_B$, which are orthonormal because $\mathbf{U}$ and $\mathbf{V}$ are orthogonal, gives

$$|\Psi\rangle = \sum_{k=0}^{\chi-1} \sigma_k |u_k\rangle_A \otimes |v_k\rangle_B \quad (1.98)$$

This is the *Schmidt decomposition*. A double sum over the product basis has collapsed to a single sum, and the SVD has done all the work. The singular values σ_k are the Schmidt coefficients, and normalisation $\langle\Psi|\Psi\rangle = 1$ gives $\sum_k \sigma_k^2 = 1$.

The reduced density matrices follow at once. Tracing out B ,

$$\rho_A = \text{Tr}_B |\Psi\rangle\langle\Psi| = \mathbf{C}\mathbf{C}^T = \mathbf{U}\mathbf{\Sigma}^2\mathbf{U}^T, \quad \rho_B = \mathbf{C}^T\mathbf{C} = \mathbf{V}\mathbf{\Sigma}^2\mathbf{V}^T, \quad (1.99)$$

which is Eqs. (1.88) and (1.89) read as physics. The two reduced density matrices have the *same* non-zero eigenvalues σ_k^2 , however different the sizes of the two parts, and their eigenvectors are the Schmidt states. The number of non-zero σ_k is the *Schmidt rank*: it equals one precisely when $|\Psi\rangle$ is a product state, so a single singular value is the signature of an unentangled state. The *entanglement entropy*

$$S_A = -\text{Tr} \rho_A \ln \rho_A = -\sum_k \sigma_k^2 \ln \sigma_k^2 \quad (1.100)$$

measures how far from a product the state is, and it is manifestly symmetric between A and B .

Truncation.

Keeping only the χ largest Schmidt coefficients gives, by Eq. (1.96), the best possible approximation to $|\Psi\rangle$ of that rank, with discarded weight

$$\epsilon_\chi = \sum_{k \geq \chi} \sigma_k^2. \quad (1.101)$$

This single quantity is the truncation error of the density-matrix renormalisation group, and χ is its bond dimension. Whether a many-body state can be compressed is therefore a question about the decay of the Schmidt spectrum, and that decay is what area laws for the entanglement entropy assert: for ground states of local, gapped Hamiltonians the entropy across a cut grows with the *boundary* of the cut rather than with the volume of the subsystem, so a modest χ suffices no matter how large the system. Matrix product

states are the systematic version of this observation, obtained by applying Eq. (1.98) at every bond of a one-dimensional lattice in turn.

The pairing model across a cut.

Take the ground state of the pairing model of Section 1.26 and split the twelve single-particle levels into the lowest six and the highest six. A seniority-zero basis state is a set of occupied levels, which factorises into its part in A and its part in B , so the coefficients form a 64×64 matrix. Its singular values are the Schmidt coefficients across the cut. Table 1.7 gives the discarded weight as a function of χ : three Schmidt states out of 64 already capture all but 5×10^{-3} of the state, and ten capture all but 8×10^{-7} . The entanglement entropy is $S_A = 0.9238$, well below the maximum $\ln 64 = 4.16$ that an arbitrary state of this size could carry.

Schmidt states kept, χ	Discarded weight ε_χ
1	4.831×10^{-1}
2	7.461×10^{-2}
3	5.448×10^{-3}
5	1.854×10^{-4}
7	7.363×10^{-6}
10	7.530×10^{-7}
15	4.225×10^{-9}

Table 1.7 Schmidt spectrum of the pairing-model ground state (12 levels, 6 pairs, $\delta = 1$, $g = 0.5$) across the cut separating the lowest six levels from the highest six. The coefficient matrix is 64×64 and the entanglement entropy is 0.9238.

```

class SchmidtDecomposition:
    """The SVD of a bipartitioned state vector.

    For  $|\Psi\rangle = \sum_{ij} C_{ij} |i\rangle_A |j\rangle_B$  the SVD  $C = U \Sigma V^T$  gives
     $|\Psi\rangle = \sum_k \sigma_k |u_k\rangle_A |v_k\rangle_B$ . The eigenvalues of both reduced
    density matrices are  $\sigma_k^2$ .
    """

    def __init__(self, C):
        self.C = np.asarray(C, dtype=float)
        self.U, self.sigma, self.Vt = np.linalg.svd(C, full_matrices=False)
        self.weights = self.sigma**2 # normalised: sum_k sigma_k^2 = 1

    def schmidt_rank(self, tol=1.0e-10):
        return int(np.sum(self.sigma > tol * self.sigma[0]))

    def entanglement_entropy(self, base=np.e):
        w = self.weights[self.weights > 1.0e-16]
        entropy = -np.sum(w * np.log(w)) / np.log(base)
        return float(entropy) if entropy > 0.0 else 0.0

    def truncation_error(self, k):
        """Weight discarded when only k Schmidt states are kept."""
        return float(np.sum(self.weights[k:]))

    def truncate(self, k):
        """The best rank-k approximation to the state, renormalised."""
        C = (self.U[:, :k] * self.sigma[:k]) @ self.Vt[:k]

```

```
return C / np.linalg.norm(C)
```

Many-body connection. The Schmidt decomposition is the bridge between linear algebra and the modern theory of many-body wave functions. The density-matrix renormalisation group truncates the Schmidt spectrum at every bond; matrix product states, projected entangled pair states and the whole family of tensor networks are built by iterating Eq. (1.98); and the entanglement entropy, computed from nothing more than a list of singular values, is what tells us in advance whether such a compression can work at all.

Quantum computing connection. The same numbers govern how hard a state is to simulate classically. A state of N qubits whose Schmidt rank across every cut is bounded by χ can be stored in $\mathcal{O}(N\chi^2)$ numbers and evolved classically; only states whose Schmidt rank grows exponentially require a quantum computer. Entanglement, quantified by Eq. (1.100), is thus simultaneously the resource that quantum algorithms exploit and the obstacle that classical ones must overcome.

1.31 Natural orbitals and occupation numbers

For two identical particles the Schmidt decomposition has a name of its own. Write the two-particle wave function in a single-particle basis $\{\phi_p\}$,

$$\Psi(x_1, x_2) = \sum_{pq} C_{pq} \phi_p(x_1) \phi_q(x_2), \quad (1.102)$$

where symmetry under exchange makes $\mathbf{C}$ symmetric (or antisymmetric). The SVD of $\mathbf{C}$, which for a symmetric matrix is its eigen-decomposition up to signs, gives the *natural expansion*

$$\Psi(x_1, x_2) = \sum_k \sigma_k \chi_k(x_1) \chi_k(x_2), \quad (1.103)$$

a single sum in terms of the transformed orbitals $\chi_k = \sum_p U_{pk} \phi_p$. These are Löwdin's *natural orbitals*, and they are the eigenvectors of the one-body density matrix

$$\rho_{pq} = \langle \Psi | a_p^\dagger a_q | \Psi \rangle, \quad \boldsymbol{\rho} \propto \mathbf{C} \mathbf{C}^\dagger = \mathbf{U} \boldsymbol{\Sigma}^2 \mathbf{U}^\dagger, \quad (1.104)$$

with eigenvalues $n_k \propto \sigma_k^2$: the *occupation numbers*. For a single Slater determinant exactly N occupation numbers equal one and the rest vanish; correlation depletes the occupied orbitals and populates the others, and the deviation of the largest n_k from unity is a direct measure of how far the state is from a mean-field one.

Table 1.8 illustrates this for two particles in a one-dimensional harmonic trap interacting through a softened Coulomb potential of strength g , solved on a grid of 40 points by full diagonalisation of the 1600-dimensional two-particle problem. At $g = 0$ the ground state is a product, there is a single natural orbital with $n_0 = 1$, and the entanglement entropy vanishes exactly. Switching on the interaction depletes n_0 and populates n_1, n_2 and so on, and the entropy rises.

Many-body connection. Natural orbitals are the optimal single-particle basis in a precise sense: for a given number of retained orbitals, no other choice gives a better representation of the state, which is Eq. (1.96) again. This is why natural-orbital truncation is the standard way of reducing a correlated basis in quantum chemistry and in nuclear structure, and why occupation numbers are reported as a diagnostic: a system whose largest occupation number is 0.99 is well described by a single determinant and Hartree-Fock theory of Chapter 2 will be a good starting point, while one whose occupation numbers are spread over several orbitals is strongly correlated and demands a multireference treatment.

g	E_0	S_A	n_0	n_1	χ for 10^{-8}
0	1.992536	0.000000	1.000000	0.000000	1
1	3.173866	0.040874	0.993528	0.006017	10
2	4.245117	0.139618	0.970732	0.027407	11
5	6.755062	0.507970	0.821806	0.168066	11

Table 1.8 Natural occupation numbers and entanglement entropy for two particles in a harmonic trap with a softened Coulomb repulsion of strength g , on a grid of 40 points. The last column gives the number of natural orbitals needed to recover all but 10^{-8} of the norm.

1.32 Low-rank factorisation of the two-body interaction

The last application concerns not the wave function but the Hamiltonian. The two-body matrix elements V_{pqrs} form a four-index array with n^4 entries, where n is the number of single-particle states, and for n of a few hundred this is already the dominant storage cost of a many-body calculation. Reading the pairs (pq) and (rs) as composite indices turns the array into an ordinary $n^2 \times n^2$ matrix,

$$V_{(pq),(rs)} = \iint dx dy \varphi_p(x) \varphi_q(x) K(x, y) \varphi_r(y) \varphi_s(y), \quad (1.105)$$

and for a Coulomb-like kernel K this matrix is symmetric and positive semidefinite. Its eigen-decomposition, equivalently its Cholesky factorisation of Section 1.15, gives

$$V_{(pq),(rs)} = \sum_{\gamma=0}^{M-1} L_{pq}^{\gamma} L_{rs}^{\gamma}, \quad L_{pq}^{\gamma} = \sqrt{\lambda_{\gamma}} U_{(pq),\gamma}. \quad (1.106)$$

The decisive fact is that the eigenvalues λ_{γ} decay rapidly, because the pair densities $\varphi_p \varphi_q$ are far from linearly independent as functions. Keeping only $M \ll n^2$ of them reduces the storage from n^4 to Mn^2 and turns the contraction of the interaction with a wave function into a sequence of matrix products.

Table 1.9 shows the decay for the twelve-Gaussian basis of Section 1.29, for which there are 144 pair indices. Nine vectors reproduce the interaction to one part in 10^6 , and ten to one part in 10^8 : a reduction of the rank by more than a factor of ten.

Target weight	Vectors M	$\ V - L^T L\ / \ V\ $
10^{-4}	7	2.1×10^{-5}
10^{-6}	9	1.7×10^{-7}
10^{-8}	10	1.2×10^{-8}

Table 1.9 Low-rank factorisation of the two-body matrix elements in a basis of twelve Gaussians, for which the pair index runs over 144 values. The leading eigenvalues of the pair matrix are 44.7, 15.6, 5.18, 1.38, 0.309, 0.055, ..., a geometric decay.

```
class TwoBodyFactorization:
    """Low-rank factorisation of a two-body interaction.

    The pair-index matrix  $V_{\{(pq), (rs)\}}$  is positive semidefinite for a
    Coulomb-like interaction, so it factorises as a sum of outer products,
     $V = \sum_{\gamma} L^{\gamma} (L^{\gamma})^T$ , with few terms.
    """

    def __init__(self, V):
```

```

self.V = np.asarray(V, dtype=float)
w, U = np.linalg.eigh(0.5 * (self.V + self.V.T))
order = np.argsort(w)[::-1]
self.eigenvalues = w[order]
self.vectors = U[:, order]

def rank_for(self, tol):
    """Smallest M such that the discarded weight is below tol."""
    w = np.maximum(self.eigenvalues, 0.0)
    tail = np.cumsum(w[::-1])[::-1]
    below = np.where(tail <= tol * np.sum(w))[0]
    return int(below[0]) if len(below) else len(w)

def cholesky_vectors(self, M):
    """The M leading vectors  $L^\gamma$ ."""
    w = np.maximum(self.eigenvalues[:M], 0.0)
    return (self.vectors[:, :M] * np.sqrt(w)).T

```

Many-body connection. Equation (1.106) is known as the Cholesky decomposition of the two-electron integrals, or, when the auxiliary vectors are chosen from a fitted basis rather than computed, as density fitting or the resolution of the identity. It is what makes large-basis coupled-cluster and many-body perturbation theory calculations feasible, since the rate-determining contractions can be performed on the factors L^γ rather than on the full four-index array.

Quantum computing connection. Applying the same idea twice gives the *double factorisation*: each matrix L_{pq}^γ is itself diagonalised, $L^\gamma = \mathbf{W}^\gamma \mathbf{d}^\gamma (\mathbf{W}^\gamma)^T$, so that the Hamiltonian becomes a short sum of squares of one-body operators. Since a one-body operator maps under the Jordan-Wigner transformation of Section 1.7 onto a small number of commuting Pauli strings, this reduces the number of terms that must be measured in a variational quantum eigensolver, and the cost of block encoding the Hamiltonian for qubitisation, from $\mathcal{O}(n^4)$ to $\mathcal{O}(n^3)$ or better. The same singular values control the truncation in both the classical and the quantum setting.

1.33 Summary and the programs

The chapter has moved from notation to algorithms. The linear algebra of Sections 1.1-1.11 supplies the language in which quantum mechanics is written: states are vectors, observables are Hermitian matrices, symmetries and basis changes are unitary transformations, and the physical content of an operator is its spectral decomposition. The numerical sections supply the means of extracting numbers from that language.

Three dividing lines are worth keeping in mind. The first separates direct from iterative methods. Direct methods – Gaussian elimination, LU, Cholesky, Householder with QL – deliver the complete answer in a fixed $\mathcal{O}(n^3)$ number of operations, and require the matrix to be stored. Iterative methods – Jacobi, Gauss-Seidel, over-relaxation, conjugate gradient, Lanczos – deliver an approximation that improves with each step, and require only the action of the matrix on a vector. The second line separates dense from structured problems: a tridiagonal system costs $\mathcal{O}(n)$ rather than $\mathcal{O}(n^3)$, and recognising structure is usually worth more than any amount of tuning.

The third line is the one drawn by the singular value decomposition, and it is of a different kind. It separates the part of a problem that carries information from the part that does not. The same list of singular values tells us the numerical rank of a matrix, how far a basis is from linear dependence, how much of a wave function survives a truncation, how entangled a state is across a cut, which single-particle orbitals are actually occupied, and how few auxiliary vectors are needed to represent a two-body interaction. Sections 1.28-1.32 are five readings of one decomposition, and the thread connecting them – Eq. (1.96), that truncating the singular values is the best approximation of its rank – is the mathematical justification for

a large part of modern many-body theory, from natural-orbital bases to the density-matrix renormalisation group.

The complete programs are collected in `doc/BookManybody/BookMaterial/Programs`:

- `direct_solvers.py` – Gaussian elimination, LU, Cholesky and the tridiagonal solver, with determinants and inverses.
- `iterative_solvers.py` – Jacobi, Gauss-Seidel, over-relaxation and the conjugate gradient method, in both a matrix-based and a matrix-free form.
- `householder.py` – Householder tridiagonalisation, the QL algorithm with implicit shifts, and the power method.
- `lanczos.py` – the Lanczos algorithm with optional full reorthogonalisation, together with the pairing-model Hamiltonian used in Table 1.3.
- `schrodinger_diagonalization.py` – Schrödinger’s equation by diagonalisation, solved both directly and by Lanczos.
- `svd.py` – the singular value decomposition: rank, the pseudoinverse, canonical orthogonalisation, the Schmidt decomposition with entanglement entropies and natural occupation numbers, and the low-rank factorisation of a two-body interaction. It also contains the two physical models used in Sections 1.30–1.32: two particles in a trap, and the pairing model bipartitioned across a cut.

Each file runs as a script and reproduces the numbers quoted in this chapter. An executable version of the same material, in which the reader can vary the parameters, is available as a Jupyter notebook in the accompanying Jupyter-book.

1.34 Exercises

Warm-up exercises

1. **Inner and outer products.** Let $\mathbf{u} = [1 - i \ 2 + i]^T$ and $\mathbf{v} = [1 \ i]^T$. (a) Compute $\mathbf{u}^\dagger \mathbf{v}$ and $\mathbf{v}^\dagger \mathbf{u}$. (b) Verify $(\mathbf{u}^\dagger \mathbf{v})^* = \mathbf{v}^\dagger \mathbf{u}$. (c) Compute the outer product $\mathbf{u}\mathbf{v}^\dagger$ and show $\text{tr}(\mathbf{u}\mathbf{v}^\dagger) = \mathbf{v}^\dagger \mathbf{u}$.
2. **Special matrix types.** For each matrix below, identify which type(s) from Table 1.1 it belongs to, and verify your answer using the element conditions in the table: (a) σ_x ; (b) σ_z ; (c) the Hadamard gate $H = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}$.
3. **Pauli algebra.** (a) Prove $\sigma_x^2 = \sigma_y^2 = \sigma_z^2 = \mathbf{I}$. (b) Compute all commutators $[\sigma_i, \sigma_j]$ for $i, j \in \{x, y, z\}$. (c) Which Pauli matrix has $|0\rangle$ and $|1\rangle$ as its eigenbasis? What are the eigenvalues?
4. **Projection operators.** (a) Verify that $\mathbf{P} = |0\rangle\langle 0|$ satisfies $\mathbf{P}^2 = \mathbf{P}$ and $\det(\mathbf{P}) = 0$. (b) Show $\mathbf{P} + \mathbf{Q} = \mathbf{I}$ and $\mathbf{P}\mathbf{Q} = 0$. (c) What does $\mathbf{P}|\psi\rangle$ equal for $|\psi\rangle = \alpha|0\rangle + \beta|1\rangle$?
5. **Commutator identities.** Prove the following for arbitrary operators $\hat{A}, \hat{B}, \hat{C}$: (a) $[\hat{A} + \hat{B}, \hat{C}] = [\hat{A}, \hat{C}] + [\hat{B}, \hat{C}]$; (b) $[\hat{A}, \hat{B}\hat{C}] = [\hat{A}, \hat{B}]\hat{C} + \hat{B}[\hat{A}, \hat{C}]$; (c) $[\hat{A}, [\hat{B}, \hat{C}]] + [\hat{B}, [\hat{C}, \hat{A}]] + [\hat{C}, [\hat{A}, \hat{B}]] = 0$ (Jacobi identity).
6. **Shared eigenstates.** Prove that if $[\hat{A}, \hat{B}] = 0$ then $\hat{A}$ and $\hat{B}$ share a common eigenbasis.
7. **Tensor products.** (a) Write out all four two-qubit basis states $|00\rangle, |01\rangle, |10\rangle, |11\rangle$ as column vectors of length 4 and verify they are mutually orthogonal. (b) Compute $C_X|00\rangle, C_X|01\rangle, C_X|10\rangle, C_X|11\rangle$. Explain the action of the CNOT gate in words.
8. **Unitary transformation.** Let $\mathbf{U}$ be a 2×2 unitary matrix and $\{|v_0\rangle, |v_1\rangle\}$ an ONB. Prove that $\{\mathbf{U}|v_0\rangle, \mathbf{U}|v_1\rangle\}$ is also an ONB.
9. **Spectral decomposition.** (a) Find the eigenvalues and normalised eigenvectors of σ_x . (b) Write the spectral decomposition $\sigma_x = \sum_i \lambda_i |i\rangle\langle i|$ explicitly. (c) Verify that $\sigma_x|\psi_a\rangle = \lambda_a|\psi_a\rangle$ for each eigenvector.

Exercise session, week 34: unitary transformations and determinants

The exercises of this session are mainly meant as reminders of the specific linear algebra elements that we shall use throughout the course.

Exercise 1: unitary transformations and orthogonality.

In many-body theories it is common to expand a new basis in terms of a known one. Well-known examples are the analytical solutions of the quantum mechanical harmonic oscillator, or the hydrogen atom with a Coulomb interaction between the electron and the nucleus. Varying the expansion coefficients leads to a variational minimum which is expected to lie closer to the true ground state of a given Hamiltonian. This is the background for mean-field theories, full configuration interaction theory and many other many-body methods, and the expansion coefficients are normally the matrix elements of a unitary or orthogonal matrix.

Assume that the matrix $\mathbf{U}$ is unitary with dimension $n \times n$ and matrix elements $u_{p\lambda}$. Using Dirac notation we represent a new vector $|\psi_p\rangle$ of length n in terms of another vector $|\phi_\lambda\rangle$ of length n through

$$|\psi_p\rangle = \mathbf{U}|\phi_\lambda\rangle. \quad (1.107)$$

It is common to assume that the basis vectors $|\phi_\lambda\rangle$ are normalised and orthogonal, and that they are the eigenvectors of an eigenvalue problem

$$\mathbf{O}|\phi_\lambda\rangle = o_\lambda|\phi_\lambda\rangle, \quad (1.108)$$

where $\mathbf{O}$ is a Hermitian matrix of dimension $n \times n$ and o_λ one of its n eigenvalues. We have thus defined our new basis by performing a unitary (or orthogonal) transformation.

- Write down the properties of unitary and orthogonal matrices.
- Write down the properties of Hermitian matrices.
- Assuming that the coefficients $u_{p\lambda}$ belong to a unitary or orthogonal transformation, show that the new basis is orthogonal and normalised as well.
- Show that the eigenvalues obtained in the new basis are the same as those of the original basis.

Exercise 2: determinants.

- Consider the determinant

$$\begin{vmatrix} \alpha_1 b_{11} + \alpha_2 b_{12} & a_{12} \\ \alpha_1 b_{21} + \alpha_2 b_{22} & a_{22} \end{vmatrix} = \alpha_1 \begin{vmatrix} b_{11} & a_{12} \\ b_{21} & a_{22} \end{vmatrix} + \alpha_2 \begin{vmatrix} b_{12} & a_{12} \\ b_{22} & a_{22} \end{vmatrix},$$

which generalises to the linearity in a column expressed by Eq. (2.12). Use it to show that, with every element a_{ij} given as a linear combination $a_{ij} = \sum_{k=1}^n b_{ik}c_{kj}$,

$$\begin{vmatrix} \sum_k b_{1k}c_{k1} & \cdots & \sum_k b_{1k}c_{kn} \\ \vdots & & \vdots \\ \sum_k b_{nk}c_{k1} & \cdots & \sum_k b_{nk}c_{kn} \end{vmatrix} = |\mathbf{C}| |\mathbf{B}|,$$

where $|\mathbf{B}|$ and $|\mathbf{C}|$ are the determinants of the $n \times n$ matrices with elements b_{ij} and c_{ij} . This is the product rule (2.13), and we shall use it repeatedly in the mean-field theories of the following chapters.

- With the new basis of the previous exercise, now specialised to a basis of single-particle functions,

$$\psi_p(x) = \sum_{\lambda} C_{p\lambda} \phi_{\lambda}(x),$$

show that the determinant in the new basis $\psi_p(x)$ can be written as (omitting the factor $1/\sqrt{N!}$, with N the number of particles)

$$\begin{vmatrix} \psi_p(x_1) & \cdots & \psi_p(x_N) \\ \vdots & \ddots & \vdots \\ \psi_t(x_1) & \cdots & \psi_t(x_N) \end{vmatrix} = \begin{vmatrix} \sum_{\lambda} C_{p\lambda} \phi_{\lambda}(x_1) & \cdots & \sum_{\lambda} C_{p\lambda} \phi_{\lambda}(x_N) \\ \vdots & \ddots & \vdots \\ \sum_{\lambda} C_{t\lambda} \phi_{\lambda}(x_1) & \cdots & \sum_{\lambda} C_{t\lambda} \phi_{\lambda}(x_N) \end{vmatrix},$$

which is nothing but $|\mathbf{C}||\mathbf{\Phi}|$, with $|\mathbf{\Phi}|$ the determinant built from the basis functions $\phi_{\lambda}(x)$.

c) Show that the new determinant differs from $|\mathbf{\Phi}|$ by a complex phase only.

Answers, week 34

Exercise 1a.

A matrix $\mathbf{U} \in \mathbb{C}^{n \times n}$ is *unitary* if

$$\mathbf{U}^{\dagger} \mathbf{U} = \mathbf{U} \mathbf{U}^{\dagger} = \mathbf{I}, \quad \text{equivalently} \quad \mathbf{U}^{-1} = \mathbf{U}^{\dagger}. \quad (1.109)$$

In terms of matrix elements this reads $\sum_k u_{ik} u_{jk}^* = \sum_k u_{ki}^* u_{kj} = \delta_{ij}$: the rows form an orthonormal set, and so do the columns. Three consequences follow at once and were used throughout Section 1.9. The inner product is preserved, $(\mathbf{U}\mathbf{x})^{\dagger}(\mathbf{U}\mathbf{y}) = \mathbf{x}^{\dagger}\mathbf{y}$, hence so is the norm; the eigenvalues all have modulus one, since $\mathbf{U}\mathbf{x} = \lambda\mathbf{x}$ with $\|\mathbf{x}\| = 1$ gives $|\lambda|^2 = \mathbf{x}^{\dagger}\mathbf{U}^{\dagger}\mathbf{U}\mathbf{x} = 1$; and $\|\mathbf{U}\| = 1$, because $|\mathbf{U}|^*|\mathbf{U}| = |\mathbf{U}^{\dagger}\mathbf{U}| = 1$. A unitary matrix is thus a rotation, and it is perfectly reversible.

A *real* unitary matrix is called *orthogonal*: $\mathbf{O}^T \mathbf{O} = \mathbf{O} \mathbf{O}^T = \mathbf{I}$ and $\mathbf{O}^{-1} = \mathbf{O}^T$, with $|\mathbf{O}| = \pm 1$. These are the matrix types listed in Table 1.1.

Exercise 1b.

A matrix is *Hermitian* if $\mathbf{A} = \mathbf{A}^{\dagger}$, that is $a_{ij} = a_{ji}^*$; a real Hermitian matrix is symmetric. As shown in Section 1.11, the eigenvalues of a Hermitian matrix are real, eigenvectors belonging to distinct eigenvalues are orthogonal, and the matrix is diagonalisable by a unitary transformation, $\mathbf{U}^{\dagger} \mathbf{A} \mathbf{U} = \text{diag}(\lambda_0, \dots, \lambda_{n-1})$. It follows that expectation values $\langle \psi | \mathbf{A} | \psi \rangle$ are real, which is why observables are represented by Hermitian operators.

Exercise 1c.

Write the transformation in components, $|\psi_p\rangle = \sum_{\lambda} u_{p\lambda} |\phi_{\lambda}\rangle$, and use $\langle \phi_{\lambda} | \phi_{\mu} \rangle = \delta_{\lambda\mu}$:

$$\langle \psi_p | \psi_q \rangle = \sum_{\lambda\mu} u_{p\lambda}^* u_{q\mu} \langle \phi_{\lambda} | \phi_{\mu} \rangle = \sum_{\lambda} u_{p\lambda}^* u_{q\lambda} = (\mathbf{U} \mathbf{U}^{\dagger})_{qp} = \delta_{pq}. \quad (1.110)$$

Orthonormality of the new basis is therefore *equivalent* to unitarity of the coefficient matrix; nothing else about $\mathbf{U}$ is needed, and nothing less would do. This is the statement proved geometrically in Eq. (2.19).

Exercise 1d.

In the new basis the operator $\mathbf{O}$ is represented by $\mathbf{O}' = \mathbf{U} \mathbf{O} \mathbf{U}^{\dagger}$. This is a similarity transformation of the kind studied in Section 1.22, and by Eq. (1.54) it leaves the spectrum unchanged: if $\mathbf{O}|\phi_{\lambda}\rangle = o_{\lambda}|\phi_{\lambda}\rangle$ then

$$\mathcal{O}'(\mathbf{U}|\phi_\lambda\rangle) = \mathbf{U}\mathbf{O}\mathbf{U}^\dagger\mathbf{U}|\phi_\lambda\rangle = o_\lambda(\mathbf{U}|\phi_\lambda\rangle). \quad (1.111)$$

The eigenvalues are the same and the eigenvectors are the rotated ones. Physically: the spectrum is a property of the operator, not of the basis we happen to write it in.

Exercise 2a.

Let $\mathbf{A} = \mathbf{B}\mathbf{C}$, so that column j of $\mathbf{A}$ is $\mathbf{a}_j = \sum_k c_{kj}\mathbf{b}_k$, a linear combination of the columns of $\mathbf{B}$. Applying the linearity (2.12) to the first column, then to the second, and so on through all n columns, gives

$$|\mathbf{A}| = \sum_{k_1=1}^n \sum_{k_2=1}^n \cdots \sum_{k_n=1}^n c_{k_1 1} c_{k_2 2} \cdots c_{k_n n} |\mathbf{b}_{k_1} \mathbf{b}_{k_2} \cdots \mathbf{b}_{k_n}|. \quad (1.112)$$

There are n^n terms, but a determinant with two equal columns vanishes, so every term in which two indices coincide drops out. Only the $n!$ terms in which $(k_1, \dots, k_n)$ is a permutation σ of $(1, \dots, n)$ survive, and reordering the columns of such a term into their natural order produces the sign $(-1)^\sigma$ together with $|\mathbf{B}|$. Hence

$$|\mathbf{A}| = \left(\sum_{\sigma} (-1)^\sigma c_{\sigma(1)1} c_{\sigma(2)2} \cdots c_{\sigma(n)n} \right) |\mathbf{B}| = |\mathbf{C}| |\mathbf{B}|, \quad (1.113)$$

where the bracket is precisely the definition (2.10) of $|\mathbf{C}|$.

Two remarks are worth making. First, the derivation is combinatorial and says nothing about how to *compute* either determinant; for that one uses the LU factorisation of Section 1.15 and Eq. (1.22), at a cost of $\mathcal{O}(n^3)$ rather than the $\mathcal{O}(n!)$ of the sum above. Second, the result immediately gives $|\mathbf{A}^{-1}| = 1/|\mathbf{A}|$ and, for a unitary matrix, $|\mathbf{U}| = 1$.

Exercise 2b.

The matrix whose determinant we want has elements

$$M_{pj} = \psi_p(x_j) = \sum_{\lambda} C_{p\lambda} \phi_{\lambda}(x_j) = (\mathbf{C}\Phi)_{pj}, \quad \Phi_{\lambda j} = \phi_{\lambda}(x_j), \quad (1.114)$$

so $\mathbf{M} = \mathbf{C}\Phi$ and part a) gives directly

$$|\mathbf{M}| = |\mathbf{C}| |\Phi|. \quad (1.115)$$

The rows of $\mathbf{M}$ are labelled by the new single-particle states and its columns by the particle coordinates, exactly as in Eq. (2.47).

Exercise 2c.

If $\mathbf{C}$ is unitary then $|\mathbf{C}|^* |\mathbf{C}| = |\mathbf{C}^\dagger \mathbf{C}| = |\mathbf{I}| = 1$, so $|\mathbf{C}| = e^{i\theta}$ for some real θ and

$$|\mathbf{M}| = e^{i\theta} |\Phi|. \quad (1.116)$$

The two determinants differ by a global phase. Since all physical predictions depend on $|\Psi|^2$ or on expectation values $\langle \Psi | \hat{O} | \Psi \rangle$, in which the phase cancels, the two describe the *same* state. This is the result exploited in Section 2.5: one may rotate the occupied orbitals among themselves at will without changing any physics, which is the freedom underlying Hartree-Fock theory and Thouless' theorem.

Chapter 2

From Linear Algebra to Many-Body Physics

Chapter 1 developed a language and a set of algorithms. This chapter puts them to work. We shall see that the central object of fermionic many-body theory, the Slater determinant, is a determinant in exactly the sense of Section 2.3; that changing the single-particle basis is a unitary transformation of the kind studied in Section 1.9; that the Hartree-Fock energy is minimised by solving a symmetric eigenvalue problem with the machinery of Sections 1.24 and 1.25; and that the obstacle which makes the whole subject difficult – the exponential growth of the Hilbert space – is the tensor product of Section 1.10.

Nothing in this chapter is new mathematics. What is new is the physics that the mathematics is asked to carry.

2.1 The many-body Hamiltonian

Throughout these lectures we assume that the interacting part of the Hamiltonian can be approximated by a two-body interaction, so the full N -particle Hamiltonian takes the form

$$\hat{H} = \hat{H}_0 + \hat{H}_I = \sum_{i=1}^N \hat{h}_0(x_i) + \sum_{i<j}^N \hat{v}(r_{ij}), \quad (2.1)$$

where x_i denotes the coordinates – position and spin – of particle i . The one-body operator splits into a kinetic and a potential part,

$$\hat{h}_0(x_i) = \hat{t}(x_i) + \hat{u}_{\text{ext}}(x_i), \quad (2.2)$$

where $\hat{u}_{\text{ext}}$ is typically a harmonic oscillator potential, the Coulomb attraction of a nucleus, or the self-consistent Hartree-Fock potential derived later in this chapter.

We shall assume that the single-particle functions ψ_α are eigenfunctions of $\hat{h}_0$,

$$\hat{h}_0(x) \psi_\alpha(x) = (\hat{t}(x) + \hat{u}_{\text{ext}}(x)) \psi_\alpha(x) = \varepsilon_\alpha \psi_\alpha(x), \quad (2.3)$$

with ε_α the *unperturbed* or non-interacting single-particle energies. In the absence of $\hat{H}_I$ the total energy would simply be the sum of the occupied ε_α . Everything that makes many-body physics hard is contained in the correction to that statement.

Equation (2.3) is an eigenvalue problem, and in practice we solve it exactly as in Section 1.27: discretise on a grid, or expand in a finite basis, and diagonalise the resulting matrix with the algorithms of Sections 1.24 and 1.25. The harmonic oscillator eigenfunctions computed there are the single-particle basis used in the numerical examples of this chapter.

Finally, the Hamiltonian (2.1) is *symmetric* under the interchange of any two particles. Writing $\hat{P}_{ij}$ for the operator that permutes particles i and j ,

$$[\hat{H}, \hat{P}_{ij}] = 0, \quad (2.4)$$

so the eigenfunctions of $\hat{H}$ can be chosen to be eigenfunctions of $\hat{P}_{ij}$ as well.

2.2 The Pauli exclusion principle and antisymmetry

Let Ψ_λ be such a simultaneous eigenfunction. Then

$$\hat{P}_{ij}\Psi_\lambda(x_1, \dots, x_i, \dots, x_j, \dots, x_N) = \beta \Psi_\lambda(x_1, \dots, x_i, \dots, x_j, \dots, x_N), \quad (2.5)$$

and since applying the permutation twice returns the original function, $\beta^2 = 1$ and $\beta = \pm 1$. For bosons $\beta = +1$; for fermions, which are our concern here, the Pauli principle requires

$$\beta = -1, \quad (2.6)$$

so the wave function changes sign under every transposition.

An arbitrary product of single-particle functions has no such symmetry, and must be projected onto the antisymmetric subspace. The projector is the *antisymmetrisation operator*

$$\hat{A} = \frac{1}{N!} \sum_p (-1)^p \hat{P}, \quad (2.7)$$

where the sum runs over all $N!$ permutations $\hat{P}$ and p is the parity of each, that is the number of transpositions required to build it. Two properties of $\hat{A}$ will do all the work in Section 2.8. First, since $\hat{H}_0$ and $\hat{H}_I$ are invariant under every permutation,

$$[\hat{H}_0, \hat{A}] = [\hat{H}_I, \hat{A}] = 0. \quad (2.8)$$

Second, $\hat{A}$ is idempotent,

$$\hat{A}^2 = \hat{A}, \quad (2.9)$$

because permuting an already antisymmetrised function reproduces it up to the sign that the prefactor $(-1)^p$ removes. A projector, in the sense of Section 1.6: $\hat{A}$ projects the full N -particle space onto its antisymmetric subspace, exactly as $\mathbf{P} = |0\rangle\langle 0|$ projects a qubit onto a computational basis state.

2.3 Determinants, and how they are actually computed

For an $n \times n$ matrix $\mathbf{A}$, the determinant is defined by summing over all $n!$ permutations of the column indices,

$$|\mathbf{A}| = |\mathbf{A}| = \sum_{i=1}^{n!} (-1)^{p_i} \hat{P}_i a_{11} a_{22} \cdots a_{nn}, \quad (2.10)$$

where $\hat{P}_i$ permutes the column indices and p_i is the number of transpositions needed to restore the natural ordering $a_{11} a_{22} \cdots a_{nn}$. For $n = 2$ the sum has two terms,

$$\begin{vmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{vmatrix} = (-1)^0 a_{11} a_{22} + (-1)^1 a_{12} a_{21}, \quad (2.11)$$

the second obtained by interchanging the column indices 1 and 2.

The resemblance between Eq. (2.10) and the antisymmetriser (2.7) is not an accident, and it is the reason determinants appear in fermionic physics at all: both are the signed sum over the symmetric group, and the determinant is simply what the antisymmetriser produces when applied to a product of single-particle functions.

Two properties we shall need.

The determinant is linear in each column separately,

$$\begin{vmatrix} a_{11} & \cdots & \sum_{k=1}^n c_k b_{1k} & \cdots \\ a_{21} & \cdots & \sum_{k=1}^n c_k b_{2k} & \cdots \\ \vdots & & \vdots & \\ a_{n1} & \cdots & \sum_{k=1}^n c_k b_{nk} & \cdots \end{vmatrix} = \sum_{k=1}^n c_k \begin{vmatrix} a_{11} & \cdots & b_{1k} & \cdots \\ a_{21} & \cdots & b_{2k} & \cdots \\ \vdots & & \vdots & \\ a_{n1} & \cdots & b_{nk} & \cdots \end{vmatrix}, \quad (2.12)$$

and if every column of $\mathbf{A}$ is a linear combination of the columns of $\mathbf{B}$ with coefficient matrix $\mathbf{C}$, that is $a_{ij} = \sum_k b_{ik} c_{kj}$ or $\mathbf{A} = \mathbf{B}\mathbf{C}$, then

$$|\mathbf{A}| = |\mathbf{B}| \cdot |\mathbf{C}|. \quad (2.13)$$

Equation (2.13) follows by applying Eq. (2.12) to each column in turn; the reader is invited to verify it for $n = 2$. It is the single most important algebraic fact in this chapter, and Section 2.5 is nothing but its physical interpretation.

A warning about Eq. (2.10).

The definition is a definition, not an algorithm. Evaluating a determinant by summing over permutations costs $\mathcal{O}(n!)$ operations, which for $n = 20$ is 2×10^{18} and for $n = 25$ exceeds the age of the universe on any machine. The LU decomposition of Section 1.15 computes the same number in $\mathcal{O}(n^3)$ operations as the product of the diagonal elements of $\mathbf{U}$, Eq. (1.22). Table 2.1 makes the comparison. Every Slater determinant evaluated in a Monte Carlo or configuration-interaction code is computed this way, never from Eq. (2.10).

n	From Eq. (2.10), $n!$	From LU, $\sim \frac{2}{3}n^3$
5	1.2×10^2	8.3×10^1
10	3.6×10^6	6.7×10^2
20	2.4×10^{18}	5.3×10^3
50	3.0×10^{64}	8.3×10^4

Table 2.1 Operation counts for evaluating an $n \times n$ determinant from the definition and from the LU decomposition of Section 1.15.

2.4 The Slater determinant

We now have everything needed to write down an antisymmetric N -particle wave function. Given N single-particle orbitals $\psi_\alpha, \psi_\beta, \dots, \psi_\sigma$, the *Slater determinant* is

$$\Phi(x_1, \dots, x_N; \alpha, \beta, \dots, \sigma) = \frac{1}{\sqrt{N!}} \begin{vmatrix} \psi_\alpha(x_1) & \psi_\alpha(x_2) & \cdots & \psi_\alpha(x_N) \\ \psi_\beta(x_1) & \psi_\beta(x_2) & \cdots & \psi_\beta(x_N) \\ \vdots & \vdots & \ddots & \vdots \\ \psi_\sigma(x_1) & \psi_\sigma(x_2) & \cdots & \psi_\sigma(x_N) \end{vmatrix}, \quad (2.14)$$

where $\alpha, \beta, \dots, \sigma$ are the quantum numbers labelling the occupied orbitals. The two defining physical properties follow from two elementary properties of determinants:

- interchanging two particles interchanges two *columns*, which reverses the sign – the wave function is antisymmetric, as Eq. (2.6) demands;
- placing two particles in the same orbital gives two identical *rows*, and the determinant vanishes – no two fermions may occupy the same single-particle state.

The Pauli principle is thus not an extra postulate bolted onto the formalism; it is a property of determinants.

An equivalent form uses the antisymmetriser directly. Writing the unsymmetrised product, the *Hartree function*,

$$\Phi_H(x_1, \dots, x_N; \alpha, \dots, \nu) = \psi_\alpha(x_1) \psi_\beta(x_2) \cdots \psi_\nu(x_N), \quad (2.15)$$

we have

$$\Phi = \frac{1}{\sqrt{N!}} \sum_p (-)^p \hat{P} \Phi_H = \sqrt{N!} \hat{A} \Phi_H. \quad (2.16)$$

This is the form used in the derivations of Section 2.8.

Link to Chapter 1. The Hartree function (2.15) is precisely a tensor product, $|\alpha\rangle \otimes |\beta\rangle \otimes \cdots \otimes |\nu\rangle$, of the kind introduced in Section 1.10. The Slater determinant is its antisymmetric projection. A tensor product of N single-particle states is the least entangled state there is, and its Schmidt rank across any cut is one; antisymmetrisation introduces exactly the correlation demanded by statistics and no more. This is the precise sense in which a Slater determinant is an *uncorrelated* state, and it is made quantitative by the occupation numbers of Section 1.31: for a single Slater determinant they are exactly 1 for the N occupied orbitals and exactly 0 for all the rest. Any deviation from those values is a measure of genuine, dynamical correlation.

The numerical companion `slaterdeterminant.py` verifies both statements. Evaluating Eq. (2.14) for three particles at (x_1, x_2, x_3) and at (x_2, x_1, x_3) gives values that cancel to machine zero, and putting two particles in the same orbital gives $\Phi = -6 \times 10^{-18}$. The eigenvalues of the density matrix $\rho_{\lambda\mu} = \sum_i C_{i\lambda}^* C_{i\mu}$ come out as $(1, 1, 1, 0, 0, 0, 0, 0)$ to machine precision.

2.5 Single-particle bases and unitary transformations

In practice one rarely knows the right orbitals in advance. The standard strategy is to expand them in a fixed, known basis $\{\phi_\lambda\}$ – harmonic oscillator functions, hydrogen-like functions, plane waves, Gaussians – and to vary the expansion coefficients,

$$\psi_p(x) = \sum_\lambda C_{p\lambda} \phi_\lambda(x). \quad (2.17)$$

The basis $\{\phi_\lambda\}$ is normally chosen as the eigenfunctions of the one-body Hamiltonian, Eq. (2.3), and is orthonormal. In Dirac notation Eq. (2.17) reads

$$|p\rangle = \sum_\lambda C_{p\lambda} |\lambda\rangle, \quad \phi_\lambda(\mathbf{r}) = \langle \mathbf{r} | \lambda \rangle. \quad (2.18)$$

Two questions arise immediately. Does the new basis remain orthonormal? And what happens to the Slater determinant? Chapter 1 has already answered both.

Orthonormality is preserved if and only if $\mathbf{C}$ is unitary.

This was proved in Section 1.9: for an orthonormal set $\mathbf{v}_j^T \mathbf{v}_i = \delta_{ij}$ and a transformation $\mathbf{w}_i = \mathbf{U} \mathbf{v}_i$,

$$\mathbf{w}_j^T \mathbf{w}_i = \mathbf{v}_j^T \underbrace{\mathbf{U}^T \mathbf{U}}_{=I} \mathbf{v}_i = \delta_{ij}, \quad (2.19)$$

so $\langle p|q \rangle = \delta_{pq}$ follows whenever $(C_{p\lambda})$ is unitary. We shall use this constantly: it is what allows us to search over single-particle bases without ever having to re-orthogonalise.

The Slater determinant picks up a phase.

Substituting Eq. (2.17) into Eq. (2.14), every entry of the determinant becomes a sum $\sum_{\lambda} C_{p\lambda} \phi_{\lambda}(x_i)$, that is, the matrix of orbitals is the product of the coefficient matrix $\mathbf{C}$ with the matrix Φ of basis functions. By the product rule (2.13),

$$\frac{1}{\sqrt{N!}} \begin{vmatrix} \sum_{\lambda} C_{p\lambda} \phi_{\lambda}(x_1) & \cdots & \sum_{\lambda} C_{p\lambda} \phi_{\lambda}(x_N) \\ \vdots & \ddots & \vdots \\ \sum_{\lambda} C_{i\lambda} \phi_{\lambda}(x_1) & \cdots & \sum_{\lambda} C_{i\lambda} \phi_{\lambda}(x_N) \end{vmatrix} = |\mathbf{C}| \cdot |\Phi|. \quad (2.20)$$

If $\mathbf{C}$ is unitary then $||\mathbf{C}|| = 1$, so the rotated Slater determinant differs from the original only by an overall phase. It describes the *same* physical state.

Many-body connection. Equation (2.20) is the algebraic foundation of a fact used everywhere in this book: one may perform any unitary rotation among the occupied orbitals without changing the many-body state, the density, or the total energy. In Hartree-Fock theory this is the freedom that makes canonical and localised orbitals equally valid descriptions of the same determinant, and it is the content of Thouless' theorem. In configuration interaction and coupled-cluster theory it is why the choice of reference orbitals is a matter of convenience rather than principle. The numerical check in `slaterdeterminant.py` makes it concrete: a random orthogonal rotation of three occupied orbitals leaves the energy unchanged to 2.7×10^{-15} , while a rotation that mixes occupied and unoccupied orbitals changes it by more than a factor of two.

Truncation, and what happens when the basis is not orthogonal.

A complete basis $\{|\lambda\rangle\}$ contains infinitely many states, and any calculation must truncate the sum in Eq. (2.17). This is the first and often the largest approximation in a many-body calculation, and its convergence has to be studied explicitly.

Furthermore, some of the most useful bases are not orthogonal at all. Gaussian basis sets in quantum chemistry, oscillator states of different oscillator length, and generator-coordinate states all have non-trivial overlap matrices $S_{ij} = \langle \phi_i | \phi_j \rangle$, and Eq. (2.19) no longer applies. The eigenvalue problem then becomes the generalised one, $\mathbf{H}\mathbf{c} = E\mathbf{S}\mathbf{c}$, treated in Section 1.29: one diagonalises $\mathbf{S}$, discards the directions whose eigenvalues lie below a threshold, and transforms with $\mathbf{X} = \mathbf{V}\mathbf{S}^{-1/2}$. The threshold is a physical choice, and Table 1.6 shows how much ill-conditioning it removes.

2.6 Particle-hole notation and the Fermi level

Before deriving the energy it is worth fixing the index conventions that will be used for the rest of the book. Given a reference Slater determinant – usually the one obtained by filling the N lowest single-particle levels – the single-particle states divide into those that are occupied in the reference and those that are not. The dividing line is the *Fermi level*, denoted F . We write

$$\begin{aligned}
i, j, k, l, \dots &\leq F \quad \text{occupied states, or } \textit{hole} \text{ states,} \\
a, b, c, d, \dots &> F \quad \text{unoccupied states, or } \textit{particle} \text{ states,} \\
p, q, r, s, \dots &\quad \text{general single-particle states.}
\end{aligned} \tag{2.21}$$

The names come from the picture in which the reference determinant is treated as a new vacuum: exciting a particle from an occupied state i to an unoccupied state a creates a particle in a and leaves a hole in i . Chapter 3 develops this into the particle-hole formalism of second quantisation, where it becomes the natural language for perturbation theory, coupled-cluster theory and the shell model. For now it is simply a convention, but it is one worth adopting at once: every equation from here on is easier to read with it.

2.7 The variational principle

Let E_0 be the exact ground-state energy of $\hat{H}$. For any normalised trial function Φ ,

$$E_0 \leq E[\Phi] = \int \Phi^* \hat{H} \Phi d\tau, \quad \int \Phi^* \Phi d\tau = 1, \tag{2.22}$$

with the shorthand $d\tau = dx_1 dx_2 \dots dx_N$. The proof is the spectral decomposition of Section 1.11 and nothing more. Expanding Φ in the exact eigenstates, $|\Phi\rangle = \sum_k c_k |\Psi_k\rangle$ with $\hat{H}|\Psi_k\rangle = E_k |\Psi_k\rangle$ and $\sum_k |c_k|^2 = 1$, gives

$$E[\Phi] = \sum_k |c_k|^2 E_k \geq E_0 \sum_k |c_k|^2 = E_0, \tag{2.23}$$

with equality only when Φ is the ground state itself.

This one inequality is the engine of most of what follows. Restricting the trial function to a single Slater determinant and minimising gives Hartree-Fock theory; allowing a linear combination of determinants gives configuration interaction; parametrising it by a quantum circuit gives the variational quantum eigensolver of Section 2.16. It is also the reason the Lanczos algorithm of Section 1.26 is so well suited to many-body problems: the lowest Ritz value is the minimum of the Rayleigh quotient over the Krylov space, and therefore an upper bound on E_0 that can only improve as the space grows.

2.8 Expectation values with a Slater determinant

We now compute $E[\Phi]$ for the Slater determinant (2.14). The calculation uses only Eqs. (2.8), (2.9) and the orthonormality of the single-particle functions, and it is worth following in full: the same manipulations recur throughout many-body theory, and second quantisation exists largely to make them automatic.

The one-body contribution.

Using the form (2.16) of the Slater determinant,

$$\int \Phi^* \hat{H}_0 \Phi d\tau = N! \int \Phi_H^* \hat{A} \hat{H}_0 \hat{A} \Phi_H d\tau. \tag{2.24}$$

Because $\hat{H}_0$ commutes with $\hat{A}$ we may move the left-hand antisymmetriser through it, and because $\hat{A}^2 = \hat{A}$ the two then collapse into one,

$$\int \Phi^* \hat{H}_0 \Phi d\tau = N! \int \Phi_H^* \hat{H}_0 \hat{A} \Phi_H d\tau. \tag{2.25}$$

Replacing $\hat{A}$ by its definition (2.7), which cancels the factor $N!$, and $\hat{H}_0$ by the sum of one-body operators,

$$\int \Phi^* \hat{H}_0 \Phi d\tau = \sum_{i=1}^N \sum_p (-)^p \int \Phi_H^* \hat{h}_0(x_i) \hat{P} \Phi_H d\tau. \quad (2.26)$$

Now the key step. The integrand factorises into single-particle integrals, and $\hat{h}_0(x_i)$ touches only one coordinate. If $\hat{P}$ permutes any two particles, at least one of the remaining factors is an overlap $\langle \mu | \nu \rangle$ between two *different* orbitals, which vanishes. Only the identity permutation survives,

$$\int \Phi^* \hat{H}_0 \Phi d\tau = \sum_{i=1}^N \int \Phi_H^* \hat{h}_0(x_i) \Phi_H d\tau, \quad (2.27)$$

and integrating out the $N - 1$ spectator coordinates, each of which contributes unity, leaves

$$\int \Phi^* \hat{H}_0 \Phi d\tau = \sum_{\mu=1}^N \int \psi_\mu^*(x) \hat{h}_0 \psi_\mu(x) dx. \quad (2.28)$$

Introducing the shorthand

$$\langle \mu | \hat{h}_0 | \mu \rangle = \int \psi_\mu^*(x) \hat{h}_0 \psi_\mu(x) dx, \quad (2.29)$$

this is simply

$$\int \Phi^* \hat{H}_0 \Phi d\tau = \sum_{\mu=1}^N \langle \mu | \hat{h}_0 | \mu \rangle. \quad (2.30)$$

The one-body energy is the sum of the diagonal matrix elements over the occupied orbitals – the result one would have guessed, now derived.

The two-body contribution.

The same three steps apply,

$$\int \Phi^* \hat{H}_I \Phi d\tau = N! \int \Phi_H^* \hat{A} \hat{H}_I \hat{A} \Phi_H d\tau = \sum_{i < j}^N \sum_p (-)^p \int \Phi_H^* \hat{v}(r_{ij}) \hat{P} \Phi_H d\tau, \quad (2.31)$$

but the conclusion is different. Because $\hat{v}(r_{ij})$ depends on *two* coordinates, the permutation that exchanges precisely those two particles no longer produces a vanishing overlap. Two permutations survive: the identity, and the transposition $\hat{P}_{ij}$. Hence

$$\int \Phi^* \hat{H}_I \Phi d\tau = \sum_{i < j}^N \int \Phi_H^* \hat{v}(r_{ij}) (1 - \hat{P}_{ij}) \Phi_H d\tau. \quad (2.32)$$

Writing out the two terms and integrating over the spectator coordinates,

$$\begin{aligned} \int \Phi^* \hat{H}_I \Phi d\tau = \frac{1}{2} \sum_{\mu=1}^N \sum_{\nu=1}^N \left[\int \psi_\mu^*(x_i) \psi_\nu^*(x_j) \hat{v}(r_{ij}) \psi_\mu(x_i) \psi_\nu(x_j) dx_i dx_j \right. \\ \left. - \int \psi_\mu^*(x_i) \psi_\nu^*(x_j) \hat{v}(r_{ij}) \psi_\nu(x_i) \psi_\mu(x_j) dx_i dx_j \right], \end{aligned} \quad (2.33)$$

where the factor $\frac{1}{2}$ compensates for the unrestricted double sum, which counts every pair twice. The first term is the *direct* or *Hartree* term: it is the classical electrostatic energy of one charge distribution in the field of another. The second is the *exchange* or *Fock* term, and it has no classical counterpart whatsoever; it exists solely because the wave function is antisymmetric, and it is the sole trace left in the energy by the transposition $\hat{P}_{ij}$ in Eq. (2.32).

Notation for the matrix elements.

Writing the single-particle functions as overlaps, $\psi_\alpha(x) = \langle x|\alpha\rangle$, we abbreviate the two integrals as

$$\langle \mu \nu | \hat{v} | \mu \nu \rangle = \int \psi_\mu^*(x_i) \psi_\nu^*(x_j) \hat{v}(r_{ij}) \psi_\mu(x_i) \psi_\nu(x_j) dx_i dx_j, \quad (2.34)$$

$$\langle \mu \nu | \hat{v} | \nu \mu \rangle = \int \psi_\mu^*(x_i) \psi_\nu^*(x_j) \hat{v}(r_{ij}) \psi_\nu(x_i) \psi_\mu(x_j) dx_i dx_j, \quad (2.35)$$

and define the *antisymmetrised* matrix element

$$\langle \mu \nu | \hat{v} | \mu \nu \rangle_{AS} = \langle \mu \nu | \hat{v} | \mu \nu \rangle - \langle \mu \nu | \hat{v} | \nu \mu \rangle. \quad (2.36)$$

Note that $\langle \mu \mu | \hat{v} | \mu \mu \rangle_{AS} = 0$: a particle does not interact with itself. The self-interaction that the Hartree term would otherwise contain is cancelled exactly by the exchange term, which is one of the reasons Hartree-Fock theory is better behaved than the Hartree theory that preceded it.

2.9 The energy as a functional of the coefficients

Collecting Eqs. (2.30) and (2.33), the expectation value of the Hamiltonian in a Slater determinant is

$$E[\Phi] = \sum_{\mu=1}^N \langle \mu | \hat{h}_0 | \mu \rangle + \frac{1}{2} \sum_{\mu=1}^N \sum_{\nu=1}^N \langle \mu \nu | \hat{v} | \mu \nu \rangle_{AS}, \quad (2.37)$$

or, in the particle-hole notation of Section 2.6,

$$E[\Phi] = \sum_{i=1}^N \langle i | \hat{h}_0 | i \rangle + \frac{1}{2} \sum_{i,j=1}^N \langle ij | \hat{v} | ij \rangle_{AS}. \quad (2.38)$$

This is a functional of the occupied orbitals. To turn it into something we can minimise we insert the expansion (2.17) and let the coefficients carry the variation,

$$E[\Psi] = \sum_{i=1}^N \sum_{\alpha\beta} C_{i\alpha}^* C_{i\beta} \langle \alpha | \hat{h}_0 | \beta \rangle + \frac{1}{2} \sum_{i,j=1}^N \sum_{\alpha\beta\gamma\delta} C_{i\alpha}^* C_{j\beta}^* C_{i\gamma} C_{j\delta} \langle \alpha\beta | \hat{v} | \gamma\delta \rangle_{AS}. \quad (2.39)$$

The Greek indices run over the full, fixed basis; the Latin indices run over the N occupied orbitals only. Everything on the right-hand side except the C 's is computed once and stored.

Minimising Eq. (2.39) with respect to $C_{i\alpha}^*$ subject to the orthonormality constraint yields the Hartree-Fock equations. We defer the derivation, but the structure of the answer can be read off immediately and is worth stating now. Define the *density matrix* and the *Fock matrix*

$$\rho_{\gamma\delta} = \sum_{i=1}^N C_{i\gamma}^* C_{i\delta}, \quad f_{\alpha\beta} = \langle \alpha | \hat{h}_0 | \beta \rangle + \sum_{\gamma\delta} \rho_{\gamma\delta} \langle \alpha\gamma | \hat{v} | \beta\delta \rangle_{AS}. \quad (2.40)$$

The condition for a stationary point is then

$$\sum_{\beta} f_{\alpha\beta} C_{i\beta} = \epsilon_i C_{i\alpha}, \quad (2.41)$$

a Hermitian eigenvalue problem of exactly the kind solved in Sections 1.24 and 1.25 – with the single complication that the matrix $\mathbf{f}$ depends, through $\boldsymbol{\rho}$, on the eigenvectors we are trying to find. One therefore guesses, builds $\mathbf{f}$, diagonalises, rebuilds and repeats until nothing changes. This is the *self-consistent field*

procedure, and it is an iteration in the sense of Section 1.17, with a matrix diagonalisation in place of a single sweep.

Summary: Hartree-Fock is Chapter 1 applied to fermions. Expand the states in an orthonormal basis (Section 1.8); change basis with a unitary transformation (Section 1.9), which leaves the determinant invariant up to a phase by the product rule (Section 2.3); minimise a quadratic form subject to a constraint, which produces a Hermitian eigenvalue problem (Section 1.11); solve it by Householder reduction followed by the QL algorithm (Sections 1.24 and 1.25); and iterate to self-consistency. Every step has already been developed.

The minimal implementation in `slaterdeterminant.py` does exactly this, using the eigenvalue solver written in Chapter 1. For spinless fermions in a one-dimensional trap with a softened Coulomb repulsion, and a basis of eight harmonic oscillator orbitals, it converges in nine iterations and lowers the energy below that of the naive determinant built from the N lowest oscillator states, as shown in Table 2.2.

N	E (lowest N orbitals)	E (Hartree-Fock)	Gain
2	2.69018342	2.66308878	-2.7×10^{-2}
3	6.46801021	6.39016133	-7.8×10^{-2}
4	11.77084712	11.62305385	-1.5×10^{-1}

Table 2.2 Self-consistent Hartree-Fock energies for N spinless fermions in a one-dimensional harmonic trap with a softened Coulomb interaction, in a basis of eight oscillator orbitals. The middle column is the energy of the determinant built from the N lowest oscillator states; the Hartree-Fock solution optimises the orbitals and always lies below it.

For $N = 2$ the exact ground state in the same truncated basis can be obtained by diagonalising the Hamiltonian in the space of antisymmetric pairs, and the difference

$$E_{\text{corr}} = E_{\text{exact}} - E_{\text{HF}} \quad (2.42)$$

defines the *correlation energy*. Here it is -3.86×10^{-3} : the system is weakly correlated, and the largest amplitude in the exact state, 0.9930, sits on the Hartree-Fock configuration. Recovering the remaining 0.7% of the amplitude is what every method in the rest of this book is built to do.

2.10 The price of the two-body matrix elements

Equation (2.39) contains a sum over four basis indices, so a calculation in a basis of n single-particle states requires n^4 two-body matrix elements. For $n = 100$ that is 10^8 numbers, for $n = 1000$ it is 10^{12} , and long before the many-body problem becomes hard the mere storage of the interaction does.

Symmetries help. For a real, local interaction the matrix elements obey

$$\langle \alpha\beta | \hat{v} | \gamma\delta \rangle = \langle \beta\alpha | \hat{v} | \delta\gamma \rangle = \langle \gamma\delta | \hat{v} | \alpha\beta \rangle, \quad (2.43)$$

which reduces the count by roughly a factor of eight, and rotational or translational invariance can reduce it much further by making most elements vanish. But the scaling remains $\mathcal{O}(n^4)$.

The structural remedy is the one developed in Section 1.32. Reading the pairs $(\alpha\gamma)$ and $(\beta\delta)$ as composite indices turns the four-index array into an ordinary matrix, which for a Coulomb-like interaction is positive semidefinite with rapidly decaying eigenvalues, so that

$$\langle \alpha\beta | \hat{v} | \gamma\delta \rangle = \sum_{\eta=0}^{M-1} L_{\alpha\gamma}^{\eta} L_{\beta\delta}^{\eta}, \quad M \ll n^2. \quad (2.44)$$

The storage drops from n^4 to Mn^2 and the contractions in Eq. (2.39) can be performed on the factors rather than on the full array. For the eight-orbital basis used above, the 64×64 pair matrix has numerical rank 15, not 64: less than a quarter of the nominal size. This is the same effect that Table 1.9 quantified, and it is what makes large-basis coupled-cluster and perturbation-theory calculations possible at all.

2.11 Evaluating Slater determinants in Monte Carlo simulations

The energy functional of Section 2.9 needs the Slater determinant once. A Monte Carlo calculation needs it several million times, and needs its gradient and its Laplacian as well. The variational and diffusion Monte Carlo methods developed later in this book propose a new configuration, accept or reject it, and repeat; the local energy requires

$$\frac{\nabla_i |\mathbf{D}|}{|\mathbf{D}|} \quad \text{and} \quad \frac{\nabla_i^2 |\mathbf{D}|}{|\mathbf{D}|} \quad (2.45)$$

for every particle i at every step. Since the cost of a single determinant is $\mathcal{O}(N^3)$ by Section 2.3, and there are $N \cdot d$ coordinates to differentiate, brute-force evaluation costs

$$Nd \cdot \mathcal{O}(N^3) \sim \mathcal{O}(d \cdot N^4) \quad (2.46)$$

per sweep. This is prohibitive, and it is entirely avoidable. The technical material of this section and the next reduces it to $\mathcal{O}(d \cdot N^2)$. It is placed here because it is nothing but linear algebra applied to the object introduced in Section 2.4; the physics that makes it necessary comes later.

The Slater matrix.

Define the $N \times N$ matrix $\mathbf{D}$ with elements

$$d_{ij} = \phi_j(\mathbf{r}_i), \quad (2.47)$$

where ϕ_j is a single-particle wave function. Rows are labelled by particles and columns by the single-particle states. The Slater determinant of Eq. (2.14) is $|\mathbf{D}|/\sqrt{N!}$.

The observation that makes everything work is this: *if we move only one particle at a time, only one row of $\mathbf{D}$ changes*. Recomputing the whole determinant then throws away almost all of the work already done. The solution is to keep track of the *inverse* of the Slater matrix.

The ratio of two determinants.

Let $\mathbf{r}^{\text{old}}$ be the current configuration and $\mathbf{r}^{\text{new}}$ the proposed one, differing only in the position of particle i . The Metropolis test needs the ratio

$$R \equiv \frac{|\mathbf{D}(\mathbf{r}^{\text{new}})|}{|\mathbf{D}(\mathbf{r}^{\text{old}})|}. \quad (2.48)$$

Expanding each determinant along row i in terms of the cofactors C_{ij} ,

$$R = \frac{\sum_{j=1}^N d_{ij}(\mathbf{r}^{\text{new}}) C_{ij}(\mathbf{r}^{\text{new}})}{\sum_{j=1}^N d_{ij}(\mathbf{r}^{\text{old}}) C_{ij}(\mathbf{r}^{\text{old}})}. \quad (2.49)$$

Now recall that the cofactor C_{ij} is the signed minor obtained by *deleting* row i and column j , so the i th row of the cofactor matrix does not depend on the i th row of $\mathbf{D}$ at all. Since the two configurations differ only in that row,

$$C_{ij}(\mathbf{r}^{\text{new}}) = C_{ij}(\mathbf{r}^{\text{old}}) \quad \forall j \in \{1, \dots, N\}. \quad (2.50)$$

Next, the inverse of a matrix is the transposed cofactor matrix divided by the determinant,

$$d_{ij}^{-1} = \frac{C_{ji}}{|\mathbf{D}|}, \quad (2.51)$$

and by definition $\sum_k d_{ik} d_{kj}^{-1} = \delta_{ij}$. Substituting Eq. (2.50) into Eq. (2.49) and replacing the cofactors using Eq. (2.51), the denominator becomes exactly unity, and we are left with

$$R = \sum_{j=1}^N d_{ij}(\mathbf{r}^{\text{new}}) d_{ji}^{-1}(\mathbf{r}^{\text{old}}) = \sum_{j=1}^N \phi_j(\mathbf{r}_i^{\text{new}}) d_{ji}^{-1}(\mathbf{r}^{\text{old}}) \quad (2.52)$$

The ratio of two $N \times N$ determinants is a single dot product between the vector of single-particle functions evaluated at the new position and the i th column of the old inverse. The cost is $\mathcal{O}(N)$ instead of $\mathcal{O}(N^3)$, and no determinant is computed at all.

Updating the inverse.

The price is that we must maintain $\mathbf{D}^{-1}$. If the move is accepted, the inverse is repaired in $\mathcal{O}(N^2)$ operations rather than recomputed in $\mathcal{O}(N^3)$. For every column $j \neq i$ form

$$S_j = \left(\mathbf{D}(\mathbf{r}^{\text{new}}) \mathbf{D}^{-1}(\mathbf{r}^{\text{old}}) \right)_{ij} = \sum_{l=1}^N d_{il}(\mathbf{r}^{\text{new}}) d_{lj}^{-1}(\mathbf{r}^{\text{old}}), \quad (2.53)$$

and update

$$d_{kj}^{-1}(\mathbf{r}^{\text{new}}) = d_{kj}^{-1}(\mathbf{r}^{\text{old}}) - \frac{S_j}{R} d_{ki}^{-1}(\mathbf{r}^{\text{old}}), \quad k = 1, \dots, N, \quad j \neq i. \quad (2.54)$$

The i th column is simply rescaled,

$$d_{ki}^{-1}(\mathbf{r}^{\text{new}}) = \frac{1}{R} d_{ki}^{-1}(\mathbf{r}^{\text{old}}), \quad k = 1, \dots, N. \quad (2.55)$$

Equations (2.54) and (2.55) are the Sherman-Morrison formula specialised to a rank-one change of a single row. Note that both contain a division by R : the algorithm degrades precisely when the proposed determinant is small, a point we return to in Section 2.14.

```
def ratio(self, i, r_new):
    """R = |D(r_new)| / |D(r_old)|, in O(N) operations."""
    row = np.array([self.s.basis.value(j, r_new)[0]
                    for j in range(self.s.N)])
    return float(row @ self.s.Dinv[:, i]), row

def accept(self, i, r_new, R=None, row=None):
    """Update D and D^{-1} after the move of particle i is accepted."""
    if row is None:
        R, row = self.ratio(i, r_new)
    Dinv = self.s.Dinv

    # S_j = (D(new) D^{-1}(old))_{ij} for every column j
    S = row @ Dinv # length N, S[i] == R
    new_inv = Dinv.copy()
    for j in range(self.s.N):
        if j != i:
            new_inv[:, j] = Dinv[:, j] - (S[j] / R) * Dinv[:, i]
    new_inv[:, i] = Dinv[:, i] / R

    self.s.Dinv = new_inv
    self.s.D[i, :] = row
    self.s.positions[i] = r_new
```

For six particles in a two-dimensional oscillator basis the ratio computed from Eq. (2.52) agrees with the brute-force quotient of two determinants to 2×10^{-16} , as `slater_update.py` confirms.

2.12 Gradients, the Laplacian and the quantum force

The same argument applies to derivatives, and this is where the real saving lies. Differentiating the Slater determinant with respect to the coordinates of particle i changes only the i th row of $\mathbf{D}$, exactly as moving particle i did. The derivation leading to Eq. (2.52) therefore goes through unchanged with d_{ij} replaced by its derivative,

$$\frac{\nabla_i |\mathbf{D}(\mathbf{r})|}{|\mathbf{D}(\mathbf{r})|} = \sum_{j=1}^N \nabla_i d_{ij}(\mathbf{r}) d_{ji}^{-1}(\mathbf{r}) = \sum_{j=1}^N \nabla_i \phi_j(\mathbf{r}_i) d_{ji}^{-1}(\mathbf{r}), \quad (2.56)$$

$$\frac{\nabla_i^2 |\mathbf{D}(\mathbf{r})|}{|\mathbf{D}(\mathbf{r})|} = \sum_{j=1}^N \nabla_i^2 d_{ij}(\mathbf{r}) d_{ji}^{-1}(\mathbf{r}) = \sum_{j=1}^N \nabla_i^2 \phi_j(\mathbf{r}_i) d_{ji}^{-1}(\mathbf{r}). \quad (2.57)$$

To obtain every derivative of the Slater determinant we therefore need only the derivatives of the *single-particle* functions and the elements of the inverse Slater matrix. Each derivative is an $\mathcal{O}(N)$ dot product, and there are $d \cdot N$ of them, so the total evaluation costs $\mathcal{O}(d \cdot N^2)$. With the $\mathcal{O}(N^2)$ update of the inverse the overall scaling is no worse, which should be compared with the $\mathcal{O}(d \cdot N^4)$ of Eq. (2.46). Table 2.3 collects the counts.

Quantity	Brute force	With the inverse
Ratio R for one move	$\mathcal{O}(N^3)$	$\mathcal{O}(N)$
Update of $\mathbf{D}^{-1}$	$\mathcal{O}(N^3)$	$\mathcal{O}(N^2)$
All gradients and Laplacians	$\mathcal{O}(dN^4)$	$\mathcal{O}(dN^2)$

Table 2.3 Operation counts for the quantities a Monte Carlo step requires, with and without maintaining the inverse Slater matrix.

Equation (2.56) is the *quantum force* (up to a factor of two) that drives the importance-sampled random walk of diffusion Monte Carlo, and Eq. (2.57) supplies the kinetic energy in the local energy. Both are therefore evaluated at every single step of the simulation.

A practical remark.

In almost all cases one has closed-form expressions for the single-particle wave functions. It is then worth differentiating them analytically and writing separate functions for $\nabla_i \phi_j$ and $\nabla_i^2 \phi_j$. Finite differences would cost both accuracy and speed, and the derivatives are needed far too often to tolerate either. The check in `slater_update.py` compares the analytic Eqs. (2.56) and (2.57) against numerical differentiation of $\ln |\mathbf{D}|$ and finds agreement to 5×10^{-10} and 9×10^{-6} respectively – the residual being the error of the finite differences, not of the formulae.

2.13 Spin factorisation of the Slater determinant

A further and quite substantial saving follows from a property of the spin degree of freedom. Consider beryllium, four electrons in the $1s$ and $2s$ orbitals. The Slater determinant reads

$$\Phi(\mathbf{r}_1, \mathbf{r}_2, \mathbf{r}_3, \mathbf{r}_4) = \frac{1}{\sqrt{4!}} \begin{vmatrix} \psi_{100\uparrow}(\mathbf{r}_1) & \psi_{100\uparrow}(\mathbf{r}_2) & \psi_{100\uparrow}(\mathbf{r}_3) & \psi_{100\uparrow}(\mathbf{r}_4) \\ \psi_{100\downarrow}(\mathbf{r}_1) & \psi_{100\downarrow}(\mathbf{r}_2) & \psi_{100\downarrow}(\mathbf{r}_3) & \psi_{100\downarrow}(\mathbf{r}_4) \\ \psi_{200\uparrow}(\mathbf{r}_1) & \psi_{200\uparrow}(\mathbf{r}_2) & \psi_{200\uparrow}(\mathbf{r}_3) & \psi_{200\uparrow}(\mathbf{r}_4) \\ \psi_{200\downarrow}(\mathbf{r}_1) & \psi_{200\downarrow}(\mathbf{r}_2) & \psi_{200\downarrow}(\mathbf{r}_3) & \psi_{200\downarrow}(\mathbf{r}_4) \end{vmatrix}. \quad (2.58)$$

As written this determinant *vanishes*, because the spatial parts of the spin-up and spin-down rows are identical: rows one and two are the same function of the spatial coordinates, and so are rows three and four. The spin labels carry the distinction, and a determinant of spatial functions alone cannot see them.

Expanding Eq. (2.58) in terms of 2×2 determinants built separately from the spin-up and spin-down orbitals gives

$$\begin{aligned} \Phi &\propto \det\uparrow(1,2) \det\downarrow(3,4) - \det\uparrow(1,3) \det\downarrow(2,4) \\ &\quad - \det\uparrow(1,4) \det\downarrow(3,2) + \det\uparrow(2,3) \det\downarrow(1,4) \\ &\quad - \det\uparrow(2,4) \det\downarrow(1,3) + \det\uparrow(3,4) \det\downarrow(1,2), \end{aligned} \quad (2.59)$$

with

$$\det\uparrow(1,2) = \frac{1}{\sqrt{2}} \begin{vmatrix} \psi_{100\uparrow}(\mathbf{r}_1) & \psi_{100\uparrow}(\mathbf{r}_2) \\ \psi_{200\uparrow}(\mathbf{r}_1) & \psi_{200\uparrow}(\mathbf{r}_2) \end{vmatrix}, \quad \det\downarrow(3,4) = \frac{1}{\sqrt{2}} \begin{vmatrix} \psi_{100\downarrow}(\mathbf{r}_3) & \psi_{100\downarrow}(\mathbf{r}_4) \\ \psi_{200\downarrow}(\mathbf{r}_3) & \psi_{200\downarrow}(\mathbf{r}_4) \end{vmatrix}. \quad (2.60)$$

The total is of course still zero; the expansion has merely reorganised the sum over spin assignments.

The Moskowitz-Kalos ansatz.

What we would like is to avoid summing over spin variables altogether, particularly when the interaction does not depend on spin. It can be shown, see Moskowitz and Kalos [2], that for the variational energy one may replace the full determinant by the single product

$$\Phi(\mathbf{r}_1, \dots, \mathbf{r}_N) \propto \det\uparrow \cdot \det\downarrow, \quad (2.61)$$

in which $\det\uparrow$ involves only the electrons with spin up and $\det\downarrow$ only those with spin down. This ansatz is *not* antisymmetric under the exchange of two electrons of opposite spin, but as long as the Hamiltonian is spin independent it yields the same expectation value for the energy as the full determinant. The reason is visible in Section 2.8: the exchange term of Eq. (2.33) connects only orbitals with the same spin, so the missing antisymmetry never enters the energy.

The saving.

The full determinant factorises,

$$|\mathbf{D}| = |\mathbf{D}|_{\uparrow} \cdot |\mathbf{D}|_{\downarrow}, \quad (2.62)$$

and the combined dimensionality of the two smaller determinants equals that of the original. The ratio then factorises too,

$$\frac{|\mathbf{D}|^{\text{new}}}{|\mathbf{D}|^{\text{old}}} = \frac{|\mathbf{D}|_{\uparrow}^{\text{new}}}{|\mathbf{D}|_{\uparrow}^{\text{old}}} \cdot \frac{|\mathbf{D}|_{\downarrow}^{\text{new}}}{|\mathbf{D}|_{\downarrow}^{\text{old}}}, \quad (2.63)$$

and when a single particle moves, one of the two factors is unchanged and cancels. Only a determinant of size $N/2$ is involved in each ratio and each inverse update, so the cost of one move drops from

$$\mathcal{O}_R(N) + \mathcal{O}_{\text{inverse}}(N^2) \quad \text{to} \quad \mathcal{O}_R(N/2) + \mathcal{O}_{\text{inverse}}(N^2/4), \quad (2.64)$$

and the cost of a full sweep over all N particles from $\mathcal{O}_R(N^2) + \mathcal{O}_{\text{inverse}}(N^3)$ to

$$\mathcal{O}_R(N^2/2) + \mathcal{O}_{\text{inverse}}(N^3/4). \quad (2.65)$$

The reduction is a constant factor, but a large one, and it is greatest when the numbers of spin-up and spin-down particles are equal so that the two factorised determinants are exactly half the size of the original.

2.14 Determinants, inverses and numerical stability

The updating scheme of Section 2.11 still needs the inverse and the determinant *once*, at the start of the simulation and whenever the accumulated error has to be reset. This is where the machinery of Chapter 1 enters directly.

The determinant from LU.

Never use the definition (2.10). The LU factorisation of Section 1.15 gives $\mathbf{PD} = \mathbf{LU}$ in $\mathcal{O}(N^3)$ operations, and the determinant follows from Eq. (1.22) as the signed product of the diagonal elements of $\mathbf{U}$,

$$|\mathbf{D}| = \pm u_{00}u_{11} \cdots u_{N-1,N-1}. \quad (2.66)$$

The same factorisation delivers the inverse for N additional back substitutions, Eq. (1.23). One factorisation therefore supplies both objects the Monte Carlo run needs.

In practice one should not store $|\mathbf{D}|$ at all but its logarithm. A product of N orbital values underflows or overflows long before N becomes physically interesting, whereas

$$\ln ||\mathbf{D}|| = \sum_{k=0}^{N-1} \ln |u_{kk}| \quad (2.67)$$

is perfectly well behaved, and the sign is tracked separately as the parity of the row interchanges. Since the Metropolis test uses only ratios, working with $\ln |\Phi|$ throughout costs nothing.

The determinant from the SVD.

The singular value decomposition of Section 1.28 gives the same number,

$$||\mathbf{D}|| = \prod_{k=0}^{N-1} \sigma_k, \quad (2.68)$$

since $\mathbf{D} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$ with $||\mathbf{U}|| = ||\mathbf{V}|| = 1$. It costs several times more than LU and is never the right way to compute a determinant on its own. It is, however, exactly the right tool when one wants to know *how close to singular* the Slater matrix is, because it returns the condition number of Eq. (1.91) in the same breath,

$$\kappa_2(\mathbf{D}) = \frac{\sigma_0}{\sigma_{N-1}}. \quad (2.69)$$

The nodal surface.

This matters because the Slater determinant vanishes on the *nodal surface*, the set of configurations at which $|\mathbf{D}| = 0$. As a configuration approaches it the matrix becomes singular, the smallest singular value goes to zero, and the condition number diverges. Table 2.4 shows what happens as two of six particles are brought together: the determinant and the smallest singular value fall in step, the condition number rises by the same factor, and the accuracy of the computed inverse degrades exactly as Eq. (1.10) predicts.

Separation	$\ \mathbf{D}\ $	$\kappa_2(\mathbf{D})$	$\ \mathbf{D}\mathbf{D}^{-1} - \mathbf{I}\ $
10^0	2.81×10^0	2.58×10^1	8.7×10^{-16}
10^{-2}	3.21×10^{-2}	1.51×10^3	5.5×10^{-14}
10^{-4}	3.20×10^{-4}	1.54×10^5	4.3×10^{-12}
10^{-6}	3.20×10^{-6}	1.54×10^7	4.7×10^{-10}

Table 2.4 Six particles in a two-dimensional oscillator basis, with two of them brought to the separation shown. Each factor of 100 in the separation costs two digits in the inverse.

When to rebuild.

Equations (2.54) and (2.55) divide by R , so every near-nodal move multiplies the error in $\mathbf{D}^{-1}$ by roughly $1/R$, and the damage is permanent: subsequent updates propagate it. Under Metropolis sampling of $|\Phi|^2$ this is rare, since the sampling density itself vanishes on the node – in a run of 5000 moves the smallest accepted $|R|$ was 0.143 and the error in the inverse never exceeded 5×10^{-14} . But rare is not never. Forcing a sequence of near-nodal moves shows the mechanism plainly:

R of the move	$\ \mathbf{D}\mathbf{D}^{-1} - \mathbf{I}\ $ after
(start)	2.2×10^{-15}
-7.0×10^{-3}	3.4×10^{-14}
1.2×10^{-4}	3.5×10^{-12}
1.2×10^{-6}	2.0×10^{-10}
1.2×10^{-8}	3.2×10^{-8}
after an LU refresh	1.7×10^{-15}

Table 2.5 Error accumulated in the updated inverse by a sequence of moves that pass close to the nodal surface, and its removal by one LU refactorisation.

The remedy is simple and universal: rebuild $\mathbf{D}^{-1}$ from an LU factorisation every few hundred accepted moves. The $\mathcal{O}(N^3)$ cost, amortised over that many $\mathcal{O}(N^2)$ updates, is negligible, and the accumulated error is reset to machine precision. A production Monte Carlo code should also monitor $\|\mathbf{D}\mathbf{D}^{-1} - \mathbf{I}\|$, or the condition number of Eq. (2.69), and refresh whenever it exceeds a threshold.

Link to Chapter 1. Three separate results from the first chapter meet here. The LU decomposition of Section 1.15 supplies the determinant and the inverse at the start of the run; the condition number of Section 1.13, computed from the singular values of Section 1.28, tells us how many digits the inverse has lost; and the observation that the accuracy of a computed inverse is governed by κ_2 rather than

by the algorithm, Eq. (1.10), explains why no cleverer updating formula could avoid the problem. The nodal surface of a fermionic wave function is, from the point of view of linear algebra, simply the locus where the Slater matrix becomes singular.

2.15 Quantum computing: the gate model and its linear-algebra foundation

A quantum computer is a well-controlled quantum many-body system, and the abstract *gate-based* model is defined entirely in the language of Chapter 1. Starting from N qubits initialised in $|\Psi\rangle^{(0)} = |00\cdots 0\rangle$, the state evolves by successive applications of unitary gates,

$$|\Psi\rangle^{(n+1)} = \hat{U}^{(n)} |\Psi\rangle^{(n)}, \quad \hat{U}^{(n)} = T \exp\left(-i \int_{t_n}^{t_{n+1}} \hat{H}(t) dt\right).$$

A *single-qubit gate* acts on qubit a via a 2×2 unitary:

$$\Psi_{i_1 \cdots i_N}^{(n+1)} = \sum_{i'_a} U_{i_a i'_a}^{(n)} \Psi_{i_1 \cdots i_{a-1} i'_a i_{a+1} \cdots i_N}^{(n)}.$$

A *two-qubit gate* acts on qubits a, b via a 4×4 unitary. Standard examples are the Pauli gates X, Y, Z of Section 1.7, the Hadamard gate

$$H = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix},$$

the phase gate $T = \begin{pmatrix} e^{i\pi/8} & 0 \\ 0 & e^{-i\pi/8} \end{pmatrix}$, and the two-qubit CNOT (controlled-NOT):

$$C_X = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{pmatrix}.$$

When qubit a is *measured*, the result $\alpha \in \{0, 1\}$ occurs with probability

$$P_\alpha = \sum_{i_1 \cdots i_a \cdots i_N} \left| \Psi_{i_1 \cdots \alpha \cdots i_N}^{(n)} \right|^2,$$

after which the state collapses to the corresponding projected state – the projection operators of Section 1.6 acting for real.

In a nutshell. Quantum computing is a restricted form of linear algebra: matrix-vector multiplications with special unitary matrices (gates) on exponentially large vectors (wave functions). The challenge – shared with classical many-body physics – is that only a tiny fraction of the 2^N -dimensional space is physically relevant; identifying that subspace is the core problem in both fields.

2.16 The exponential-dimension problem and why it matters

The state of N spin- $\frac{1}{2}$ particles requires 2^N complex numbers to describe, as the tensor-product construction of Section 1.10 already made clear. For $N \approx 300$, $2^{300} \gg 10^{82}$, the estimated number of atoms in the observable universe. The same counting applies to the fermionic problem: the number of Slater determinants

that can be formed by distributing N particles over n single-particle states is $\binom{n}{N}$, which for $n = 40$ and $N = 20$ is already 1.4×10^{11} .

How do physicists actually compute?

1. **Clever approximations:** Bethe ansatz, perturbation theory, mean-field approximations such as the Hartree-Fock theory derived above.
2. **Computer simulations:** DFT, quantum Monte Carlo, FCI, coupled cluster, DMRG.
3. **Restricting to tractable subsystems:** systems where a single-particle basis is a good approximation.

Chapter 1 supplied two of the reasons this is not hopeless. The first is that we rarely want the whole spectrum: the Lanczos algorithm of Section 1.26 extracts the lowest eigenvalues of a matrix of dimension 10^9 using only matrix-vector products, and the pairing-model example there converged to machine precision in forty of them. The second is that physically relevant states are compressible: the Schmidt spectrum of a ground state decays rapidly (Section 1.30), and the Eckart-Young theorem, Eq. (1.96), guarantees that truncating it is the best approximation of its rank. Density-matrix renormalisation group and tensor-network methods are built on precisely that observation.

For a quantum circuit of depth D acting on N qubits, the accessible subspace has dimension $\mathcal{O}(DN)$, vastly smaller than 2^N . For $N = 100$ and $D = 10^6$, this subspace still has 10^{20} times fewer degrees of freedom than the full Hilbert space. The question is whether the *physically relevant* states lie in this reachable subspace – a central open problem in quantum many-body theory and quantum computing alike (see, e.g., LaRocca et al.).

The variational quantum eigensolver (VQE) tackles the problem by using a parametrised quantum circuit as an ansatz for the ground state and minimising $\langle \Psi(\boldsymbol{\theta}) | \hat{H} | \Psi(\boldsymbol{\theta}) \rangle$ classically over the circuit parameters. It is the quantum-computing form of the variational principle (2.22) that underlies Hartree-Fock and coupled-cluster theory: the bound $E(\boldsymbol{\theta}) \geq E_0$ holds for every choice of parameters, and the quality of the result is decided entirely by whether the ansatz can reach the true ground state.

2.17 Exercises

Warm-up exercises

1. **Determinant product rule.** For 2×2 matrices $\mathbf{B}$ and $\mathbf{C}$, define $a_{ij} = \sum_k b_{ik} c_{kj}$. Show by direct calculation that $\det(\mathbf{A}) = \det(\mathbf{B}) \det(\mathbf{C})$.
2. **Slater determinant:** $N = 2$. Write the Slater determinant for two electrons in orbitals ψ_α and ψ_β . (a) Show it is antisymmetric under particle exchange. (b) Show it vanishes if $\psi_\alpha = \psi_\beta$.
3. **Antisymmetriser.** Verify for $N = 2$ that $\hat{A}^2 = \hat{A}$ and that $[\hat{H}_0, \hat{A}] = 0$ when $\hat{H}_0 = \hat{h}_0(x_1) + \hat{h}_0(x_2)$.
4. **Basis change of the Slater determinant.** With the expansion $\psi_p = \sum_\lambda C_{p\lambda} \phi_\lambda$ and a 2×2 unitary $\mathbf{C}$, show that the two-particle Slater determinant in the new basis equals $\det(\mathbf{C}) \det(\boldsymbol{\Phi})$ and that $|\det(\mathbf{C})| = 1$ leaves the physics unchanged.
5. **Hartree-Fock energy.** For a two-electron system with orbitals ψ_α and ψ_β , write out the HF energy functional (2.37) explicitly, identifying the kinetic, direct (Hartree), and exchange (Fock) contributions.
6. **Variational principle.** Show that $E[\boldsymbol{\Phi}] \geq E_0$ for any normalised trial state $\boldsymbol{\Phi}$, where E_0 is the exact ground state energy. Use the spectral decomposition of $\hat{H}$ in its eigenbasis.
7. **Exponential dimension.** How many complex numbers are needed to store the full N -particle wave function for $N = 10, 20, 50$? A modern GPU can perform roughly 10^{14} floating-point operations per second. Estimate how long a single matrix-vector multiplication in the $N = 50$ space would take.
8. **One-qubit gates numerically.** Write a short Python function that represents $|0\rangle$ and $|1\rangle$ as NumPy arrays and applies each Pauli matrix, the Hadamard gate H , and the phase gate T in turn. Verify the action of X (spin flip), Z (phase flip) and $H|0\rangle = (|0\rangle + |1\rangle)/\sqrt{2}$ (creation of a superposition state).
9. **Pauli gate as a quantum NOT.** Explain why σ_x is the quantum analogue of the classical NOT gate. What is the difference between the classical and quantum cases when the input is a superposition $\alpha|0\rangle + \beta|1\rangle$?

10. **VQE and the variational principle.** The variational quantum eigensolver minimises $E(\boldsymbol{\theta}) = \langle \Psi(\boldsymbol{\theta}) | \hat{H} | \Psi(\boldsymbol{\theta}) \rangle$ over circuit parameters $\boldsymbol{\theta}$. (a) Explain why $E(\boldsymbol{\theta}) \geq E_0$ for any $\boldsymbol{\theta}$. (b) Describe the connection to the Hartree-Fock variational principle. (c) What determines the quality of the approximation in VQE compared to FCI?
11. **The antisymmetriser as a projector.** Using Eq. (2.7), verify for $N = 2$ that $\hat{A}^2 = \hat{A}$ and that $[\hat{H}_0, \hat{A}] = 0$ when $\hat{H}_0 = \hat{h}_0(x_1) + \hat{h}_0(x_2)$. Explain why $\hat{A}$ is a projection operator in the sense of Section 1.6, and identify the subspace it projects onto.
12. **Why only two permutations survive.** In the derivation of Eq. (2.32) it was argued that only the identity and the transposition $\hat{P}_{ij}$ contribute to the two-body expectation value. Write out the case $N = 3$ explicitly, listing all six permutations, and show that four of them give a vanishing integral. Then explain why the corresponding argument for the one-body operator leaves only one surviving permutation.
13. **Self-interaction.** Show from Eq. (2.36) that $\langle \mu\mu | \hat{v} | \mu\mu \rangle_{\text{AS}} = 0$, and explain in words what would go wrong in Eq. (2.37) if the exchange term were omitted.
14. **Occupation numbers of a Slater determinant.** For the determinant (2.14) the density matrix is $\rho_{\lambda\mu} = \sum_i C_{i\lambda}^* C_{i\mu}$. (a) Show that $\boldsymbol{\rho}^2 = \boldsymbol{\rho}$ when the coefficients are orthonormal. (b) Conclude that the eigenvalues – the natural occupation numbers of Section 1.31 – can only be 0 or 1. (c) What does it mean physically when a computed occupation number comes out as 0.93?
15. **Cost of a determinant.** Reproduce Table 2.1. For $n = 20$, estimate how long the evaluation of a single determinant would take from Eq. (2.10) on a machine performing 10^{12} operations per second, and compare with the LU route of Section 1.15.
16. **Two-body storage.** How many two-body matrix elements $\langle \alpha\beta | \hat{v} | \gamma\delta \rangle$ are there for $n = 50$, $n = 200$ and $n = 1000$ single-particle states? How much memory does each case require in double precision? Using Eq. (2.44) with $M = 5n$, how much is saved?
17. **Hartree-Fock numerically.** Using `slaterdeterminant.py`, reproduce Table 2.2. Then vary the interaction strength and plot the correlation energy (2.42) and the largest natural occupation number against it. At what point would you no longer trust a single-determinant description?
18. **The ratio R by hand.** Take $N = 2$ and a Slater matrix $\mathbf{D}$ with elements $d_{ij} = \phi_j(\mathbf{r}_i)$. Move particle 1 and compute the ratio R both from the definition (2.48), by evaluating two 2×2 determinants, and from Eq. (2.52). Verify that they agree, and count the floating point operations each route requires.
19. **Why the cofactor row does not change.** Explain in your own words why $C_{ij}(\mathbf{r}^{\text{new}}) = C_{ij}(\mathbf{r}^{\text{old}})$ when only particle i has moved, Eq. (2.50). Where in the derivation of Eq. (2.52) would the argument fail if two particles were moved at once?
20. **Updating the inverse.** Implement Eqs. (2.53)–(2.55) for a 3×3 Slater matrix and check the result against `numpy.linalg.inv`. Confirm that the update costs $\mathcal{O}(N^2)$ operations by counting the loops.
21. **Derivatives of the determinant.** Using `slater_update.py`, compare the analytic gradient and Laplacian ratios of Eqs. (2.56) and (2.57) with numerical differentiation of $\ln|\mathbf{D}|$. Why does the Laplacian agree to fewer digits than the gradient?
22. **The nodal surface as an ill-conditioned matrix.** Reproduce Table 2.4. Plot $\ln \kappa_2(\mathbf{D})$ against $\ln(\text{separation})$ and explain the slope. Using Eq. (1.10), predict how many digits of $\mathbf{D}^{-1}$ survive at a separation of 10^{-8} , and check your prediction.
23. **When to refresh.** Design a rule for deciding when to rebuild $\mathbf{D}^{-1}$ from an LU factorisation. Estimate the fraction of the total run time that the refreshes consume for $N = 50$ if you rebuild every 500 accepted moves, and compare with the cost of doing so at every move.
24. **Spin factorisation.** For $N = 4$ with two spin-up and two spin-down particles, write out Eq. (2.59) and verify that the full determinant of Eq. (2.58) vanishes. Then show, using the structure of Eq. (2.33), that the factorised ansatz (2.61) gives the same energy for a spin-independent Hamiltonian.

Exercise session, week 35: Slater determinants and a first FCI model

Exercise 1: expectation values with a Slater determinant.

Consider the fermion Slater determinant used as an ansatz for a quantum-mechanical state,

$$\Phi_{\lambda}^{AS}(x_1 x_2 \dots x_N; \alpha_1 \alpha_2 \dots \alpha_N) = \frac{1}{\sqrt{N!}} \sum_p (-1)^p \hat{P} \prod_{i=1}^N \psi_{\alpha_i}(x_i), \quad (2.70)$$

where $\hat{P}$ permutes the particles and the sum runs over all $N!$ permutations. The number of particles equals the number of available single-particle states $\alpha_1 \alpha_2 \dots \alpha_N$, and the single-particle basis ψ_{α_i} is orthonormal.

- Write out Φ^{AS} for $N = 3$.
- Show that

$$\int dx_1 dx_2 \dots dx_N \left| \Phi_{\lambda}^{AS}(x_1 x_2 \dots x_N; \alpha_1 \alpha_2 \dots \alpha_N) \right|^2 = 1.$$

- Define a general one-body operator $\hat{F} = \sum_i^N \hat{f}(x_i)$ and a general two-body operator $\hat{G} = \sum_{i>j}^N \hat{g}(x_i, x_j)$, with $\hat{g}$ invariant under the interchange of the coordinates of particles i and j . Calculate the matrix elements for a two-particle Slater determinant,

$$\langle \Phi_{\alpha_1 \alpha_2}^{AS} | \hat{F} | \Phi_{\alpha_1 \alpha_2}^{AS} \rangle \quad \text{and} \quad \langle \Phi_{\alpha_1 \alpha_2}^{AS} | \hat{G} | \Phi_{\alpha_1 \alpha_2}^{AS} \rangle.$$

Which properties do you expect these operators to have, in addition to permutation symmetry?

- Repeat the calculation for a general N -particle Slater determinant.

Exercise 2: a three-level model.

We now consider a simple three-level problem, our first and very simple model of a many-fermion system and of what we shall later call full configuration interaction theory. The particles are fermions. The single-particle states are labelled by the quantum number p and each can accommodate up to two particles, so every single-particle state is doubly degenerate; one may think of the two as spin up and spin down. The spacing between the doubly degenerate states is constant with value d , the first state has energy d , and there are three available states $p = 1, 2, 3$, as in Figure 2.1.




Fig. 2.1 Schematic plot of the possible single-particle levels with double degeneracy. The filled circles indicate occupied states and the spacing between the levels is constant. Four possible two-particle states are shown; the two leftmost are the configurations retained in parts b) and c), in which both particles occupy the same level.

- How many two-particle Slater determinants can we construct in this space?
- Limit the system to the two lowest orbits, $p = 1$ and $p = 2$, and two particles. Write the Hamiltonian as $\hat{H} = \hat{H}_0 + \hat{H}_I$, with the one-body part satisfying

$$\hat{h}_0 \psi_{p\sigma} = p d \psi_{p\sigma},$$

where σ is the spin quantum number. Assume that the only two-particle states that can exist are those in which both particles occupy the same level p , and that all two-particle matrix elements of $\hat{H}_I$ have the constant value $-g$. Show that the Hamiltonian matrix can be written as

$$\begin{pmatrix} 2d - g & -g \\ -g & 4d - g \end{pmatrix},$$

and find its eigenvalues and eigenvectors. What is the admixture of the state with two particles in $p = 2$ to the wave function with two particles in $p = 1$? Discuss the result in terms of a linear combination of Slater determinants.

- c) Add the possibility that the two particles occupy $p = 3$, and find the Hamiltonian matrix, the eigenvalues and the eigenvectors. We still insist that only two-particle states with both particles in the same level are allowed. The 3×3 matrix may be diagonalised numerically.

This small model catches several birds with one stone. It shows how linear combinations of Slater determinants are built and how the coefficients are interpreted as admixtures to a given state. The two-particle states above $p = 1$ are excitations of the ground-state configuration, and the reliability of the single-determinant ansatz depends on the ratio of the interaction strength g to the level spacing d . Finally, the model is a schematic ansatz for pairing correlations, and hence for superfluidity and superconductivity in fermionic systems; it is the two-level ancestor of the pairing model used in Section 1.26.

Answers, week 35

Exercise 1a.

For $N = 3$ the six permutations give

$$\begin{aligned} \Phi_\lambda^{AS} = \frac{1}{\sqrt{6}} & \left(\psi_{\alpha_1}(x_1) \psi_{\alpha_2}(x_2) \psi_{\alpha_3}(x_3) - \psi_{\alpha_1}(x_2) \psi_{\alpha_2}(x_1) \psi_{\alpha_3}(x_3) \right. \\ & + \psi_{\alpha_1}(x_2) \psi_{\alpha_2}(x_3) \psi_{\alpha_3}(x_1) - \psi_{\alpha_1}(x_3) \psi_{\alpha_2}(x_2) \psi_{\alpha_3}(x_1) \\ & \left. + \psi_{\alpha_1}(x_3) \psi_{\alpha_2}(x_1) \psi_{\alpha_3}(x_2) - \psi_{\alpha_1}(x_1) \psi_{\alpha_2}(x_3) \psi_{\alpha_3}(x_2) \right), \end{aligned} \quad (2.71)$$

which is the expansion of the 3×3 determinant of Eq. (2.14) along its rows.

Exercise 1b.

Introduce the antisymmetriser (2.7) and the Hartree function (2.15), so that $\Phi_\lambda^{AS} = \sqrt{N!} \hat{A} \Phi_H$. Then

$$\langle \Phi_\lambda^{AS} | \Phi_\lambda^{AS} \rangle = N! \int \Phi_H^* \hat{A}^\dagger \hat{A} \Phi_H d\tau = N! \int \Phi_H^* \hat{A} \Phi_H d\tau, \quad (2.72)$$

where we used that $\hat{A}$ is Hermitian and idempotent, Eq. (2.9). Writing out the antisymmetriser cancels the factor $N!$,

$$\langle \Phi_\lambda^{AS} | \Phi_\lambda^{AS} \rangle = \sum_p (-1)^p \int \Phi_H^* \hat{P} \Phi_H d\tau. \quad (2.73)$$

The permutation operator acts on only one of the two Hartree functions, so any $\hat{P}$ other than the identity produces at least one overlap $\langle \alpha_i | \alpha_j \rangle$ with $i \neq j$, which vanishes by orthonormality. Only the identity survives and $\langle \Phi_\lambda^{AS} | \Phi_\lambda^{AS} \rangle = \int \Phi_H^* \Phi_H d\tau = 1$.

Exercise 1c.

For two particles the same manipulation gives

$$\langle \Phi_{\alpha_1 \alpha_2}^{AS} | \hat{F} | \Phi_{\alpha_1 \alpha_2}^{AS} \rangle = \sum_p (-)^p \int \psi_{\alpha_1}^*(x_1) \psi_{\alpha_2}^*(x_2) [\hat{f}(x_1) + \hat{f}(x_2)] \hat{P} \psi_{\alpha_1}(x_1) \psi_{\alpha_2}(x_2) d\tau. \quad (2.74)$$

As in b), a permuted Hartree function makes the integral vanish, so only the identity contributes. In each remaining term the one-body operator acts on a single coordinate; integrating out the other gives unity, and

$$\langle \Phi_{\alpha_1 \alpha_2}^{AS} | \hat{F} | \Phi_{\alpha_1 \alpha_2}^{AS} \rangle = \langle \alpha_1 | \hat{f} | \alpha_1 \rangle + \langle \alpha_2 | \hat{f} | \alpha_2 \rangle, \quad (2.75)$$

with the shorthand of Eq. (2.29).

For the two-body operator there is only one term, $\hat{G} = \hat{g}(x_1, x_2)$, and now the transposition survives as well, since $\hat{g}$ depends on both coordinates. The two permutations give

$$\begin{aligned} \langle \Phi_{\alpha_1 \alpha_2}^{AS} | \hat{G} | \Phi_{\alpha_1 \alpha_2}^{AS} \rangle &= \int \psi_{\alpha_1}^*(x_1) \psi_{\alpha_2}^*(x_2) \hat{g}(x_1, x_2) \psi_{\alpha_1}(x_1) \psi_{\alpha_2}(x_2) d\tau \\ &\quad - \int \psi_{\alpha_1}^*(x_1) \psi_{\alpha_2}^*(x_2) \hat{g}(x_1, x_2) \psi_{\alpha_1}(x_2) \psi_{\alpha_2}(x_1) d\tau \\ &= \langle \alpha_1 \alpha_2 | \hat{g} | \alpha_1 \alpha_2 \rangle_{AS}. \end{aligned} \quad (2.76)$$

The first term is the direct or Hartree term, the second the exchange or Fock term, exactly as in Eq. (2.33).

Beyond permutation symmetry we expect both operators to be *Hermitian*, since they represent physical observables and must have real eigenvalues. Permutation symmetry is forced on us because the particles are indistinguishable, and it implies $[\hat{A}, \hat{F}] = [\hat{A}, \hat{G}] = 0$, which is what allowed the manipulations above.

Exercise 1d.

For general N the steps are the same. For the one-body operator,

$$\langle \Phi^{AS} | \hat{F} | \Phi^{AS} \rangle = \sum_{i=1}^N \sum_p (-)^p \int \Phi_H^* \hat{f}(x_i) \hat{P} \Phi_H d\tau = \sum_{i=1}^N \langle \alpha_i | \hat{f} | \alpha_i \rangle. \quad (2.77)$$

For the two-body operator, each pair (x_i, x_j) receives two contributing permutations with opposite signs, giving the factor $(1 - \hat{P}_{ij})$ of Eq. (2.32). Replacing the restricted sum $i < j$ by an unrestricted double sum with a compensating $\frac{1}{2}$, and noting that the $i = j$ terms cancel identically, leaves

$$\langle \Phi^{AS} | \hat{G} | \Phi^{AS} \rangle = \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \langle \alpha_i \alpha_j | \hat{g} | \alpha_i \alpha_j \rangle_{AS}. \quad (2.78)$$

Equations (2.77) and (2.78) together are Eq. (2.37), the Hartree-Fock energy functional.

Exercise 2a.

With three levels, each doubly degenerate, there are six single-particle states and hence $\binom{6}{2} = 15$ two-particle Slater determinants in general. If we count only the level labels and ignore spin, three states with two distinct levels and three with both particles in the same level give six configurations. Under the restriction of parts b) and c), that both particles occupy the *same* level, exactly three Slater determinants remain, one for each p .

Exercise 2b.

Denote by $|\Phi_p\rangle$ the state with both particles in level p . Since $\psi_{p\sigma}$ is an eigenstate of $\hat{h}_0$ with eigenvalue pd , the one-body part contributes $\langle\Phi_p|\hat{H}_0|\Phi_p\rangle = 2\langle p|\hat{h}_0|p\rangle = 2pd$, using Eq. (2.77) with $N = 2$. Every matrix element of $\hat{H}_I$ equals $-g$. With $|\Phi_1\rangle$ and $|\Phi_2\rangle$ as basis,

$$\mathbf{H} = \begin{pmatrix} 2d-g & -g \\ -g & 4d-g \end{pmatrix}, \quad (2.79)$$

which is the requested matrix. This is the matrix representation of an operator in a chosen basis, exactly as in Section 1.21: the Slater determinants are not eigenstates of $\hat{H}$, but they form a basis in which the eigenstates can be expanded, $|\Psi_i\rangle = c_{i0}|\Phi_1\rangle + c_{i1}|\Phi_2\rangle$.

The characteristic polynomial follows from the trace $6d - 2g$ and the determinant $(2d - g)(4d - g) - g^2 = 8d^2 - 6dg$,

$$\lambda^2 + (2g - 6d)\lambda + 8d^2 - 6dg = 0, \quad (2.80)$$

and solving it gives

$$\lambda_{\pm} = 3d - g \pm \sqrt{d^2 + g^2}. \quad (2.81)$$

Substituting back, the unnormalised eigenvectors are

$$|\Psi_{\mp}\rangle = \frac{d \pm \sqrt{d^2 + g^2}}{g} |\Phi_1\rangle + |\Phi_2\rangle, \quad (2.82)$$

or, with the dimensionless ratio $\gamma \equiv d/g$,

$$|\Psi_{\mp}\rangle = \left(\gamma \pm \sqrt{1 + \gamma^2} \right) |\Phi_1\rangle + |\Phi_2\rangle. \quad (2.83)$$

The single parameter γ measures the strength of the interaction against the level spacing. For $g \rightarrow 0$ ($\gamma \rightarrow \infty$) the ground state becomes pure $|\Phi_1\rangle$: the single Slater determinant is exact. As g grows the admixture of $|\Phi_2\rangle$ grows with it, and at $\gamma = 0$ the two determinants enter with equal weight. This is correlation, in its simplest possible form, and it is the same story told quantitatively by the occupation numbers of Section 1.31.

Exercise 2c.

Adding $p = 3$ changes nothing structurally: the diagonal element is $2pd - g = 6d - g$ and the off-diagonal elements remain $-g$, so

$$\mathbf{H} = \begin{pmatrix} 2d-g & -g & -g \\ -g & 4d-g & -g \\ -g & -g & 6d-g \end{pmatrix} = g \begin{pmatrix} 2\gamma-1 & -1 & -1 \\ -1 & 4\gamma-1 & -1 \\ -1 & -1 & 6\gamma-1 \end{pmatrix}. \quad (2.84)$$

The eigenvalues may be found symbolically, or numerically with the Householder and QL solver of Sections 1.24 and 1.25 – three lines of Python. For $\gamma \gg 1$ the ground state is again dominated by $|\Phi_1\rangle$ with admixtures of order $1/\gamma$ from the two excited configurations. Continuing this construction to L levels and n pairs gives exactly the pairing model of Table 1.3, whose dimension $\binom{L}{n}$ grows fast enough that the Lanczos algorithm becomes the only practical route.

Exercise session, week 36: Condon-Slater rules and second quantisation

Exercise 1: Condon-Slater rules for expectation values.

Consider three N -particle Slater determinants $|SD\rangle$, $|SD_i^j\rangle$ and $|SD_{ij}^{kl}\rangle$, where $|SD_i^j\rangle$ differs from $|SD\rangle$ by a single single-particle state – the state ψ_i has been replaced by ψ_j – and $|SD_{ij}^{kl}\rangle$ differs by two. Define, as in week 35, a general one-body operator $\hat{F} = \sum_i^N \hat{f}(x_i)$ and a general two-body operator $\hat{G} = \sum_{i>j}^N \hat{g}(x_i, x_j)$ with $\hat{g}$ invariant under interchange of two coordinates. The single-particle states ψ_i are not necessarily eigenstates of $\hat{f}$.

- Find $\langle SD|\hat{F}|SD\rangle$ and $\langle SD|\hat{G}|SD\rangle$.
- Find $\langle SD|\hat{F}|SD_i^j\rangle$ and $\langle SD|\hat{G}|SD_i^j\rangle$.
- Find $\langle SD|\hat{F}|SD_{ij}^{kl}\rangle$ and $\langle SD|\hat{G}|SD_{ij}^{kl}\rangle$. What happens to the two-body operator for a transition $\langle SD|\hat{G}|SD_{ijk}^{lmn}\rangle$, in which the determinant on the right differs by more than two single-particle states?

Exercise 2: a first encounter with second quantisation.

- Show that the density of particles at $\mathbf{x}$,

$$n(\mathbf{x}) = N \int d\mathbf{x}_2 \cdots d\mathbf{x}_N |\Psi_{AS}(\mathbf{x}, \mathbf{x}_2, \dots, \mathbf{x}_N)|^2,$$

can be written in terms of the single-particle states as

$$n(\mathbf{x}) = \sum_k |\psi_k(\mathbf{x})|^2.$$

- Calculate the matrix elements $\langle \alpha_1 \alpha_2 | \hat{F} | \alpha_1 \alpha_2 \rangle$ and $\langle \alpha_1 \alpha_2 | \hat{G} | \alpha_1 \alpha_2 \rangle$ with

$$\begin{aligned} |\alpha_1 \alpha_2\rangle &= a_{\alpha_1}^\dagger a_{\alpha_2}^\dagger |0\rangle, \\ \hat{F} &= \sum_{\alpha\beta} \langle \alpha | f | \beta \rangle a_\alpha^\dagger a_\beta, \\ \hat{G} &= \frac{1}{2} \sum_{\alpha\beta\gamma\delta} \langle \alpha\beta | g | \gamma\delta \rangle a_\alpha^\dagger a_\beta^\dagger a_\delta a_\gamma, \end{aligned}$$

where the matrix elements are the integrals of Eqs. (2.29), (2.34) and (2.35). Compare with exercise 1c) of week 35.

Answers, week 36

Throughout we assume the two determinants have been brought to *maximum coincidence*: the common single-particle states appear in the same order in both, so that no additional sign from reordering rows is left over. If they have not, every result below acquires the factor $(-1)^\sigma$ of the permutation needed to bring them into that form.

Exercise 1a: the diagonal elements.

These are the results already derived in Section 2.8 and in exercise 1d) of week 35,

$$\langle SD|\hat{F}|SD\rangle = \sum_{i=1}^N \langle i|\hat{f}|i\rangle, \quad \langle SD|\hat{G}|SD\rangle = \frac{1}{2} \sum_{i,j=1}^N \langle ij|\hat{g}|ij\rangle_{AS}. \quad (2.85)$$

The one-body operator sees each occupied orbital once; the two-body operator sees each occupied pair once, with the exchange term subtracted.

Exercise 1b: one differing state.

Write $|SD_i^j\rangle$ for the determinant in which ψ_i has been replaced by ψ_j , with j not among the states occupied in $|SD\rangle$. Using $\Phi = \sqrt{N!}\hat{A}\Phi_H$ and the same argument as in Section 2.8, a term survives only if every orbital not acted on by the operator overlaps with an *identical* orbital on the other side. Since exactly one orbital differs, the one-body operator must be the one that connects ψ_i to ψ_j , and every other factor contributes unity:

$$\langle SD|\hat{F}|SD_i^j\rangle = \langle i|\hat{f}|j\rangle. \quad (2.86)$$

A single term, not a sum: the one-body operator can repair exactly one mismatch.

For the two-body operator, one of the two coordinates must carry the $i \rightarrow j$ replacement while the other runs over all the orbitals common to both determinants. Both the direct and the exchange contraction survive, so

$$\langle SD|\hat{G}|SD_i^j\rangle = \sum_{k=1}^N \langle ik|\hat{g}|jk\rangle_{AS} = \sum_{k=1}^N [\langle ik|\hat{g}|jk\rangle - \langle ik|\hat{g}|kj\rangle], \quad (2.87)$$

where the sum runs over the states occupied in $|SD\rangle$. The term $k = i$ contributes nothing, since $\langle ii|\hat{g}|ji\rangle_{AS} = 0$ by the antisymmetry of Eq. (2.36).

Equation (2.87) is worth pausing on. It says that a two-body interaction, evaluated between determinants differing by one orbital, behaves like an *effective one-body* operator obtained by summing over the occupied states. That sum is precisely the interaction part of the Fock matrix of Eq. (2.40); the Hartree-Fock condition $\langle SD|\hat{H}|SD_i^j\rangle = 0$, that no single excitation couples to the reference, is Brillouin's theorem.

Exercise 1c: two differing states, and beyond.

The one-body operator can repair only one mismatch at a time, so with two differing orbitals at least one overlap $\langle i|k\rangle$ between different states is always left over and

$$\langle SD|\hat{F}|SD_{ij}^{kl}\rangle = 0. \quad (2.88)$$

The two-body operator can repair exactly two, and both coordinates are used up doing so. Nothing is left to sum over, and

$$\langle SD|\hat{G}|SD_{ij}^{kl}\rangle = \langle ij|\hat{g}|kl\rangle_{AS} = \langle ij|\hat{g}|kl\rangle - \langle ij|\hat{g}|lk\rangle. \quad (2.89)$$

Finally, if the determinants differ by three or more single-particle states,

$$\langle SD|\hat{G}|SD_{ijk}^{lmn}\rangle = 0, \quad (2.90)$$

because a two-body operator acts on at most two coordinates and therefore cannot bridge three mismatches; at least one vanishing overlap always survives.

Equations (2.85)–(2.90) are the *Condon-Slater rules*, and the pattern is simple enough to memorise: an n -body operator connects Slater determinants differing by at most n single-particle states. Their consequence is the reason large-scale diagonalisation is feasible at all. A configuration-interaction Hamiltonian matrix contains only one- and two-body operators, so all matrix elements between determinants differing by three or more orbitals vanish identically. The matrix is therefore extremely *sparse*: of the $\binom{n}{N}^2$ possible elements only a tiny fraction are non-zero, which is exactly the structure the Lanczos algorithm of Section 1.26 was designed to exploit, and the reason the Hamiltonian need never be stored.

Exercise 2a: the density.

Expand the Slater determinant along the row belonging to the coordinate $\mathbf{x}$,

$$\Psi_{AS}(\mathbf{x}, \mathbf{x}_2, \dots, \mathbf{x}_N) = \frac{1}{\sqrt{N}} \sum_{k=1}^N (-)^{k+1} \psi_k(\mathbf{x}) \Psi_{AS}^{(k)}(\mathbf{x}_2, \dots, \mathbf{x}_N), \quad (2.91)$$

where $\Psi_{AS}^{(k)}$ is the normalised $(N-1)$ -particle Slater determinant built from all the occupied orbitals *except* ψ_k . Because two such determinants with $k \neq k'$ differ by one orbital, they are orthogonal,

$$\int d\mathbf{x}_2 \dots d\mathbf{x}_N \Psi_{AS}^{(k)*} \Psi_{AS}^{(k')} = \delta_{kk'}. \quad (2.92)$$

Squaring Eq. (2.91) and integrating therefore kills every cross term and leaves

$$\int d\mathbf{x}_2 \dots d\mathbf{x}_N |\Psi_{AS}|^2 = \frac{1}{N} \sum_{k=1}^N |\psi_k(\mathbf{x})|^2, \quad (2.93)$$

so that

$$n(\mathbf{x}) = N \cdot \frac{1}{N} \sum_{k=1}^N |\psi_k(\mathbf{x})|^2 = \sum_{k=1}^N |\psi_k(\mathbf{x})|^2. \quad (2.94)$$

The density of a Slater determinant is the sum of the densities of its occupied orbitals, with no interference between them. Integrating gives $\int n(\mathbf{x}) d\mathbf{x} = N$, as it must. Note that Eq. (2.94) is the diagonal of the one-body density matrix of Eq. (1.104) in coordinate space, whose eigenvalues we already know to be 0 and 1 for a single determinant.

Exercise 2b: the same results from second quantisation.

Start from the one-body operator. Acting with a_β on the two-particle state and using $\{a_\alpha, a_\beta^\dagger\} = \delta_{\alpha\beta}$,

$$a_\beta a_{\alpha_1}^\dagger a_{\alpha_2}^\dagger |0\rangle = \delta_{\beta\alpha_1} a_{\alpha_2}^\dagger |0\rangle - \delta_{\beta\alpha_2} a_{\alpha_1}^\dagger |0\rangle, \quad (2.95)$$

so that

$$a_\alpha^\dagger a_\beta |\alpha_1 \alpha_2\rangle = \delta_{\beta\alpha_1} a_\alpha^\dagger a_{\alpha_2}^\dagger |0\rangle - \delta_{\beta\alpha_2} a_\alpha^\dagger a_{\alpha_1}^\dagger |0\rangle. \quad (2.96)$$

Projecting on $\langle \alpha_1 \alpha_2 | = \langle 0 | a_{\alpha_2} a_{\alpha_1}$ and using $\langle \alpha_1 \alpha_2 | \alpha \alpha_2 \rangle = \delta_{\alpha\alpha_1}$ and $\langle \alpha_1 \alpha_2 | \alpha \alpha_1 \rangle = -\delta_{\alpha\alpha_2}$ gives

$$\langle \alpha_1 \alpha_2 | a_\alpha^\dagger a_\beta | \alpha_1 \alpha_2 \rangle = \delta_{\alpha\alpha_1} \delta_{\beta\alpha_1} + \delta_{\alpha\alpha_2} \delta_{\beta\alpha_2}, \quad (2.97)$$

and therefore

$$\langle \alpha_1 \alpha_2 | \hat{F} | \alpha_1 \alpha_2 \rangle = \langle \alpha_1 | f | \alpha_1 \rangle + \langle \alpha_2 | f | \alpha_2 \rangle. \quad (2.98)$$

For the two-body operator, $a_\delta a_\gamma$ annihilates the state unless $\{\gamma, \delta\} = \{\alpha_1, \alpha_2\}$, and likewise $a_\alpha^\dagger a_\beta^\dagger$ must recreate the same pair. The four surviving index assignments give, after collecting the signs from the anti-commutators,

$$\langle \alpha_1 \alpha_2 | \hat{G} | \alpha_1 \alpha_2 \rangle = \langle \alpha_1 \alpha_2 | g | \alpha_1 \alpha_2 \rangle - \langle \alpha_1 \alpha_2 | g | \alpha_2 \alpha_1 \rangle = \langle \alpha_1 \alpha_2 | g | \alpha_1 \alpha_2 \rangle_{AS}. \quad (2.99)$$

The factor $\frac{1}{2}$ in the definition of $\hat{G}$ is cancelled by the two equivalent assignments of the pair.

Comparison.

Equations (2.98) and (2.99) are *identical* to the results (2.75) and (2.76) obtained in week 35 by integrating over permutations of the Hartree function. This is the point of the exercise. The first-quantised route required us to write out the antisymmetriser, sort through $N!$ permutations, and argue term by term about which overlaps vanish; the second-quantised route required only the anticommutation relations $\{a_\alpha, a_\beta^\dagger\} = \delta_{\alpha\beta}$. Antisymmetry, which was an explicit and increasingly painful bookkeeping task in Section 2.8, has been built into the algebra of the operators themselves.

This is why the next chapter is devoted to second quantisation. The Slater determinant of Eq. (2.14) becomes a string of creation operators acting on the vacuum, the Condon-Slater rules of exercise 1 become elementary consequences of Wick's theorem, and calculations that are barely feasible by hand for $N = 3$ become mechanical for arbitrary N .

Chapter 3

Second quantization

Chapter 2 ended with a warning. Every expectation value there was obtained by writing out the antisymmetriser, sorting through $N!$ permutations, and arguing term by term about which overlaps vanish. For two particles this is an exercise; for the eight-operator strings of Section 2.8 it is already unpleasant; for the excited determinants that perturbation theory and coupled-cluster theory require it is impossible. Exercise 2 of the week-36 session made the point sharply: the same matrix elements fall out of the anticommutation relations in three lines.

This chapter develops that observation into a formalism. Antisymmetry, which was an explicit and increasingly costly bookkeeping task in Chapter 2, is built once and for all into the algebra of the creation and annihilation operators. The Slater determinant of Eq. (2.14) becomes a string of creation operators acting on the vacuum, the Condon-Slater rules of Section 2.6 become consequences of Wick's theorem, and calculations that were barely feasible by hand for $N = 3$ become mechanical for arbitrary N .

We introduce the time-independent operators $a_\alpha^\dagger$ and a_α which create and annihilate, respectively, a particle in the single-particle state φ_α . We define the fermion creation operator $a_\alpha^\dagger$

$$a_\alpha^\dagger |0\rangle \equiv |\alpha\rangle, \quad (3.1)$$

and

$$a_\alpha^\dagger |\alpha_1 \dots \alpha_n\rangle_{\text{AS}} \equiv |\alpha \alpha_1 \dots \alpha_n\rangle_{\text{AS}} \quad (3.2)$$

In Eq. (3.1) the operator $a_\alpha^\dagger$ acts on the vacuum state $|0\rangle$, which does not contain any particles. Alternatively, we could define a closed-shell nucleus or atom as our new vacuum, but then we need to introduce the particle-hole formalism, see the discussion to come.

In Eq. (3.2) $a_\alpha^\dagger$ acts on an antisymmetric n -particle state and creates an antisymmetric $(n+1)$ -particle state, where the one-body state φ_α is occupied, under the condition that $\alpha \neq \alpha_1, \alpha_2, \dots, \alpha_n$. It follows that we can express an antisymmetric state as the product of the creation operators acting on the vacuum state.

$$|\alpha_1 \dots \alpha_n\rangle_{\text{AS}} = a_{\alpha_1}^\dagger a_{\alpha_2}^\dagger \dots a_{\alpha_n}^\dagger |0\rangle \quad (3.3)$$

It is easy to derive the commutation and anticommutation rules for the fermionic creation operators $a_\alpha^\dagger$. Using the antisymmetry of the states (3.3)

$$|\alpha_1 \dots \alpha_i \dots \alpha_k \dots \alpha_n\rangle_{\text{AS}} = -|\alpha_1 \dots \alpha_k \dots \alpha_i \dots \alpha_n\rangle_{\text{AS}} \quad (3.4)$$

we obtain

$$a_{\alpha_i}^\dagger a_{\alpha_k}^\dagger = -a_{\alpha_k}^\dagger a_{\alpha_i}^\dagger \quad (3.5)$$

Using the Pauli principle

$$|\alpha_1 \dots \alpha_i \dots \alpha_i \dots \alpha_n\rangle_{\text{AS}} = 0 \quad (3.6)$$

it follows that

$$a_{\alpha_i}^\dagger a_{\alpha_i}^\dagger = 0. \quad (3.7)$$

If we combine Eqs. (3.5) and (3.7), we obtain the well-known anti-commutation rule

$$a_\alpha^\dagger a_\beta^\dagger + a_\beta^\dagger a_\alpha^\dagger \equiv \{a_\alpha^\dagger, a_\beta^\dagger\} = 0 \quad (3.8)$$

The hermitian conjugate of $a_\alpha^\dagger$ is

$$a_\alpha = (a_\alpha^\dagger)^\dagger \quad (3.9)$$

If we take the hermitian conjugate of Eq. (3.8), we arrive at

$$\{a_\alpha, a_\beta\} = 0 \quad (3.10)$$

What is the physical interpretation of the operator a_α and what is the effect of a_α on a given state $|\alpha_1 \alpha_2 \dots \alpha_n\rangle_{AS}$? Consider the following matrix element

$$\langle \alpha_1 \alpha_2 \dots \alpha_n | a_\alpha | \alpha'_1 \alpha'_2 \dots \alpha'_m \rangle \quad (3.11)$$

where both sides are antisymmetric. We distinguish between two cases. The first (1) is when $\alpha \in \{\alpha_i\}$. Using the Pauli principle of Eq. (3.6) it follows

$$\langle \alpha_1 \alpha_2 \dots \alpha_n | a_\alpha = 0 \quad (3.12)$$

The second (2) case is when $\alpha \notin \{\alpha_i\}$. It follows that an hermitian conjugation

$$\langle \alpha_1 \alpha_2 \dots \alpha_n | a_\alpha = \langle \alpha \alpha_1 \alpha_2 \dots \alpha_n | \quad (3.13)$$

Eq. (3.13) holds for case (1) since the lefthand side is zero due to the Pauli principle. We write Eq. (3.11) as

$$\langle \alpha_1 \alpha_2 \dots \alpha_n | a_\alpha | \alpha'_1 \alpha'_2 \dots \alpha'_m \rangle = \langle \alpha_1 \alpha_2 \dots \alpha_n | \alpha \alpha'_1 \alpha'_2 \dots \alpha'_m \rangle \quad (3.14)$$

Here we must have $m = n + 1$ if Eq. (3.14) has to be trivially different from zero.

For the last case, the minus and plus signs apply when the sequence $\alpha, \alpha_1, \alpha_2, \dots, \alpha_n$ and $\alpha'_1, \alpha'_2, \dots, \alpha'_{n+1}$ are related to each other via even and odd permutations. If we assume that $\alpha \notin \{\alpha_i\}$ we obtain

$$\langle \alpha_1 \alpha_2 \dots \alpha_n | a_\alpha | \alpha'_1 \alpha'_2 \dots \alpha'_{n+1} \rangle = 0 \quad (3.15)$$

when $\alpha \in \{\alpha'_i\}$. If $\alpha \notin \{\alpha'_i\}$, we obtain

$$a_\alpha | \underbrace{\alpha'_1 \alpha'_2 \dots \alpha'_{n+1}}_{\neq \alpha} \rangle = 0 \quad (3.16)$$

and in particular

$$a_\alpha |0\rangle = 0 \quad (3.17)$$

If $\{\alpha \alpha_i\} = \{\alpha'_i\}$, performing the right permutations, the sequence $\alpha, \alpha_1, \alpha_2, \dots, \alpha_n$ is identical with the sequence $\alpha'_1, \alpha'_2, \dots, \alpha'_{n+1}$. This results in

$$\langle \alpha_1 \alpha_2 \dots \alpha_n | a_\alpha | \alpha \alpha_1 \alpha_2 \dots \alpha_n \rangle = 1 \quad (3.18)$$

and thus

$$a_\alpha | \alpha \alpha_1 \alpha_2 \dots \alpha_n \rangle = | \alpha_1 \alpha_2 \dots \alpha_n \rangle \quad (3.19)$$

The action of the operator a_α from the left on a state vector is to remove one particle in the state α . If the state vector does not contain the single-particle state α , the outcome of the operation is zero. The operator a_α is normally called for a destruction or annihilation operator.

The next step is to establish the commutator algebra of $a_\alpha^\dagger$ and a_β .

The action of the anti-commutator $\{a_\alpha^\dagger, a_\alpha\}$ on a given n -particle state is

$$\begin{aligned}
a_\alpha^\dagger a_\alpha \underbrace{|\alpha_1 \alpha_2 \dots \alpha_n\rangle}_{\neq \alpha} &= 0 \\
a_\alpha a_\alpha^\dagger \underbrace{|\alpha_1 \alpha_2 \dots \alpha_n\rangle}_{\neq \alpha} &= a_\alpha \underbrace{|\alpha \alpha_1 \alpha_2 \dots \alpha_n\rangle}_{\neq \alpha} = \underbrace{|\alpha_1 \alpha_2 \dots \alpha_n\rangle}_{\neq \alpha}
\end{aligned} \tag{3.20}$$

if the single-particle state α is not contained in the state.

If it is present we arrive at

$$\begin{aligned}
a_\alpha^\dagger a_\alpha |\alpha_1 \alpha_2 \dots \alpha_k \alpha \alpha_{k+1} \dots \alpha_{n-1}\rangle &= a_\alpha^\dagger a_\alpha (-1)^k |\alpha \alpha_1 \alpha_2 \dots \alpha_{n-1}\rangle \\
&= (-1)^k |\alpha \alpha_1 \alpha_2 \dots \alpha_{n-1}\rangle = |\alpha_1 \alpha_2 \dots \alpha_k \alpha \alpha_{k+1} \dots \alpha_{n-1}\rangle \\
a_\alpha a_\alpha^\dagger |\alpha_1 \alpha_2 \dots \alpha_k \alpha \alpha_{k+1} \dots \alpha_{n-1}\rangle &= 0
\end{aligned} \tag{3.21}$$

From Eqs. (3.20) and (3.21) we arrive at

$$\{a_\alpha^\dagger, a_\alpha\} = a_\alpha^\dagger a_\alpha + a_\alpha a_\alpha^\dagger = 1 \tag{3.22}$$

The action of $\{a_\alpha^\dagger, a_\beta\}$, with $\alpha \neq \beta$ on a given state yields three possibilities. The first case is a state vector which contains both α and β , then either α or β and finally none of them.

The first case results in

$$\begin{aligned}
a_\alpha^\dagger a_\beta |\alpha \beta \alpha_1 \alpha_2 \dots \alpha_{n-2}\rangle &= 0 \\
a_\beta a_\alpha^\dagger |\alpha \beta \alpha_1 \alpha_2 \dots \alpha_{n-2}\rangle &= 0
\end{aligned} \tag{3.23}$$

while the second case gives

$$\begin{aligned}
a_\alpha^\dagger a_\beta \underbrace{|\beta \alpha_1 \alpha_2 \dots \alpha_{n-1}\rangle}_{\neq \alpha} &= \underbrace{|\alpha \alpha_1 \alpha_2 \dots \alpha_{n-1}\rangle}_{\neq \alpha} \\
a_\beta a_\alpha^\dagger \underbrace{|\beta \alpha_1 \alpha_2 \dots \alpha_{n-1}\rangle}_{\neq \alpha} &= a_\beta \underbrace{|\alpha \beta \beta \alpha_1 \alpha_2 \dots \alpha_{n-1}\rangle}_{\neq \alpha} \\
&= - \underbrace{|\alpha \alpha_1 \alpha_2 \dots \alpha_{n-1}\rangle}_{\neq \alpha}
\end{aligned} \tag{3.24}$$

Finally if the state vector does not contain α and β

$$\begin{aligned}
a_\alpha^\dagger a_\beta \underbrace{|\alpha_1 \alpha_2 \dots \alpha_n\rangle}_{\neq \alpha, \beta} &= 0 \\
a_\beta a_\alpha^\dagger \underbrace{|\alpha_1 \alpha_2 \dots \alpha_n\rangle}_{\neq \alpha, \beta} &= a_\beta \underbrace{|\alpha \alpha_1 \alpha_2 \dots \alpha_n\rangle}_{\neq \alpha, \beta} = 0
\end{aligned} \tag{3.25}$$

For all three cases we have

$$\{a_\alpha^\dagger, a_\beta\} = a_\alpha^\dagger a_\beta + a_\beta a_\alpha^\dagger = 0, \quad \alpha \neq \beta \tag{3.26}$$

We can summarize our findings in Eqs. (3.22) and (3.26) as

$$\{a_\alpha^\dagger, a_\beta\} = \delta_{\alpha\beta} \tag{3.27}$$

with $\delta_{\alpha\beta}$ is the Kroenecker δ -symbol.

The properties of the creation and annihilation operators can be summarized as (for fermions)

$$a_\alpha^\dagger |0\rangle \equiv |\alpha\rangle,$$

and

$$a_{\alpha}^{\dagger}|\alpha_1 \dots \alpha_n\rangle_{AS} \equiv |\alpha \alpha_1 \dots \alpha_n\rangle_{AS}.$$

from which follows

$$|\alpha_1 \dots \alpha_n\rangle_{AS} = a_{\alpha_1}^{\dagger} a_{\alpha_2}^{\dagger} \dots a_{\alpha_n}^{\dagger} |0\rangle.$$

The hermitian conjugate has the following properties

$$a_{\alpha} = (a_{\alpha}^{\dagger})^{\dagger}.$$

Finally we found

$$a_{\alpha} |\underbrace{\alpha'_1 \alpha'_2 \dots \alpha'_{n+1}}_{\neq \alpha}\rangle = 0, \quad \text{in particular } a_{\alpha} |0\rangle = 0,$$

and

$$a_{\alpha} |\alpha \alpha_1 \alpha_2 \dots \alpha_n\rangle = |\alpha_1 \alpha_2 \dots \alpha_n\rangle,$$

and the corresponding commutator algebra

$$\{a_{\alpha}^{\dagger}, a_{\beta}^{\dagger}\} = \{a_{\alpha}, a_{\beta}\} = 0 \quad \{a_{\alpha}^{\dagger}, a_{\beta}\} = \delta_{\alpha\beta}.$$

3.1 One-body operators in second quantization

The result we are about to derive was obtained once already, in Section 2.8, by integrating over permutations of the Hartree function. It is worth comparing the two routes as we go: what took a page of careful argument there follows here from the anticommutation relations alone.

A very useful operator is the so-called number-operator. Most physics cases we will study in this text conserve the total number of particles. The number operator is therefore a useful quantity which allows us to test that our many-body formalism conserves the number of particles. In for example (d, p) or (p, d) reactions it is important to be able to describe quantum mechanical states where particles get added or removed. A creation operator $a_{\alpha}^{\dagger}$ adds one particle to the single-particle state α of a give many-body state vector, while an annihilation operator a_{α} removes a particle from a single-particle state α .

Let us consider an operator proportional with $a_{\alpha}^{\dagger} a_{\beta}$ and $\alpha = \beta$. It acts on an n -particle state resulting in

$$a_{\alpha}^{\dagger} a_{\alpha} |\alpha_1 \alpha_2 \dots \alpha_n\rangle = \begin{cases} 0 & \alpha \notin \{\alpha_i\} \\ |\alpha_1 \alpha_2 \dots \alpha_n\rangle & \alpha \in \{\alpha_i\} \end{cases} \quad (3.28)$$

Summing over all possible one-particle states we arrive at

$$\left(\sum_{\alpha} a_{\alpha}^{\dagger} a_{\alpha} \right) |\alpha_1 \alpha_2 \dots \alpha_n\rangle = n |\alpha_1 \alpha_2 \dots \alpha_n\rangle \quad (3.29)$$

The operator

$$\hat{N} = \sum_{\alpha} a_{\alpha}^{\dagger} a_{\alpha} \quad (3.30)$$

is called the number operator since it counts the number of particles in a give state vector when it acts on the different single-particle states. It acts on one single-particle state at the time and falls therefore under category one-body operators. Next we look at another important one-body operator, namely $\hat{H}_0$ and study its operator form in the occupation number representation.

We want to obtain an expression for a one-body operator which conserves the number of particles. Here we study the one-body operator for the kinetic energy plus an eventual external one-body potential. The action of this operator on a particular n -body state with its pertinent expectation value has already been studied in coordinate space. In coordinate space the operator reads

$$\hat{H}_0 = \sum_i \hat{h}_0(x_i) \quad (3.31)$$

and the anti-symmetric n -particle Slater determinant is defined as

$$\Phi(x_1, x_2, \dots, x_n, \alpha_1, \alpha_2, \dots, \alpha_n) = \frac{1}{\sqrt{n!}} \sum_p (-1)^p \hat{P} \psi_{\alpha_1}(x_1) \psi_{\alpha_2}(x_2) \dots \psi_{\alpha_n}(x_n).$$

Defining

$$\hat{h}_0(x_i) \psi_{\alpha_i}(x_i) = \sum_{\alpha'_i} \psi_{\alpha'_i}(x_i) \langle \alpha'_i | \hat{h}_0 | \alpha_i \rangle \quad (3.32)$$

we can easily evaluate the action of $\hat{H}_0$ on each product of one-particle functions in Slater determinant. From Eq. (3.32) we obtain the following result without permuting any particle pair

$$\begin{aligned} & \left(\sum_i \hat{h}_0(x_i) \right) \psi_{\alpha_1}(x_1) \psi_{\alpha_2}(x_2) \dots \psi_{\alpha_n}(x_n) \\ = & \sum_{\alpha'_1} \langle \alpha'_1 | \hat{h}_0 | \alpha_1 \rangle \psi_{\alpha'_1}(x_1) \psi_{\alpha_2}(x_2) \dots \psi_{\alpha_n}(x_n) \\ + & \sum_{\alpha'_2} \langle \alpha'_2 | \hat{h}_0 | \alpha_2 \rangle \psi_{\alpha_1}(x_1) \psi_{\alpha'_2}(x_2) \dots \psi_{\alpha_n}(x_n) \\ + & \dots \\ + & \sum_{\alpha'_n} \langle \alpha'_n | \hat{h}_0 | \alpha_n \rangle \psi_{\alpha_1}(x_1) \psi_{\alpha_2}(x_2) \dots \psi_{\alpha'_n}(x_n) \end{aligned} \quad (3.33)$$

If we interchange particles 1 and 2 we obtain

$$\begin{aligned} & \left(\sum_i \hat{h}_0(x_i) \right) \psi_{\alpha_1}(x_2) \psi_{\alpha_2}(x_1) \dots \psi_{\alpha_n}(x_n) \\ = & \sum_{\alpha'_2} \langle \alpha'_2 | \hat{h}_0 | \alpha_2 \rangle \psi_{\alpha_1}(x_2) \psi_{\alpha'_2}(x_1) \dots \psi_{\alpha_n}(x_n) \\ + & \sum_{\alpha'_1} \langle \alpha'_1 | \hat{h}_0 | \alpha_1 \rangle \psi_{\alpha'_1}(x_2) \psi_{\alpha_2}(x_1) \dots \psi_{\alpha_n}(x_n) \\ + & \dots \\ + & \sum_{\alpha'_n} \langle \alpha'_n | \hat{h}_0 | \alpha_n \rangle \psi_{\alpha_1}(x_2) \psi_{\alpha_2}(x_1) \dots \psi_{\alpha'_n}(x_n) \end{aligned} \quad (3.34)$$

We can continue by computing all possible permutations. We rewrite also our Slater determinant in its second quantized form and skip the dependence on the quantum numbers x_i . Summing up all contributions and taking care of all phases $(-1)^p$ we arrive at

$$\begin{aligned} \hat{H}_0 |\alpha_1, \alpha_2, \dots, \alpha_n\rangle = & \sum_{\alpha'_1} \langle \alpha'_1 | \hat{h}_0 | \alpha_1 \rangle |\alpha'_1 \alpha_2 \dots \alpha_n\rangle \\ + & \sum_{\alpha'_2} \langle \alpha'_2 | \hat{h}_0 | \alpha_2 \rangle |\alpha_1 \alpha'_2 \dots \alpha_n\rangle \\ + & \dots \\ + & \sum_{\alpha'_n} \langle \alpha'_n | \hat{h}_0 | \alpha_n \rangle |\alpha_1 \alpha_2 \dots \alpha'_n\rangle \end{aligned} \quad (3.35)$$

In Eq. (3.35) we have expressed the action of the one-body operator of Eq. (3.31) on the n -body state in its second quantized form. This equation can be further manipulated if we use the properties of the creation and annihilation operator on each primed quantum number, that is

$$|\alpha_1 \alpha_2 \dots \alpha'_k \dots \alpha_n\rangle = a_{\alpha'_k}^\dagger a_{\alpha_k} |\alpha_1 \alpha_2 \dots \alpha_k \dots \alpha_n\rangle \quad (3.36)$$

Inserting this in the right-hand side of Eq. (3.35) results in

$$\begin{aligned} \hat{H}_0 |\alpha_1 \alpha_2 \dots \alpha_n\rangle &= \sum_{\alpha'_1} \langle \alpha'_1 | \hat{h}_0 | \alpha_1 \rangle a_{\alpha'_1}^\dagger a_{\alpha_1} |\alpha_1 \alpha_2 \dots \alpha_n\rangle \\ &+ \sum_{\alpha'_2} \langle \alpha'_2 | \hat{h}_0 | \alpha_2 \rangle a_{\alpha'_2}^\dagger a_{\alpha_2} |\alpha_1 \alpha_2 \dots \alpha_n\rangle \\ &+ \dots \\ &+ \sum_{\alpha'_n} \langle \alpha'_n | \hat{h}_0 | \alpha_n \rangle a_{\alpha'_n}^\dagger a_{\alpha_n} |\alpha_1 \alpha_2 \dots \alpha_n\rangle \\ &= \sum_{\alpha, \beta} \langle \alpha | \hat{h}_0 | \beta \rangle a_\alpha^\dagger a_\beta |\alpha_1 \alpha_2 \dots \alpha_n\rangle \end{aligned} \quad (3.37)$$

In the number occupation representation or second quantization we get the following expression for a one-body operator which conserves the number of particles

$$\hat{H}_0 = \sum_{\alpha, \beta} \langle \alpha | \hat{h}_0 | \beta \rangle a_\alpha^\dagger a_\beta \quad (3.38)$$

Obviously, $\hat{H}_0$ can be replaced by any other one-body operator which preserved the number of particles. The structure of the operator is therefore not limited to say the kinetic or single-particle energy only.

The operator $\hat{H}_0$ takes a particle from the single-particle state β to the single-particle state α with a probability for the transition given by the expectation value $\langle \alpha | \hat{h}_0 | \beta \rangle$.

It is instructive to verify Eq. (3.38) by computing the expectation value of $\hat{H}_0$ between two single-particle states

$$\langle \alpha_1 | \hat{h}_0 | \alpha_2 \rangle = \sum_{\alpha, \beta} \langle \alpha | \hat{h}_0 | \beta \rangle \langle 0 | a_{\alpha_1} a_\alpha^\dagger a_\beta a_{\alpha_2}^\dagger | 0 \rangle \quad (3.39)$$

Using the commutation relations for the creation and annihilation operators we have

$$a_{\alpha_1} a_\alpha^\dagger a_\beta a_{\alpha_2}^\dagger = (\delta_{\alpha\alpha_1} - a_\alpha^\dagger a_{\alpha_1}) (\delta_{\beta\alpha_2} - a_{\alpha_2}^\dagger a_\beta), \quad (3.40)$$

which results in

$$\langle 0 | a_{\alpha_1} a_\alpha^\dagger a_\beta a_{\alpha_2}^\dagger | 0 \rangle = \delta_{\alpha\alpha_1} \delta_{\beta\alpha_2} \quad (3.41)$$

and

$$\langle \alpha_1 | \hat{h}_0 | \alpha_2 \rangle = \sum_{\alpha, \beta} \langle \alpha | \hat{h}_0 | \beta \rangle \delta_{\alpha\alpha_1} \delta_{\beta\alpha_2} = \langle \alpha_1 | \hat{h}_0 | \alpha_2 \rangle \quad (3.42)$$

3.2 Two-body operators in second quantization

The antisymmetrised matrix elements $\langle pq | \hat{v} | rs \rangle_{AS}$ that appear below are exactly those of Eq. (2.36), and the n^4 of them are the storage problem discussed in Section 2.10. Nothing in this chapter changes that count; what changes is the ease with which they are contracted.

Let us now derive the expression for our two-body interaction part, which also conserves the number of particles. We can proceed in exactly the same way as for the one-body operator. In the coordinate representation our two-body interaction part takes the following expression

$$\hat{H}_I = \sum_{i < j} V(x_i, x_j) \quad (3.43)$$

where the summation runs over distinct pairs. The term V can be an interaction model for the nucleon-nucleon interaction or the interaction between two electrons. It can also include additional two-body interaction terms.

The action of this operator on a product of two single-particle functions is defined as

$$V(x_i, x_j) \psi_{\alpha_k}(x_i) \psi_{\alpha_l}(x_j) = \sum_{\alpha'_k \alpha'_l} \psi'_{\alpha'_k}(x_i) \psi'_{\alpha'_l}(x_j) \langle \alpha'_k \alpha'_l | \hat{v} | \alpha_k \alpha_l \rangle \quad (3.44)$$

We can now let $\hat{H}_I$ act on all terms in the linear combination for $|\alpha_1 \alpha_2 \dots \alpha_n\rangle$. Without any permutations we have

$$\begin{aligned} & \left(\sum_{i < j} V(x_i, x_j) \right) \psi_{\alpha_1}(x_1) \psi_{\alpha_2}(x_2) \dots \psi_{\alpha_n}(x_n) \\ = & \sum_{\alpha'_1 \alpha'_2} \langle \alpha'_1 \alpha'_2 | \hat{v} | \alpha_1 \alpha_2 \rangle \psi'_{\alpha'_1}(x_1) \psi'_{\alpha'_2}(x_2) \dots \psi_{\alpha_n}(x_n) \\ + & \dots \\ + & \sum_{\alpha'_1 \alpha'_n} \langle \alpha'_1 \alpha'_n | \hat{v} | \alpha_1 \alpha_n \rangle \psi'_{\alpha'_1}(x_1) \psi_{\alpha_2}(x_2) \dots \psi'_{\alpha'_n}(x_n) \\ + & \dots \\ + & \sum_{\alpha'_2 \alpha'_n} \langle \alpha'_2 \alpha'_n | \hat{v} | \alpha_2 \alpha_n \rangle \psi_{\alpha_1}(x_1) \psi'_{\alpha'_2}(x_2) \dots \psi'_{\alpha'_n}(x_n) \\ + & \dots \end{aligned} \quad (3.45)$$

where on the rhs we have a term for each distinct pairs.

For the other terms on the rhs we obtain similar expressions and summing over all terms we obtain

$$\begin{aligned} H_I |\alpha_1 \alpha_2 \dots \alpha_n\rangle = & \sum_{\alpha'_1, \alpha'_2} \langle \alpha'_1 \alpha'_2 | \hat{v} | \alpha_1 \alpha_2 \rangle |\alpha'_1 \alpha'_2 \dots \alpha_n\rangle \\ + & \dots \\ + & \sum_{\alpha'_1, \alpha'_n} \langle \alpha'_1 \alpha'_n | \hat{v} | \alpha_1 \alpha_n \rangle |\alpha'_1 \alpha_2 \dots \alpha'_n\rangle \\ + & \dots \\ + & \sum_{\alpha'_2, \alpha'_n} \langle \alpha'_2 \alpha'_n | \hat{v} | \alpha_2 \alpha_n \rangle |\alpha_1 \alpha'_2 \dots \alpha'_n\rangle \\ + & \dots \end{aligned} \quad (3.46)$$

We introduce second quantization via the relation

$$\begin{aligned} & a_{\alpha'_k}^\dagger a_{\alpha'_l}^\dagger a_{\alpha_l} a_{\alpha_k} |\alpha_1 \alpha_2 \dots \alpha_k \dots \alpha_l \dots \alpha_n\rangle \\ = & (-1)^{k-1} (-1)^{l-2} a_{\alpha'_k}^\dagger a_{\alpha'_l}^\dagger a_{\alpha_l} a_{\alpha_k} |\alpha_k \alpha_l \underbrace{\alpha_1 \alpha_2 \dots \alpha_n}_{\neq \alpha_k, \alpha_l}\rangle \\ = & (-1)^{k-1} (-1)^{l-2} |\alpha'_k \alpha'_l \underbrace{\alpha_1 \alpha_2 \dots \alpha_n}_{\neq \alpha'_k, \alpha'_l}\rangle \\ = & |\alpha_1 \alpha_2 \dots \alpha'_k \dots \alpha'_l \dots \alpha_n\rangle \end{aligned} \quad (3.47)$$

Inserting this in (3.46) gives

$$\begin{aligned}
H_I |\alpha_1 \alpha_2 \dots \alpha_n\rangle &= \sum_{\alpha'_1, \alpha'_2} \langle \alpha'_1 \alpha'_2 | \hat{v} | \alpha_1 \alpha_2 \rangle a_{\alpha'_1}^\dagger a_{\alpha'_2}^\dagger a_{\alpha_2} a_{\alpha_1} |\alpha_1 \alpha_2 \dots \alpha_n\rangle \\
&+ \dots \\
&= \sum_{\alpha'_1, \alpha'_n} \langle \alpha'_1 \alpha'_n | \hat{v} | \alpha_1 \alpha_n \rangle a_{\alpha'_1}^\dagger a_{\alpha'_n}^\dagger a_{\alpha_n} a_{\alpha_1} |\alpha_1 \alpha_2 \dots \alpha_n\rangle \\
&+ \dots \\
&= \sum_{\alpha'_2, \alpha'_n} \langle \alpha'_2 \alpha'_n | \hat{v} | \alpha_2 \alpha_n \rangle a_{\alpha'_2}^\dagger a_{\alpha'_n}^\dagger a_{\alpha_n} a_{\alpha_2} |\alpha_1 \alpha_2 \dots \alpha_n\rangle \\
&+ \dots \\
&= \sum'_{\alpha, \beta, \gamma, \delta} \langle \alpha \beta | \hat{v} | \gamma \delta \rangle a_\alpha^\dagger a_\beta^\dagger a_\delta a_\gamma |\alpha_1 \alpha_2 \dots \alpha_n\rangle \quad (3.48)
\end{aligned}$$

Here we let $\sum'$ indicate that the sums running over α and β run over all single-particle states, while the summations γ and δ run over all pairs of single-particle states. We wish to remove this restriction and since

$$\langle \alpha \beta | \hat{v} | \gamma \delta \rangle = \langle \beta \alpha | \hat{v} | \delta \gamma \rangle \quad (3.49)$$

we get

$$\sum_{\alpha \beta} \langle \alpha \beta | \hat{v} | \gamma \delta \rangle a_\alpha^\dagger a_\beta^\dagger a_\delta a_\gamma = \sum_{\alpha \beta} \langle \beta \alpha | \hat{v} | \delta \gamma \rangle a_\alpha^\dagger a_\beta^\dagger a_\delta a_\gamma \quad (3.50)$$

$$= \sum_{\alpha \beta} \langle \beta \alpha | \hat{v} | \delta \gamma \rangle a_\beta^\dagger a_\alpha^\dagger a_\gamma a_\delta \quad (3.51)$$

where we have used the anti-commutation rules.

Changing the summation indices α and β in (3.51) we obtain

$$\sum_{\alpha \beta} \langle \alpha \beta | \hat{v} | \gamma \delta \rangle a_\alpha^\dagger a_\beta^\dagger a_\delta a_\gamma = \sum_{\alpha \beta} \langle \alpha \beta | \hat{v} | \delta \gamma \rangle a_\alpha^\dagger a_\beta^\dagger a_\gamma a_\delta \quad (3.52)$$

From this it follows that the restriction on the summation over γ and δ can be removed if we multiply with a factor $\frac{1}{2}$, resulting in

$$\hat{H}_I = \frac{1}{2} \sum_{\alpha \beta \gamma \delta} \langle \alpha \beta | \hat{v} | \gamma \delta \rangle a_\alpha^\dagger a_\beta^\dagger a_\delta a_\gamma \quad (3.53)$$

where we sum freely over all single-particle states α, β, γ and δ .

With this expression we can now verify that the second quantization form of $\hat{H}_I$ in Eq. (3.53) results in the same matrix between two anti-symmetrized two-particle states as its corresponding coordinate space representation. We have

$$\langle \alpha_1 \alpha_2 | \hat{H}_I | \beta_1 \beta_2 \rangle = \frac{1}{2} \sum_{\alpha \beta \gamma \delta} \langle \alpha \beta | \hat{v} | \gamma \delta \rangle \langle 0 | a_{\alpha_2} a_{\alpha_1} a_\alpha^\dagger a_\beta^\dagger a_\delta a_\gamma a_{\beta_1}^\dagger a_{\beta_2}^\dagger | 0 \rangle. \quad (3.54)$$

Using the commutation relations we get

$$\begin{aligned}
& a_{\alpha_2} a_{\alpha_1} a_\alpha^\dagger a_\beta^\dagger a_\delta a_\gamma a_{\beta_1}^\dagger a_{\beta_2}^\dagger \\
= & a_{\alpha_2} a_{\alpha_1} a_\alpha^\dagger a_\beta^\dagger (a_\delta \delta_{\gamma \beta_1} a_{\beta_2}^\dagger - a_\delta a_{\beta_1}^\dagger a_\gamma a_{\beta_2}^\dagger) \\
= & a_{\alpha_2} a_{\alpha_1} a_\alpha^\dagger a_\beta^\dagger (\delta_{\gamma \beta_1} \delta_{\delta \beta_2} - \delta_{\gamma \beta_1} a_{\beta_2}^\dagger a_\delta - a_\delta a_{\beta_1}^\dagger \delta_{\gamma \beta_2} + a_\delta a_{\beta_1}^\dagger a_{\beta_2}^\dagger a_\gamma) \\
= & a_{\alpha_2} a_{\alpha_1} a_\alpha^\dagger a_\beta^\dagger (\delta_{\gamma \beta_1} \delta_{\delta \beta_2} - \delta_{\gamma \beta_1} a_{\beta_2}^\dagger a_\delta \\
& - \delta_{\delta \beta_1} \delta_{\gamma \beta_2} + \delta_{\gamma \beta_2} a_{\beta_1}^\dagger a_\delta + a_\delta a_{\beta_1}^\dagger a_{\beta_2}^\dagger a_\gamma) \quad (3.55)
\end{aligned}$$

The vacuum expectation value of this product of operators becomes

$$\begin{aligned}
 & \langle 0 | a_{\alpha_2} a_{\alpha_1} a_{\alpha}^{\dagger} a_{\beta}^{\dagger} a_{\delta} a_{\gamma} a_{\beta_1}^{\dagger} a_{\beta_2}^{\dagger} | 0 \rangle \\
 = & (\delta_{\gamma\beta_1} \delta_{\delta\beta_2} - \delta_{\delta\beta_1} \delta_{\gamma\beta_2}) \langle 0 | a_{\alpha_2} a_{\alpha_1} a_{\alpha}^{\dagger} a_{\beta}^{\dagger} | 0 \rangle \\
 = & (\delta_{\gamma\beta_1} \delta_{\delta\beta_2} - \delta_{\delta\beta_1} \delta_{\gamma\beta_2}) (\delta_{\alpha\alpha_1} \delta_{\beta\alpha_2} - \delta_{\beta\alpha_1} \delta_{\alpha\alpha_2})
 \end{aligned} \tag{3.56}$$

Insertion of Eq. (3.56) in Eq. (3.54) results in

$$\begin{aligned}
 \langle \alpha_1 \alpha_2 | \hat{H}_I | \beta_1 \beta_2 \rangle &= \frac{1}{2} [\langle \alpha_1 \alpha_2 | \hat{v} | \beta_1 \beta_2 \rangle - \langle \alpha_1 \alpha_2 | \hat{v} | \beta_2 \beta_1 \rangle \\
 &\quad - \langle \alpha_2 \alpha_1 | \hat{v} | \beta_1 \beta_2 \rangle + \langle \alpha_2 \alpha_1 | \hat{v} | \beta_2 \beta_1 \rangle] \\
 &= \langle \alpha_1 \alpha_2 | \hat{v} | \beta_1 \beta_2 \rangle - \langle \alpha_1 \alpha_2 | \hat{v} | \beta_2 \beta_1 \rangle \\
 &= \langle \alpha_1 \alpha_2 | \hat{v} | \beta_1 \beta_2 \rangle_{AS}.
 \end{aligned} \tag{3.57}$$

The two-body operator can also be expressed in terms of the anti-symmetrized matrix elements we discussed previously as

$$\begin{aligned}
 \hat{H}_I &= \frac{1}{2} \sum_{\alpha\beta\gamma\delta} \langle \alpha\beta | \hat{v} | \gamma\delta \rangle a_{\alpha}^{\dagger} a_{\beta}^{\dagger} a_{\delta} a_{\gamma} \\
 &= \frac{1}{4} \sum_{\alpha\beta\gamma\delta} [\langle \alpha\beta | \hat{v} | \gamma\delta \rangle - \langle \alpha\beta | \hat{v} | \delta\gamma \rangle] a_{\alpha}^{\dagger} a_{\beta}^{\dagger} a_{\delta} a_{\gamma} \\
 &= \frac{1}{4} \sum_{\alpha\beta\gamma\delta} \langle \alpha\beta | \hat{v} | \gamma\delta \rangle_{AS} a_{\alpha}^{\dagger} a_{\beta}^{\dagger} a_{\delta} a_{\gamma}
 \end{aligned} \tag{3.58}$$

The factors in front of the operator, either $\frac{1}{4}$ or $\frac{1}{2}$ tells whether we use antisymmetrized matrix elements or not.

We can now express the Hamiltonian operator for a many-fermion system in the occupation basis representation as

$$H = \sum_{\alpha,\beta} \langle \alpha | \hat{t} + \hat{u}_{\text{ext}} | \beta \rangle a_{\alpha}^{\dagger} a_{\beta} + \frac{1}{4} \sum_{\alpha\beta\gamma\delta} \langle \alpha\beta | \hat{v} | \gamma\delta \rangle_{AS} a_{\alpha}^{\dagger} a_{\beta}^{\dagger} a_{\delta} a_{\gamma}. \tag{3.59}$$

This is the form we will use in the rest of these lectures, assuming that we work with anti-symmetrized two-body matrix elements.

3.3 Particle-hole formalism

The index conventions used here $-i, j, k, l$ for hole states below the Fermi level, a, b, c, d for particle states above it, and p, q, r, s for either – were introduced in Section 2.6. We now give them their proper setting: the reference determinant becomes a new vacuum, and the operators are redefined with respect to it.

Second quantization is a useful and elegant formalism for constructing many-body states and quantum mechanical operators. One can express and translate many physical processes into simple pictures such as Feynman diagrams. Expectation values of many-body states are also easily calculated. However, although the equations are seemingly easy to set up, from a practical point of view, that is the solution of Schroedinger's equation, there is no particular gain. The many-body equation is equally hard to solve, irrespective of representation. The cliché that there is no free lunch brings us down to earth again. Note however that a transformation to a particular basis, for cases where the interaction obeys specific symmetries, can ease the solution of Schroedinger's equation.

But there is at least one important case where second quantization comes to our rescue. It is namely easy to introduce another reference state than the pure vacuum $|0\rangle$, where all single-particle states are active. With many particles present it is often useful to introduce another reference state than the vacuum state $|0\rangle$.

We will label this state $|c\rangle$ (c for core) and as we will see it can reduce considerably the complexity and thereby the dimensionality of the many-body problem. It allows us to sum up to infinite order specific many-body correlations. The particle-hole representation is one of these handy representations.

In the original particle representation these states are products of the creation operators $a_{\alpha_i}^\dagger$ acting on the true vacuum $|0\rangle$. Following Eq. (3.3) we have

$$|\alpha_1 \alpha_2 \dots \alpha_{n-1} \alpha_n\rangle = a_{\alpha_1}^\dagger a_{\alpha_2}^\dagger \dots a_{\alpha_{n-1}}^\dagger a_{\alpha_n}^\dagger |0\rangle \quad (3.60)$$

$$|\alpha_1 \alpha_2 \dots \alpha_{n-1} \alpha_n \alpha_{n+1}\rangle = a_{\alpha_1}^\dagger a_{\alpha_2}^\dagger \dots a_{\alpha_{n-1}}^\dagger a_{\alpha_n}^\dagger a_{\alpha_{n+1}}^\dagger |0\rangle \quad (3.61)$$

$$|\alpha_1 \alpha_2 \dots \alpha_{n-1}\rangle = a_{\alpha_1}^\dagger a_{\alpha_2}^\dagger \dots a_{\alpha_{n-1}}^\dagger |0\rangle \quad (3.62)$$

If we use Eq. (3.60) as our new reference state, we can simplify considerably the representation of this state

$$|c\rangle \equiv |\alpha_1 \alpha_2 \dots \alpha_{n-1} \alpha_n\rangle = a_{\alpha_1}^\dagger a_{\alpha_2}^\dagger \dots a_{\alpha_{n-1}}^\dagger a_{\alpha_n}^\dagger |0\rangle \quad (3.63)$$

The new reference states for the $n+1$ and $n-1$ states can then be written as

$$|\alpha_1 \alpha_2 \dots \alpha_{n-1} \alpha_n \alpha_{n+1}\rangle = (-1)^n a_{\alpha_{n+1}}^\dagger |c\rangle \equiv (-1)^n |\alpha_{n+1}\rangle_c \quad (3.64)$$

$$|\alpha_1 \alpha_2 \dots \alpha_{n-1}\rangle = (-1)^{n-1} a_{\alpha_n} |c\rangle \equiv (-1)^{n-1} |\alpha_{n-1}\rangle_c \quad (3.65)$$

The first state has one additional particle with respect to the new vacuum state $|c\rangle$ and is normally referred to as a one-particle state or one particle added to the many-body reference state. The second state has one particle less than the reference vacuum state $|c\rangle$ and is referred to as a one-hole state. When dealing with a new reference state it is often convenient to introduce new creation and annihilation operators since we have from Eq. (3.65)

$$a_\alpha |c\rangle \neq 0 \quad (3.66)$$

since α is contained in $|c\rangle$, while for the true vacuum we have $a_\alpha |0\rangle = 0$ for all α .

The new reference state leads to the definition of new creation and annihilation operators which satisfy the following relations

$$b_\alpha |c\rangle = 0 \quad (3.67)$$

$$\{b_\alpha^\dagger, b_\beta^\dagger\} = \{b_\alpha, b_\beta\} = 0 \quad (3.68)$$

$$\{b_\alpha^\dagger, b_\beta\} = \delta_{\alpha\beta} \quad (3.68)$$

We assume also that the new reference state is properly normalized

$$\langle c|c\rangle = 1 \quad (3.69)$$

The physical interpretation of these new operators is that of so-called quasiparticle states. This means that a state defined by the addition of one extra particle to a reference state $|c\rangle$ may not necessarily be interpreted as one particle coupled to a core. We define now new creation operators that act on a state α creating a new quasiparticle state

$$b_\alpha^\dagger |c\rangle = \begin{cases} a_\alpha^\dagger |c\rangle = |\alpha\rangle, & \alpha > F \\ a_\alpha |c\rangle = |\alpha^{-1}\rangle, & \alpha \leq F \end{cases} \quad (3.70)$$

where F is the Fermi level representing the last occupied single-particle orbit of the new reference state $|c\rangle$.

The annihilation is the hermitian conjugate of the creation operator

$$b_\alpha = (b_\alpha^\dagger)^\dagger,$$

resulting in

$$b_{\alpha}^{\dagger} = \begin{cases} a_{\alpha}^{\dagger} & \alpha > F \\ a_{\alpha} & \alpha \leq F \end{cases} \quad b_{\alpha} = \begin{cases} a_{\alpha} & \alpha > F \\ a_{\alpha}^{\dagger} & \alpha \leq F \end{cases} \quad (3.71)$$

With the new creation and annihilation operator we can now construct many-body quasiparticle states, with one-particle-one-hole states, two-particle-two-hole states etc in the same fashion as we previously constructed many-particle states. We can write a general particle-hole state as

$$|\beta_1 \beta_2 \dots \beta_{n_p} \gamma_1^{-1} \gamma_2^{-1} \dots \gamma_{n_h}^{-1}\rangle \equiv \underbrace{b_{\beta_1}^{\dagger} b_{\beta_2}^{\dagger} \dots b_{\beta_{n_p}}^{\dagger}}_{>F} \underbrace{b_{\gamma_1}^{\dagger} b_{\gamma_2}^{\dagger} \dots b_{\gamma_{n_h}}^{\dagger}}_{\leq F} |c\rangle \quad (3.72)$$

We can now rewrite our one-body and two-body operators in terms of the new creation and annihilation operators. The number operator becomes

$$\hat{N} = \sum_{\alpha} a_{\alpha}^{\dagger} a_{\alpha} = \sum_{\alpha > F} b_{\alpha}^{\dagger} b_{\alpha} + n_c - \sum_{\alpha \leq F} b_{\alpha}^{\dagger} b_{\alpha} \quad (3.73)$$

where n_c is the number of particle in the new vacuum state $|c\rangle$. The action of $\hat{N}$ on a many-body state results in

$$\hat{N} |\beta_1 \beta_2 \dots \beta_{n_p} \gamma_1^{-1} \gamma_2^{-1} \dots \gamma_{n_h}^{-1}\rangle = (n_p + n_c - n_h) |\beta_1 \beta_2 \dots \beta_{n_p} \gamma_1^{-1} \gamma_2^{-1} \dots \gamma_{n_h}^{-1}\rangle \quad (3.74)$$

Here $n = n_p + n_c - n_h$ is the total number of particles in the quasi-particle state of Eq. (3.72). Note that $\hat{N}$ counts the total number of particles present

$$N_{qp} = \sum_{\alpha} b_{\alpha}^{\dagger} b_{\alpha}, \quad (3.75)$$

gives us the number of quasi-particles as can be seen by computing

$$N_{qp} |\beta_1 \beta_2 \dots \beta_{n_p} \gamma_1^{-1} \gamma_2^{-1} \dots \gamma_{n_h}^{-1}\rangle = (n_p + n_h) |\beta_1 \beta_2 \dots \beta_{n_p} \gamma_1^{-1} \gamma_2^{-1} \dots \gamma_{n_h}^{-1}\rangle \quad (3.76)$$

where $n_{qp} = n_p + n_h$ is the total number of quasi-particles.

We express the one-body operator $\hat{H}_0$ in terms of the quasi-particle creation and annihilation operators, resulting in

$$\begin{aligned} \hat{H}_0 = & \sum_{\alpha\beta > F} \langle \alpha | \hat{h}_0 | \beta \rangle b_{\alpha}^{\dagger} b_{\beta} + \sum_{\alpha > F, \beta \leq F} \left[\langle \alpha | \hat{h}_0 | \beta \rangle b_{\alpha}^{\dagger} b_{\beta}^{\dagger} + \langle \beta | \hat{h}_0 | \alpha \rangle b_{\beta} b_{\alpha} \right] \\ & + \sum_{\alpha \leq F} \langle \alpha | \hat{h}_0 | \alpha \rangle - \sum_{\alpha\beta \leq F} \langle \beta | \hat{h}_0 | \alpha \rangle b_{\alpha}^{\dagger} b_{\beta} \end{aligned} \quad (3.77)$$

The first term gives contribution only for particle states, while the last one contributes only for holestates. The second term can create or destroy a set of quasi-particles and the third term is the contribution from the vacuum state $|c\rangle$.

Before we continue with the expressions for the two-body operator, we introduce a nomenclature we will use for the rest of this text. It is inspired by the notation used in quantum chemistry. We reserve the labels $i, j, k, \dots$ for hole states and $a, b, c, \dots$ for states above F , viz. particle states. This means also that we will skip the constraint $\leq F$ or $> F$ in the summation symbols. Our operator $\hat{H}_0$ reads now

$$\begin{aligned} \hat{H}_0 = & \sum_{ab} \langle a | \hat{h} | b \rangle b_a^{\dagger} b_b + \sum_{ai} \left[\langle a | \hat{h} | i \rangle b_a^{\dagger} b_i^{\dagger} + \langle i | \hat{h} | a \rangle b_i b_a \right] \\ & + \sum_i \langle i | \hat{h} | i \rangle - \sum_{ij} \langle j | \hat{h} | i \rangle b_i^{\dagger} b_j \end{aligned} \quad (3.78)$$

The two-particle operator in the particle-hole formalism is more complicated since we have to translate four indices $\alpha\beta\gamma\delta$ to the possible combinations of particle and hole states. When performing the commutator algebra we can regroup the operator in five different terms

$$\hat{H}_I = \hat{H}_I^{(a)} + \hat{H}_I^{(b)} + \hat{H}_I^{(c)} + \hat{H}_I^{(d)} + \hat{H}_I^{(e)} \quad (3.79)$$

Using anti-symmetrized matrix elements, the term $\hat{H}_I^{(a)}$ is

$$\hat{H}_I^{(a)} = \frac{1}{4} \sum_{abcd} \langle ab|\hat{V}|cd\rangle b_a^\dagger b_b^\dagger b_d b_c \quad (3.80)$$

The next term $\hat{H}_I^{(b)}$ reads

$$\hat{H}_I^{(b)} = \frac{1}{4} \sum_{abci} \left(\langle ab|\hat{V}|ci\rangle b_a^\dagger b_b^\dagger b_i^\dagger b_c + \langle ai|\hat{V}|cb\rangle b_a^\dagger b_i b_b b_c \right) \quad (3.81)$$

This term conserves the number of quasiparticles but creates or removes a three-particle-one-hole state. For $\hat{H}_I^{(c)}$ we have

$$\begin{aligned} \hat{H}_I^{(c)} = & \frac{1}{4} \sum_{abij} \left(\langle ab|\hat{V}|ij\rangle b_a^\dagger b_b^\dagger b_j^\dagger b_i^\dagger + \langle ij|\hat{V}|ab\rangle b_a b_b b_j b_i \right) + \\ & \frac{1}{2} \sum_{abij} \langle ai|\hat{V}|bj\rangle b_a^\dagger b_j^\dagger b_b b_i + \frac{1}{2} \sum_{abi} \langle ai|\hat{V}|bi\rangle b_a^\dagger b_b. \end{aligned} \quad (3.82)$$

The first line stands for the creation of a two-particle-two-hole state, while the second line represents the creation to two one-particle-one-hole pairs while the last term represents a contribution to the particle single-particle energy from the hole states, that is an interaction between the particle states and the hole states within the new vacuum state. The fourth term reads

$$\begin{aligned} \hat{H}_I^{(d)} = & \frac{1}{4} \sum_{aijk} \left(\langle ai|\hat{V}|jk\rangle b_a^\dagger b_k^\dagger b_j^\dagger b_i + \langle ji|\hat{V}|ak\rangle b_k^\dagger b_j b_i b_a \right) + \\ & \frac{1}{4} \sum_{aij} \left(\langle ai|\hat{V}|ji\rangle b_a^\dagger b_j^\dagger + \langle ji|\hat{V}|ai\rangle - \langle ji|\hat{V}|ia\rangle b_j b_a \right). \end{aligned} \quad (3.83)$$

The terms in the first line stand for the creation of a particle-hole state interacting with hole states, we will label this as a two-hole-one-particle contribution. The remaining terms are a particle-hole state interacting with the holes in the vacuum state. Finally we have

$$\hat{H}_I^{(e)} = \frac{1}{4} \sum_{ijkl} \langle kl|\hat{V}|ij\rangle b_i^\dagger b_j^\dagger b_l b_k + \frac{1}{2} \sum_{ijk} \langle ij|\hat{V}|kj\rangle b_k^\dagger b_i + \frac{1}{2} \sum_{ij} \langle ij|\hat{V}|ij\rangle \quad (3.84)$$

The first terms represents the interaction between two holes while the second stands for the interaction between a hole and the remaining holes in the vacuum state. It represents a contribution to single-hole energy to first order. The last term collects all contributions to the energy of the ground state of a closed-shell system arising from hole-hole correlations.

3.4 Summarizing and defining a normal-ordered Hamiltonian

Everything collected here has an exact counterpart in Chapter 2. The Slater determinant is Eq. (2.14); the antisymmetrised two-body matrix element is Eq. (2.36); and the energy functional that follows from the normal-ordered Hamiltonian is Eq. (2.37), whose minimisation gives the Hartree-Fock equations (2.41). The difference is that the derivations there ran over permutations, while here they will run over contractions.

$$\Phi_{AS}(\alpha_1, \dots, \alpha_A; x_1, \dots, x_A) = \frac{1}{\sqrt{A}} \sum_{\hat{P}} (-1)^P \prod_{i=1}^A \psi_{\alpha_i}(x_i),$$

which is equivalent with $|\alpha_1 \dots \alpha_A\rangle = a_{\alpha_1}^\dagger \dots a_{\alpha_A}^\dagger |0\rangle$. We have also

$$a_p^\dagger |0\rangle = |p\rangle, \quad a_p |q\rangle = \delta_{pq} |0\rangle$$

$$\delta_{pq} = \{a_p, a_q^\dagger\},$$

and

$$0 = \{a_p^\dagger, a_q\} = \{a_p, a_q\} = \{a_p^\dagger, a_q^\dagger\}$$

$$|\Phi_0\rangle = |\alpha_1 \dots \alpha_A\rangle, \quad \alpha_1, \dots, \alpha_A \leq \alpha_F$$

$$\{a_p^\dagger, a_q\} = \delta_{pq}, p, q \leq \alpha_F$$

$$\{a_p, a_q^\dagger\} = \delta_{pq}, p, q > \alpha_F$$

with $i, j, \dots \leq \alpha_F$, $a, b, \dots > \alpha_F$, $p, q, \dots$ – any

$$a_i |\Phi_0\rangle = |\Phi_i\rangle, \quad a_a^\dagger |\Phi_0\rangle = |\Phi^a\rangle$$

and

$$a_i^\dagger |\Phi_0\rangle = 0 \quad a_a |\Phi_0\rangle = 0$$

The one-body operator is defined as

$$\hat{F} = \sum_{pq} \langle p | \hat{f} | q \rangle a_p^\dagger a_q$$

while the two-body operator is defined as

$$\hat{V} = \frac{1}{4} \sum_{pqrs} \langle pq | \hat{v} | rs \rangle_{AS} a_p^\dagger a_q^\dagger a_s a_r$$

where we have defined the antisymmetric matrix elements

$$\langle pq | \hat{v} | rs \rangle_{AS} = \langle pq | \hat{v} | rs \rangle - \langle pq | \hat{v} | sr \rangle.$$

We can also define a three-body operator

$$\hat{V}_3 = \frac{1}{36} \sum_{pqrstu} \langle pqr | \hat{v}_3 | stu \rangle_{AS} a_p^\dagger a_q^\dagger a_r^\dagger a_u a_t a_s$$

with the antisymmetrized matrix element

$$\langle pqr | \hat{v}_3 | stu \rangle_{AS} = \langle pqr | \hat{v}_3 | stu \rangle + \langle pqr | \hat{v}_3 | tus \rangle + \langle pqr | \hat{v}_3 | ust \rangle - \langle pqr | \hat{v}_3 | sut \rangle - \langle pqr | \hat{v}_3 | tsu \rangle - \langle pqr | \hat{v}_3 | uts \rangle. \quad (3.85)$$

3.5 Normal ordering

The manipulations of the previous sections all follow the same pattern. Take a product of creation and annihilation operators, push the annihilation operators to the right using the anticommutation relations until they meet the vacuum and give zero, and collect the Kronecker deltas that fall out along the way. For two operators this is immediate. For the four operators of an overlap between two-particle states it is already an exercise in care, and for the eight or more operators that arise in perturbation theory and coupled-cluster theory it is hopeless. Wick's theorem organises the whole calculation once and for all. We need two definitions first.

A product of creation and annihilation operators is in *normal-ordered* form when all creation operators stand to the left of all annihilation operators. For an arbitrary product we define the normal-ordered product

$$N[xyz \cdots w] = (-1)^\kappa \left(\text{the same operators, all creation operators to the left,} \right) \quad (3.86)$$

where κ counts the interchanges of neighbouring operators needed to bring the product into that form. Every interchange of two fermion operators costs a minus sign, and the factor $(-1)^\kappa$ records them. Two elementary cases are

$$N[a_p a_q^\dagger] = -a_q^\dagger a_p, \quad N[a_p^\dagger a_q] = a_p^\dagger a_q, \quad (3.87)$$

the first requiring one interchange and the second none. The ordering of the creation operators among themselves, and of the annihilation operators among themselves, is a matter of convention as long as the same sign rule is applied consistently.

The one property that makes normal ordering useful follows immediately from $a_p|0\rangle = 0$ and $\langle 0|a_p^\dagger = 0$:

$$\boxed{\langle 0|N[xyz \cdots w]|0\rangle = 0} \quad \text{for every non-empty product.} \quad (3.88)$$

Whatever else it contains, a normal-ordered product has an annihilation operator on the far right or a creation operator on the far left, and either one kills the vacuum. A normal-ordered product therefore contributes nothing to a vacuum expectation value. The whole problem of evaluating $\langle 0|xyz \cdots w|0\rangle$ reduces to finding the pieces that are *not* normal-ordered.

3.6 Contractions

The *contraction* of two operators is defined as the difference between their product and their normal-ordered product,

$$\overline{xy} \equiv xy - N[xy], \quad (3.89)$$

and we indicate it by a line joining the two operators. Since the normal-ordered piece has vanishing vacuum expectation value by Eq. (3.88), a contraction is just a number,

$$\overline{xy} = \langle 0|xy|0\rangle. \quad (3.90)$$

There are four possible pairs of operators, and only one of them survives:

$$\overline{a_p a_q^\dagger} = \langle 0|a_p a_q^\dagger|0\rangle = \delta_{pq}, \quad (3.91)$$

$$\overline{a_p a_q} = \langle 0|a_p a_q|0\rangle = 0, \quad (3.92)$$

$$\overline{a_q^\dagger a_p} = \langle 0|a_q^\dagger a_p|0\rangle = 0, \quad (3.93)$$

$$\overline{a_p^\dagger a_q^\dagger} = \langle 0|a_p^\dagger a_q^\dagger|0\rangle = 0. \quad (3.94)$$

Only an annihilation operator standing to the *left* of a creation operator gives a non-vanishing contraction. This single fact carries all the bookkeeping that follows: it is the statement that the vacuum contains no particles, and every zero in the calculations below can be traced back to it.

When the two contracted operators are not adjacent, the contraction carries the sign of the interchanges needed to bring them together. For a set of contractions the total sign is

$$(-1)^v, \quad v = \text{number of crossings of the contraction lines,} \quad (3.95)$$

since each crossing corresponds to one interchange of neighbouring fermion operators. We shall use this rule constantly, and it is verified numerically for arbitrary strings in `wick.py`.

3.7 Two operators

We now build the theorem up from the simplest cases, using nothing more than the anticommutation relations. For two operators the content of the theorem is the definition (3.89) rearranged,

$$\begin{aligned} & \square \\ xy &= N[xy] + xy. \end{aligned} \quad (3.96)$$

Take $x = a_p$ and $y = a_q^\dagger$. The anticommutation relation $\{a_p, a_q^\dagger\} = \delta_{pq}$ gives

$$a_p a_q^\dagger = \delta_{pq} - a_q^\dagger a_p = \delta_{pq} + N[a_p a_q^\dagger], \quad (3.97)$$

where the second step is Eq. (3.87). Comparing with Eq. (3.96) identifies the contraction as δ_{pq} , which is Eq. (3.91). Taking the vacuum expectation value and using Eq. (3.88) recovers the overlap of two single-particle states,

$$\langle p|q \rangle = \langle 0|a_p a_q^\dagger|0 \rangle = \delta_{pq}. \quad (3.98)$$

The three remaining combinations $a_p a_q$, $a_p^\dagger a_q$ and $a_p^\dagger a_q^\dagger$ are already normal-ordered up to a sign, so their contractions vanish and so do their vacuum expectation values.

3.8 Three operators

An odd number of operators can never be paired up completely. Whatever contractions we carry out, at least one operator is always left over, and a single uncontracted operator inside a vacuum expectation value gives

$$\langle 0|a_p|0 \rangle = \langle 0|a_p^\dagger|0 \rangle = 0. \quad (3.99)$$

Hence

$$\langle 0|xyz|0 \rangle = 0 \quad (3.100)$$

for every choice of the three operators, as one may also check directly by anticommuting. Take $x = a_p$, $y = a_q^\dagger$, $z = a_r^\dagger$:

$$a_p a_q^\dagger a_r^\dagger = (\delta_{pq} - a_q^\dagger a_p) a_r^\dagger = \delta_{pq} a_r^\dagger - a_q^\dagger (\delta_{pr} - a_r^\dagger a_p), \quad (3.101)$$

and every surviving term contains a single unpaired creation operator acting on the vacuum, so the vacuum expectation value vanishes.

The result is worth stating explicitly because it is used constantly: *any* vacuum expectation value of an odd number of creation and annihilation operators is zero. The same argument reappears in the proof of the theorem, where it disposes of half the terms in the induction step.

3.9 Four operators

Four operators is the first genuinely interesting case, and it is the one that appears whenever we compute the overlap of two two-particle states,

$$|rs \rangle = a_r^\dagger a_s^\dagger |0 \rangle, \quad |pq \rangle = a_p^\dagger a_q^\dagger |0 \rangle, \quad \langle rs|pq \rangle = \langle 0|a_s a_r a_p^\dagger a_q^\dagger|0 \rangle. \quad (3.102)$$

By anticommutation.

We evaluate it the elementary way first, so that the answer can be checked against the theorem afterwards. Anticommuting a_r past $a_p^\dagger$ with $a_r a_p^\dagger = \delta_{rp} - a_p^\dagger a_r$,

$$a_s a_r a_p^\dagger a_q^\dagger = a_s (\delta_{rp} - a_p^\dagger a_r) a_q^\dagger = \delta_{rp} a_s a_q^\dagger - a_s a_p^\dagger a_r a_q^\dagger. \quad (3.103)$$

The first term is handled by Eq. (3.97),

$$\delta_{rp} a_s a_q^\dagger = \delta_{rp} (\delta_{sq} - a_q^\dagger a_s) \longrightarrow \delta_{rp} \delta_{sq}, \quad (3.104)$$

the arrow denoting that we take the vacuum expectation value and drop the normal-ordered remainder. For the second term we anticommute twice more,

$$-a_s a_p^\dagger a_r a_q^\dagger = -a_s a_p^\dagger (\delta_{rq} - a_q^\dagger a_r) = -\delta_{rq} a_s a_p^\dagger + a_s a_p^\dagger a_q^\dagger a_r, \quad (3.105)$$

and the last term has an annihilation operator on the far right, so it dies on the vacuum. The surviving piece is

$$-\delta_{rq} a_s a_p^\dagger = -\delta_{rq} (\delta_{sp} - a_p^\dagger a_s) \longrightarrow -\delta_{rq} \delta_{sp}. \quad (3.106)$$

Collecting the two contributions,

$$\boxed{\langle rs|pq\rangle = \langle 0|a_s a_r a_p^\dagger a_q^\dagger|0\rangle = \delta_{rp} \delta_{sq} - \delta_{sp} \delta_{rq}} \quad (3.107)$$

Two-particle states are orthonormal up to precisely the antisymmetry expected of fermions: the overlap is the antisymmetrised product of single-particle overlaps.

By contractions.

The same answer can be read off without any algebra at all. Four operators can be paired in three ways. The first pairs the two annihilation operators with each other and the two creation operators with each other,

$$\overbrace{a_s a_r} \overbrace{a_p^\dagger a_q^\dagger} = 0, \quad (3.108)$$

which vanishes by Eqs. (3.92) and (3.94). The second is the nested pairing, with no crossings and hence a plus sign,

$$\overbrace{a_s a_r} \overbrace{a_p^\dagger a_q^\dagger} = +\delta_{sq} \delta_{rp}. \quad (3.109)$$

The third has the two lines crossing once, and therefore a minus sign by Eq. (3.95); we draw one line above and one below,

$$\overbrace{a_s a_r} \overbrace{a_p^\dagger a_q^\dagger} = -\delta_{sp} \delta_{rq}. \quad (3.110)$$

Adding the three gives Eq. (3.107) again.

Three pairings, one of which vanishes on inspection and two of which are read straight off: this is the economy the theorem buys. The program `wick.py` carries out both routes for arbitrary strings and confirms that they agree. For the eight-operator strings of Section 3.13 the anticommutation recursion takes 129 steps, while the theorem enumerates 105 pairings of which only a handful survive.

3.10 Wick's theorem

The pattern of the previous three sections generalises. *Wick's theorem* states that an arbitrary product of creation and annihilation operators can be written as the normal-ordered product plus the sum of all possible contractions,

$$\begin{aligned}
 xyz \cdots w &= N[xyz \cdots w] \\
 &+ \sum_{(1)} N[xyz \cdots w] \quad (\text{all single contractions}) \\
 &+ \sum_{(2)} N[xyz \cdots w] \quad (\text{all double contractions}) \\
 &+ \cdots + \sum_{[M/2]} N[xyz \cdots w], \tag{3.111}
 \end{aligned}$$

where $\sum_{(m)}$ runs over all ways of choosing m disjoint pairs of operators to contract, each term carrying the sign (3.95), and $[M/2]$ is the largest number of disjoint pairs that M operators admit.

The theorem becomes a computational tool the moment we take a vacuum expectation value. Every term in Eq. (3.111) that still contains an uncontracted operator is normal-ordered and therefore vanishes by Eq. (3.88). Only the *fully* contracted terms survive:

$$\langle 0 | xyz \cdots w | 0 \rangle = \sum_{[M/2]} N[xyz \cdots w] \quad (\text{fully contracted terms only}) \tag{3.112}$$

For M even there are $(M-1)!!$ such terms – one, three, fifteen and one hundred and five for $M = 2, 4, 6, 8$ – and most of them vanish on inspection because one of their contractions is of the type (3.92)–(3.94). For M odd there are none, which is Eq. (3.100) again.

The practical consequence is the one anticipated in Section 2.16. Because an n -body operator can connect only Slater determinants differing by at most n single-particle states, a configuration-interaction Hamiltonian built from one- and two-body operators is extremely sparse: of the $\binom{n}{N}^2$ possible matrix elements only a vanishing fraction are non-zero. That sparsity is precisely what the Lanczos algorithm of Section 1.26 exploits, and it is why the Hamiltonian never has to be stored.

3.11 Proof of Wick's theorem

The proof is by induction on the number of operators. The case $M = 2$ is Eq. (3.96), established directly from the anticommutation relation. The induction step requires one preparatory result, which does all the real work.

A lemma: appending one operator on the right.

Let $x, y, z, \dots, w$ be a set of operators and let Ω be one further creation or annihilation operator. Then

$$N[xyz \cdots w] \Omega = N[xyz \cdots w \Omega] + \sum_{(1)\Omega} N[xyz \cdots w \Omega], \tag{3.113}$$

where $\sum_{(1)\Omega}$ runs over all single contractions *involving* Ω , that is over the contractions of Ω with each of $x, y, z, \dots, w$ in turn.

Proof of the lemma.

There are four cases to consider, and three of them are immediate.

- (i) Ω is an annihilation operator. Then $N[xyz \cdots w] \Omega$ is already in normal-ordered form, since appending an annihilation operator on the far right cannot disturb the ordering, and so the left-hand side equals $N[xyz \cdots w \Omega]$. At the same time every contraction with Ω vanishes: a non-zero contraction requires an annihilation operator to the *left* of a creation operator, and here Ω is itself an annihilation operator standing to the right of everything. Both extra terms are absent and the lemma holds trivially.
- (ii) All of $x, y, \dots, w$ are annihilation operators and Ω is an annihilation operator. This is a special case of (i), and the lemma is again trivially satisfied.
- (iii) Ω is a creation operator. Here we need only prove the lemma when all of $x, y, \dots, w$ are annihilation operators. The reason is that a creation operator inside the normal-ordered product already stands to the left of Ω 's eventual position, and its contraction with Ω vanishes by Eq. (3.94),

$$\overline{a_i^\dagger a_j^\dagger} = 0.$$

Creation operators in the string thus contribute neither to the reordering nor to the sum of contractions, and may be carried along untouched.

- (iv) It remains to anticommute Ω leftwards through the annihilation operators $x, y, \dots, w$, one at a time. Each step uses

$$a_j \Omega = \delta_{j\Omega} - \Omega a_j,$$

and produces two pieces: a Kronecker delta, which is precisely the contraction of a_j with Ω carrying the sign of the operators Ω has had to cross, and a term in which Ω has moved one place to the left. Iterating until Ω has passed all of them, the accumulated delta terms are exactly the sum $\sum_{(1)\Omega} N[xyz \cdots w \Omega]$, and the final term, in which Ω stands to the left of every annihilation operator, is exactly $N[xyz \cdots w \Omega]$. This is Eq. (3.113). $\square$

The lemma at work.

For $M = 1$ the lemma reads

$$N[a_1] a_\ell^\dagger = N[a_1 a_\ell^\dagger] + N[\overline{a_1 a_\ell^\dagger}] = -\langle 0 | a_\ell^\dagger a_1 | 0 \rangle + \delta_{1\ell}, \quad (3.114)$$

which is just Eq. (3.97) once more. For $M = 2$,

$$a_0 N[a_1] a_\ell^\dagger = N[a_0 a_1 a_\ell^\dagger] + N[\overline{a_0 a_1 a_\ell^\dagger}] + N[\overline{a_0 a_1 a_\ell^\dagger}], \quad (3.115)$$

that is $N[a_0 a_1 a_\ell^\dagger] + \delta_{1\ell} N[a_0] + \delta_{0\ell} N[a_1]$ up to the signs of the crossings. Generally,

$$a_0 N[a_1 a_2 \cdots a_M] a_\ell^\dagger = \sum_{(1)\Omega} N[a_0 a_1 \cdots a_M a_\ell^\dagger] + N[a_0 a_1 \cdots a_M a_\ell^\dagger]. \quad (3.116)$$

The induction step.

Assume Wick's theorem, Eq. (3.111), holds for a product of M operators. Multiply both sides from the right by one further operator Ω and apply the lemma to every term. The first term gives

$$N[xyz \cdots w] \Omega = N[xyz \cdots w \Omega] + \sum_{(1)\Omega} N[xyz \cdots w \Omega], \quad (3.117)$$

the m -fold contracted terms give

$$\sum_{(m)} N[xyz \cdots w] \Omega = \sum_{(m) \neq \Omega} N[xyz \cdots w \Omega] + \sum_{(m) \Omega} N[xyz \cdots w \Omega], \quad (3.118)$$

where $\sum_{(m) \neq \Omega}$ collects the m -fold contractions that do not involve Ω and $\sum_{(m) \Omega}$ those in which one of the m contractions ends on Ω . But an m -fold contraction of the $M+1$ operators either involves Ω or it does not, so the two sums recombine,

$$\sum_{(m) \neq \Omega} + \sum_{(m) \Omega} = \sum_{(m)} \quad \text{over all } M+1 \text{ operators.} \quad (3.119)$$

Collecting the terms by the number of contractions therefore reproduces Eq. (3.111) with M replaced by $M+1$,

$$xyz \cdots w \Omega = N[xyz \cdots w \Omega] + \sum_{(1)} N[\cdots] + \sum_{(2)} N[\cdots] + \cdots + \sum_{[(M+1)/2]} N[\cdots]. \quad (3.120)$$

This completes the induction. $\square$

A remark on parity.

If $M+1$ is odd, the last sum $\sum_{[(M+1)/2]}$ still leaves one operator uncontracted, since an odd number of operators cannot be paired up completely. On taking the vacuum expectation value these terms vanish by Eq. (3.99), and we recover the result of Section 3.8: vacuum expectation values of odd strings are zero. For $M+1$ even the last sum consists of the fully contracted terms, and Eq. (3.112) follows.

3.12 Wick's generalised theorem

In practice we rarely meet a bare product of operators. What we meet is a product of groups, each of which is *already* in normal-ordered form: the bra, the operator, and the ket. Wick's generalised theorem states that for such a product

$$N[A_1 A_2 \cdots] N[B_1 B_2 \cdots] N[C_1 C_2 \cdots] = N[A_1 A_2 \cdots B_1 B_2 \cdots C_1 C_2 \cdots] + \sum_{\text{all}} N[A_1 A_2 \cdots C_1 C_2 \cdots], \quad (3.121)$$

where the sum runs over all contractions *between operators belonging to different normal-ordered groups*. Contractions within a single group are absent, precisely because that group is already normal-ordered and its internal contractions vanish by Eqs. (3.92)–(3.94).

This is a substantial saving. In the examples of the next section the naive count of pairings runs into the hundreds, while the number of contractions between different groups is a handful.

3.13 Applications: matrix elements between two-particle states

We close with the two calculations that motivated the whole construction. Both use the generalised theorem, and in both the bra, the operator and the ket are separately normal-ordered.

A one-body operator.

With

$$\hat{O}^{(1)} = \sum_{pq} \langle p | \hat{\delta}^{(1)} | q \rangle a_p^\dagger a_q, \quad (3.122)$$

and the two-particle states $|ij\rangle = a_i^\dagger a_j^\dagger |0\rangle$, $\langle ij| = \langle 0| a_j a_i$, we need

$$\langle ij | \hat{O}^{(1)} | k\ell \rangle = \sum_{pq} \langle p | \hat{\delta}^{(1)} | q \rangle \langle 0 | \underbrace{a_j a_i}_{\text{bra}} \underbrace{a_p^\dagger a_q}_{\text{operator}} \underbrace{a_k^\dagger a_\ell^\dagger}_{\text{ket}} | 0 \rangle. \quad (3.123)$$

Only contractions between different groups contribute, and each must join an annihilation operator to a creation operator standing to its right. One such term is

$$\overbrace{a_j a_i a_p^\dagger a_q a_k^\dagger a_\ell^\dagger}^{\text{contraction}} = \delta_{ip} \delta_{qk} \delta_{j\ell}, \quad (3.124)$$

which contributes $\delta_{j\ell} \langle i | \hat{\delta}^{(1)} | k \rangle$ after the sums over p and q are carried out. Collecting the four such terms gives the familiar result

$$\langle ij | \hat{O}^{(1)} | k\ell \rangle = \delta_{j\ell} \langle i | \hat{\delta}^{(1)} | k \rangle - \delta_{jk} \langle i | \hat{\delta}^{(1)} | \ell \rangle - \delta_{i\ell} \langle j | \hat{\delta}^{(1)} | k \rangle + \delta_{ik} \langle j | \hat{\delta}^{(1)} | \ell \rangle. \quad (3.125)$$

The matrix element vanishes whenever the two Slater determinants differ by more than one single-particle state, since then no assignment of the deltas can be satisfied. This is the Condon-Slater rule of Section 2.6, now obtained without a single explicit permutation.

The two-body interaction.

With the normal-ordered interaction

$$\hat{H}_I = \frac{1}{2} \sum_{pqrs} \langle pq | v | rs \rangle a_p^\dagger a_q^\dagger a_s a_r, \quad (3.126)$$

the diagonal matrix element in the state $|ij\rangle$ is

$$\langle ij | \hat{H}_I | ij \rangle = \frac{1}{2} \sum_{pqrs} \langle pq | v | rs \rangle \langle 0 | \underbrace{a_j a_i}_{\text{bra}} \underbrace{a_p^\dagger a_q^\dagger a_s a_r}_{\text{operator}} \underbrace{a_i^\dagger a_j^\dagger}_{\text{ket}} | 0 \rangle. \quad (3.127)$$

This is an eight-operator string. Enumerating its 105 pairings by hand would be unpleasant, but the generalised theorem restricts us to contractions between the three groups, and of those only four survive. Two of them are

$$\overbrace{a_j a_i a_p^\dagger a_q^\dagger a_s a_r a_i^\dagger a_j^\dagger}^{\text{contraction}} = \delta_{ip} \delta_{jq} \delta_{ri} \delta_{sj}, \quad (3.128)$$

giving $\langle ij | v | ij \rangle$, and the one obtained by exchanging the roles of r and s , which crosses one more line and therefore enters with the opposite sign, giving $-\langle ij | v | ji \rangle$. The two remaining terms are these two with i and j interchanged; they are equal to the first two and cancel the factor $\frac{1}{2}$. The result is

$$\boxed{\langle ij | \hat{H}_I | ij \rangle = \langle ij | v | ij \rangle - \langle ij | v | ji \rangle = \langle ij | v | ij \rangle_{\text{AS}}} \quad (3.129)$$

which is exactly the antisymmetrised matrix element obtained in Chapter 2 by writing out the antisymmetriser and sorting through permutations. The two derivations agree, as they must; the point is that this one required no permutations at all, only the observation that a contraction is non-zero when an annihilation operator stands to the left of a creation operator.

The program `wick.py` performs every calculation of these sections in both ways – by brute-force anticommutation and by summing contractions with their signs – and verifies that they agree, for strings

of two, four, six and eight operators. It also reproduces Eq. (3.129) by enumerating which of the (p, q, r, s) assignments survive.

3.14 Operators in second quantization

What follows is the bridge from the formalism to a running code. The Hamiltonian matrix built here is the sparse, never-stored matrix that the Lanczos algorithm of Section 1.26 was designed for: only the product $\mathbf{H}\mathbf{v}$ is needed, and the Condon-Slater rules of Section 3.13 say which of its elements can be non-zero. The bit representation below is how that product is evaluated in practice.

In the build-up of a shell-model or FCI code that is meant to tackle large dimensionalities is the action of the Hamiltonian $\hat{H}$ on a Slater determinant represented in second quantization as

$$|\alpha_1 \dots \alpha_n\rangle = a_{\alpha_1}^\dagger a_{\alpha_2}^\dagger \dots a_{\alpha_n}^\dagger |0\rangle.$$

The time consuming part stems from the action of the Hamiltonian on the above determinant,

$$\left(\sum_{\alpha\beta} \langle \alpha | t + u | \beta \rangle a_\alpha^\dagger a_\beta + \frac{1}{4} \sum_{\alpha\beta\gamma\delta} \langle \alpha\beta | \hat{v} | \gamma\delta \rangle a_\alpha^\dagger a_\beta^\dagger a_\delta a_\gamma \right) a_{\alpha_1}^\dagger a_{\alpha_2}^\dagger \dots a_{\alpha_n}^\dagger |0\rangle.$$

A practically useful way to implement this action is to encode a Slater determinant as a bit pattern.

Assume that we have at our disposal n different single-particle orbits $\alpha_0, \alpha_2, \dots, \alpha_{n-1}$ and that we can distribute among these orbits $N \leq n$ particles.

A Slater determinant can then be coded as an integer of n bits. As an example, if we have $n = 16$ single-particle states $\alpha_0, \alpha_1, \dots, \alpha_{15}$ and $N = 4$ fermions occupying the states $\alpha_3, \alpha_6, \alpha_{10}$ and α_{13} we could write this Slater determinant as

$$\Phi_A = a_{\alpha_3}^\dagger a_{\alpha_6}^\dagger a_{\alpha_{10}}^\dagger a_{\alpha_{13}}^\dagger |0\rangle.$$

The unoccupied single-particle states have bit value 0 while the occupied ones are represented by bit state 1. In the binary notation we would write this 16 bits long integer as

$$\begin{array}{cccccccccccccccc} \alpha_0 & \alpha_1 & \alpha_2 & \alpha_3 & \alpha_4 & \alpha_5 & \alpha_6 & \alpha_7 & \alpha_8 & \alpha_9 & \alpha_{10} & \alpha_{11} & \alpha_{12} & \alpha_{13} & \alpha_{14} & \alpha_{15} \\ 0 & 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 \end{array}$$

which translates into the decimal number

$$2^3 + 2^6 + 2^{10} + 2^{13} = 9288.$$

We can thus encode a Slater determinant as a bit pattern.

With N particles that can be distributed over n single-particle states, the total number of Slater determinants (and defining thereby the dimensionality of the system) is

$$\dim(\mathcal{H}) = \binom{n}{N}.$$

The total number of bit patterns is 2^n .

We assume again that we have at our disposal n different single-particle orbits $\alpha_0, \alpha_2, \dots, \alpha_{n-1}$ and that we can distribute among these orbits $N \leq n$ particles. The ordering among these states is important as it defines the order of the creation operators. We will write the determinant

$$\Phi_A = a_{\alpha_3}^\dagger a_{\alpha_6}^\dagger a_{\alpha_{10}}^\dagger a_{\alpha_{13}}^\dagger |0\rangle,$$

in a more compact way as

$$\Phi_{3,6,10,13} = |0001001000100100\rangle.$$

The action of a creation operator is thus

$$a_{\alpha_4}^\dagger \Phi_{3,6,10,13} = a_{\alpha_4}^\dagger |0001001000100100\rangle = a_{\alpha_4}^\dagger a_{\alpha_3}^\dagger a_{\alpha_6}^\dagger a_{\alpha_{10}}^\dagger a_{\alpha_{13}}^\dagger |0\rangle,$$

which becomes

$$-a_{\alpha_3}^\dagger a_{\alpha_4}^\dagger a_{\alpha_6}^\dagger a_{\alpha_{10}}^\dagger a_{\alpha_{13}}^\dagger |0\rangle = -|0001101000100100\rangle.$$

Similarly

$$a_{\alpha_6}^\dagger \Phi_{3,6,10,13} = a_{\alpha_6}^\dagger |0001001000100100\rangle = a_{\alpha_6}^\dagger a_{\alpha_3}^\dagger a_{\alpha_6}^\dagger a_{\alpha_{10}}^\dagger a_{\alpha_{13}}^\dagger |0\rangle,$$

which becomes

$$-a_{\alpha_4}^\dagger (a_{\alpha_6}^\dagger)^2 a_{\alpha_{10}}^\dagger a_{\alpha_{13}}^\dagger |0\rangle = 0!$$

This gives a simple recipe:

- If one of the bits b_j is 1 and we act with a creation operator on this bit, we return a null vector
- If $b_j = 0$, we set it to 1 and return a sign factor $(-1)^l$, where l is the number of bits set before bit j .

Consider the action of $a_{\alpha_2}^\dagger$ on various slater determinants:

$$\begin{aligned} a_{\alpha_2}^\dagger \Phi_{00111} &= a_{\alpha_2}^\dagger |00111\rangle = 0 \times |00111\rangle \\ a_{\alpha_2}^\dagger \Phi_{01011} &= a_{\alpha_2}^\dagger |01011\rangle = (-1) \times |01111\rangle \\ a_{\alpha_2}^\dagger \Phi_{01101} &= a_{\alpha_2}^\dagger |01101\rangle = 0 \times |01101\rangle \\ a_{\alpha_2}^\dagger \Phi_{01110} &= a_{\alpha_2}^\dagger |01110\rangle = 0 \times |01110\rangle \\ a_{\alpha_2}^\dagger \Phi_{10011} &= a_{\alpha_2}^\dagger |10011\rangle = (-1) \times |10111\rangle \\ a_{\alpha_2}^\dagger \Phi_{10101} &= a_{\alpha_2}^\dagger |10101\rangle = 0 \times |10101\rangle \\ a_{\alpha_2}^\dagger \Phi_{10110} &= a_{\alpha_2}^\dagger |10110\rangle = 0 \times |10110\rangle \\ a_{\alpha_2}^\dagger \Phi_{11001} &= a_{\alpha_2}^\dagger |11001\rangle = (+1) \times |11101\rangle \\ a_{\alpha_2}^\dagger \Phi_{11010} &= a_{\alpha_2}^\dagger |11010\rangle = (+1) \times |11110\rangle \end{aligned}$$

What is the simplest way to obtain the phase when we act with one annihilation(creation) operator on the given Slater determinant representation?

We have an SD representation

$$\Phi_\Lambda = a_{\alpha_0}^\dagger a_{\alpha_3}^\dagger a_{\alpha_6}^\dagger a_{\alpha_{10}}^\dagger a_{\alpha_{13}}^\dagger |0\rangle,$$

in a more compact way as

$$\Phi_{0,3,6,10,13} = |1001001000100100\rangle.$$

The action of

$$a_{\alpha_4}^\dagger a_{\alpha_0} \Phi_{0,3,6,10,13} = a_{\alpha_4}^\dagger |0001001000100100\rangle = a_{\alpha_4}^\dagger a_{\alpha_3}^\dagger a_{\alpha_6}^\dagger a_{\alpha_{10}}^\dagger a_{\alpha_{13}}^\dagger |0\rangle,$$

which becomes

$$-a_{\alpha_3}^\dagger a_{\alpha_4}^\dagger a_{\alpha_6}^\dagger a_{\alpha_{10}}^\dagger a_{\alpha_{13}}^\dagger |0\rangle = -|0001101000100100\rangle.$$

The action

$$a_{\alpha_0} \Phi_{0,3,6,10,13} = |0001001000100100\rangle,$$

can be obtained by subtracting the logical sum (AND operation) of $\Phi_{0,3,6,10,13}$ and a word which represents only α_0 , that is

$$|1000000000000000\rangle,$$

from $\Phi_{0,3,6,10,13} = |1001001000100100\rangle$.

This operation gives $|0001001000100100\rangle$.

Similarly, we can form $a_{\alpha_4}^\dagger a_{\alpha_0} \Phi_{0,3,6,10,13}$, say, by adding $|0000100000000000\rangle$ to $a_{\alpha_0} \Phi_{0,3,6,10,13}$, first checking that their logical sum is zero in order to make sure that orbital α_4 is not already occupied.

It is trickier however to get the phase $(-1)^l$. One possibility is as follows

- Let S_1 be a word that represents the 1-bit to be removed and all others set to zero.

In the previous example $S_1 = |1000000000000000\rangle$

- Define S_2 as the similar word that represents the bit to be added, that is in our case

$S_2 = |0000100000000000\rangle$.

- Compute then $S = S_1 - S_2$, which here becomes

$$S = |0111000000000000\rangle$$

- Perform then the logical AND operation of S with the word containing

$$\Phi_{0,3,6,10,13} = |1001001000100100\rangle,$$

which results in $|0001000000000000\rangle$. Counting the number of 1-bits gives the phase. Here you need however an algorithm for bitcounting. Several efficient ones available.

Chapter 4

Physical systems and models

The three preceding chapters built a formalism. This one puts it to work on four systems.

The models collected here share a property that makes them invaluable: each is simple enough to be solved exactly, at least for small numbers of particles, yet each contains a genuine many-body mechanism. The Lipkin model has a collective mode driven by a two-body interaction that scatters pairs across a gap; the pairing model has superfluid correlations; the Fermi-Hubbard model has the competition between kinetic delocalisation and on-site repulsion that drives the Mott transition; and the Calogero-Sutherland model is exactly solvable for arbitrary particle number, with a ground state that is a pure Jastrow function. Because the exact answer is available, these models are the benchmarks against which every approximate method in the rest of this book is measured.

They also share a structure. Every one of them is written as

$$\hat{H} = \sum_{pq} \langle p | \hat{h}_0 | q \rangle a_p^\dagger a_q + \frac{1}{4} \sum_{pqrs} \langle pq | \hat{v} | rs \rangle_{AS} a_p^\dagger a_q^\dagger a_s a_r, \quad (4.1)$$

the general second-quantised Hamiltonian of Chapter 3, with a particular and very restricted choice of matrix elements. What differs between them is which single-particle states there are and which matrix elements are allowed to be non-zero. The symmetries that follow from those restrictions are what make the models tractable, and identifying them is the first step in every case.

Throughout, the detailed algebra is relegated to the exercises at the end of the chapter, where full solutions are given. The main text states each model, identifies its symmetries, and reports what the exact solution looks like. All numbers quoted are produced by `models.py`.

4.1 The Lipkin model

The Lipkin-Meshkov-Glick model [3] describes N fermions distributed over two levels, each of degeneracy Ω . The levels carry the quantum number $\sigma = \pm 1$: the upper has $\sigma = +1$ and energy $+\varepsilon/2$, the lower has $\sigma = -1$ and energy $-\varepsilon/2$. The substates of each level are labelled $p = 1, \dots, \Omega$, so a single-particle state is $|p\sigma\rangle = a_{p\sigma}^\dagger |0\rangle$.

The Hamiltonian is

$$\hat{H} = \hat{H}_0 + \hat{H}_1 + \hat{H}_2, \quad (4.2)$$

with

$$\hat{H}_0 = \frac{1}{2} \varepsilon \sum_{p\sigma} \sigma a_{p\sigma}^\dagger a_{p\sigma}, \quad (4.3)$$

$$\hat{H}_1 = \frac{1}{2} V \sum_{pp'\sigma} a_{p\sigma}^\dagger a_{p'\sigma}^\dagger a_{p'-\sigma} a_{p-\sigma}, \quad (4.4)$$

$$\hat{H}_2 = \frac{1}{2} W \sum_{pp'\sigma} a_{p\sigma}^\dagger a_{p'-\sigma}^\dagger a_{p'\sigma} a_{p-\sigma}. \quad (4.5)$$

The operator $\hat{H}_1$ moves a *pair* of particles from one level to the other, with strength V ; $\hat{H}_2$ is a spin-exchange term of strength W , which moves a pair from $(p\sigma, p' - \sigma)$ to $(p - \sigma, p'\sigma)$. Neither changes the substate labels p , and this is the crucial restriction: the interaction is completely degenerate in p .

Quasispin.

The degeneracy in p means the model has an $SU(2)$ symmetry. Define the *quasispin* operators

$$\hat{J}_\pm = \sum_p a_{p\pm}^\dagger a_{p\mp}, \quad \hat{J}_z = \frac{1}{2} \sum_{p\sigma} \sigma a_{p\sigma}^\dagger a_{p\sigma}, \quad \hat{J}^2 = \hat{J}_+ \hat{J}_- + \hat{J}_z^2 - \hat{J}_z. \quad (4.6)$$

Exercise 4.6 shows, using only the anticommutation relations of Chapter 3, that these satisfy the angular-momentum algebra

$$[\hat{J}_z, \hat{J}_\pm] = \pm \hat{J}_\pm, \quad [\hat{J}_+, \hat{J}_-] = 2\hat{J}_z, \quad [\hat{J}^2, \hat{J}_\pm] = [\hat{J}^2, \hat{J}_z] = 0, \quad (4.7)$$

and that the Hamiltonian collapses to

$$\boxed{\hat{H} = \varepsilon \hat{J}_z + \frac{1}{2} V (\hat{J}_+^2 + \hat{J}_-^2) + \frac{1}{2} W (-\hat{N} + \hat{J}_+ \hat{J}_- + \hat{J}_- \hat{J}_+)} \quad (4.8)$$

with $\hat{N} = \sum_{p\sigma} a_{p\sigma}^\dagger a_{p\sigma}$ the number operator. Every trace of the substate index p has disappeared.

The consequence is decisive. Since $[\hat{H}, \hat{J}^2] = 0$ and $[\hat{H}, \hat{N}] = 0$, both J and N are good quantum numbers, and the Hamiltonian is block diagonal in J . A Hilbert space of dimension $\binom{2\Omega}{N}$ collapses into blocks of dimension $2J + 1$. For $N = 4$ particles at half filling the largest block is $J = 2$, of dimension five: a 70-dimensional problem has become a 5×5 matrix.

The $J = 2$ block.

The states are labelled $|J, J_z\rangle$ with $J_z = -J, \dots, J$. The lowest, with all four spins down, is

$$|2, -2\rangle = a_{1-}^\dagger a_{2-}^\dagger a_{3-}^\dagger a_{4-}^\dagger |0\rangle, \quad (4.9)$$

and the remaining four follow by applying $\hat{J}_+$ with the standard normalisations

$$\hat{J}_\pm |J, J_z\rangle = \sqrt{J(J+1) - J_z(J_z \pm 1)} |J, J_z \pm 1\rangle. \quad (4.10)$$

Working out the matrix elements of Eq. (4.8) – again Exercise 4.6 – gives for $\varepsilon = 2$, $V = -1/3$, $W = -1/4$

$$\mathbf{H}_{J=2}^{(1)} = \begin{bmatrix} -4 & 0 & -\sqrt{6}/3 & 0 & 0 \\ 0 & -11/4 & 0 & -1 & 0 \\ -\sqrt{6}/3 & 0 & -1 & 0 & -\sqrt{6}/3 \\ 0 & -1 & 0 & 5/4 & 0 \\ 0 & 0 & -\sqrt{6}/3 & 0 & 4 \end{bmatrix}, \quad (4.11)$$

whose eigenvalues are -4.21288 , -2.98607 , -0.91914 , 1.48607 and 4.13201 . The ground state is

$$|\psi_0\rangle = 0.96735|2, -2\rangle + 0.25221|2, 0\rangle + 0.02507|2, 2\rangle, \quad E_0 = -4.21288. \quad (4.12)$$

Note the structure: the interaction connects only states differing by two units of J_z , because $\hat{J}_\pm^2$ moves a pair. The matrix therefore splits further into even and odd J_z .

Increasing the interaction to $V = -4/3$, $W = -1$ gives eigenvalues -7.75122 , -7.47214 , -1.55581 , 1.47214 , 5.30704 and the ground state

$$|\psi_0\rangle = 0.64268|2, -2\rangle + 0.73816|2, 0\rangle + 0.20515|2, 2\rangle, \quad E_0 = -7.75122. \quad (4.13)$$

In the first case the unperturbed configuration $|2, -2\rangle$ carries 94% of the probability and a single Slater determinant is a good approximation. In the second the interaction matrix elements are comparable to the single-particle spacing, $|2, 0\rangle$ has become the dominant component, and no single determinant will do. This is the same story the occupation numbers of Section 1.31 told, in a model small enough to see the whole of it.

Why this model matters. The Lipkin model was invented to test many-body methods, and it still is. Because the exact answer is a 5×5 eigenvalue problem while the interaction can be tuned from weak to strong, it exposes exactly where an approximation breaks down: Hartree-Fock theory undergoes a spurious phase transition at a critical coupling, the random-phase approximation predicts an excitation energy that goes imaginary there, and coupled-cluster theory tracks the exact result much further. Every one of these statements can be checked against Eq. (4.12) and its analogues.

Mapping to qubits.

The Lipkin model is a natural first target for quantum algorithms, and how one maps it onto qubits is worth a moment. There are three obvious schemes.

The *occupation mapping* assigns one qubit to each single-particle state (p, σ) , the qubit being $|1\rangle$ or $|0\rangle$ according as that state is occupied or empty. It costs 2Ω qubits.

The *level mapping* exploits the fact that there are only two levels. Restricting to half filling, $N = \Omega$, each doublet $((p, +1), (p, -1))$ holds exactly one particle, so one qubit per doublet suffices: $|0\rangle$ for the particle in the upper state and $|1\rangle$ for the lower. This costs Ω qubits, half of the occupation mapping. It has a second advantage: moving a pair changes the state of two doublets, so any operator in the Hamiltonian is at most a two-qubit gate.

The *spin mapping* goes further and assigns a qubit to each state $|J, J_z\rangle$. For a given J there are only $2J + 1$ of these, so with $J_{\max} = N/2$ the largest register is $N + 1$ qubits. The price is that one must run a separate circuit for each J and take the lowest of the resulting energies.

Under the level mapping the quasispin operators become sums of Pauli matrices. Writing $\hat{J}_z = \sum_p j_z^{(p)}$ and $\hat{J}_\pm = \sum_p j_\pm^{(p)}$ with

$$j_z^{(p)} = \frac{1}{2} \sum_\sigma \sigma a_{p\sigma}^\dagger a_{p\sigma}, \quad j_\pm^{(p)} = a_{p\pm}^\dagger a_{p\mp}, \quad (4.14)$$

each doublet operator is a single-qubit operator, and the correspondence with the Pauli matrices of Section 1.7 is

$$j_z^{(p)} \longleftrightarrow \frac{1}{2} Z_p, \quad j_\pm^{(p)} \longleftrightarrow \frac{1}{2} (X_p \pm iY_p). \quad (4.15)$$

The Hamiltonian (4.8) therefore becomes a sum of Pauli strings of weight at most two – exactly the form a variational quantum eigensolver requires. The circuits themselves belong to the quantum-computing part of this book; what matters here is that the mapping costs Ω qubits and two-qubit interactions only, which is a direct consequence of the $SU(2)$ structure identified above.

4.2 The pairing model

The pairing model has appeared repeatedly already – as the test case for the Lanczos algorithm in Section 1.26, for the Schmidt decomposition in Section 1.30, and as the three-level exercise of week 35. We now set it out properly.

The single-particle space consists of L doubly degenerate, equally spaced levels labelled $p = 1, 2, \dots, L$ with spin $\sigma = \pm 1$. The Hamiltonian is

$$\hat{H} = \hat{H}_0 + \hat{V}, \quad \hat{H}_0 = \xi \sum_{p\sigma} (p-1) a_{p\sigma}^\dagger a_{p\sigma}, \quad \hat{V} = -\frac{1}{2} g \sum_{pq} a_{p+}^\dagger a_{p-}^\dagger a_{q-} a_{q+}. \quad (4.16)$$

The spacing ξ is set to unity without loss of generality. The interaction has a single term and a single strength g : it destroys a pair of particles in level q with opposite spins and creates one in level p . Nothing else happens. In the language of Eq. (4.1), only the matrix elements $\langle p+p- | \hat{V} | q+q- \rangle$ are non-zero, and they are all equal.

Symmetries.

Exercise 4.6 shows that both $\hat{H}_0$ and $\hat{V}$ commute with the spin projection and the total spin,

$$\hat{S}_z = \frac{1}{2} \sum_{p\sigma} \sigma a_{p\sigma}^\dagger a_{p\sigma}, \quad \hat{S}^2 = \hat{S}_z^2 + \frac{1}{2} (\hat{S}_+ \hat{S}_- + \hat{S}_- \hat{S}_+), \quad \hat{S}_\pm = \sum_p a_{p\pm}^\dagger a_{p\mp}, \quad (4.17)$$

so the Hamiltonian is block diagonal in S and S_z . We work in the $S = 0$ block. There it is natural to introduce the pair operators

$$\hat{P}_p^+ = a_{p+}^\dagger a_{p-}^\dagger, \quad \hat{P}_p^- = a_{p-} a_{p+}, \quad (4.18)$$

in terms of which

$$\hat{H} = \sum_{p\sigma} (p-1) a_{p\sigma}^\dagger a_{p\sigma} - \frac{1}{2} g \sum_{pq} \hat{P}_p^+ \hat{P}_q^-. \quad (4.19)$$

The pair creation operators commute among themselves, $[\hat{P}_p^+, \hat{P}_q^+] = 0$, which is why they behave like bosons and why the model describes a superfluid. They are *not* bosons, however: $(\hat{P}_p^+)^2 = 0$, the Pauli principle asserting itself.

The seniority-zero space.

Because the interaction only moves pairs, it never breaks one. If we start from a state in which every occupied level carries a full pair, we stay in that space forever. A basis state is then simply a choice of which n of the L levels carry a pair, and the dimension is $\binom{L}{n}$ rather than $\binom{2L}{2n}$. For $L = 4$ and $n = 2$ this is 6 instead of 28; for $L = 12$ and $n = 6$ it is 924 instead of 2704156. This is the space in which Table 1.3 was computed.

In that basis the Hamiltonian is easy to write down: the diagonal is $2 \sum_{p \in \text{occ}} (p-1) - \frac{1}{2} gn$, and any two configurations differing by the position of a single pair are connected by $-\frac{1}{2} g$.

Table 4.1 gives the exact ground-state energy for four particles in four levels as a function of g , together with the energy of the unperturbed configuration. At $g = 0$ the reference determinant is exact. For small $|g|$ the correlation energy grows quadratically, as second-order perturbation theory requires, and for $g = 1$ it has reached 36% of the reference energy.

Why this model matters. The pairing model is the schematic version of superfluidity and superconductivity, and it is the standard proving ground for the hierarchy of many-body methods. Because the

g	E_{ref}	E_0 (exact)	E_{corr}
−1.00	3.000000	2.77987014	−0.22012986
−0.50	2.500000	2.43688426	−0.06311574
0.00	2.000000	2.00000000	0.00000000
0.50	1.500000	1.41677428	−0.08322572
1.00	1.000000	0.63554847	−0.36445153

Table 4.1 The pairing model with four doubly degenerate levels and two pairs, $\xi = 1$. E_{ref} is the energy of the configuration with both pairs in the two lowest levels; E_0 is the exact ground-state energy in the six-dimensional seniority-zero space.

exact answer is available for any g , one can plot Hartree-Fock, second- and third-order perturbation theory, coupled cluster with doubles, and full configuration interaction on the same axes and watch each one fail at a different coupling. The second midterm of FYS4480 does exactly this, and the exercises at the end of this chapter set up the ingredients.

4.3 The Fermi-Hubbard model

The two models so far live in an abstract single-particle space with no geometry. The Fermi-Hubbard model introduces a lattice, and with it the competition that dominates the physics of correlated electrons: kinetic energy wants to delocalise the particles, the interaction wants to keep them apart.

Electrons hop between neighbouring sites of a lattice and pay an energy U whenever two of them, necessarily of opposite spin, occupy the same site,

$$\hat{H} = - \sum_{\langle ij \rangle \sigma} t_{ij} \left(c_{i\sigma}^\dagger c_{j\sigma} + \text{h.c.} \right) + U \sum_i \hat{n}_{i\uparrow} \hat{n}_{i\downarrow} \quad (4.20)$$

with $\hat{n}_{i\sigma} = c_{i\sigma}^\dagger c_{i\sigma}$. This is again Eq. (4.1): the one-body part is the hopping matrix, and the only non-vanishing two-body matrix element is the on-site one.

The single-particle basis is now the set of lattice sites, so the creation operators carry a position label. Fermionic antisymmetry is enforced by the Jordan-Wigner string: $c_i^\dagger c_j$ acting on an occupation-number state picks up the parity of the occupied sites strictly between i and j , which is exactly the phase computed by the bit-counting algorithm at the end of Chapter 3.

Symmetries and sector dimensions.

The Hamiltonian conserves $N_\uparrow$ and $N_\downarrow$ separately, so we may diagonalise in a fixed sector of dimension

$$\dim = \binom{N}{N_\uparrow} \binom{N}{N_\downarrow}. \quad (4.21)$$

Better still, the hopping term acts on the two spin species independently, so the sector factorises as a tensor product and the kinetic part of the Hamiltonian is

$$\mathbf{K}_\uparrow \otimes \mathbf{I}_\downarrow + \mathbf{I}_\uparrow \otimes \mathbf{K}_\downarrow, \quad (4.22)$$

the tensor-product construction of Section 1.10 used in earnest. The interaction is diagonal in the occupation basis. On a periodic chain translational invariance gives a further reduction into momentum sectors.

From metal to Mott insulator.

Table 4.2 follows the half-filled four-site ring as U grows. At $U = 0$ the model is exact free fermions: the single-particle energies are $\varepsilon_k = -2t \cos k$ with $k = 0, \pm\pi/2, \pi$, and filling the two lowest levels for each spin gives $E_0 = -4t$, which the diagonalisation reproduces. As U increases, the double occupancy falls steadily – the electrons avoid each other – and the ground-state energy rises towards zero.

U/t	E_0/t	$E_0/(Nt)$	$\sum_i \langle n_{i\uparrow} n_{i\downarrow} \rangle$
0	-4.00000000	-1.000000	1.000000
2	-2.82842712	-0.707107	0.453082
4	-2.10274848	-0.525687	0.287325
8	-1.32023496	-0.330059	0.129992
16	-0.72286432	-0.180716	0.042033

Table 4.2 The half-filled four-site Hubbard ring, $N_\uparrow = N_\downarrow = 2$, sector dimension 36. The last column is the total double occupancy in the ground state.

The strong-coupling limit.

When $U \gg t$ the doubly occupied states are pushed far up in energy and can be eliminated perturbatively. What remains is a model of localised spins,

$$\hat{H}_{tJ} = \hat{P}_G \left[- \sum_{\langle ij \rangle \sigma} t_{ij} \left(c_{i\sigma}^\dagger c_{j\sigma} + \text{h.c.} \right) \right] \hat{P}_G + \sum_{\langle ij \rangle} J_{ij} \left[\mathbf{S}_i \cdot \mathbf{S}_j - \frac{1}{4} \hat{P}_G \hat{n}_i \hat{n}_j \hat{P}_G \right], \quad J_{ij} = \frac{4t_{ij}^2}{U}, \quad (4.23)$$

where $\hat{P}_G$ is the Gutzwiller projector removing all doubly occupied sites. At half filling there is no empty site to hop into, the kinetic term is frozen, and the t - J model reduces to the Heisenberg antiferromagnet.

This is a prediction we can check. For the four-site ring the Heisenberg part contributes $-2J$ and the constant $-\frac{1}{4}n_i n_j$ term on four bonds contributes $-J$, so $E_0 \rightarrow -3J = -12t^2/U$. Table 4.3 shows the ratio E_0/J converging to -3 .

U/t	E_0/t	$E_0/J, \quad J = 4t^2/U$
8	-1.3202349583	-2.64047
16	-0.7228643224	-2.89146
32	-0.3714104223	-2.97128
64	-0.1870445461	-2.99271
128	-0.0936928521	-2.99817

Table 4.3 The half-filled four-site Hubbard ring at strong coupling. The ratio converges to -3 , confirming the mapping onto the Heisenberg model with $J = 4t^2/U$.

Why this model matters. The Hubbard model is the minimal model of strong correlation. It is believed to contain the essential physics of the cuprate superconductors, it is the standard benchmark for quantum Monte Carlo, tensor-network and machine-learning wave functions, and it is among the first

targets for quantum simulation with cold atoms and with quantum computers. It is also the model in which the exponential wall of Section 2.16 bites hardest: the sector dimension (4.21) for a 4×4 lattice at half filling is already 1.66×10^8 , and the two-dimensional model remains unsolved.

4.4 The Calogero-Sutherland model

The three models above are solvable because a symmetry cuts the Hilbert space down to a manageable size. The Calogero-Sutherland model [4,5] is solvable for a different and much rarer reason: it is *integrable*, and its ground state is known in closed form for any number of particles.

The Calogero model describes N particles on a line in a harmonic trap, interacting through an inverse-square potential,

$$\hat{H} = -\frac{\hbar^2}{2m} \sum_{i=1}^N \frac{\partial^2}{\partial x_i^2} + \frac{m\omega^2}{2} \sum_{i=1}^N x_i^2 + \sum_{i<j} \frac{g}{(x_i - x_j)^2}. \quad (4.24)$$

It is conventional to parametrise the coupling as

$$g = \frac{\hbar^2}{m} \lambda(\lambda - 1), \quad (4.25)$$

so that λ measures the interaction strength. The Sutherland model is the same system on a ring of circumference L , with the inverse square replaced by the chord distance,

$$\hat{H} = -\frac{\hbar^2}{2m} \sum_i \frac{\partial^2}{\partial x_i^2} + \frac{\hbar^2}{m} \left(\frac{\pi}{L}\right)^2 \sum_{i<j} \frac{\lambda(\lambda - 1)}{\sin^2(\pi(x_i - x_j)/L)}. \quad (4.26)$$

The exact ground state.

In units $\hbar = m = 1$ the ground state of Eq. (4.24) is

$$\Psi_0(x_1, \dots, x_N) = \prod_{i<j} |x_i - x_j|^\lambda \exp\left(-\frac{\omega}{2} \sum_i x_i^2\right) \quad (4.27)$$

with energy

$$E_0 = \omega \left[\frac{N}{2} + \frac{\lambda N(N-1)}{2} \right]. \quad (4.28)$$

Exercise 4.6 verifies this by direct substitution; the inverse-square form of the potential is precisely what is needed for the cross terms generated by the Laplacian to cancel.

The factor $\prod_{i<j} |x_i - x_j|^\lambda$ is a *Jastrow factor*, and it is the reason this model deserves a place in a book about many-body methods. It is a correlated wave function of exactly the type that variational Monte Carlo optimises numerically – but here it is exact, for every N and every λ . The model is therefore the sharpest available test of a Jastrow ansatz: any method that cannot reproduce Eq. (4.27) is missing something structural.

Two limits are instructive. At $\lambda = 0$ the coupling g vanishes and the particles are free bosons in a trap, $E_0 = \omega N/2$. At $\lambda = 1$ the coupling vanishes again, but now

$$\prod_{i<j} |x_i - x_j| \quad (4.29)$$

is the absolute value of the Vandermonde determinant of Section 1.20, which is precisely the Slater determinant of N free fermions in a harmonic trap. The Calogero model interpolates continuously between bosons and fermions, and for non-integer λ it describes particles obeying *fractional statistics*.

A numerical check.

For $N = 2$ the centre-of-mass and relative motions separate. With $x = x_1 - x_2$ the relative Hamiltonian is

$$\hat{h}_{\text{rel}} = -\frac{d^2}{dx^2} + \frac{\omega^2}{4}x^2 + \frac{\lambda(\lambda-1)}{x^2}, \quad (4.30)$$

whose exact ground-state energy is $\omega(\lambda + \frac{1}{2})$. Discretising on a grid and diagonalising with the machinery of Section 1.27 reproduces this to better than 10^{-5} for λ between 1 and 3, as `models.py` confirms. The inverse-square barrier keeps the wave function away from the origin, so the boundary condition $\Psi \rightarrow 0$ as $x \rightarrow 0$ is automatic for $\lambda > 0$.

Why this model matters. The Calogero-Sutherland model sits at the junction of several subjects. Its ground state is a Laughlin-type wave function, $\prod_{i<j}(z_i - z_j)^m e^{-\sum |z_i|^2/4}$, connecting it to the fractional quantum Hall effect; its excitation spectrum is that of free particles with fractional statistics; and the probability density $|\Psi_0|^2 \propto \prod_{i<j} |x_i - x_j|^{2\lambda} e^{-\omega \sum x_i^2}$ is exactly the eigenvalue distribution of a random matrix ensemble, with $\beta = 2\lambda$ the Dyson index. For our purposes it is above all a benchmark with a known answer for arbitrary N – something none of the other models in this chapter can offer.

4.5 What the four models have in common

Table 4.4 collects the four. The pattern is worth stating explicitly, because it is the pattern of the whole subject.

Model	Single-particle space	Interaction	Symmetry used	Exactly solvable
Lipkin	two levels, degeneracy Ω	pair transfer, spin exchange	$SU(2)$ quasispin	yes, block by block
Pairing	L doubly degenerate levels	pair transfer	S, S_z , seniority	yes, small L
Hubbard	lattice sites	on-site U	$N_\uparrow, N_\downarrow$, momentum	1D only
Calogero	continuum, N particles	inverse square	integrability	yes, any N

Table 4.4 The four models of this chapter. Each is the general Hamiltonian (4.1) with a drastically restricted set of two-body matrix elements, and in each case a symmetry reduces the dimension of the problem.

In every case the route is the same. Write the Hamiltonian in second quantisation; identify the operators that commute with it; use them to block diagonalise; and solve the blocks. The first step is Chapter 3, the second is a commutator calculation of the kind Wick's theorem makes routine, the third is the statement that a symmetry gives a basis in which the matrix is block diagonal (Section 1.11), and the fourth is the eigenvalue problem of Sections 1.24 and 1.26.

What separates these models from a realistic system is only the last step of the reduction. In a real nucleus or a real molecule the symmetries reduce the dimension by orders of magnitude but not to five; the matrix is sparse but enormous; and one must fall back on the approximate methods that occupy the rest of this book. The value of the models in this chapter is that they let us watch those approximations succeed and fail against an answer we already know.

4.6 Exercises

Exercise 1: the Lipkin model in the quasispin basis

- Show that the quasispin operators (4.6) obey the angular-momentum algebra (4.7), using only the anti-commutation relations of Chapter 3.
- Show that the Lipkin Hamiltonian (4.2) can be rewritten as Eq. (4.8).
- Show that $[\hat{H}, \hat{J}^2] = 0$ and $[\hat{H}, \hat{N}] = 0$, and explain what this implies for the structure of the Hamiltonian matrix.
- Construct the five states $|2, J_z\rangle$ for four particles, and compute the matrix (4.11). Diagonalise it and compare with Eq. (4.12).

Answer to a).

Insert the definitions and use $\{a_l, a_k\} = \{a_l^\dagger, a_k^\dagger\} = 0$ and $\{a_l^\dagger, a_k\} = \delta_{lk}$. For the first commutator,

$$[\hat{J}_z, \hat{J}_\pm] = \frac{1}{2} \sum_{pp'\sigma} \sigma \left(a_{p\sigma}^\dagger a_{p\sigma} a_{p'\pm}^\dagger a_{p'\mp} - a_{p'\pm}^\dagger a_{p'\mp} a_{p\sigma}^\dagger a_{p\sigma} \right). \quad (4.31)$$

Anticommute the operators of the second product into the order of the first. Using $a_{p'\mp} a_{p\sigma}^\dagger = \delta_{pp'} \delta_{\mp\sigma} - a_{p\sigma}^\dagger a_{p'\mp}$ and then $a_{p\sigma} a_{p'\pm}^\dagger = \delta_{pp'} \delta_{\pm\sigma} - a_{p'\pm}^\dagger a_{p\sigma}$, the four-operator terms cancel in pairs and what survives is

$$[\hat{J}_z, \hat{J}_\pm] = \frac{1}{2} \sum_{p\sigma} \sigma \left(\delta_{\pm\sigma} a_{p\sigma}^\dagger a_{p\mp} - \delta_{\mp\sigma} a_{p\pm}^\dagger a_{p\sigma} \right) = \pm \frac{1}{2} \sum_p \left(a_{p\pm}^\dagger a_{p\mp} + a_{p\pm}^\dagger a_{p\mp} \right) = \pm \hat{J}_\pm, \quad (4.32)$$

where the two Kronecker deltas each pick out one value of σ , and the factor $\sigma = \pm 1$ supplies the sign. The same manipulation gives $[\hat{J}_+, \hat{J}_-] = 2\hat{J}_z$, and $[\hat{J}^2, \hat{J}_\pm] = [\hat{J}^2, \hat{J}_z] = 0$ then follows from the definition $\hat{J}^2 = \hat{J}_+ \hat{J}_- + \hat{J}_z^2 - \hat{J}_z$ exactly as for ordinary angular momentum.

Answer to b).

The one-body part is immediate: comparing Eqs. (4.3) and (4.6) gives $\hat{H}_0 = \epsilon \hat{J}_z$.

For $\hat{H}_1$, first anticommute the two annihilation operators, $a_{p'-\sigma} a_{p-\sigma} = -a_{p-\sigma} a_{p'-\sigma}$, and then move $a_{p-\sigma}$ past $a_{p'\sigma}^\dagger$ using $a_{p-\sigma} a_{p'\sigma}^\dagger = \delta_{pp'} \delta_{-\sigma\sigma} - a_{p'\sigma}^\dagger a_{p-\sigma}$. Since $\delta_{-\sigma\sigma} = 0$, the delta term drops out and

$$\hat{H}_1 = \frac{1}{2} V \sum_{pp'\sigma} a_{p\sigma}^\dagger a_{p-\sigma} a_{p'\sigma}^\dagger a_{p'-\sigma}. \quad (4.33)$$

The two factors are now independent sums over p and p' . Writing out $\sigma = \pm$ separately,

$$\hat{H}_1 = \frac{1}{2} V \left[\left(\sum_p a_{p+}^\dagger a_{p-} \right) \left(\sum_{p'} a_{p'+}^\dagger a_{p'-} \right) + \left(\sum_p a_{p-}^\dagger a_{p+} \right) \left(\sum_{p'} a_{p'-}^\dagger a_{p'+} \right) \right] = \frac{1}{2} V (\hat{J}_+^2 + \hat{J}_-^2). \quad (4.34)$$

For $\hat{H}_2$ the same two steps give

$$\hat{H}_2 = \frac{1}{2} W \left(- \sum_{p\sigma} a_{p\sigma}^\dagger a_{p\sigma} + \sum_{pp'\sigma} a_{p\sigma}^\dagger a_{p-\sigma} a_{p'-\sigma}^\dagger a_{p'\sigma} \right), \quad (4.35)$$

where this time $\delta_{-\sigma-\sigma} = 1$ and the delta term survives as $-\hat{N}$. Splitting the remaining sum over σ as before identifies it as $\hat{J}_+ \hat{J}_- + \hat{J}_- \hat{J}_+$, so $\hat{H}_2 = \frac{1}{2} W (-\hat{N} + \hat{J}_+ \hat{J}_- + \hat{J}_- \hat{J}_+)$. Adding the three pieces gives Eq. (4.8).

Answer to c).

Written as in Eq. (4.8), $\hat{H}$ is built entirely from $\hat{J}_z$, $\hat{J}_\pm$ and $\hat{N}$. Since $\hat{J}^2$ commutes with each of $\hat{J}_z$ and $\hat{J}_\pm$ by part a), and with $\hat{N}$ trivially, it commutes with $\hat{H}$. For $\hat{N}$: the one-body term conserves particle number, and every term of the interaction contains two creation and two annihilation operators, so $[\hat{H}, \hat{N}] = 0$.

Both J and N are therefore good quantum numbers. In a basis $|J, J_z\rangle$ the Hamiltonian matrix is block diagonal, one block for each J , of dimension $2J + 1$. Within a block, $\hat{J}_z$ is diagonal and $\hat{J}_\pm^2$ connects J_z to $J_z \pm 2$, so the block splits once more into even and odd J_z . A problem of nominal dimension $\binom{2\Omega}{N}$ has been reduced to a set of matrices of dimension at most $N/2 + 1$.

Answer to d).

Start from $|2, -2\rangle$, Eq. (4.9), and apply Eq. (4.10). With $J = 2$ and $J_z = -2$ the factor is $\sqrt{6-2} = 2$, so

$$|2, -1\rangle = \frac{1}{2}\hat{J}_+|2, -2\rangle = \frac{1}{2}\left(a_{1+}^\dagger a_{2-}^\dagger a_{3-}^\dagger a_{4-}^\dagger + a_{1-}^\dagger a_{2+}^\dagger a_{3-}^\dagger a_{4-}^\dagger + a_{1-}^\dagger a_{2-}^\dagger a_{3+}^\dagger a_{4-}^\dagger + a_{1-}^\dagger a_{2-}^\dagger a_{3-}^\dagger a_{4+}^\dagger\right)|0\rangle, \quad (4.36)$$

one term for each substate that can be flipped. Applying $\hat{J}_+$ again with the factor $\sqrt{6-0} = \sqrt{6}$ gives $|2, 0\rangle$ as the normalised sum of the six ways of flipping two of the four substates, and so on up to $|2, 2\rangle$ with all four spins up.

The matrix elements follow from Eq. (4.8). The diagonal is

$$\langle 2, J_z | \hat{H} | 2, J_z \rangle = \epsilon J_z + \frac{W}{2} \left(-N + 2[J(J+1) - J_z^2] \right), \quad (4.37)$$

since $\hat{J}_+ \hat{J}_- + \hat{J}_- \hat{J}_+ = 2(\hat{J}^2 - \hat{J}_z^2)$. The off-diagonal elements come from $\frac{1}{2}V(\hat{J}_+^2 + \hat{J}_-^2)$,

$$\langle 2, J_z + 2 | \hat{H} | 2, J_z \rangle = \frac{V}{2} \sqrt{J(J+1) - J_z(J_z+1)} \sqrt{J(J+1) - (J_z+1)(J_z+2)}. \quad (4.38)$$

With $\epsilon = 2$, $V = -1/3$, $W = -1/4$ and $N = 4$ this gives Eq. (4.11); for instance the $(-2, 0)$ element is $\frac{1}{2}V\sqrt{4}\sqrt{6} = -\sqrt{6}/3$, and the $J_z = -1$ diagonal element is $-2 + \frac{1}{2}(-\frac{1}{4})(-4 + 2[6-1]) = -2 - \frac{3}{4} = -\frac{11}{4}$. Diagonalising reproduces Eq. (4.12).

Exercise 2: the pairing model

- Show that $\hat{H}_0$ and $\hat{V}$ of Eq. (4.16) commute with $\hat{S}_z$ and $\hat{S}^2$ of Eq. (4.17).
- Show that the Hamiltonian can be written in the pair form (4.19), and that the pair creation operators commute among themselves while $(\hat{P}_p^+)^2 = 0$.
- Construct the Hamiltonian matrix for four particles in the four lowest levels, with no broken pairs and $S = 0$, and find its eigenvalues as a function of g .

Answer to a).

$\hat{H}_0$ is a sum of number operators $a_{p\sigma}^\dagger a_{p\sigma}$, and $\hat{S}_z$ is a weighted sum of the same operators, so the two commute term by term. For $\hat{V}$, note that $a_{p+}^\dagger a_{p-}^\dagger$ creates one particle with $\sigma = +1$ and one with $\sigma = -1$, changing S_z by $\frac{1}{2}(+1) + \frac{1}{2}(-1) = 0$; likewise $a_{q-} a_{q+}$ destroys such a pair. Hence $[\hat{V}, \hat{S}_z] = 0$.

For $\hat{S}^2$ it is enough to show $[\hat{V}, \hat{S}_\pm] = 0$. The operator $\hat{S}_+ = \sum_p a_{p+}^\dagger a_{p-}$ acting on a fully paired level gives zero, because the state $|p+\rangle$ is already occupied and $(a_{p+}^\dagger)^2 = 0$; acting on an empty level it also gives zero. Since $\hat{V}$ maps the fully paired space into itself, $\hat{S}_+ \hat{V}$ and $\hat{V} \hat{S}_+$ both annihilate it, and the commutator

vanishes. Both S and S_z are therefore good quantum numbers and the Hamiltonian is block diagonal in them.

Answer to b).

Insert the definitions (4.18) into Eq. (4.16):

$$-\frac{1}{2}g \sum_{pq} a_{p+}^\dagger a_{p-}^\dagger a_{q-} a_{q+} = -\frac{1}{2}g \sum_{pq} \hat{P}_p^+ \hat{P}_q^-, \quad (4.39)$$

which is Eq. (4.19).

For the commutator, $\hat{P}_p^+ \hat{P}_q^+ = a_{p+}^\dagger a_{p-}^\dagger a_{q+}^\dagger a_{q-}^\dagger$. Moving $a_{q+}^\dagger a_{q-}^\dagger$ to the front requires four interchanges of neighbouring creation operators, each contributing a minus sign, so the total sign is $(-1)^4 = +1$ and $[\hat{P}_p^+, \hat{P}_q^+] = 0$. The pair operators therefore behave like bosons under exchange.

They are not bosons, however. Since $(a_{p+}^\dagger)^2 = 0$ by the Pauli principle,

$$(\hat{P}_p^+)^2 = a_{p+}^\dagger a_{p-}^\dagger a_{p+}^\dagger a_{p-}^\dagger = -\left(a_{p+}^\dagger\right)^2 \left(a_{p-}^\dagger\right)^2 = 0. \quad (4.40)$$

A level can hold one pair but not two. This is the hard-core constraint that distinguishes the pairing model from a gas of free bosons, and it is what makes the model non-trivial.

Answer to c).

With four levels and two pairs there are $\binom{4}{2} = 6$ basis states, which we label by the occupied levels: (12), (13), (14), (23), (24), (34). The diagonal element of a configuration is $2 \sum_{p \in \text{occ}} (p-1) - g$, the last term coming from the $p = q$ terms of the interaction with two pairs present. Every pair of configurations differing by the position of one pair is connected by $-g/2$; configurations differing in both pairs are not connected. Ordering the basis as above,

$$\mathbf{H} = \begin{bmatrix} 2-g & -g/2 & -g/2 & -g/2 & -g/2 & 0 \\ -g/2 & 4-g & -g/2 & -g/2 & 0 & -g/2 \\ -g/2 & -g/2 & 6-g & 0 & -g/2 & -g/2 \\ -g/2 & -g/2 & 0 & 6-g & -g/2 & -g/2 \\ -g/2 & 0 & -g/2 & -g/2 & 8-g & -g/2 \\ 0 & -g/2 & -g/2 & -g/2 & -g/2 & 10-g \end{bmatrix}. \quad (4.41)$$

Diagonalising for a range of g gives Table 4.1. Note that the (13) and (23) configurations are degenerate at $g = 0$ – both have unperturbed energy 6 – and that the correlation energy is negative for both signs of g , since the interaction can always lower the energy by mixing configurations.

Exercise 3: the Hubbard model

- Show that the Hubbard Hamiltonian (4.20) conserves $N_\uparrow$ and $N_\downarrow$ separately, and hence S_z .
- Diagonalise the $U = 0$ Hubbard ring analytically and verify the entry $E_0 = -4t$ of Table 4.2.
- Derive the effective spin Hamiltonian in the limit $U \gg t$ at half filling, and predict the ratio E_0/J for the four-site ring.

Answer to a).

The hopping term contains $c_{i\sigma}^\dagger c_{j\sigma}$ with the *same* σ on both operators, so it conserves the number of particles of each spin separately. The interaction is a product of number operators and conserves everything. Hence $[\hat{H}, \hat{N}_\uparrow] = [\hat{H}, \hat{N}_\downarrow] = 0$, and since $\hat{S}_z = \frac{1}{2}(\hat{N}_\uparrow - \hat{N}_\downarrow)$, also $[\hat{H}, \hat{S}_z] = 0$. This is what allows the diagonalisation to be carried out in a fixed sector of dimension (4.21).

Answer to b).

At $U = 0$ the Hamiltonian is a sum of independent one-body problems. On a ring of N sites with periodic boundary conditions the hopping matrix is diagonalised by plane waves $c_{k\sigma}^\dagger = N^{-1/2} \sum_j e^{ikj} c_{j\sigma}^\dagger$ with $k = 2\pi n/N$, giving

$$\epsilon_k = -2t \cos k. \quad (4.42)$$

For $N = 4$ the allowed momenta are $k = 0, \pm\pi/2, \pi$ with energies $-2t, 0, 0, +2t$. Filling the two lowest for each spin gives $-2t + 0 = -2t$ per spin species, hence $E_0 = -4t$, which is the first line of Table 4.2. The degeneracy at $k = \pm\pi/2$ means the $U = 0$ ground state is not unique; switching on U lifts the degeneracy.

Answer to c).

At half filling and $U \gg t$ every site carries exactly one electron in the ground state, and states with a doubly occupied site lie an energy U higher. Treat the hopping as a perturbation. A single hop creates a doubly occupied site and an empty one, costing U ; hopping back restores a singly occupied configuration. Second-order perturbation theory therefore gives an energy gain $-2t^2/U$ for each such virtual excursion, and the process is possible only if the two spins involved are antiparallel, since the Pauli principle forbids two electrons of the same spin on one site.

Collecting the terms gives the Heisenberg Hamiltonian

$$\hat{H}_{\text{eff}} = J \sum_{\langle ij \rangle} \left(\mathbf{s}_i \cdot \mathbf{s}_j - \frac{\hat{n}_i \hat{n}_j}{4} \right), \quad J = \frac{4t^2}{U}, \quad (4.43)$$

which is the half-filled limit of the t - J model, Eq. (4.23). For the four-site ring the Heisenberg ground-state energy is $-2J$, and the constant $-\frac{1}{4}n_i n_j$ contributes $-\frac{1}{4}$ per bond on four bonds, that is $-J$. Hence

$$E_0 \longrightarrow -3J = -\frac{12t^2}{U}, \quad \text{so} \quad \frac{E_0}{J} \longrightarrow -3, \quad (4.44)$$

which is what Table 4.3 confirms numerically.

Exercise 4: the Calogero ground state

- Verify by direct substitution that Eq. (4.27) is an eigenstate of Eq. (4.24) with energy (4.28).
- Show that at $\lambda = 1$ the ground state is the Slater determinant of N free fermions in a harmonic trap.
- Separate the $N = 2$ problem and verify the relative energy $\omega(\lambda + \frac{1}{2})$.

Answer to a).

Work in units $\hbar = m = 1$ and write $\Psi_0 = e^{-\Phi}$ with

$$\Phi = \frac{\omega}{2} \sum_i x_i^2 - \lambda \sum_{i < j} \ln |x_i - x_j|. \quad (4.45)$$

For any such form,

$$-\frac{1}{2} \sum_i \frac{\partial^2 \Psi_0}{\partial x_i^2} = \frac{1}{2} \left[\sum_i \frac{\partial^2 \Phi}{\partial x_i^2} - \sum_i \left(\frac{\partial \Phi}{\partial x_i} \right)^2 \right] \Psi_0. \quad (4.46)$$

The derivatives are

$$\frac{\partial \Phi}{\partial x_i} = \omega x_i - \lambda \sum_{j \neq i} \frac{1}{x_i - x_j}, \quad \frac{\partial^2 \Phi}{\partial x_i^2} = \omega + \lambda \sum_{j \neq i} \frac{1}{(x_i - x_j)^2}. \quad (4.47)$$

Squaring the first,

$$\sum_i \left(\frac{\partial \Phi}{\partial x_i} \right)^2 = \omega^2 \sum_i x_i^2 - 2\lambda \omega \sum_i \sum_{j \neq i} \frac{x_i}{x_i - x_j} + \lambda^2 \sum_i \sum_{j \neq i} \sum_{k \neq i} \frac{1}{(x_i - x_j)(x_i - x_k)}. \quad (4.48)$$

Two simplifications now do all the work. In the second term, pairing (i, j) with (j, i) gives $x_i/(x_i - x_j) + x_j/(x_j - x_i) = 1$, so the double sum equals $N(N-1)/2$ times two, that is $-2\lambda \omega \cdot N(N-1)/2$. In the third term the pieces with $j = k$ give $\lambda^2 \sum_{i \neq j} (x_i - x_j)^{-2}$, while the pieces with $j \neq k$ cancel identically by the three-term identity

$$\frac{1}{(x_i - x_j)(x_i - x_k)} + \frac{1}{(x_j - x_i)(x_j - x_k)} + \frac{1}{(x_k - x_i)(x_k - x_j)} = 0. \quad (4.49)$$

Assembling everything, the inverse-square terms combine as

$$\frac{1}{2} \left[\lambda \sum_{i \neq j} \frac{1}{(x_i - x_j)^2} - \lambda^2 \sum_{i \neq j} \frac{1}{(x_i - x_j)^2} \right] = -\lambda(\lambda - 1) \sum_{i < j} \frac{1}{(x_i - x_j)^2}, \quad (4.50)$$

which cancels the interaction term of the Hamiltonian exactly; the $\omega^2 x_i^2$ terms cancel the trap; and what is left is the constant

$$E_0 = \frac{1}{2} [N\omega + \lambda \omega N(N-1)] = \omega \left[\frac{N}{2} + \frac{\lambda N(N-1)}{2} \right], \quad (4.51)$$

which is Eq. (4.28). The cancellation in Eq. (4.50) works only because the interaction falls off as the inverse *square*; this is the sense in which the model is fine-tuned to be solvable.

Answer to b).

At $\lambda = 1$ the coupling $g = \lambda(\lambda - 1)$ vanishes and the particles are free, but the wave function

$$\Psi_0 = \prod_{i < j} |x_i - x_j| e^{-\omega \sum_i x_i^2 / 2} \quad (4.52)$$

is not the bosonic ground state. Up to an overall sign it is the Vandermonde determinant of Section 1.20,

$$\prod_{i < j} (x_j - x_i) = \det \left[x_i^{j-1} \right], \quad (4.53)$$

and multiplying each column by the Gaussian and taking suitable linear combinations turns $x_i^{j-1} e^{-\omega x_i^2 / 2}$ into the Hermite functions of Section 1.27. The result is precisely the Slater determinant of the N lowest harmonic oscillator orbitals. The energy check is immediate: Eq. (4.28) with $\lambda = 1$ gives $E_0 = \omega N^2 / 2 = \omega \sum_{n=0}^{N-1} (n + \frac{1}{2})$, the sum of the N lowest oscillator energies.

For non-integer λ the wave function is neither symmetric nor antisymmetric, and the particles obey fractional statistics interpolating between the two.

Answer to c).

Introduce $X = (x_1 + x_2)/2$ and $x = x_1 - x_2$. The kinetic energy separates as $-\frac{1}{4}\partial_X^2 - \partial_x^2$ and the trap as $\omega^2 X^2 + \frac{1}{4}\omega^2 x^2$, so with $\hbar = m = 1$

$$\hat{H} = \left[-\frac{1}{4} \frac{\partial^2}{\partial X^2} + \omega^2 X^2 \right] + \left[-\frac{\partial^2}{\partial x^2} + \frac{\omega^2}{4} x^2 + \frac{\lambda(\lambda-1)}{x^2} \right]. \quad (4.54)$$

The centre-of-mass part is a harmonic oscillator with ground-state energy $\omega/2$. The relative part is Eq. (4.30); its ground state is $x^\lambda e^{-\omega x^2/4}$ with energy $\omega(\lambda + \frac{1}{2})$, as substitution confirms. The total is $\omega/2 + \omega(\lambda + \frac{1}{2}) = \omega(1 + \lambda)$, which agrees with Eq. (4.28) for $N = 2$. Solving the relative problem numerically on a grid, as in Section 1.27, reproduces $\omega(\lambda + \frac{1}{2})$ to better than 10^{-5} ; see `models.py`.

Exercise 5: numerical explorations

- Using `models.py`, plot the Lipkin ground-state energy and the weight of $|2, -2\rangle$ in the ground state as functions of V at fixed ε and W . At roughly what value of V/ε does the single-configuration picture break down?
- Repeat the pairing calculation of Table 4.1 for six and eight levels. How does the correlation energy at $g = 1$ scale with the number of levels?
- Compute the double occupancy of the Hubbard ring as a function of U/t and identify the crossover. Compare a six-site ring with the four-site one.
- For the pairing model with twelve levels and six pairs, compare the exact ground-state energy from full diagonalisation with the Lanczos result of Table 1.3, and with the Schmidt truncation of Table 1.7. How many Schmidt states are needed for the energy to be correct to six digits?

Part II

Wave function based methods

Chapter 5

Full configuration interaction theory

We will start our discussions of many-body methods with what is called full configuration interaction theory, or just FCI. The reason for this is simple. If the effective degrees of freedom which are relevant for describing a physical system give a sufficiently good reproduction of the static properties (and later we will add dynamical properties) of a many-body system, FCI gives us all relevant eigenpairs (eigenvalues and eigenstates) and allows us to compute other expectation values than the energy levels.

Although it would have been natural to start with mean-field methods (FCI is often labelled as a post Hartree-Fock method), we have chosen a slightly different route. The reason for this is that FCI, within a restricted effective space, provides us with the exact eigenvalues and can thus serve as a benchmark for the approximative methods discussed in the coming chapters, from mean field methods to the Tamm-Dancoff (TDA) and random-phase (RPA) approximations, variants of many-body perturbation theory and coupled-cluster theory, Greens' function approaches and other. We will use FCI to illustrate how these other methods can be interpreted in terms of specific truncations/approximations of FCI.

We start thus with, what this authour likes to claim as the simplest possible approach to many-body problems, an approach which involves us being able to set up a many-body Hamiltonian matrix (within an effective and reduced Hilbert space), and then simply diagonalize. The overarching principles and the underlying mathematics are straightforward, although the actual implementations may easily lead to endless frustrations and sleepless nights. Let us get started.

5.1 Slater determinants as a basis

The simplest many-body wave function is a product,

$$\Psi(x_1, x_2, \dots, x_N) \approx \phi_1(x_1)\phi_2(x_2) \cdots \phi_N(x_N), \quad (5.1)$$

because we are really only good at thinking about one particle at a time. Product functions are easy to work with: if the single-particle states are orthonormal the products are too, and matrix elements factorise into one-dimensional integrals. The price is the complete absence of correlations, which must then be built up by superposing many, many products. This is the trade-off that runs through the whole subject – a compact representation with difficult integrals, or easy integrals and very many states.

Antisymmetry turns the product into a determinant, and Chapter 2 showed that the two elementary properties of determinants – a sign change under the interchange of two columns, and vanishing when two rows coincide – are exactly the Pauli principle:

- no two particles at the same place (two equal columns);
- no two particles in the same state (two equal rows).

Beyond $N = 4$ the coordinate-space determinant becomes unwieldy, which is why Chapter 3 replaced it by the occupation representation. There a Slater determinant is a string of creation operators acting on the

vacuum, each index appearing at most once because $a_i^\dagger a_i^\dagger = 0$, and antisymmetry is carried by the algebra rather than by the notation.

We take that formalism as given. Writing the reference determinant as

$$|\Phi_0\rangle = \left(\prod_{i \leq F} \hat{a}_i^\dagger \right) |0\rangle, \quad (5.2)$$

the excited determinants are

$$|\Phi_i^a\rangle = \hat{a}_a^\dagger \hat{a}_i |\Phi_0\rangle, \quad |\Phi_{ij}^{ab}\rangle = \hat{a}_a^\dagger \hat{a}_b^\dagger \hat{a}_j \hat{a}_i |\Phi_0\rangle, \quad (5.3)$$

and in general

$$|\Phi_{ijk\dots}^{abc\dots}\rangle = \hat{a}_a^\dagger \hat{a}_b^\dagger \hat{a}_c^\dagger \dots \hat{a}_k \hat{a}_j \hat{a}_i |\Phi_0\rangle, \quad (5.4)$$

with the hole indices i, j, k below the Fermi level and the particle indices a, b, c above it, following the convention of Section 2.6. We call such a determinant an np - nh excitation, or a *configuration*.

5.2 The configuration-interaction expansion

The exact ground state is an arbitrary superposition of all of these,

$$|\Psi_0\rangle = C_0 |\Phi_0\rangle + \sum_{ai} C_i^a |\Phi_i^a\rangle + \sum_{abij} C_{ij}^{ab} |\Phi_{ij}^{ab}\rangle + \dots = (C_0 + \hat{C}) |\Phi_0\rangle, \quad (5.5)$$

where we have introduced the *correlation operator*

$$\hat{C} = \sum_{ai} C_i^a \hat{a}_a^\dagger \hat{a}_i + \sum_{abij} C_{ij}^{ab} \hat{a}_a^\dagger \hat{a}_b^\dagger \hat{a}_j \hat{a}_i + \dots \quad (5.6)$$

The normalisation of $|\Psi_0\rangle$ is at our disposal, and C_0 is by hypothesis non-zero, so we may set $C_0 = 1$ with proportional changes in all the other coefficients. With this *intermediate normalisation*,

$$\langle \Psi_0 | \Phi_0 \rangle = \langle \Phi_0 | \Phi_0 \rangle = 1, \quad |\Psi_0\rangle = (1 + \hat{C}) |\Phi_0\rangle. \quad (5.7)$$

It is convenient to compress the notation. Writing H for a set of $0, 1, \dots, n$ hole indices and P for a set of $0, 1, \dots, n$ particle indices,

$$|\Psi_0\rangle = \sum_{PH} C_H^P |\Phi_H^P\rangle = \left(\sum_{PH} C_H^P \hat{A}_H^P \right) |\Phi_0\rangle, \quad (5.8)$$

with $\hat{A}_H^P$ the operator that creates the corresponding excitation. Unit normalisation reads $\sum_{PH} |C_H^P|^2 = 1$, and the energy is

$$E = \langle \Psi_0 | \hat{H} | \Psi_0 \rangle = \sum_{PP'HH'} C_H^{*P} \langle \Phi_H^P | \hat{H} | \Phi_{H'}^{P'} \rangle C_{H'}^{P'}. \quad (5.9)$$

If the expansion runs over *all* determinants that can be built from the chosen single-particle basis, this is *full configuration interaction*. The only approximation left is the truncation of the single-particle space itself.

5.3 The eigenvalue problem

Equation (5.9) is solved by diagonalisation, and it is worth seeing why the diagonalisation is not merely convenient but is the variational answer. We minimise the energy subject to normalisation,

$$\delta [\langle \Psi_0 | \hat{H} | \Psi_0 \rangle - \lambda \langle \Psi_0 | \Psi_0 \rangle] = 0, \quad (5.10)$$

with λ a Lagrange multiplier. Carrying out the variation,

$$\sum_{P'H'} \left\{ \delta[C_H^{*P}] \langle \Phi_H^P | \hat{H} | \Phi_{H'}^{P'} \rangle C_{H'}^{P'} + C_H^{*P} \langle \Phi_H^P | \hat{H} | \Phi_{H'}^{P'} \rangle \delta[C_{H'}^{P'}] - \lambda (\delta[C_H^{*P}] C_{H'}^{P'} + C_H^{*P} \delta[C_{H'}^{P'}]) \right\} = 0. \quad (5.11)$$

Since $\delta[C_H^{*P}]$ and $\delta[C_{H'}^{P'}]$ are complex conjugates, it is necessary and sufficient that the quantities multiplying $\delta[C_H^{*P}]$ vanish, giving

$$\sum_{P'H'} \langle \Phi_H^P | \hat{H} | \Phi_{H'}^{P'} \rangle C_{H'}^{P'} = \lambda C_H^P \quad (5.12)$$

for every set PH . Multiplying by C_H^{*P} and summing identifies $\lambda = E$, so the multiplier is the energy and

$$\boxed{\sum_{P'H'} \langle \Phi_H^P | \hat{H} - E | \Phi_{H'}^{P'} \rangle C_{H'}^{P'} = 0} \quad (5.13)$$

This is the full CI equation. The same result follows more directly by writing $(\hat{H} - E)|\Psi_0\rangle = 0$ and projecting successively onto every $\langle \Phi_H^P |$.

Equation (5.13) is a Hermitian eigenvalue problem, and the diagonalisation is a Rayleigh-Ritz variational calculation in the space spanned by the chosen determinants: the lowest eigenvalue is the best upper bound to E_0 obtainable in that space, and it can only improve as the space grows. This is exactly the property that made the Lanczos algorithm of Section 1.26 so well suited to the problem – the lowest Ritz value there is the same variational bound, computed in a Krylov space instead of a space of determinants.

Everything now depends on being able to build and diagonalise the matrix $\langle \Phi_H^P | \hat{H} | \Phi_{H'}^{P'} \rangle$. The matrix elements are supplied by the Condon-Slater rules of Section 3.13, and the diagonalisation by the algorithms of Chapter 1: Householder and QL for small matrices, Lanczos for large ones.

5.4 The structure of the Hamiltonian matrix

Suppose we have six fermions below the Fermi level, so that at most $6p$ - $6h$ excitations are possible. Ordering the determinants by excitation level, the Hamiltonian matrix has the block structure

	0p0h	1p1h	2p2h	3p3h	4p4h	5p5h	6p6h
0p0h	×	×	×	0	0	0	0
1p1h	×	×	×	×	0	0	0
2p2h	×	×	×	×	×	0	0
3p3h	0	×	×	×	×	×	0
4p4h	0	0	×	×	×	×	×
5p5h	0	0	0	×	×	×	×
6p6h	0	0	0	0	×	×	×

(5.14)

The zeros are not accidental. A two-body operator connects Slater determinants differing by at most two single-particle states, which is the last of the Condon-Slater rules of Section 3.13. The matrix is therefore *banded* in the excitation level, with bandwidth two, and overwhelmingly sparse: for six particles the fraction of non-zero blocks is already small, and it falls as the space grows. Chapter 1 called this the structure that Lanczos exploits; here we see where it comes from.

In a Hartree-Fock basis.

If the single-particle states are the solutions of the Hartree-Fock equations rather than the eigenstates of $\hat{h}_0$, one more set of elements vanishes. Brillouin's theorem, met in Section 3.13, states that

$$\langle \Phi_0 | \hat{H} | \Phi_i^a \rangle = 0, \quad (5.15)$$

so the reference determinant does not couple to the singles at all,

	$0p0h$	$1p1h$	$2p2h$	$3p3h$	$4p4h$	$5p5h$	$6p6h$
$0p0h$	$\tilde{\times}$	0	$\tilde{\times}$	0	0	0	0
$1p1h$	0	$\tilde{\times}$	$\tilde{\times}$	$\tilde{\times}$	0	0	0
$2p2h$	$\tilde{\times}$	$\tilde{\times}$	$\tilde{\times}$	$\tilde{\times}$	$\tilde{\times}$	0	0
$3p3h$	0	$\tilde{\times}$	$\tilde{\times}$	$\tilde{\times}$	$\tilde{\times}$	$\tilde{\times}$	0
$4p4h$	0	0	$\tilde{\times}$	$\tilde{\times}$	$\tilde{\times}$	$\tilde{\times}$	$\tilde{\times}$
$5p5h$	0	0	0	$\tilde{\times}$	$\tilde{\times}$	$\tilde{\times}$	$\tilde{\times}$
$6p6h$	0	0	0	0	$\tilde{\times}$	$\tilde{\times}$	$\tilde{\times}$

(5.16)

where the tildes are a reminder that a unitary transformation of the single-particle basis has changed every matrix element. The correlation energy is then carried entirely by the doubles at leading order, which is the starting point for both many-body perturbation theory and coupled-cluster theory.

The programme `fci.py` builds this structure explicitly. For the pairing model with four levels and four particles the full space has $\binom{8}{4} = 70$ determinants, distributed as 1, 16, 36, 16, 1 over the excitation levels 0 to 4, and the computed block table is

	$0p0h$	$1p1h$	$2p2h$	$3p3h$	$4p4h$
$0p0h$	$\times$	0	$\times$	0	0
$1p1h$	0	$\times$	0	$\times$	0
$2p2h$	$\times$	0	$\times$	0	$\times$
$3p3h$	0	$\times$	0	$\times$	0
$4p4h$	0	0	$\times$	0	$\times$

(5.17)

Even the odd blocks are empty here, because the pairing interaction moves particles two at a time and can never change the number of broken pairs. The symmetry has done more work than Brillouin's theorem would.

5.5 The exponential wall

Full configuration interaction is exact, and that is precisely its problem. With N particles distributed over n single-particle states the number of Slater determinants is

$$\dim \mathcal{H} = \binom{n}{N}, \quad (5.18)$$

which is the count of bit patterns with N bits set out of n , in the representation of Section 3.14. This is a combinatorial expression, but combinatorial growth turns into exponential growth as soon as the particle number scales with the size of the basis. At half filling, $N = n/2$, Stirling's formula gives

$$\binom{n}{n/2} \sim \frac{2^n}{\sqrt{\pi n/2}} \sim \frac{2^n}{\sqrt{n}}, \quad (5.19)$$

so the dimension doubles every time one single-particle state is added. The program `fci.py` tabulates both sides:

n	$\binom{n}{n/2}$	$2^n/\sqrt{n}$
10	2.520×10^2	3.238×10^2
20	1.848×10^5	2.345×10^5
40	1.378×10^{11}	1.738×10^{11}
80	1.075×10^{23}	1.352×10^{23}
160	9.205×10^{46}	1.155×10^{47}

The estimate is accurate to within a factor of order one over five decades of n and twenty-five orders of magnitude of dimension.

A realistic example.

Consider oxygen-16 and let the neutrons occupy the four lowest major oscillator shells, the $0s$, $0p$, $1s0d$ and $1p0f$ shells, which together hold $2 + 6 + 12 + 20 = 40$ single-particle states for each species. The eight neutrons can be distributed in

$$\binom{40}{8} = 7.690 \times 10^7 \quad (5.20)$$

ways, and since the protons can be distributed independently the full proton-neutron space has

$$\left(\binom{40}{8}\right)^2 = 5.914 \times 10^{15} \quad (5.21)$$

Slater determinants. Rotational and parity symmetry reduce this by orders of magnitude, but adding a single further major shell, or relaxing the truncation, restores the growth immediately. Direct diagonalisation by Householder and QL, Sections 1.24 and 1.25, is limited to matrices of dimension around 10^5 because it stores and modifies the whole matrix. Lanczos, Section 1.26, needs only the product $\mathbf{H}\mathbf{v}$ and reaches perhaps 10^{10} on a large machine. Beyond that no rearrangement of the linear algebra helps: the expansion itself must be truncated.

The same wall in the models of Chapter 4.

Each of the models of the previous chapter meets Eq. (5.19) in its own way.

- The *Lipkin model* with N particles in two N -fold degenerate levels has $\binom{2N}{N}$ determinants in total and $\binom{N}{N/2} \sim 2^N/\sqrt{N}$ in the maximal-quasispin multiplet. That the SU(2) structure of Section 4.1 collapses the problem to an $(N+1)$ -dimensional matrix is a symmetry accident, not a general escape.
- The *pairing model* of Section 4.2 has $\binom{2n}{n}$ determinants for n doubly degenerate levels, of which only $\binom{n}{n/2}$ have seniority zero. Again a symmetry – here the conservation of the number of broken pairs – takes the square root of the dimension, and again it is special to the model.
- The *Hubbard model* of Section 4.3 has a local Hilbert space of dimension four, so L sites give 4^L states in total. At fixed particle number this becomes $\binom{L}{N_\uparrow}\binom{L}{N_\downarrow}$, Eq. (4.21), which is again exponential at any extensive filling.
- A *spin chain*, the limit of the Hubbard model at large U , has no combinatorial disguise at all: $\mathcal{H} = \bigotimes_{i=1}^L \mathbb{C}^2$ has dimension 2^L by construction.

The pattern is worth stating plainly. Combinatorial counting and tensor-product structure are two faces of the same growth; the tensor product is fundamental, the binomial coefficient is what survives after a conservation law has been imposed. A conserved quantity can remove a factor, but never the exponential.

Why the wall is not the end of the story. The dimension of $\mathcal{H}$ counts the coefficients an *arbitrary* state would need. Physical ground states are not arbitrary. For local Hamiltonians with a gap, the entanglement entropy across a cut grows with the area of the cut rather than with the volume it encloses, and states obeying such an area law are described to high accuracy by a number of parameters polynomial in L . The Schmidt decomposition of Section 1.30 is the one-cut version of this statement, and the singular values decaying rapidly is exactly what makes the truncation work. Applying it at every bond of a chain gives a matrix product state,

$$|\psi\rangle = \sum_{s_1 \dots s_L} A_1^{s_1} A_2^{s_2} \dots A_L^{s_L} |s_1 s_2 \dots s_L\rangle, \quad (5.22)$$

whose cost is polynomial in the bond dimension D ; this is the density-matrix renormalisation group. The same idea appears elsewhere in this book in different clothing: a neural-network quantum state replaces the exponentially many amplitudes by a polynomially parameterised function of the occupation pattern, and quantum computing replaces them by n qubits whose Hilbert space is exponential by construction. Full CI is the reference against which all of these compressions are measured, which is why we treat it first.

5.6 Full CI for the pairing model

The pairing Hamiltonian of Section 4.2,

$$\hat{H} = \xi \sum_{p\sigma} (p-1) \hat{a}_{p\sigma}^\dagger \hat{a}_{p\sigma} - \frac{1}{2} g \sum_{pq} \hat{a}_{p+}^\dagger \hat{a}_{p-}^\dagger \hat{a}_{q-} \hat{a}_{q+}, \quad (5.23)$$

is the ideal first test. It is small enough to diagonalise in the full space, and Chapter 4 already gave the exact answer in the seniority-zero subspace, so the two calculations can be checked against each other.

Take four doubly degenerate levels and four particles, $\xi = 1$. The full space has $\binom{8}{4} = 70$ determinants, distributed over the excitation levels as 1, 16, 36, 16, 1. The seniority-zero subspace used in Chapter 4 has only $\binom{4}{2} = 6$ states. Diagonalising the 70×70 matrix at $g = 1$ gives $E_0 = 0.63554847$, identical to the value in Table 4.1 to all digits shown. The determinants with broken pairs are present in the basis but carry no amplitude in the ground state, because the interaction cannot create a broken pair from a paired configuration. The symmetry that we imposed by hand in Chapter 4 is recovered automatically by the diagonalisation.

This is the general situation. A symmetry of $\hat{H}$ block-diagonalises the FCI matrix, and one may either exploit it, gaining a smaller matrix and a labelled spectrum, or ignore it, gaining a simpler code at the cost of diagonalising a larger matrix. Production codes exploit it; the program accompanying this chapter does not, precisely so that the two answers can be compared.

5.7 Full CI for the Hubbard model

The Hubbard model of Section 4.3 is the harder test, and the more instructive one. We take a ring of $L = 6$ sites at half filling, $N_\uparrow = N_\downarrow = 3$, and work in the momentum basis, where the hopping term is diagonal,

$$\hat{H} = \sum_{k\sigma} \epsilon_k \hat{c}_{k\sigma}^\dagger \hat{c}_{k\sigma} + \frac{U}{L} \sum_{kk'q} \hat{c}_{k+q\uparrow}^\dagger \hat{c}_{k'-q\downarrow}^\dagger \hat{c}_{k'\downarrow} \hat{c}_{k\uparrow}, \quad \epsilon_k = -2t \cos k, \quad (5.24)$$

with $k = 2\pi m/L$. In this basis the non-interacting ground state is a single determinant – the Fermi sea – so the machinery of Section 5.2 applies directly with $|\Phi_0\rangle$ the filled Fermi sea. The dimension is $\binom{6}{3}^2 = 400$, with the excitation levels populated as 1, 18, 99, 164, 99, 18, 1.

The interesting quantity is how much of the correlation energy survives truncation as U/t grows:

U/t	E_{ref}	CIS	CID	CISD	FCI
1	-6.500000	-6.50000000	-6.59980159	-6.59989778	-6.60115829
2	-5.000000	-5.00000000	-5.38877703	-5.39026458	-5.40945685
4	-2.000000	-2.00000000	-3.41192866	-3.43018837	-3.66870618
8	4.000000	2.48338852	-0.36779217	-1.05692183	-2.04813089

The doubles recover 98.66% of the correlation energy at $U/t = 1$ but only 72.22% at $U/t = 8$. Two features deserve comment.

First, the reference does not couple to the singles at all: the interaction conserves total momentum, and a single excitation changes it, so $\langle \Phi_0 | \hat{H} | \Phi_i^a \rangle = 0$ exactly, for every U . This is Brillouin's theorem, Eq. (5.15), enforced here by a symmetry rather than by a variational condition – the plane-wave determinant is the Hartree-Fock solution of a translationally invariant problem. The CIS matrix is therefore block diagonal, and its lowest eigenvalue is the smaller of the reference energy and the lowest eigenvalue of the singles block alone. For $U/t \leq 4$ the reference wins and CIS reproduces it exactly. At $U/t = 8$ it does not: the singles block, strongly mixed by the large interaction, has an eigenvalue at 2.483, below the reference at 4.000. The number in the table is that eigenvalue, and the state it belongs to is orthogonal to $|\Phi_0\rangle$. What CIS reports as a ground state is by then not a correlated version of the reference at all.

Second, the deterioration with U/t is the signature of *strong correlation*. At small U the ground state is a single determinant with small corrections and any method built on that picture works. At large U the ground state is a superposition of very many determinants with comparable weight, the reference is a poor starting point, and a truncation at any fixed excitation level fails. This is where the tensor-network and quantum-computing approaches of the later chapters become attractive, and it is why the Hubbard model at intermediate U/t remains a benchmark problem.

5.8 Truncating the expansion

Since the full expansion is out of reach for all but the smallest systems, we truncate it. The natural truncation is by excitation level: keep all determinants up to n particle-hole pairs and discard the rest. This gives the familiar hierarchy

$$|\Psi_{\text{CIS}}\rangle = (1 + \hat{C}_1) |\Phi_0\rangle, \quad (5.25)$$

$$|\Psi_{\text{CID}}\rangle = (1 + \hat{C}_2) |\Phi_0\rangle, \quad (5.26)$$

$$|\Psi_{\text{CISD}}\rangle = (1 + \hat{C}_1 + \hat{C}_2) |\Phi_0\rangle, \quad (5.27)$$

and so on to CISDT, CISDTQ, until at $\hat{C}_N$ we are back at full CI. Each of these is again a variational calculation in a subspace, so each lowest eigenvalue is an upper bound to E_0 and the bounds decrease monotonically as the truncation is relaxed. The cost, for a basis of n states and N particles, is roughly $\binom{N}{k} \binom{n-N}{k}$ determinants at excitation level k , that is polynomial of degree $2k$ – which is why CISD, at $\mathcal{O}(n^6)$, is affordable and CISDTQ, at $\mathcal{O}(n^{10})$, is not.

For the pairing model with four levels and four particles the truncated spaces have dimensions 1, 17, 53, 69, 70, and the energies are

g	E_{ref}	CIS	CID	CISD	FCI
0.25	1.750000	1.75000000	1.73050700	1.73050700	1.73045985
0.50	1.500000	1.50000000	1.41759645	1.41759645	1.41677428
1.00	1.000000	1.00000000	0.64890655	0.64890655	0.63554847
2.00	0.000000	0.00000000	-1.35208670	-1.35208670	-1.48965216

The singles change nothing at all, and CISD coincides with CID: the pairing interaction moves pairs and cannot break them, so it has no matrix element between the reference and a $1p\text{-}1h$ state, nor between a $1p\text{-}1h$ state and the reference through any intermediate. The doubles carry essentially the whole correlation energy – 99.8% at $g = 0.25$, still 96.3% at $g = 1$, but only 90.8% at $g = 2$. The pattern is the same as in the Hubbard model: truncation at fixed excitation level works while the reference is good and degrades as the coupling grows.

5.9 Size consistency

There is a defect of truncated CI that is more serious than the loss of accuracy, because it does not go away as the basis is improved. Consider two identical systems A and B , infinitely far apart and non-interacting. The Hamiltonian is $\hat{H}_A + \hat{H}_B$ and the exact ground state is the product $|\Psi_A\rangle|\Psi_B\rangle$, so the exact energy must satisfy

$$E(AB) = E(A) + E(B) = 2E(A). \quad (5.28)$$

A method that reproduces this is called *size consistent*.

Full CI is size consistent, since the product state lies in the full space. Truncated CI is not. The reason is immediate: if A and B are each doubly excited, the combined state is a *quadruple* excitation of the combined reference, and CID has no room for it. The product of two doubles is a quadruple, and truncating at a fixed excitation level is not a multiplicatively closed operation.

The program `fci.py` demonstrates this on two non-interacting copies of a two-level, two-particle pairing system:

g	$2E(A)$	$E(AB), \text{FCI}$	$E(AB), \text{CID}$	CID error
0.25	-0.26556444	-0.26556444	-0.26550480	5.96×10^{-5}
0.50	-0.56155281	-0.56155281	-0.56066017	8.93×10^{-4}
1.00	-1.23606798	-1.23606798	-1.22474487	1.13×10^{-2}

Full CI passes to machine precision; CID fails, and the failure grows with the coupling. With M subsystems instead of two the error grows with M , so the energy per particle computed with truncated CI drifts as the system is made larger. For an extended system – nuclear matter, the electron gas, a large molecule – this is fatal.

The repair is to make the ansatz multiplicative rather than additive. Replacing $1 + \hat{C}$ by $\exp(\hat{T})$ with $\hat{T} = \hat{T}_1 + \hat{T}_2$ gives

$$|\Psi_{\text{CCSD}}\rangle = e^{\hat{T}_1 + \hat{T}_2} |\Phi_0\rangle = (1 + \hat{T}_1 + \hat{T}_2 + \frac{1}{2}\hat{T}_2^2 + \dots) |\Phi_0\rangle, \quad (5.29)$$

in which the quadruple excitation $\frac{1}{2}\hat{T}_2^2$ appears automatically, with its coefficient fixed by the doubles rather than determined independently. For two non-interacting subsystems $e^{\hat{T}} = e^{\hat{T}_A} e^{\hat{T}_B}$ and size consistency is exact. This is coupled-cluster theory, and Eq. (5.29) is the reason it, rather than truncated CI, is the workhorse of modern many-body physics. We shall develop it properly in a later chapter; the point here is that its central idea is a response to a defect that full CI lets us see clearly.

5.10 Truncations and the other many-body methods

We can now do what the introduction to this chapter promised, and read the standard many-body methods as truncations of Eq. (5.13). Write the projected equations explicitly. Multiplying $(\hat{H} - E)|\Psi_0\rangle = 0$ from the left by $\langle\Phi_0|$ and by $\langle\Phi_H^P|$ gives

$$E = \langle \Phi_0 | \hat{H} | \Phi_0 \rangle + \sum_{P'H' \neq 0} \langle \Phi_0 | \hat{H} | \Phi_{H'}^{P'} \rangle C_{H'}^{P'}, \quad (5.30)$$

$$(E - \langle \Phi_H^P | \hat{H} | \Phi_H^P \rangle) C_H^P = \sum_{P'H' \neq PH} \langle \Phi_H^P | \hat{H} | \Phi_{H'}^{P'} \rangle C_{H'}^{P'}. \quad (5.31)$$

The first equation says that the correlation energy comes entirely from the coupling of the reference to the doubles, since only those have a non-zero matrix element with $|\Phi_0\rangle$. The second is a set of coupled equations for the coefficients, which can be solved by iteration: guess the C 's, compute the right-hand side, update, repeat. Every method in the remaining chapters is a prescription for which terms of Eq. (5.31) to keep and how to solve it.

- **Hartree-Fock theory** keeps only C_0 and instead varies the single-particle basis so that $\langle \Phi_0 | \hat{H} | \Phi_i^a \rangle = 0$ – the Brillouin condition of Eq. (5.15). It is the optimisation of the reference rather than of the expansion, and it fixes the basis in which every subsequent truncation is made. In the language of Eq. (5.16), Hartree-Fock chooses the unitary transformation that empties the $0p0h$ - $1p1h$ block.
- **The Tamm-Dancoff approximation** restricts the space to the $1p$ - $1h$ block alone and diagonalises it. It is CIS applied to excited states, and it gives the excitation spectrum of the reference at the cheapest non-trivial level.
- **The random-phase approximation** enlarges TDA by allowing the ground state to contain $2p$ - $2h$ correlations, so that the excitation operator carries both $1p$ - $1h$ and $1h$ - $1p$ pieces. It is a partial summation of the doubles block, performed to all orders in the coupling but only for a selected class of terms.
- **Many-body perturbation theory** solves Eq. (5.31) by expanding in powers of the interaction rather than by iterating to convergence. Setting E on the left-hand side to the reference energy and $C_{H'}^{P'}$ on the right to zero gives the first-order coefficients; substituting these back gives the second-order energy, and so on. The order in perturbation theory is the number of times the right-hand side has been substituted into itself.
- **Coupled-cluster theory** keeps the same truncation in excitation level as CISD but writes the wave function multiplicatively, Eq. (5.29). The equations look like Eq. (5.31) with additional nonlinear terms, and it is those nonlinear terms that restore size consistency.

Seen this way the methods are not a list of unrelated techniques but a set of answers to one question: given that we cannot keep all of Eq. (5.5), which parts do we keep, and do we keep them additively or multiplicatively? Full CI is the fixed point that lets us measure each answer. Whenever a model is small enough for Eq. (5.13) to be solved exactly – the pairing model, the Lipkin model, a short Hubbard ring – we should solve it, and use the result to calibrate the approximation before applying it to a system where no calibration is possible.

5.11 Practical considerations

Three remarks on implementation, which anticipate the code discussed in the companion notebook.

Never store the matrix.

The FCI matrix is sparse, with the band structure of Eq. (5.14), and for interesting dimensions it cannot be stored. The Lanczos algorithm of Section 1.26 needs only the product $\mathbf{H}\mathbf{v}$, and that product is computed by looping over the determinants and over the non-zero matrix elements of $\hat{H}$, using the bit representation of Section 3.14 to find which determinant each term produces and with which sign. This is the single most important structural choice in an FCI code.

Reorthogonalise.

The Lanczos vectors lose orthogonality as soon as an eigenvalue has converged, and spurious copies of converged eigenvalues appear. Section 1.26 discussed the remedies; for FCI the usual choice is full or selective reorthogonalisation, at the price of storing the Lanczos vectors.

Use the symmetries.

Angular momentum, parity, particle number, total momentum and spin projection all commute with $\hat{H}$ and block-diagonalise the matrix. Working in a single symmetry block reduces the dimension by orders of magnitude and, just as importantly, labels the states. The pairing calculation of Section 5.6 illustrated both options – 70 determinants without the symmetry, 6 with it, the same ground-state energy.

5.12 Summary

Full configuration interaction is the statement that a many-body state is a superposition of all Slater determinants that the single-particle basis allows, together with the observation that finding the best superposition is a Hermitian eigenvalue problem, Eq. (5.13). Everything else in this chapter follows from that. The Condon-Slater rules make the matrix banded in the excitation level; a Hartree-Fock basis empties one more block; the combinatorics of Eq. (5.18) makes the matrix too large to build beyond a few tens of single-particle states; and the truncations that this forces upon us are precisely the approximations that the following chapters elevate into methods.

The program `fci.py` in `BookMaterial/Programs/` carries out every calculation quoted above: the pairing model in the full 70-dimensional space and in the seniority-zero subspace, the block structure of the Hamiltonian, the CIS/CID/CISD/FCI comparison for pairing and for the Hubbard ring, the size-consistency test, and the dimension tables of Section 5.5. The accompanying notebook `fci.ipynb` runs the same code cell by cell.

5.13 Exercises

Exercise 1: counting determinants

- Show that the number of Slater determinants with N particles in n single-particle states that are k particle-hole excitations of a given reference is $\binom{N}{k} \binom{n-N}{k}$, and verify that $\sum_k \binom{N}{k} \binom{n-N}{k} = \binom{n}{N}$.
- Evaluate this for $n = 8$, $N = 4$ and compare with the distribution 1, 16, 36, 16, 1 quoted in Section 5.4.
- Use Stirling's formula to derive Eq. (5.19), and determine the largest n for which $\binom{n}{n/2}$ stays below 10^{10} .

Answer to a).

Choosing which k of the N occupied states to empty can be done in $\binom{N}{k}$ ways, and which k of the $n - N$ empty states to fill in $\binom{n-N}{k}$ ways; the two choices are independent. Every determinant with N particles arises in this way exactly once, for exactly one value of k , namely the number of occupied states it does not share with the reference. Summing over k must therefore give the total, which is Vandermonde's identity

$$\sum_k \binom{N}{k} \binom{n-N}{k} = \sum_k \binom{N}{N-k} \binom{n-N}{k} = \binom{n}{N}. \quad (5.32)$$

Answer to b).

With $n = 8$ and $N = 4$ the products are $\binom{4}{k}^2 = 1, 16, 36, 16, 1$ for $k = 0, \dots, 4$, summing to $70 = \binom{8}{4}$. This is exactly the distribution that `fci.py` finds by classifying the 70 determinants of the pairing model by excitation level.

Answer to c).

Stirling's formula $m! \sim \sqrt{2\pi m} (m/e)^m$ gives

$$\binom{n}{n/2} = \frac{n!}{[(n/2)!]^2} \sim \frac{\sqrt{2\pi n} (n/e)^n}{2\pi (n/2) [(n/2e)^{n/2}]^2} = \frac{2^n}{\sqrt{\pi n/2}}, \quad (5.33)$$

which is Eq. (5.19) up to the constant $\sqrt{\pi/2} \approx 1.25$. Solving $\binom{n}{n/2} = 10^{10}$ numerically: $\binom{36}{18} = 9.08 \times 10^9$ and $\binom{38}{19} = 3.49 \times 10^{10}$, so $n = 36$ is the last even value below the limit. Every further pair of single-particle states multiplies the dimension by about four.

Exercise 2: which blocks vanish

Let $\hat{H} = \hat{H}_0 + \hat{H}_I$ with $\hat{H}_I$ a two-body operator.

- Show that $\langle \Phi_H^P | \hat{H} | \Phi_{H'}^{P'} \rangle = 0$ whenever the two determinants differ by more than two single-particle states, and deduce the band structure of Eq. (5.14).
- Suppose $\hat{H}$ also contains a genuine three-body force. How does the band structure change, and what does that do to the cost of a matrix-vector product?
- For the pairing Hamiltonian, Eq. (5.23), show that every block connecting an even to an odd excitation level vanishes, as in Eq. (5.17).

Answer to a).

Write both determinants in the occupation representation of Section 3.14. A two-body operator is a product of two creation and two annihilation operators, so it can change at most two occupation numbers from one to zero and two from zero to one. If the two determinants differ in more than four occupation numbers – that is, by more than two single-particle states – then in the Wick expansion of Section 3.13 at least one operator is left uncontracted in every term, and the matrix element is zero. Since an $np-nh$ and an $mp-mh$ determinant differ by at least $|n-m|$ single-particle states, the element vanishes for $|n-m| > 2$. The matrix is therefore banded with bandwidth two in the excitation level, which is Eq. (5.14).

Answer to b).

A three-body operator changes up to three occupations of each kind, so the bandwidth becomes three and the $0p0h$ row now reaches the $3p3h$ block. The number of non-zero elements per row grows from $\mathcal{O}(N^2(n-N)^2)$ to $\mathcal{O}(N^3(n-N)^3)$, so a matrix-vector product costs a further factor $N(n-N)$ – typically two to three orders of magnitude. This is why three-body forces, though physically important in nuclear physics, are usually normal-ordered with respect to the reference determinant and then truncated at the two-body level, a procedure whose justification rests on Section 3.5.

Answer to c).

The interaction $-\frac{1}{2}g \sum_{pq} \hat{a}_{p+}^\dagger \hat{a}_{p-}^\dagger \hat{a}_{q-} \hat{a}_{q+}$ removes a complete pair from level q and creates a complete pair in level p . Define the seniority s as the number of unpaired particles. Both the one-body term and the interaction leave s unchanged, so $[\hat{H}, \hat{s}] = 0$ and $\hat{H}$ is block diagonal in seniority. An excitation of odd order relative to the paired reference must leave at least two particles unpaired, so it has $s \geq 2$, whereas the reference and all even excitations reached from it by moving whole pairs have $s = 0$. Every even-odd matrix element therefore connects different seniorities and vanishes. Numerically the 70×70 matrix splits into a six-dimensional $s = 0$ block, a $s = 2$ block and so on, and the ground state lies in the first.

Exercise 3: full CI for the pairing model

Take four doubly degenerate levels with single-particle energies $\varepsilon_p = (p-1)\xi$, $p = 1, \dots, 4$, four particles, $\xi = 1$, and the Hamiltonian of Eq. (5.23).

- Set up the six seniority-zero configurations and write down the 6×6 Hamiltonian matrix.
- Diagonalise it for $g = -1, -0.5, 0, 0.5, 1$ and reproduce Table 4.1.
- Repeat the calculation in the full 70-dimensional space and confirm that the ground-state energy is unchanged. Inspect the eigenvector and verify that every determinant with a broken pair has zero amplitude.
- Plot the correlation energy against g and show that it is quadratic for small $|g|$, as second-order perturbation theory requires.

Answer to a).

Label a seniority-zero configuration by the pair of levels that are occupied. With four levels and two pairs there are $\binom{4}{2} = 6$ of them,

$$|12\rangle, |13\rangle, |14\rangle, |23\rangle, |24\rangle, |34\rangle, \quad (5.34)$$

where $|pq\rangle$ has pairs in levels p and q . The diagonal element of $|pq\rangle$ is $2\xi[(p-1) + (q-1)] - g$: the factor two because each level holds two particles, and $-g$ because the $p = q$ terms of the interaction give $-\frac{1}{2}g$ for each of the two occupied pairs. A single application of the interaction moves *one* pair, so two configurations are connected by $-\frac{1}{2}g$ if and only if they share exactly one occupied level; the three pairs of configurations that share no level, $|12\rangle$ – $|34\rangle$, $|13\rangle$ – $|24\rangle$ and $|14\rangle$ – $|23\rangle$, are not connected at all. With $\xi = 1$ and the ordering of Eq. (5.34),

$$\mathbf{H} = \begin{pmatrix} 2-g & -g/2 & -g/2 & -g/2 & -g/2 & 0 \\ -g/2 & 4-g & -g/2 & -g/2 & 0 & -g/2 \\ -g/2 & -g/2 & 6-g & 0 & -g/2 & -g/2 \\ -g/2 & -g/2 & 0 & 6-g & -g/2 & -g/2 \\ -g/2 & 0 & -g/2 & -g/2 & 8-g & -g/2 \\ 0 & -g/2 & -g/2 & -g/2 & -g/2 & 10-g \end{pmatrix}. \quad (5.35)$$

Answer to b).

Diagonalising Eq. (5.35) reproduces Table 4.1; at $g = 1$ the lowest eigenvalue is 0.63554847 and the reference energy is $E_{\text{ref}} = H_{11} = 1$.

Answer to c).

The class `PairingFCI` in `fci.py` builds the full space with no seniority restriction. Its dimension is $\binom{8}{4} = 70$, the excitation levels are populated as 1, 16, 36, 16, 1, and the lowest eigenvalue at $g = 1$ is 0.63554847, identical to b) to every digit. Projecting the ground-state eigenvector onto the 64 determinants with at least one broken pair gives zero to machine precision, as Exercise 2c) predicts.

Answer to d).

The reference is $|12\rangle$. Second-order perturbation theory in the interaction gives

$$E^{(2)} = \sum_{\Phi \neq \Phi_0} \frac{|\langle \Phi_0 | \hat{H}_I | \Phi \rangle|^2}{E_0^{(0)} - E_\Phi^{(0)}} = \frac{g^2}{4} \left(\frac{1}{-2} + \frac{1}{-4} + \frac{1}{-4} + \frac{1}{-6} \right) = -\frac{7}{24}g^2, \quad (5.36)$$

using the four configurations reachable from $|12\rangle$ by moving one pair, with unperturbed excitation energies 2, 4, 4, 6. A log-log plot of $|E_{\text{corr}}|$ against g has slope 2 for $|g| \lesssim 0.2$ and bends away above that, where the higher orders take over. At $g = 0.25$, Eq. (5.36) gives -0.01823 against the exact -0.01954 ; at $g = 0.5$ it gives -0.0729 against -0.0832 , and at $g = 1$ only -0.2917 against -0.3645 .

Exercise 4: truncations and size consistency

- Using the code of the previous exercise, restrict the basis to excitation levels ≤ 1 and ≤ 2 and reproduce the CIS and CID columns of Section 5.8. Explain why CIS returns exactly the reference energy.
- Build a system of two identical, non-interacting two-level pairing systems. Show that full CI gives exactly twice the energy of one subsystem while CID does not, and tabulate the error for $g = 0.25, 0.5, 1$.
- Identify the configuration that CID is missing, and show that including quadruple excitations restores the exact result.
- Show that the ansatz $\exp(\hat{T}_2)|\Phi_0\rangle$, with the same number of independent parameters as CID, is size consistent for this system.

Answer to a).

In a truncated space the eigenvalue problem is Eq. (5.13) restricted to the corresponding rows and columns. The dimensions for the cuts at 0, 1, 2, 3, 4 are 1, 17, 53, 69, 70. CIS returns the reference energy because $\langle \Phi_0 | \hat{H} | \Phi_i^a \rangle = 0$: a single excitation breaks a pair and so changes the seniority, which the pairing Hamiltonian conserves (Exercise 2c). The reference is therefore decoupled from the entire singles block and remains an exact eigenstate of the truncated matrix. Its energy is also the lowest diagonal element, and the singles block lies far above it for the couplings considered, so it is the CIS ground state.

Answer to b).

Let each subsystem have two levels at 0 and ξ , holding one pair. The combined system then has four levels with energies 0, ξ , 0, ξ and two pairs, one per subsystem, with the interaction acting only within a subsystem. The class `NonInteractingCopies` gives

g	$2E(A)$	$E(AB)_{\text{FCI}}$	$E(AB)_{\text{CID}}$	CID error
0.25	-0.26556444	-0.26556444	-0.26550480	5.96×10^{-5}
0.50	-0.56155281	-0.56155281	-0.56066017	8.93×10^{-4}
1.00	-1.23606798	-1.23606798	-1.22474487	1.13×10^{-2}

(5.37)

Full CI is exact to machine precision; the CID error grows by roughly an order of magnitude each time g is doubled.

Answer to c).

Writing $\hat{A}_2^X$ for the operator that excites the pair of subsystem X , the exact product state is

$$(1 + c\hat{A}_2^A)(1 + c\hat{A}_2^B)|\Phi_0\rangle = (1 + c\hat{A}_2^A + c\hat{A}_2^B + c^2\hat{A}_2^A\hat{A}_2^B)|\Phi_0\rangle. \quad (5.38)$$

The last term excites both subsystems at once and is a $4p$ - $4h$ excitation of the combined reference; it is absent from CID. Adding the quadruples – which for this small system means going to the full space – recovers the exact energy, and the coefficient of the quadruple comes out equal to c^2 , the product of the two double amplitudes.

Answer to d).

With $\hat{T}_2 = t(\hat{A}_2^A + \hat{A}_2^B)$ the two terms commute, since they act on disjoint sets of orbitals, so

$$e^{\hat{T}_2} = e^{t\hat{A}_2^A} e^{t\hat{A}_2^B} = (1 + t\hat{A}_2^A)(1 + t\hat{A}_2^B), \quad (5.39)$$

the last step because $(\hat{A}_2^X)^2 = 0$ – one cannot excite the same pair twice. The exponential ansatz therefore generates exactly the product state of Eq. (5.38) with one parameter per subsystem, the same count as CID, and the energy is additive by construction. This is coupled cluster with doubles in miniature, and it shows that the cure for size inconsistency costs nothing in parameters – only a change from a linear to a multiplicative ansatz.

Exercise 5: the Hubbard ring by full CI

Take a ring of $L = 6$ sites at half filling with $t = 1$, in the momentum basis of Eq. (5.24).

- Verify that the dimension is $\binom{6}{3}^2 = 400$ and that the excitation levels are populated as 1, 18, 99, 164, 99, 18, 1.
- Compute the exact ground-state energy for $U/t = 1, 2, 4, 8$ and compare with the CID result.
- Show that $\langle\Phi_0|\hat{H}|\Phi_i^a\rangle = 0$ for every single excitation, and explain this using conservation of total momentum.
- At $U/t = 8$ the CIS energy lies *below* the reference energy even though the reference does not couple to the singles. Explain how this is possible and what it says about the reliability of a variational bound.
- Compare the exact energies at large U with the Heisenberg limit of Section 4.3.

Answer to a).

There are $\binom{6}{3} = 20$ ways to place three electrons of each spin, and the two species are independent, so the dimension is 400. Counting for each determinant the number of occupied momenta outside the Fermi sea, separately for each spin, and adding the two counts gives 1, 18, 99, 164, 99, 18, 1, which sums to 400. The distribution is symmetric because a kp - kh excitation of the Fermi sea is a $(6-k)p$ - $(6-k)h$ excitation of the completely inverted determinant.

Answer to b).

The values are those of Section 5.7. Relative to the reference, the doubles recover 98.66%, 94.95%, 84.61% and 72.22% of the correlation energy at $U/t = 1, 2, 4, 8$.

Answer to c).

The interaction is $\frac{U}{L} \sum_{kk'q} \hat{c}_{k+q\uparrow}^\dagger \hat{c}_{k'-q\downarrow}^\dagger \hat{c}_{k'\downarrow} \hat{c}_{k\uparrow}$, in which the momenta of the two created operators sum to those of the two annihilated ones; total crystal momentum is therefore conserved. A single excitation $i \rightarrow a$ changes the total momentum by $k_a - k_i \neq 0$ and so cannot be reached from the Fermi sea. The one-body part is diagonal in k and cannot connect them either. Evaluating the couplings numerically gives zero to machine precision for every U . Note that this is Brillouin's theorem, Eq. (5.15), obtained here from a symmetry rather than from a variational condition: for a translationally invariant problem the plane-wave determinant is the Hartree-Fock solution.

Answer to d).

Because the coupling vanishes, the CIS matrix is block diagonal: a 1×1 block holding E_{ref} and an 18×18 singles block. Its lowest eigenvalue is the smaller of the two numbers. At $U/t \leq 4$ the reference wins and CIS reproduces it exactly. At $U/t = 8$ the singles block, strongly mixed by the large interaction, has an eigenvalue at 2.483, which lies below $E_{\text{ref}} = 4.000$, and the variational principle selects the corresponding state – one that is exactly orthogonal to $|\Phi_0\rangle$. The number is a legitimate upper bound to E_0 , but it is not a correlated correction to the reference; the calculation has silently changed which state it describes. The moral is that a lower variational energy does not by itself mean a better description of the state one set out to compute, and that the overlap with the reference should always be monitored alongside the energy.

Answer to e).

For $U/t \gg 1$ the half-filled Hubbard model maps onto a Heisenberg chain with $J = 4t^2/U$, Eq. (4.23), whose ground-state energy per site tends to $-J(\ln 2 - \frac{1}{4})$ in the thermodynamic limit. At $U/t = 8$ and $L = 6$ this estimate gives about -1.3 against the exact -2.05 . The discrepancy is expected: the mapping neglects terms of order t^4/U^3 , the ring is short, and $U/t = 8$ is not yet deep in the strong-coupling regime. Repeating the comparison at $U/t = 20$ and 40 shows the two converging.

Exercise 6: the exponential wall in practice

- A shell-model code stores the Hamiltonian matrix in double precision. For a dimension of 10^6 with on average 10^3 non-zero elements per row, and a 4-byte column index per element, how much memory does the sparse matrix need? Compare with storing the full matrix.
- How large a single-particle space can be treated at half filling if a Lanczos code can handle 10^{10} determinants?
- A matrix product state with bond dimension D on L sites has about LD^2d parameters, with d the local dimension. For a spin chain with $d = 2$, $L = 100$ and $D = 200$, compare this with 2^L .
- Discuss what must be true of the physical state for the comparison in c) to be a fair one.

Answer to a).

The sparse matrix holds $10^6 \times 10^3 = 10^9$ elements at $8 + 4 = 12$ bytes each, or 12 GB – large but routine. The full matrix would need 8×10^{12} bytes, 8 TB, and diagonalising it by Householder and QL would take $\mathcal{O}(10^{18})$ floating-point operations. This is the practical content of Section 5.11: the dense algorithms of Sections 1.24 and 1.25 are unusable here, and even the sparse matrix is best not stored at all but regenerated on the fly inside the matrix-vector product.

Answer to b).

From Eq. (5.19), $2^n / \sqrt{n} = 10^{10}$ gives $n \approx 38$, and the exact condition $\binom{n}{n/2} \leq 10^{10}$ gives $n = 36$. A Lanczos calculation at the limit of current hardware therefore reaches roughly 36–40 single-particle states at half filling. Away from half filling the binomial coefficient is much smaller and the reach is greater, which is why calculations near closed shells are so much easier than calculations in the middle of a shell.

Answer to c).

The matrix product state has about $100 \times 200^2 \times 2 = 8 \times 10^6$ parameters, against $2^{100} = 1.27 \times 10^{30}$ amplitudes for a general state – a compression by twenty-four orders of magnitude.

Answer to d).

The comparison is fair only if the state of interest is actually representable at that bond dimension. The Schmidt spectrum across every bond must decay fast enough that keeping D singular values costs little, which is the truncation error of the Eckart-Young theorem in Section 1.30 applied bond by bond. This holds for the ground states of gapped one-dimensional local Hamiltonians, which obey an area law. For critical chains the required D grows polynomially with L ; for volume-law states – generic excited states, or the state produced by a quantum quench after a time proportional to L – no polynomial D suffices and the exponential wall stands undiminished. The same caveat applies to every other compressed representation, including the neural-network and quantum-circuit ansatzes of the later chapters: each buys its efficiency by restricting the class of states it can describe, and full configuration interaction remains the only method that makes no such restriction at all.

Part III

Stochastic methods

Part IV
Machine Learning based approaches

Part V
Quantum Computing

References

1. Gene H. Golub and Charles F. Van Loan. *Matrix Computations*. Johns Hopkins University Press, Baltimore, 3rd edition, 1996.
2. Jules W. Moskowitz and Malvin H. Kalos. A new look at correlation energy in atomic and molecular systems. *International Journal of Quantum Chemistry*, 20(5):1107–1119, 1981.
3. H. J. Lipkin, N. Meshkov, and A. J. Glick. Validity of many-body approximation methods for a solvable model. *Nuclear Physics*, 62:188–198, 1965.
4. Francesco Calogero. Solution of the one-dimensional n -body problems with quadratic and/or inversely quadratic pair potentials. *Journal of Mathematical Physics*, 12:419–436, 1971.
5. Bill Sutherland. Quantum many-body problem in one dimension: ground state. *Journal of Mathematical Physics*, 12:246–250, 1971.

Index

- Calogero-Sutherland model, 101
- Cholesky decomposition
 - of two-body integrals, 38
- Cholesky factorisation, 15
- condition number, 11
- configuration interaction, 111
 - truncated, 117
- conjugate gradient method, 17
- contraction, 84

- density fitting, 38
- determinant
 - from LU decomposition, 58
 - from singular values, 58

- Eckart-Young theorem, 34
- eigenvalue problem, 21
- exponential wall, 114

- Francis' algorithm, 24

- Gauss-Seidel method, 17
- Gaussian elimination, 12

- Hilbert space
 - dimension, 114
- Householder's method, 22
- Hubbard model, 99

- iterative methods, 16

- Jacobi's method, 17

- Krylov space, 26

- Lanczos' algorithm, 26
- Lipkin model, 95
- Lippmann-Schwinger equation, 20

- LU decomposition, 14

- natural orbitals, 37
- norm
 - matrix, 11
 - vector, 10
- normal ordering, 83

- occupation numbers, 37

- pairing model, 98
- pivoting, 13
- power method, 22
- pseudoinverse, 33

- QL algorithm, 24

- rank, 33

- Schmidt decomposition, 34
- Schrodinger equation
 - by diagonalisation, 29
- similarity transformation, 21
- singular value decomposition, 31
- size consistency, 118
- Slater determinant
 - efficient evaluation, 54
 - spin factorisation, 57
- spline interpolation, 19
- successive over-relaxation, 17

- tridiagonal matrix, 16

- Vandermonde matrix, 20

- Wick's theorem, 87
 - generalised, 89